Gov 50: 10. Summarizing Bivariate Relationships

Matthew Blackwell

Harvard University

Roadmap

1. Z-scores and standardization

2. Correlation

3. Writing our own functions

1/ Z-scores and standardization

COVID vaccination rates and votes

```
library(tidyverse)
library(gov50data)
covid_votes
```

```
## # A tibble: 3,114 x 8
##     fips  county state one_dose_5plus_pct one_dose_65plus_pct
##     <chr> <chr>  <chr>              <dbl>               <dbl>
##  1 26039 Crawf~ MI                  55.7                77.3
##  2 40015 Caddo~ OK                  83.3                95
##  3 17007 Boone~ IL                  71.1                94.5
##  4 12055 Highl~ FL                  68.9                93.7
##  5 34029 Ocean~ NJ                  71                  95
##  6 01067 Henry~ AL                  58.5                85.5
##  7 27037 Dakot~ MN                  81                  95
##  8 27115 Pine ~ MN                  56.5                85
##  9 51750 Radfo~ VA                  41.5                73.8
## 10 22009 Avoye~ LA                  59.7                80.1
## # i 3,104 more rows
## # i 3 more variables: booster_5plus_pct <dbl>,
## #   dem_pct_2000 <dbl>, dem_pct_2020 <dbl>
```

Is 60% vaccinated a lot?
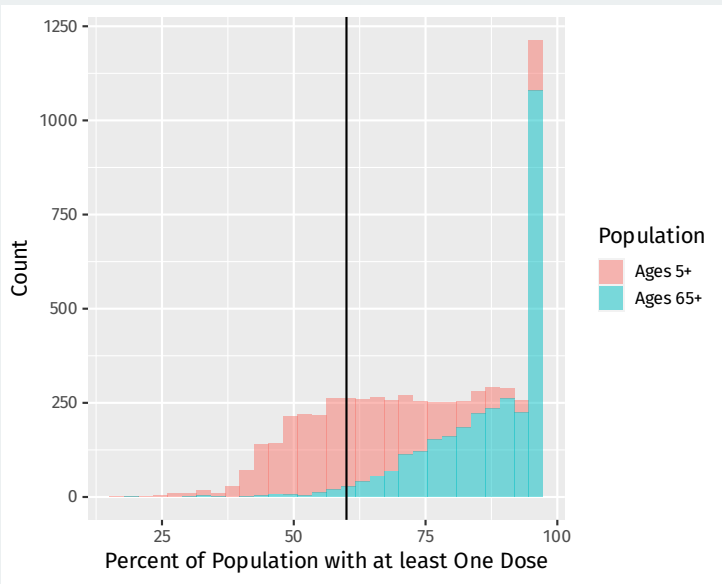
How large is large?

- How large 60% vaccinated is depends on the distribution!

How large is large?

- How large 60% vaccinated is depends on the distribution!
  - Clear to see from the histogram

How large is large?

- How large 60% vaccinated is depends on the distribution!
  - Clear to see from the histogram
  - Middling for the 5+ group, but very low for the 65+ group.

How large is large?

- How large 60% vaccinated is depends on the distribution!
  - Clear to see from the histogram
  - Middling for the 5+ group, but very low for the 65+ group.

- Can we transform the values of our variables to be **common units**?

How large is large?

- How large 60% vaccinated is depends on the distribution!
  - Clear to see from the histogram
  - Middling for the 5+ group, but very low for the 65+ group.

- Can we transform the values of our variables to be **common units**?

- Yes, with two transformations:

How large is large?

- How large 60% vaccinated is depends on the distribution!
  - Clear to see from the histogram
  - Middling for the 5+ group, but very low for the 65+ group.

- Can we transform the values of our variables to be **common units**?

- Yes, with two transformations:
 - **Centering**: subtract the mean of the variable from each value.

How large is large?

- How large 60% vaccinated is depends on the distribution!
  - Clear to see from the histogram
  - Middling for the 5+ group, but very low for the 65+ group.

- Can we transform the values of our variables to be **common units**?

- Yes, with two transformations:
  - **Centering**: subtract the mean of the variable from each value.
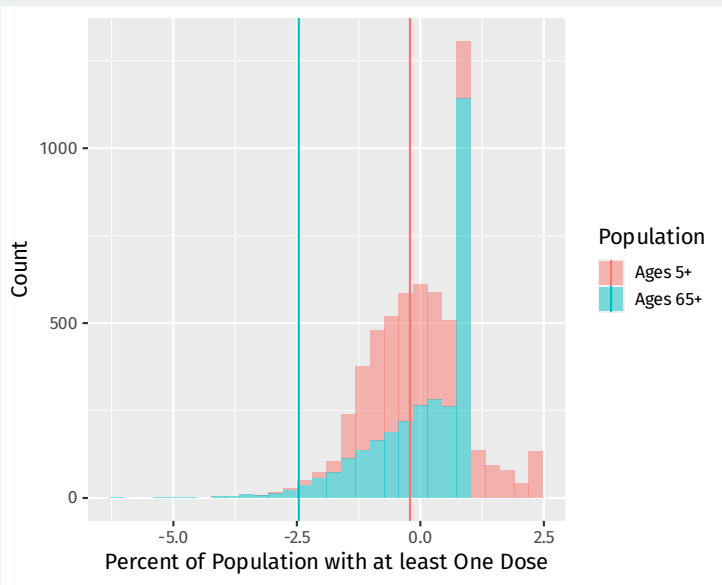 - **Scaling**: dividing deviations from the mean by the standard deviation.

Original distributions
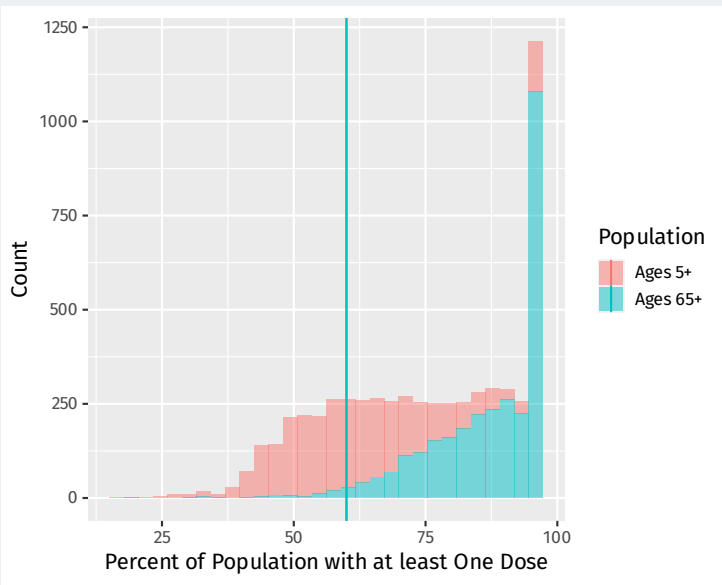
Centered distributions
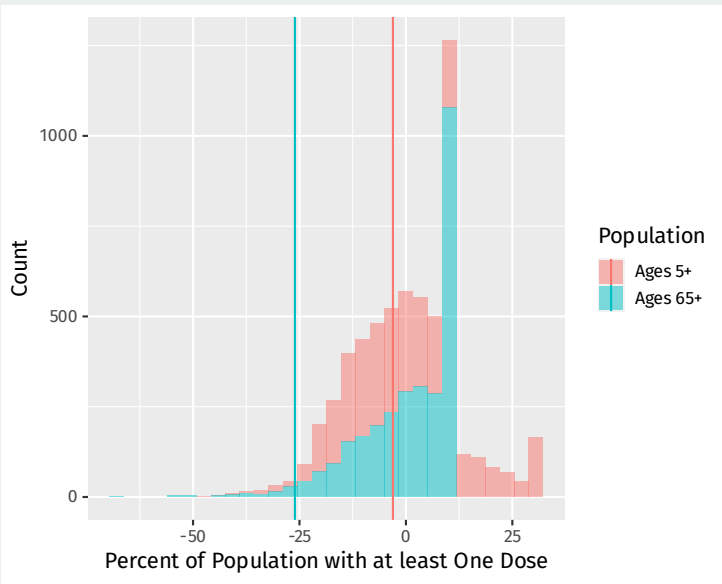
Centered and scaled distributions

Z-scores

- Centering tells us immediately if a value is above or below the mean.

Z-scores

- Centering tells us immediately if a value is above or below the mean.

- Scaling tells us how many standard deviations away from the mean it is.

Z-scores

- Centering tells us immediately if a value is above or below the mean.

- Scaling tells us how many standard deviations away from the mean it is.

- Combine them with the **z-score** transformation:

$$\text{z-score of } x_i = \frac{x_i - \text{mean of } x}{\text{standard deviation of } x}$$

Z-scores

- Centering tells us immediately if a value is above or below the mean.

- Scaling tells us how many standard deviations away from the mean it is.

- Combine them with the **z-score** transformation:

$$\text{z-score of } x_i = \frac{x_i - \text{mean of } x}{\text{standard deviation of } x}$$

- Useful heuristic: data more than 3 SDs away from mean are rare.

z-score example

```
covid_votes |>
  mutate(one_dose_centered = one_dose_5plus_pct -
          mean(one_dose_5plus_pct, na.rm = TRUE)) |>
  select(fips:state, one_dose_5plus_pct, one_dose_centered)
```

```
## # A tibble: 3,114 x 5
##    fips  county   state one_dose_5plus_pct one_dose_centered
##    <chr> <chr>    <chr>              <dbl>             <dbl>
## 1  26039 Crawfor~ MI                  55.7             -7.35
## 2  40015 Caddo C~ OK                  83.3             20.2
## 3  17007 Boone C~ IL                  71.1              8.05
## 4  12055 Highlan~ FL                  68.9              5.85
## 5  34029 Ocean C~ NJ                  71                7.95
## 6  01067 Henry C~ AL                  58.5             -4.55
## 7  27037 Dakota ~ MN                  81               17.9
## 8  27115 Pine Co~ MN                  56.5             -6.55
## 9  51750 Radford~ VA                  41.5            -21.6
## 10 22009 Avoyell~ LA                  59.7             -3.35
## # i 3,104 more rows
```

z-score example

```
covid_votes |>
  mutate(
    one_dose_z =
      (one_dose_5plus_pct - mean(one_dose_5plus_pct, na.rm = TRUE)) /
      sd(one_dose_5plus_pct, na.rm = TRUE))   |>
  select(fips:state, one_dose_5plus_pct, one_dose_z)
```

```
## # A tibble: 3,114 x 5
##     fips  county         state one_dose_5plus_pct one_dose_z
##     <chr> <chr>          <chr>              <dbl>      <dbl>
##  1 26039 Crawford County MI                   55.7     -0.508
##  2 40015 Caddo County    OK                   83.3      1.40
##  3 17007 Boone County    IL                   71.1      0.556
##  4 12055 Highlands Coun~ FL                   68.9      0.404
##  5 34029 Ocean County    NJ                   71        0.549
##  6 01067 Henry County    AL                   58.5     -0.314
##  7 27037 Dakota County   MN                   81        1.24
##  8 27115 Pine County     MN                   56.5     -0.452
##  9 51750 Radford city    VA                   41.5     -1.49
## 10 22009 Avoyelles Pari~ LA                   59.7     -0.231
## # i 3,104 more rows
```

2/ Correlation

Correlation

- How do variables move together on average?

Correlation

- How do variables move together on average?

- When x_i is big, what is y_i likely to be?

Correlation

- How do variables move together on average?

- When x_i is big, what is y_i likely to be?

  - Positive correlation: when x_i is big, y_i is also big

Correlation

- How do variables move together on average?

- When x_i is big, what is y_i likely to be?

  - Positive correlation: when x_i is big, y_i is also big
  - Negative correlation: when x_i is big, y_i is small

Correlation

- How do variables move together on average?

- When x_i is big, what is y_i likely to be?

  - Positive correlation: when x_i is big, y_i is also big
  - Negative correlation: when x_i is big, y_i is small
  - High magnitude of correlation: data cluster tightly around a line.

Correlation

- How do variables move together on average?

- When x_i is big, what is y_i likely to be?

    - Positive correlation: when x_i is big, y_i is also big
    - Negative correlation: when x_i is big, y_i is small
    - High magnitude of correlation: data cluster tightly around a line.

- The technical definition of the **correlation coefficient**:

$$\frac{1}{n-1} \sum_{i=1}^{n} \left[(\text{z-score for } x_i) \times (\text{z-score for } y_i) \right]$$

Correlation

- How do variables move together on average?

- When x_i is big, what is y_i likely to be?

    - Positive correlation: when x_i is big, y_i is also big
    - Negative correlation: when x_i is big, y_i is small
    - High magnitude of correlation: data cluster tightly around a line.

- The technical definition of the **correlation coefficient**:

$$\frac{1}{n-1} \sum_{i=1}^{n} \left[(\text{z-score for } x_i) \times (\text{z-score for } y_i) \right]$$

- Interpretation:

Correlation

- How do variables move together on average?

- When x_i is big, what is y_i likely to be?

  - Positive correlation: when x_i is big, y_i is also big
  - Negative correlation: when x_i is big, y_i is small
  - High magnitude of correlation: data cluster tightly around a line.

- The technical definition of the **correlation coefficient**:

$$\frac{1}{n-1} \sum_{i=1}^{n} [(\text{z-score for } x_i) \times (\text{z-score for } y_i)]$$

- Interpretation:

  - Correlation is between -1 and 1

Correlation

- How do variables move together on average?

- When x_i is big, what is y_i likely to be?

  - Positive correlation: when x_i is big, y_i is also big
  - Negative correlation: when x_i is big, y_i is small
  - High magnitude of correlation: data cluster tightly around a line.

- The technical definition of the **correlation coefficient**:

$$\frac{1}{n-1} \sum_{i=1}^{n} [(\text{z-score for } x_i) \times (\text{z-score for } y_i)]$$

- Interpretation:

  - Correlation is between -1 and 1
 - Correlation of 0 means no linear association.

Correlation

- How do variables move together on average?

- When x_i is big, what is y_i likely to be?

    - Positive correlation: when x_i is big, y_i is also big
    - Negative correlation: when x_i is big, y_i is small
    - High magnitude of correlation: data cluster tightly around a line.

- The technical definition of the **correlation coefficient**:

$$\frac{1}{n-1} \sum_{i=1}^{n} [(\text{z-score for } x_i) \times (\text{z-score for } y_i)]$$

- Interpretation:

    - Correlation is between -1 and 1
    - Correlation of 0 means no linear association.
    - Positive correlations $\rightsquigarrow$ positive associations.

Correlation

- How do variables move together on average?

- When x_i is big, what is y_i likely to be?

    - Positive correlation: when x_i is big, y_i is also big
    - Negative correlation: when x_i is big, y_i is small
    - High magnitude of correlation: data cluster tightly around a line.

- The technical definition of the **correlation coefficient**:

$$\frac{1}{n-1} \sum_{i=1}^{n} \left[(\text{z-score for } x_i) \times (\text{z-score for } y_i)\right]$$

- Interpretation:

    - Correlation is between -1 and 1
    - Correlation of 0 means no linear association.
    - Positive correlations $\rightsquigarrow$ positive associations.
    - Negative correlations $\rightsquigarrow$ negative associations.

Correlation

- How do variables move together on average?

- When x_i is big, what is y_i likely to be?

    - Positive correlation: when x_i is big, y_i is also big
    - Negative correlation: when x_i is big, y_i is small
    - High magnitude of correlation: data cluster tightly around a line.

- The technical definition of the **correlation coefficient**:

$$\frac{1}{n-1} \sum_{i=1}^{n} [(\text{z-score for } x_i) \times (\text{z-score for } y_i)]$$

- Interpretation:

    - Correlation is between -1 and 1
    - Correlation of 0 means no linear association.
    - Positive correlations $\rightsquigarrow$ positive associations.
    - Negative correlations $\rightsquigarrow$ negative associations.
    - Closer to -1 or 1 means stronger association.

Correlation intuition
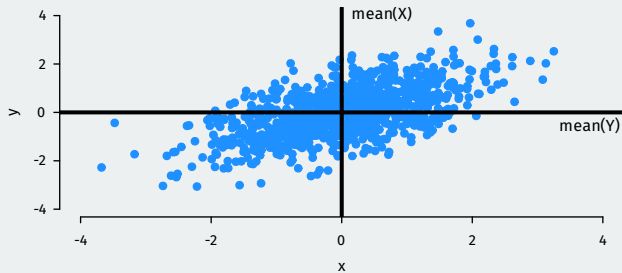
Correlation intuition
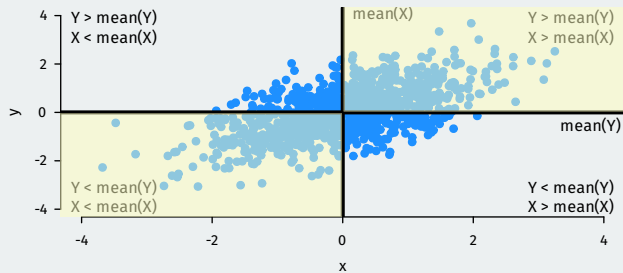
- Large values of X tend to occur with large values of Y:
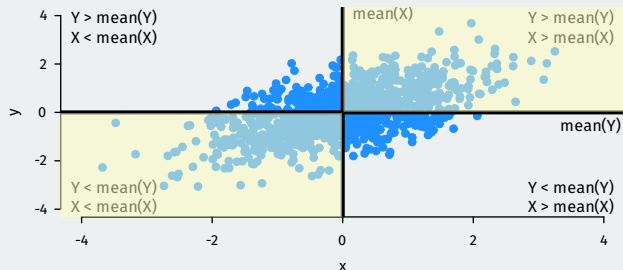
Correlation intuition
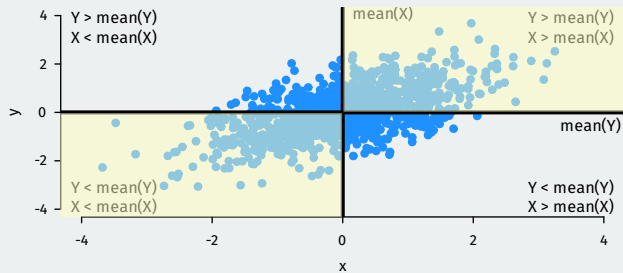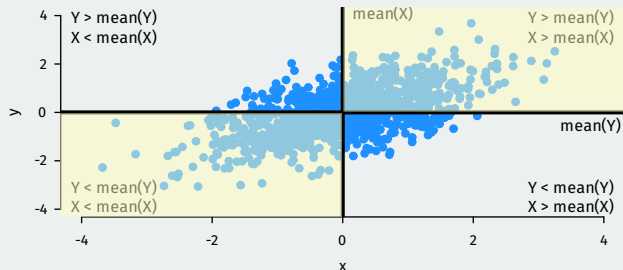

- Large values of X tend to occur with large values of Y:
  - (z-score for x_i) $\times$ (z-score for y_i) = (pos. num.) $\times$ (pos. num) = $+$

Correlation intuition
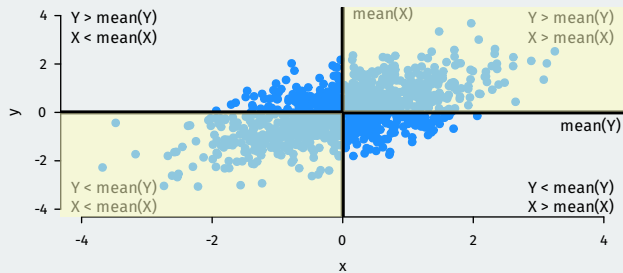
- Large values of X tend to occur with large values of Y:

  - (z-score for x_i) $\times$ (z-score for y_i) = (pos. num.) $\times$ (pos. num) = $+$

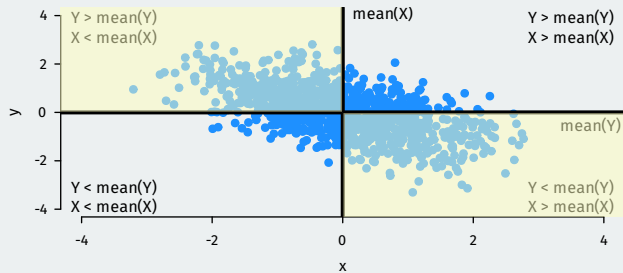- Small values of X tend to occur with small values of Y:

Correlation intuition

- Large values of X tend to occur with large values of Y:
    - (z-score for x_i) × (z-score for y_i) = (pos. num.) × (pos. num) = +

- Small values of X tend to occur with small values of Y:
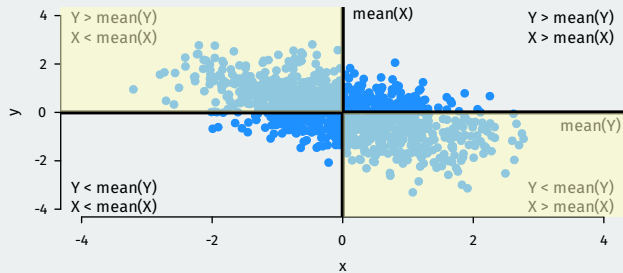   - (z-score for x_i) × (z-score for y_i) = (neg. num.) × (neg. num) = +

Correlation intuition

- Large values of X tend to occur with large values of Y:
  - (z-score for x_i) $\times$ (z-score for y_i) = (pos. num.) $\times$ (pos. num.) = $+$
- Small values of X tend to occur with small values of Y:
  - (z-score for x_i) $\times$ (z-score for y_i) = (neg. num.) $\times$ (neg. num.) = $+$
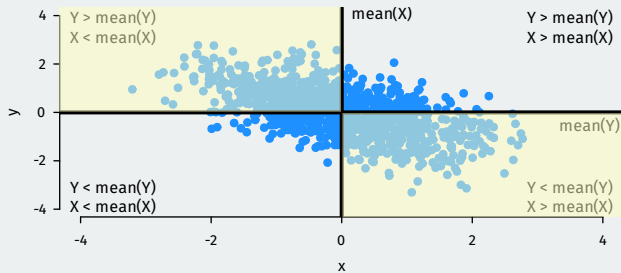- If these dominate $\rightsquigarrow$ positive correlation.

Correlation intuition

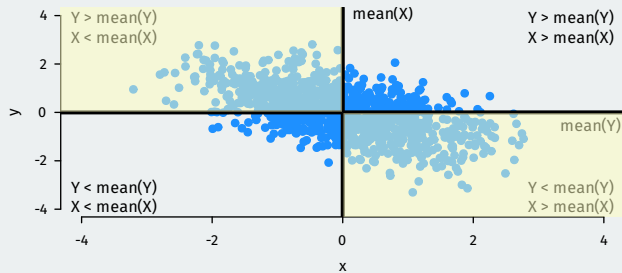- Large values of X tend to occur with small values of Y:
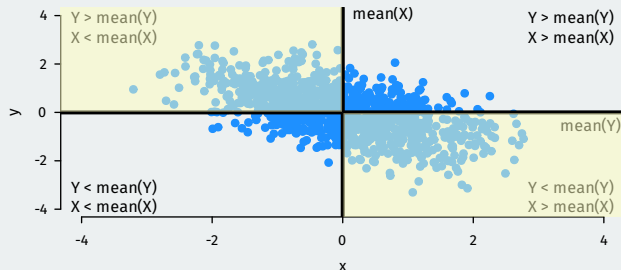
- Large values of X tend to occur with small values of Y:

  - (z-score for x_i) $\times$ (z-score for y_i) = (pos. num.) $\times$ (neg. num) = $-$

Correlation intuition

- Large values of X tend to occur with small values of Y:
  - (z-score for x_i) $\times$ (z-score for y_i) = (pos. num.) $\times$ (neg. num) $= -$
- Small values of X tend to occur with large values of Y:

Correlation intuition

- Large values of X tend to occur with small values of Y:
  - (z-score for x_i) × (z-score for y_i) = (pos. num.) × (neg. num) = −

- Small values of X tend to occur with large values of Y:
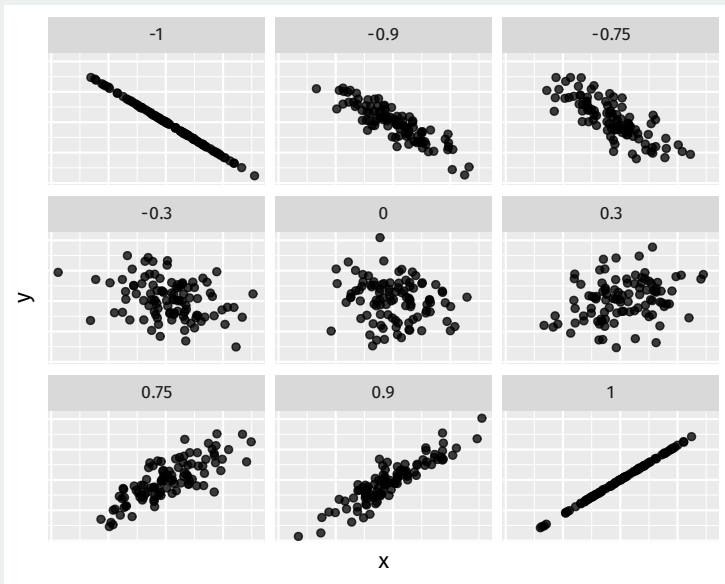 - (z-score for x_i) × (z-score for y_i) = (neg. num.) × (pos. num) = −

Correlation intuition

- Large values of X tend to occur with small values of Y:
    - (z-score for x_i) $\times$ (z-score for y_i) = (pos. num.) $\times$ (neg. num) = $-$
- Small values of X tend to occur with large values of Y:
    - (z-score for x_i) $\times$ (z-score for y_i) = (neg. num.) $\times$ (pos. num) = $-$
- If these dominate ⤳ negative correlation.

Correlation examples

Properties of correlation coefficient

- Correlation measures **linear** association.

Properties of correlation coefficient

- Correlation measures **linear** association.

- Order doesn't matter: `cor(x,y) = cor(y,x)`

Properties of correlation coefficient

- Correlation measures **linear** association.

- Order doesn't matter: `cor(x,y) = cor(y,x)`

- Not affected by changes of scale:

Properties of correlation coefficient

- Correlation measures **linear** association.

- Order doesn't matter: `cor(x,y) = cor(y,x)`

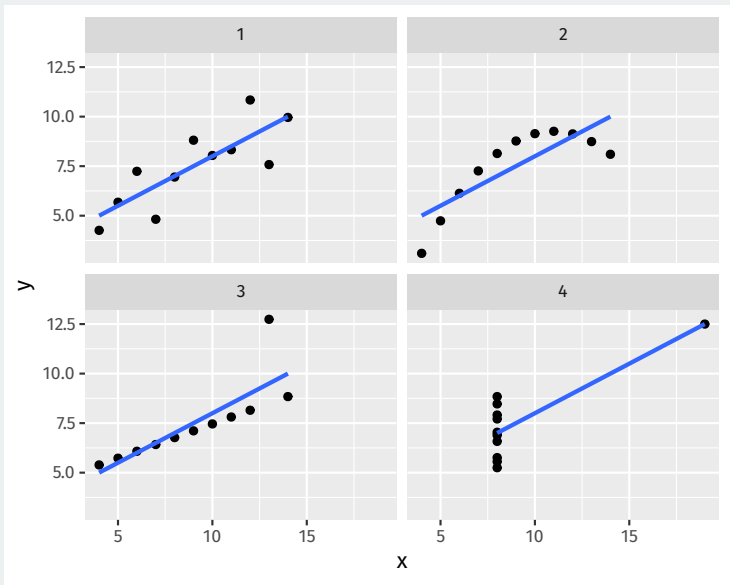- Not affected by changes of scale:

    - `cor(x,y) = cor(ax+b, cy+d)`

Properties of correlation coefficient

- Correlation measures **linear** association.

- Order doesn't matter: `cor(x,y) = cor(y,x)`

- Not affected by changes of scale:
    - `cor(x,y) = cor(ax+b, cy+d)`
    - Celsius vs. Fahreneheit; dollars vs. pesos; cm vs. in.

All 4 relationships have 0.816 correlation

Correlation in R

Use the cor() function:

```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2020)
```

```
## [1] NA
```

Correlation in R

Use the `cor()` function:

```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2020)
```

```
## [1] NA
```

Missing values: set the `use = "pairwise"` → available case analysis
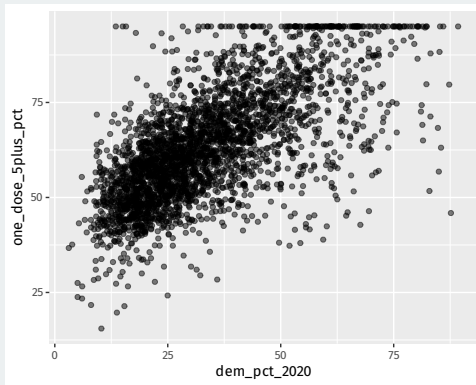
```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2020,
    use = "pairwise")
```

```
## [1] 0.666
```

Comparing correlations

```
covid_votes |>
  ggplot(aes(x = dem_pct_2020, y = one_dose_5plus_pct)) +
  geom_point(alpha = 0.5)
```



```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2020,
    use = "pairwise")
```

```
## [1] 0.666
```

Comparing correlations

```
covid_votes |>
  ggplot(aes(x = dem_pct_2000, y = one_dose_5plus_pct)) +
  geom_point(alpha = 0.5)
```



```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2000,
    use = "pairwise")
```
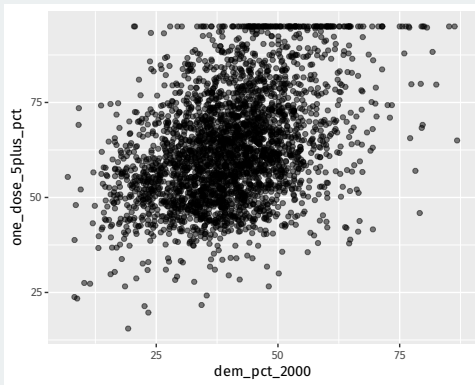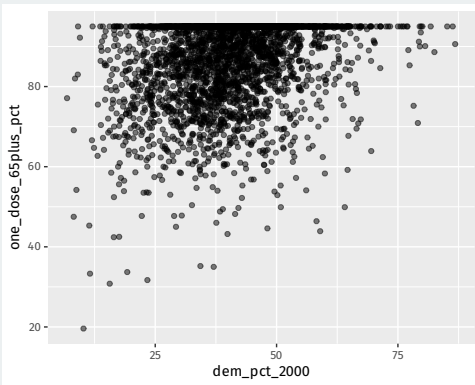
```
## [1] 0.394
```

Comparing correlations

```r
covid_votes |>
  ggplot(aes(x = dem_pct_2000, y = one_dose_65plus_pct)) +
  geom_point(alpha = 0.5)
```



```r
cor(covid_votes$one_dose_65plus_pct, covid_votes$dem_pct_2000,
    use = "pairwise")
```

```
## [1] 0.263
```

3/ Writing our own functions

Why write functions?

Copy-pasting code tedious and prone to failure:

```
covid_votes |>
  mutate(
    one_dose_5p_z =
      (one_dose_5plus_pct - mean(one_dose_5plus_pct, na.rm = TRUE)) /
      sd(one_dose_5plus_pct, na.rm = TRUE),
    one_dose_65p_z =
      (one_dose_65plus_pct - mean(one_dose_65plus_pct, na.rm = TRUE)) /
      sd(one_dose_65plus_pct, na.rm = TRUE),
    booster_z =
      (booster_5plus_pct - mean(booster_5plus_pct, na.rm = TRUE)) /
      sd(booster_5plus_pct, na.rm = TRUE),
    dem_pct_2000_z =
      (dem_pct_2000 - mean(dem_pct_2000, na.rm = TRUE)) /
      sd(dem_pct_2000, na.rm = TRUE),
    dem_pct_2020_z =
      (dem_pct_2020 - mean(dem_pct_2020, na.rm = TRUE)) /
      sd(dem_pct_2020, na.rm = TRUE)
  )
```

Writing a new function

Notice that all of the mutations follow the same template:

```
( █ - mean( █ , na.rm = TRUE)) / sd( █ , na.rm = TRUE)
```

Only one thing varies: the column of data, represented with █

Components of a function

We create functions like so:

```
name <- function(arguments) {
  body
}
```

We create functions like so:

```
name <- function(arguments) {
  body
}
```

Three components:

1. **Name**: the name of the function that we'll use to call it. Maybe
   z_score?

Components of a function

We create functions like so:

```
name <- function(arguments) {
  body
}
```

Three components:

1. **Name**: the name of the function that we'll use to call it. Maybe `z_score`?
2. **Arguments**: things that we want to vary across calls of our function. We'll use x.

Components of a function

We create functions like so:

```
name <- function(arguments) {
  body
}
```

Three components:

1. **Name**: the name of the function that we'll use to call it. Maybe
   z_score?
2. **Arguments**: things that we want to vary across calls of our function.
   We'll use x.
3. **Body**: the code that the function performs.

Our first function

Convert our template to a function:

```r
z_score <- function(x) {
  (x - mean(x, na.rm = TRUE)) / sd(x, na.rm = TRUE)
}
```

Convert our template to a function:

```
z_score <- function(x) {
  (x - mean(x, na.rm = TRUE)) / sd(x, na.rm = TRUE)
}
```

Check that it seems to work:

```
z_score(c(1,2, 3, 4, 5))
```

```
## [1] -1.265 -0.632  0.000  0.632  1.265
```

Cleaning up our code

```
covid_votes |>
  mutate(
    one_dose_5p_z = z_score(one_dose_5plus_pct),
    one_dose_65p_z = z_score(one_dose_65plus_pct),
    booster_z = z_score(booster_5plus_pct),
    dem_pct_2000_z = z_score(dem_pct_2000),
    dem_pct_2020_z = z_score(dem_pct_2020)
  )
```

`across()` **function**

If we want to replace our variables with z-scores, we can use the `across()` function to perform many mutations at once:

```
covid_votes |>
  mutate(across(one_dose_5plus_pct:dem_pct_2020, z_score))
```

```
## # A tibble: 3,114 x 8
##     fips  county state one_dose_5plus_pct one_dose_65plus_pct
##     <chr> <chr>  <chr>              <dbl>               <dbl>
##  1 26039 Crawf~ MI                -0.508              -0.829
##  2 40015 Caddo~ OK                 1.40                0.843
##  3 17007 Boone~ IL                 0.556               0.795
##  4 12055 Highl~ FL                 0.404               0.720
##  5 34029 Ocean~ NJ                 0.549               0.843
##  6 01067 Henry~ AL                -0.314              -0.0545
##  7 27037 Dakot~ MN                 1.24                0.843
##  8 27115 Pine ~ MN                -0.452              -0.102
##  9 51750 Radfo~ VA                -1.49               -1.16
## 10 22009 Avoye~ LA                -0.231              -0.564
## # i 3,104 more rows
## # i 3 more variables: booster_5plus_pct <dbl>,
## #   dem_pct_2000 <dbl>, dem_pct_2020 <dbl>
```

Alternative approach

We could also target all the numeric variables:

```
covid_votes |>
  mutate(across(where(is.numeric), z_score))
```

```
## # A tibble: 3,114 x 8
##     fips  county state one_dose_5plus_pct one_dose_65plus_pct
##     <chr> <chr>  <chr>              <dbl>               <dbl>
##  1 26039 Crawf~ MI                -0.508              -0.829
##  2 40015 Caddo~ OK                 1.40                0.843
##  3 17007 Boone~ IL                 0.556               0.795
##  4 12055 Highl~ FL                 0.404               0.720
##  5 34029 Ocean~ NJ                 0.549               0.843
##  6 01067 Henry~ AL                -0.314              -0.0545
##  7 27037 Dakot~ MN                 1.24                0.843
##  8 27115 Pine ~ MN                -0.452              -0.102
##  9 51750 Radfo~ VA                -1.49               -1.16
## 10 22009 Avoye~ LA                -0.231              -0.564
## # i 3,104 more rows
## # i 3 more variables: booster_5plus_pct <dbl>,
## #   dem_pct_2000 <dbl>, dem_pct_2020 <dbl>
```

Alternative approach

We could also target only the first dose variables:

```
covid_votes |>
  mutate(across(starts_with("one_dose"), z_score))
```

```
## # A tibble: 3,114 x 8
##     fips  county state one_dose_5plus_pct one_dose_65plus_pct
##    <chr> <chr>  <chr>              <dbl>               <dbl>
##  1 26039 Crawf~ MI                -0.508              -0.829
##  2 40015 Caddo~ OK                 1.40                0.843
##  3 17007 Boone~ IL                 0.556               0.795
##  4 12055 Highl~ FL                 0.404               0.720
##  5 34029 Ocean~ NJ                 0.549               0.843
##  6 01067 Henry~ AL                -0.314              -0.0545
##  7 27037 Dakot~ MN                 1.24                0.843
##  8 27115 Pine ~ MN                -0.452              -0.102
##  9 51750 Radfo~ VA                -1.49               -1.16
## 10 22009 Avoye~ LA                -0.231              -0.564
## # i 3,104 more rows
## # i 3 more variables: booster_5plus_pct <dbl>,
## #   dem_pct_2000 <dbl>, dem_pct_2020 <dbl>
```

Adding arguments to our function

What if we want to be able to control na.rm in the calls to mean() and sd() in our z_score function? Add an argument!

```
z_score2 <- function(x, na.rm = FALSE) {
  (x - mean(x, na.rm = na.rm)) / sd(x, na.rm = na.rm)
}
```

```
head(z_score2(covid_votes$one_dose_5plus_pct))
```

```
## [1] NA NA NA NA NA NA
```

```
head(z_score2(covid_votes$one_dose_5plus_pct, na.rm = TRUE))
```

```
## [1] -0.508  1.398  0.556  0.404  0.549 -0.314
```