

Gov 50: 10. Summarizing Bivariate Relationships

Matthew Blackwell

Harvard University

Roadmap

1. Z-scores and standardization
2. Correlation
3. Writing our own functions


1/ Z-scores and standardization

COVID vaccination rates and votes

```
library(tidyverse)
library(gov50data)
covid_votes
```

```
## # A tibble: 3,114 x 8
##   fips county state one_dose_5plus_pct one_dose_65plus_pct
##   <chr> <chr> <chr>           <dbl>           <dbl>
## 1 26039 Crawf~ MI             55.7             77.3
## 2 40015 Caddo~ OK             83.3             95
## 3 17007 Boone~ IL             71.1            94.5
## 4 12055 Highl~ FL             68.9            93.7
## 5 34029 Ocean~ NJ             71              95
## 6 01067 Henry~ AL             58.5            85.5
## 7 27037 Dakot~ MN             81              95
## 8 27115 Pine ~ MN             56.5             85
## 9 51750 Radfo~ VA             41.5            73.8
## 10 22009 Avoye~ LA             59.7            80.1
## # i 3,104 more rows
## # i 3 more variables: booster_5plus_pct <dbl>,
## #   dem_pct_2000 <dbl>, dem_pct_2020 <dbl>
```


Is 60% vaccinated a lot?



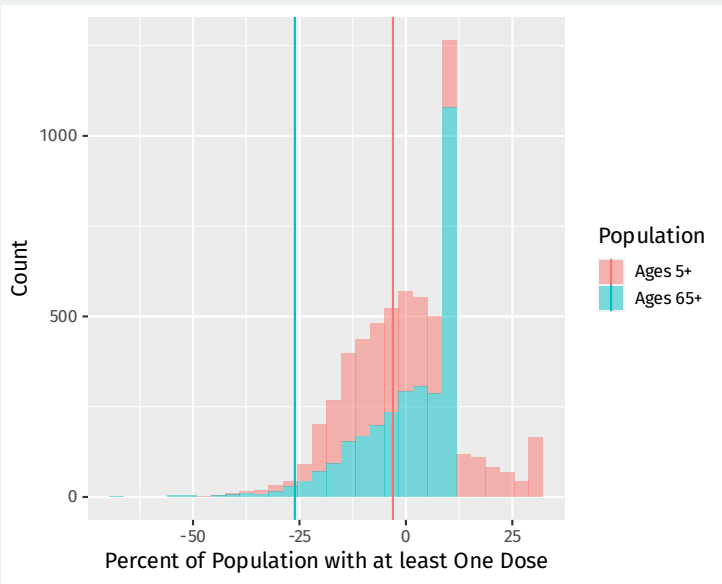
How large is large?

- How large 60% vaccinated is depends on the distribution!
 - Clear to see from the histogram
 - Middling for the 5+ group, but very low for the 65+ group.
- Can we transform the values of our variables to be **common units**?
- Yes, with two transformations:
 - **Centering**: subtract the mean of the variable from each value.
 - **Scaling**: dividing deviations from the mean by the standard deviation.


Original distributions



Centered distributions



Centered and scaled distributions



Z-scores

- Centering tells us immediately if a value is above or below the mean.
- Scaling tells us how many standard deviations away from the mean it is.
- Combine them with the **z-score** transformation:

$$\text{z-score of } x_i = \frac{x_i - \text{mean of } x}{\text{standard deviation of } x}$$

- Useful heuristic: data more than 3 SDs away from mean are rare.

z-score example

```
covid_votes |>
  mutate(one_dose_centered = one_dose_5plus_pct -
         mean(one_dose_5plus_pct, na.rm = TRUE)) |>
  select(fips:state, one_dose_5plus_pct, one_dose_centered)
```

```
## # A tibble: 3,114 x 5
##   fips  county  state one_dose_5plus_pct one_dose_centered
##   <chr> <chr>   <chr>           <dbl>           <dbl>
## 1 26039 Crawfor~ MI             55.7            -7.35
## 2 40015 Caddo C~ OK             83.3            20.2
## 3 17007 Boone C~ IL             71.1             8.05
## 4 12055 Highlan~ FL             68.9             5.85
## 5 34029 Ocean C~ NJ              71             7.95
## 6 01067 Henry C~ AL             58.5            -4.55
## 7 27037 Dakota ~ MN              81            17.9
## 8 27115 Pine Co~ MN             56.5            -6.55
## 9 51750 Radford~ VA             41.5           -21.6
## 10 22009 Avoyell~ LA             59.7            -3.35
## # i 3,104 more rows
```

z-score example

```
covid_votes |>
  mutate(
    one_dose_z =
      (one_dose_5plus_pct - mean(one_dose_5plus_pct, na.rm = TRUE)) /
      sd(one_dose_5plus_pct, na.rm = TRUE)) |>
  select(fips:state, one_dose_5plus_pct, one_dose_z)
```

```
## # A tibble: 3,114 x 5
##   fips county          state one_dose_5plus_pct one_dose_z
##   <chr> <chr>          <chr>          <dbl>         <dbl>
## 1 26039 Crawford County MI             55.7         -0.508
## 2 40015 Caddo County    OK             83.3           1.40
## 3 17007 Boone County   IL             71.1           0.556
## 4 12055 Highlands Coun~ FL             68.9           0.404
## 5 34029 Ocean County   NJ             71             0.549
## 6 01067 Henry County   AL             58.5          -0.314
## 7 27037 Dakota County  MN             81             1.24
## 8 27115 Pine County   MN             56.5          -0.452
## 9 51750 Radford city   VA             41.5          -1.49
## 10 22009 Avoyelles Pari~ LA             59.7          -0.231
## # i 3,104 more rows
```

2/ Correlation


Correlation

- How do variables move together on average?
- When x_i is big, what is y_i likely to be?
 - Positive correlation: when x_i is big, y_i is also big
 - Negative correlation: when x_i is big, y_i is small
 - High magnitude of correlation: data cluster tightly around a line.
- The technical definition of the **correlation coefficient**:


$$\frac{1}{n-1} \sum_{i=1}^n [(z\text{-score for } x_i) \times (z\text{-score for } y_i)]$$

- Interpretation:
 - Correlation is between -1 and 1
 - Correlation of 0 means no linear association.
 - Positive correlations \rightsquigarrow positive associations.
 - Negative correlations \rightsquigarrow negative associations.
 - Closer to -1 or 1 means stronger association.

Correlation intuition




Correlation intuition




- Large values of X tend to occur with large values of Y :
 - $(\text{z-score for } x_i) \times (\text{z-score for } y_i) = (\text{pos. num.}) \times (\text{pos. num.}) = +$
- Small values of X tend to occur with small values of Y :
 - $(\text{z-score for } x_i) \times (\text{z-score for } y_i) = (\text{neg. num.}) \times (\text{neg. num.}) = +$
- If these dominate \rightsquigarrow positive correlation.

Correlation intuition



- Large values of X tend to occur with small values of Y :
 - $(\text{z-score for } x_i) \times (\text{z-score for } y_i) = (\text{pos. num.}) \times (\text{neg. num.}) = -$
- Small values of X tend to occur with large values of Y :
 - $(\text{z-score for } x_i) \times (\text{z-score for } y_i) = (\text{neg. num.}) \times (\text{pos. num.}) = -$
- If these dominate \rightsquigarrow negative correlation.


Correlation examples



Properties of correlation coefficient

- Correlation measures **linear** association.
- Order doesn't matter: $\text{cor}(x, y) = \text{cor}(y, x)$
- Not affected by changes of scale:
 - $\text{cor}(x, y) = \text{cor}(ax+b, cy+d)$
 - Celsius vs. Fahrenheit; dollars vs. pesos; cm vs. in.

All 4 relationships have 0.816 correlation



Correlation in R

Use the `cor()` function:

```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2020)
```

```
## [1] NA
```


Missing values: set the `use = "pairwise"` → available case analysis

```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2020,  
    use = "pairwise")
```

```
## [1] 0.666
```

Comparing correlations

```
covid_votes |>  
  ggplot(aes(x = dem_pct_2020, y = one_dose_5plus_pct)) +  
  geom_point(alpha = 0.5)
```




```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2020,  
    use = "pairwise")
```

```
## [1] 0.666
```

Comparing correlations

```
covid_votes |>  
  ggplot(aes(x = dem_pct_2000, y = one_dose_5plus_pct)) +  
  geom_point(alpha = 0.5)
```




```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2000,  
    use = "pairwise")
```

```
## [1] 0.394
```

Comparing correlations

```
covid_votes |>  
  ggplot(aes(x = dem_pct_2000, y = one_dose_65plus_pct)) +  
  geom_point(alpha = 0.5)
```



```
cor(covid_votes$one_dose_65plus_pct, covid_votes$dem_pct_2000,  
    use = "pairwise")
```

```
## [1] 0.263
```


3/ Writing our own functions

Why write functions?

Copy-pasting code tedious and prone to failure:

```
covid_votes |>
  mutate(
    one_dose_5p_z =
      (one_dose_5plus_pct - mean(one_dose_5plus_pct, na.rm = TRUE)) /
      sd(one_dose_5plus_pct, na.rm = TRUE),
    one_dose_65p_z =
      (one_dose_65plus_pct - mean(one_dose_65plus_pct, na.rm = TRUE)) /
      sd(one_dose_65plus_pct, na.rm = TRUE),
    booster_z =
      (booster_5plus_pct - mean(booster_5plus_pct, na.rm = TRUE)) /
      sd(booster_5plus_pct, na.rm = TRUE),
    dem_pct_2000_z =
      (dem_pct_2000 - mean(dem_pct_2000, na.rm = TRUE)) /
      sd(dem_pct_2000, na.rm = TRUE),
    dem_pct_2020_z =
      (dem_pct_2020 - mean(dem_pct_2020, na.rm = TRUE)) /
      sd(dem_pct_2020, na.rm = TRUE)
  )
```

Writing a new function

Notice that all of the mutations follow the same template:

```
(█ - mean(█, na.rm = TRUE)) / sd(█, na.rm = TRUE)
```

Only one thing varies: the column of data, represented with `█`

Components of a function

We create functions like so:

```
name <- function(arguments) {  
  body  
}
```

Three components:

1. **Name:** the name of the function that we'll use to call it. Maybe `z_score`?
2. **Arguments:** things that we want to vary across calls of our function. We'll use `x`.
3. **Body:** the code that the function performs.

Our first function

Convert our template to a function:

```
z_score <- function(x) {  
  (x - mean(x, na.rm = TRUE)) / sd(x, na.rm = TRUE)  
}
```

Check that it seems to work:

```
z_score(c(1,2, 3, 4, 5))
```

```
## [1] -1.265 -0.632  0.000  0.632  1.265
```

Cleaning up our code

```
covid_votes |>
  mutate(
    one_dose_5p_z = z_score(one_dose_5plus_pct),
    one_dose_65p_z = z_score(one_dose_65plus_pct),
    booster_z = z_score(booster_5plus_pct),
    dem_pct_2000_z = z_score(dem_pct_2000),
    dem_pct_2020_z = z_score(dem_pct_2020)
  )
```

across() function

If we want to replace our variables with z-scores, we can use the `across()` function to perform many mutations at once:

```
covid_votes |>
  mutate(across(one_dose_5plus_pct:dem_pct_2020, z_score))
```

```
## # A tibble: 3,114 x 8
##   fips county state one_dose_5plus_pct one_dose_65plus_pct
##   <chr> <chr>  <chr>          <dbl>          <dbl>
## 1 26039 Crawf~ MI             -0.508         -0.829
## 2 40015 Caddo~ OK              1.40          0.843
## 3 17007 Boone~ IL              0.556          0.795
## 4 12055 Highl~ FL              0.404          0.720
## 5 34029 Ocean~ NJ              0.549          0.843
## 6 01067 Henry~ AL             -0.314        -0.0545
## 7 27037 Dakot~ MN              1.24          0.843
## 8 27115 Pine ~ MN             -0.452        -0.102
## 9 51750 Radfo~ VA             -1.49         -1.16
## 10 22009 Avoye~ LA             -0.231        -0.564
## # i 3,104 more rows
## # i 3 more variables: booster_5plus_pct <dbl>,
## #   dem_pct_2000 <dbl>, dem_pct_2020 <dbl>
```

Alternative approach

We could also target all the numeric variables:

```
covid_votes |>  
  mutate(across(where(is.numeric), z_score))
```

```
## # A tibble: 3,114 x 8  
##   fips county state one_dose_5plus_pct one_dose_65plus_pct  
##   <chr> <chr> <chr>           <dbl>           <dbl>  
## 1 26039 Crawf~ MI             -0.508           -0.829  
## 2 40015 Caddo~ OK              1.40            0.843  
## 3 17007 Boone~ IL              0.556            0.795  
## 4 12055 Highl~ FL              0.404            0.720  
## 5 34029 Ocean~ NJ              0.549            0.843  
## 6 01067 Henry~ AL             -0.314          -0.0545  
## 7 27037 Dakot~ MN              1.24            0.843  
## 8 27115 Pine ~ MN             -0.452          -0.102  
## 9 51750 Radfo~ VA             -1.49           -1.16  
## 10 22009 Avoye~ LA             -0.231          -0.564  
## # i 3,104 more rows  
## # i 3 more variables: booster_5plus_pct <dbl>,  
## #   dem_pct_2000 <dbl>, dem_pct_2020 <dbl>
```


Alternative approach

We could also target only the first dose variables:

```
covid_votes |>  
  mutate(across(starts_with("one_dose"), z_score))
```

```
## # A tibble: 3,114 x 8  
##   fips county state one_dose_5plus_pct one_dose_65plus_pct  
##   <chr> <chr> <chr>          <dbl>          <dbl>  
## 1 26039 Crawf~ MI          -0.508          -0.829  
## 2 40015 Caddo~ OK           1.40           0.843  
## 3 17007 Boone~ IL           0.556           0.795  
## 4 12055 Highl~ FL           0.404           0.720  
## 5 34029 Ocean~ NJ           0.549           0.843  
## 6 01067 Henry~ AL          -0.314          -0.0545  
## 7 27037 Dakot~ MN           1.24           0.843  
## 8 27115 Pine ~ MN          -0.452          -0.102  
## 9 51750 Radfo~ VA          -1.49           -1.16  
## 10 22009 Avoye~ LA          -0.231          -0.564  
## # i 3,104 more rows  
## # i 3 more variables: booster_5plus_pct <dbl>,  
## #   dem_pct_2000 <dbl>, dem_pct_2020 <dbl>
```

Adding arguments to our function

What if we want to be able to control `na.rm` in the calls to `mean()` and `sd()` in our `z_score` function? Add an argument!

```
z_score2 <- function(x, na.rm = FALSE) {  
  (x - mean(x, na.rm = na.rm)) / sd(x, na.rm = na.rm)  
}
```

```
head(z_score2(covid_votes$one_dose_5plus_pct))
```

```
## [1] NA NA NA NA NA NA
```

```
head(z_score2(covid_votes$one_dose_5plus_pct, na.rm = TRUE))
```

```
## [1] -0.508  1.398  0.556  0.404  0.549 -0.314
```