

Gov 50: 17. Sampling Distributions

Matthew Blackwell

Harvard University

Roadmap

1. Poll example
2. Random variables and probability distributions
3. Sampling distribution
4. Normal variables and the Central Limit Theorem

1/ Poll example

How popular is Joe Biden?



- What proportion of the public approves of Biden's job as president?
- Latest Gallup poll:
 - Sept 1st-16th
 - 812 adult Americans
 - Telephone interviews
 - Approve (42%), Disapprove (56%)

Poll in our framework

- **Population:** adults 18+ living in 50 US states and DC.
- **Population parameter:** population proportion of all US adults that approve of Biden.
 - Census: not possible.
- **Sample:** random digit dialing phone numbers (cell and landline).
- **Point estimate:** sample proportion that approve of Biden

2/ Random variables and probability distributions

Random variables

Random variables are numerical summaries of chance processes:

$$X_i = \begin{cases} 1 & \text{if respondent } i \text{ supports Biden,} \\ 0 & \text{otherwise} \end{cases}$$

With a simple random sample, chance of $X_i = 1$ is equal to the population proportion of people that support Biden.

Types of random variables

- **Discrete:** X can take a finite (or countably infinite) number of values.
 - Number of heads in 5 coin flips
 - Sampled senator is a woman ($X = 1$) or not ($X = 0$)
 - Number of battle deaths in a civil war
- **Continuous:** X can take any real value (usually within an interval).
 - GDP per capita (average income) in a country.
 - Share of population that approves of Biden.
 - Amount of time spent on a website.

Probability distributions

Probability distributions tell us the chances of different values of a r.v. occurring

Discrete variables: like a frequency barplot for the population distribution.


Continuous variables: like a continuous version of population histogram.

Discrete probability distribution

We can use the `y = ..prop..` aesthetic to get a barplot with proportions instead of count to show us the chance/probability of selecting a first-year student:

```
library(gov50data)
class_years |>
  mutate(first_year = as.numeric(year == "First-Year")) |>
  ggplot(aes(x = first_year)) +
  geom_bar(mapping = aes(y = ..prop..), width = 0.1)
```

Discrete probability distribution



Midwest data

```
library(ggplot2)
midwest
```


```
## # A tibble: 437 x 28
##       PID county    state  area poptotal popdensity popwhite
##   <int> <chr>    <chr> <dbl>    <int>      <dbl>    <int>
## 1   561 ADAMS      IL    0.052   66090    1271.    63917
## 2   562 ALEXANDER IL    0.014   10626     759     7054
## 3   563 BOND       IL    0.022   14991     681.    14477
## 4   564 BOONE      IL    0.017   30806    1812.    29344
## 5   565 BROWN      IL    0.018    5836     324.     5264
## 6   566 BUREAU     IL    0.05    35688     714.    35157
## 7   567 CALHOUN    IL    0.017    5322     313.     5298
## 8   568 CARROLL    IL    0.027   16805     622.    16519
## 9   569 CASS       IL    0.024   13437     560.    13384
## 10  570 CHAMPAIGN IL    0.058  173025    2983.   146506
## # i 427 more rows
## # i 21 more variables: popblack <int>, popamerindian <int>,
## #   popasian <int>, popother <int>, percwhite <dbl>,
## #   percblack <dbl>, percamerindian <dbl>, percasian <dbl>,
## #   percother <dbl>, popadults <int>, perchsd <dbl>,
## #   percollege <dbl>, percprof <dbl>,
```

Continuous probability distribution

We can use the `y = ..density..` to create a **density histogram** instead of a count histogram so that the area of the histogram boxes are equal to the chance of randomly selecting a unit in that bin:

```
midwest |>
  ggplot(aes(x = percollege)) +
  geom_histogram(aes(y = ..density..), binwidth = 1)
```

Continuous probability distribution



Why density?

Histograms with **density** on the y-axis are drawn so that the area of each box is equal to the proportion of units in the sample in that horizontal bin.

Easier to compare distributions across sample sizes.

Sum up all the area = 1 (but heights can go above 1)

3/ Sampling distribution

Key properties of sums and means

Suppose X_1, X_2, \dots, X_n is a simple random sample from a population distribution with mean μ (“mu”) and variance σ^2 (“sigma squared”)

Sample mean: $\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$

$$\bar{X}_n = \frac{X_1 + X_2 + \dots + X_n}{n}$$

...

\bar{X}_n is a random variable with a distribution!!


Sample means/proportions distribution

Sampling distributions are the probability distributions of an estimator like \bar{X}_n


When we have access to the full population, we can approximate the sampling distribution with repeated sampling.

```
library(infer)
midwest |>
  rep_slice_sample(n = 50, reps = 100) |>
  group_by(replicate) |>
  summarize(`Average Percent College` = mean(percollege)) |>
  ggplot(aes(x = `Average Percent College`)) +
  geom_histogram(mapping = aes(y = ..density..), binwidth = 0.5) +
  coord_cartesian(xlim = c(14, 23), ylim = c(0, 0.7)) +
  labs(title = "100 Repititions") +
  stat_function(fun = dnorm, args = c(mean(midwest$percollege), sd(midwest$percollege)),
    color = "indianred1", size = 1.5, xlim = c(14, 23))
```


100 Repititions




1,000 Repititions



10,000 Repititions



10,000 Repititions



Sampling distribution of the sample mean

Suppose X_1, X_2, \dots, X_n is a simple random sample from a population distribution with mean μ and variance σ^2 .

Expected value of the distribution of \bar{X}_n is the population mean, μ .

Standard error of the distribution of \bar{X}_n is approximately σ/\sqrt{n} :

$$SE \approx \frac{\text{population standard deviation}}{\sqrt{\text{sample size}}}$$


Unbiasedness

An estimator is **unbiased** when its expected value across repeated samples equals the population parameter of interest.

Sample mean of a simple random sample is **unbiased** for the population mean, $\mathbb{E}[\bar{X}_n] = \mu$

An estimator that isn't unbiased is called **biased**.

Precision vs accuracy



Law of large numbers


Law of large numbers

Let X_1, \dots, X_n be a simple random sample from a population with mean μ and finite variance σ^2 . Then, \bar{X}_n converges to μ as n gets large.

- Probability of \bar{X}_n being “far away” from μ goes to 0 as n gets big.
- The distribution of sample mean “collapses” to population mean.
- Can see this from the SE of \bar{X}_n : $SE = \sigma/\sqrt{n}$.
- Not necessarily true with a biased sample!


4/ Normal variables and the Central Limit Theorem

Normal random variable



- A **normal distribution** has a PDF that is the classic “bell-shaped” curve.
 - Extremely ubiquitous in statistics.
 - An r.v. is more likely to be in the center, rather than the tails.
- Three key properties of this PDF:
 - **Unimodal**: one peak at the mean.
 - **Symmetric** around the mean.
 - **Everywhere positive**: any real value can possibly occur.

Normal distribution



- A normal distribution can be affected by two values:
 - **mean/expected value** usually written as μ
 - **variance** written as σ^2 (standard deviation is σ)
 - Written $X \sim N(\mu, \sigma^2)$.
- **Standard normal distribution:** mean 0 and standard deviation 1.


Central limit theorem

Central limit theorem

Let X_1, \dots, X_n be a simple random sample from a population with mean μ and finite variance σ^2 . Then, \bar{X}_n will be approximately distributed $N(\mu, \sigma^2/n)$ in large samples.


- “Sample means tend to be normally distributed as samples get large.”
- \rightsquigarrow we know (an approx. of) the entire probability distribution of \bar{X}_n
 - Approximation is better as n goes up.
 - Does not depend on the distribution of X_i !

Empirical Rule for the Normal Distribution




- If $X \sim N(\mu, \sigma^2)$, then:

Empirical Rule for the Normal Distribution




- If $X \sim N(\mu, \sigma^2)$, then:
 - $\approx 68\%$ of the distribution of X is within 1 SD of the mean.

Empirical Rule for the Normal Distribution




- If $X \sim N(\mu, \sigma^2)$, then:
 - $\approx 68\%$ of the distribution of X is within 1 SD of the mean.
 - $\approx 95\%$ of the distribution of X is within 2 SDs of the mean.

Empirical Rule for the Normal Distribution



- If $X \sim N(\mu, \sigma^2)$, then:
 - $\approx 68\%$ of the distribution of X is within 1 SD of the mean.
 - $\approx 95\%$ of the distribution of X is within 2 SDs of the mean.
 - $\approx 99.7\%$ of the distribution of X is within 3 SDs of the mean.
- CLT + empirical rule: we'll know the rough distribution of estimation errors we should expect.

Where are we going?



We only get 1 sample. Can we learn about the population from that sample?