

# Gov 50: 1. Introduction

Matthew Blackwell

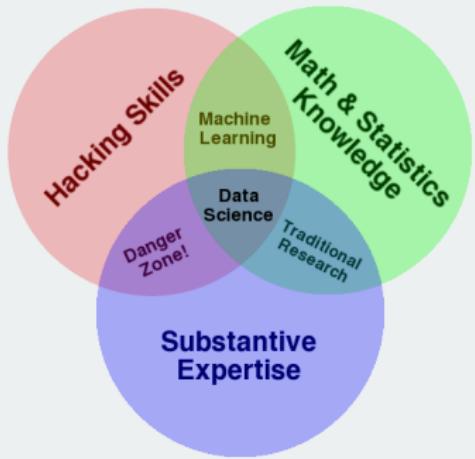
Harvard University

# Roadmap

1. Welcome and Motivation
2. Course Details

# 1/ Welcome and Motivation

# What is data science?



- **Data science:** wrangling, visualizing, and analyzing data to understand the world
- Who does data science? Tech companies, non-tech companies, nonprofits, governments.

# Glassdoor's No. 3 best job in the U.S. has seen job growth surge 480%

BY MEGHAN MALAS

March 08, 2022, 1:12 PM

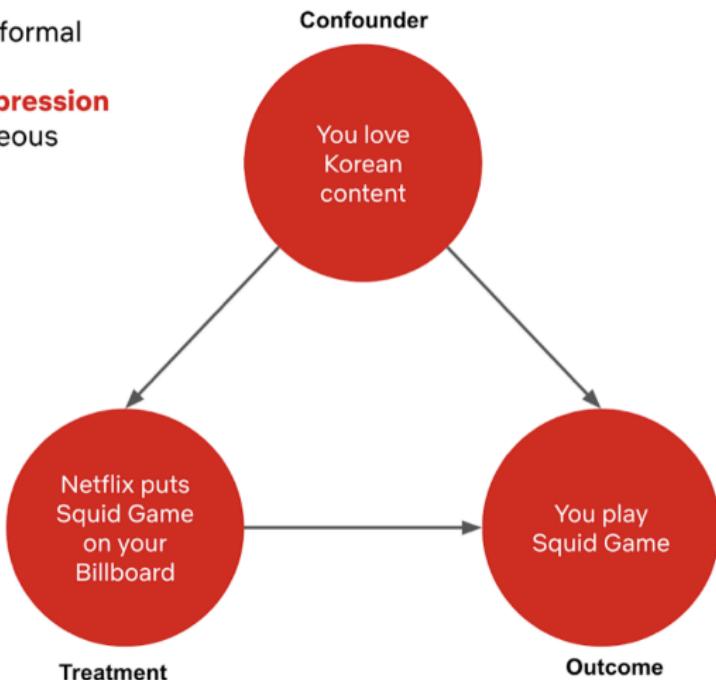


A COMMUTER BOARDS A BAY AREA RAPID TRANSIT (BART) TRAIN IN THE NEW MONTGOMERY STATION IN SAN FRANCISCO, CALIFORNIA, AS SEEN IN MARCH 2022. (PHOTOGRAPHER: DAVID PAUL MORRIS—BLOOMBERG/GETTY IMAGES)

**What problems are data scientists working on?**

# Causality

**Causal Inference** provides formal tools to tease out the true **incremental** value of an **impression** for each profile: Heterogeneous Treatment Effect (**HTE**)



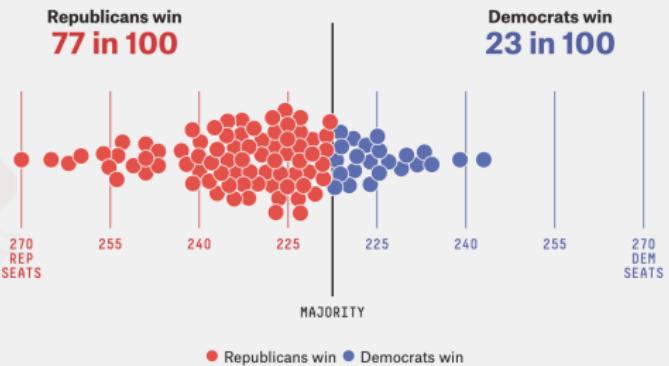
Compared to machine learning, causal inference allows us to build a robust framework that controls for confounders in order to estimate the true incremental impact to members

# Prediction

UPDATED 17 MINUTES AGO

## Republicans are *favored* to win the House

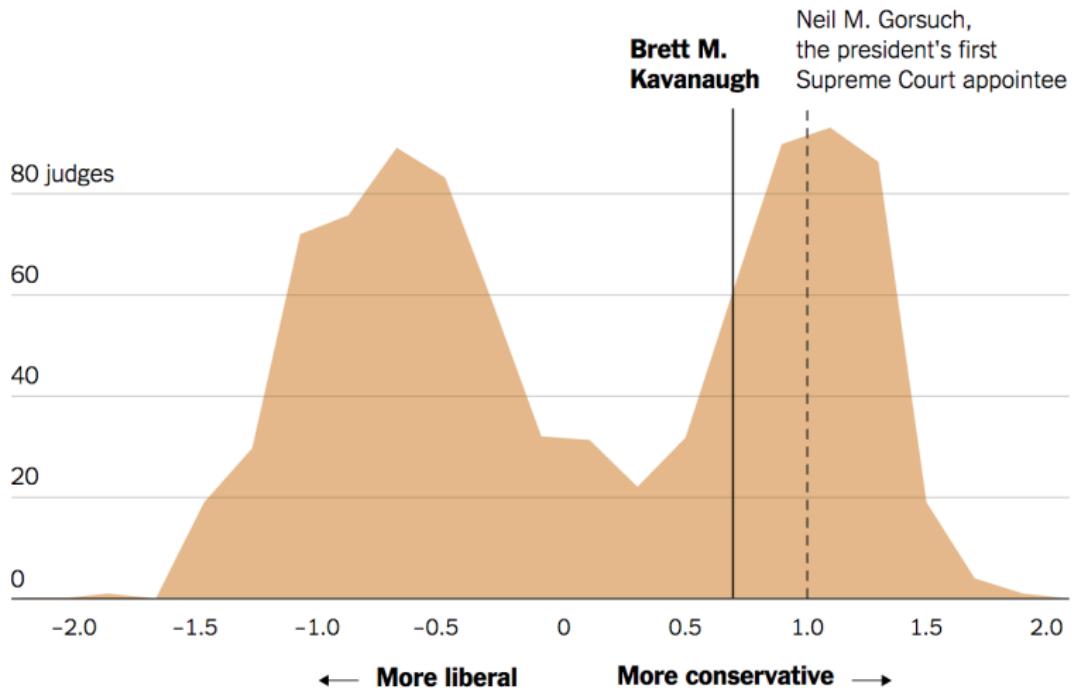
The Deluxe version of our model simulates the election 40,000 times to see which party wins the House most often. This sample of 100 outcomes gives you an idea of the range of scenarios the model considers possible.



We use numbers to express uncertainty. Upset wins are surprising but not impossible.

# Measurement

## How Kavanaugh's Ideology Compares With Other Federal Judges



Based on the campaign finance scores of all current and former federal district and court of appeals judges nominated since 1980. | Source: [Database on Ideology, Money in Politics, and Elections](#); Adam Bonica, Stanford University Department of Political Science; Maya Sen, Harvard University, Kennedy School of Government; Adam Chilton and Kyle Rozema, University of Chicago Law School.

# Understanding the socioeconomic world

HOME SEARCH

The New York Times

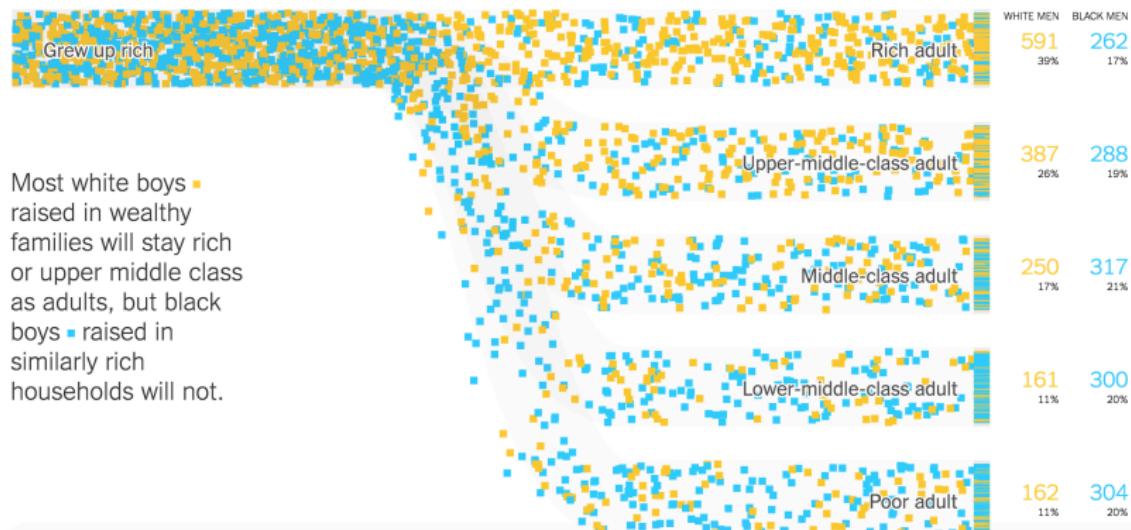
SHARE

## Extensive Data Shows Punishing Reach of Racism for Black Boys

By EMILY BADGER, CLAIRE CAIN MILLER, ADAM PEARCE and KEVIN QUEALY MARCH 19, 2018

Follow the lives of 5,734 boys who grew up in rich families ...

...and see where they end up as adults:



Adult outcomes reflect household incomes in 2014 and 2015.

# Making government work better

The screenshot shows the CityScore website. At the top, there's a navigation bar with links for 'CITY OF BOSTON' (Mayor Martin J. Walsh), 'PUBLIC NOTICES', 'PAY AND APPLY', 'FEEDBACK', and 'TRANSLATE'. Below the navigation is a large graphic with the word 'CITY SCORE' in large letters, with a small registered trademark symbol. The 'O' in 'SCORE' is replaced by a circular seal featuring a building and text. Below this graphic, the word 'CITYSCORE' is written in a bold, sans-serif font. A sub-headline below it reads: 'CityScore is an initiative designed to inform the Mayor and city managers about the overall health of the City at a moment's notice by aggregating key performance metrics into one number. Here we will provide you with an overview of the CityScore tool and data, but more importantly we will show you how we are using CityScore to make improvements across the City.' Underneath this, there's a section titled 'THE SCORE' with a horizontal scale from 0.0 to 1.0. The current score is 1.02, indicated by a central dot. Arrows on either side of the scale point left and right, with explanatory text: '< Scores below 1.0 indicate that performance is below the target.' and '> Scores above 1.0 indicate that performance is exceeding the target.'

Topic	Day	Week	Month	QTR
311 CALL CENTER PERFORMANCE			0.94	0.93
CODE ENFORCEMENT ON-TIME %	1.23	1.24	1.24	1.24
CODE ENFORCEMENT TRASH COLLECTION	1.25	1.18	1.15	1.10
GRAFFITI ON-TIME %	0.33	0.13	0.20	0.27
MISSSED TRASH ON-TIME %	1.21	1.20	1.20	1.20
PARKS MAINTENANCE ON-TIME %	0.82	0.83	0.90	0.88
POTHOLE ON-TIME %	1.25	0.88	0.69	0.67
SIGN INSTALLATION ON-TIME %	1.00	0.23	0.24	0.49
SIGNAL REPAIR ON-TIME %	1.25	1.25	1.10	1.09
STREETLIGHT ON-TIME %	0.55	0.60	0.56	0.54
TREE MAINTENANCE ON-TIME %	1.19	1.18	1.17	1.13
ON-TIME PERMIT REVIEWS	0.88	0.88	0.85	0.81
LIBRARY USERS	1.51	1.42	1.43	1.42
BPS ATTENDANCE				0.90
BFD RESPONSE TIME	0.91	0.94	0.94	0.94
BFD INCIDENTS	1.03	0.92	0.93	0.94
EMS RESPONSE TIME	0.90	0.84	0.84	0.84
PART 1 CRIMES	2.26	1.48	1.41	1.40

# Combining art and data to inform

Who Gets Miscounted In The Census ?

Black

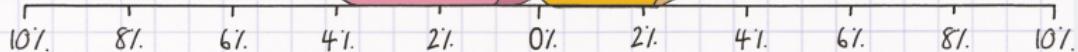


Hispanic

Native  
American

Asian

White

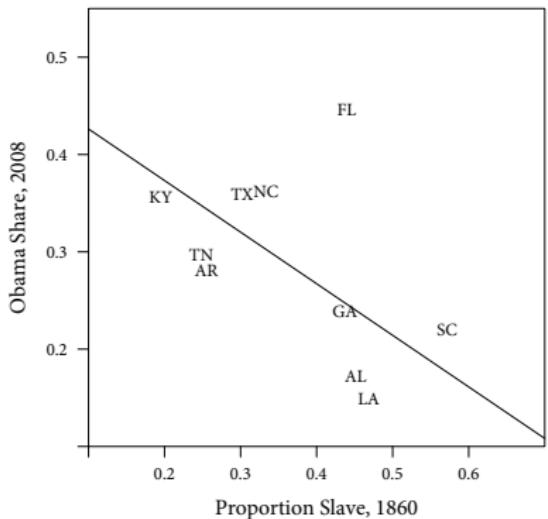


%. NOT COUNTED

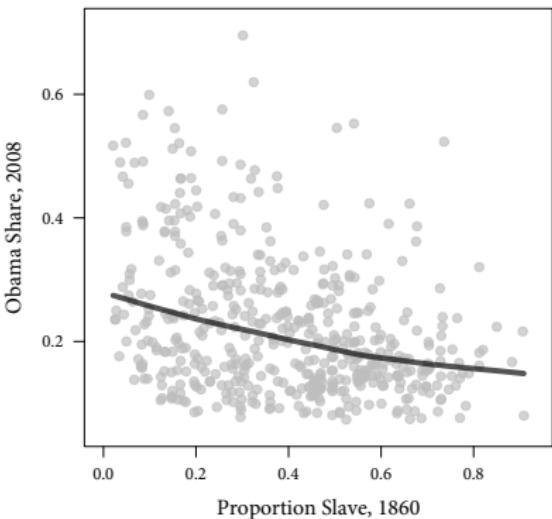
%. DOUBLE COUNTED OR FOUND

# Understanding how the past matters

States

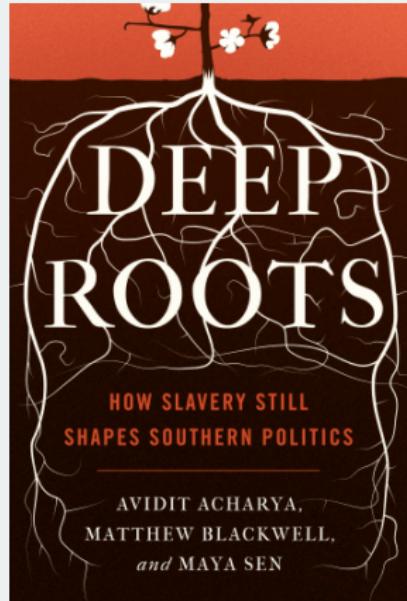
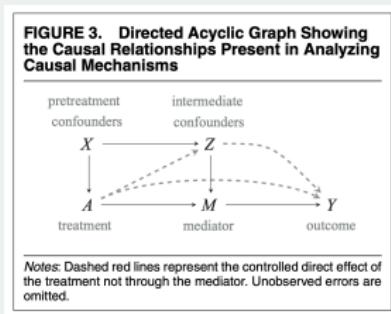


Counties



# **2/** Course Details

# About me



# What will you learn in this class?

- Summarize and visualize data
- Wrangle messy data into tidy forms
- Evaluate claims about causality
- Be able to use linear regression to analyze data
- Understand uncertainty in data analysis and how to quantify it
- Use professional tools like R, RStudio, git, and GitHub

# Teaching philosophy

- Deliberate pacing and tons of support.
- Emphasize intuition and computational approaches over mathematical equations.
- Practice, practice, practice.

# Pep talk, part I



Hadley Wickham (chief  
data scientist at RStudio)

*It's easy when you start out programming to get really frustrated and think, "Oh it's me, I'm really stupid," or, "I'm not made out to program." But, that is absolutely not the case. Everyone gets frustrated. I still get frustrated occasionally when writing R code. It's just a natural part of programming. So, it happens to everyone and gets less and less over time. Don't blame yourself. Just take a break, do something fun, and then come back and try again later.*

# Pep talk, part II



**Hadley Wickham**

@hadleywickham

...

The only way to write good code is to write tons of shitty code first. Feeling shame about bad code stops you from getting to good code

10:11 AM · Apr 17, 2015 · Echofon

---

892 Retweets 55 Quote Tweets 1,144 Likes

# Should I take this course?

- Prerequisites: **NONE** (no prior coding, statistics, data science)
- Gov 50 fulfills Gov methods requirement, data science track, and QRD
- Material useful to students interested in political science, sociology, economics, public policy, health policy, and many other fields in the social sciences.

# Class meetings

- Lectures:
  - Broad coverage of the course material.
  - Coding demonstrations (follow along with your laptop!)
  - Slides/videos will be posted to Canvas shortly before lecture.
- Section:
  - Guided practice through problems and concepts led by our amazing TFs.
  - Material in section will closely mirror assignments.
- Optional speaker series with industry data scientists, TBA!

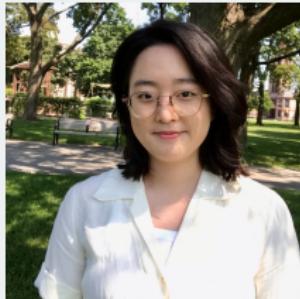
# Teaching fellows



Angelo Dagonel



Dorothy Manevich



Sooahn Shin



Dominic Valentino

# Computing

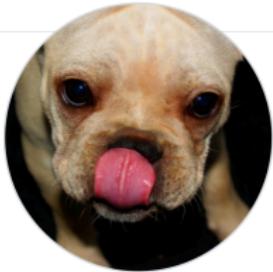
- We'll use the R statistical environment to analyze data
  - It's free
  - Extremely popular for data analysis
  - Academics, 538, NYT, Facebook, Google, Twitter, nonprofits, governments all use R.
  - Huge benefit to your resume to have R skills.
- Interface with R via a program called RStudio
- Problem Set 0 on the website helps get everything installed.
- Lots of help in section, study halls, office hours.

# git and GitHub



- Other core tools: git and GitHub
  - **Version control system:** an archive of project versions.
  - Allows you to revert back to old versions easily
  - Makes collaboration much more manageable.
- Will feel very odd at first, but you git used to it
- Why learn this now?
  - Knowing git/GitHub is a huge plus for data jobs.
  - Your GitHub profile can showcase your amazing new skills with data!

# Sample GitHub profile



Mara Averick  
batpigandme

[Follow](#)

At 742 followers · 136 following

[@tidyverse](#)  
[Missoula, MT](#)  
[@dataandme](#)

Achievements



Beta [Send feedback](#)

Overview Repositories 227 Projects Packages Stars 808

batpigandme / README.ind

Hieeeee! I'm Mara. 🐶

```
> # I'm the tidyverse developer advocate
> tidyverse::tidyverse_logo()

```

[Followers 742](#) [Follow @dataandme 47k](#)

Popular repositories

**night-owlish** Public

An RStudio, tmThemes, and Ace editor adaptation of @sdraa's Night Owl VS Code theme...

1 CSS 125 49

**tverse\_contributing** Public

\*contributing to the tidyverse\* from rstudio::conf(2018)

23

# GovCodes workshops

- Gov department providing supplemental GovCodes workshops to provide additional computing practice.
- First meeting: tomorrow! Be on the lookout for a sign-up email.
  - Topic: getting everything installed and working on your computer!
  - Good to attend if Problem Set 0 is giving you trouble.

# Textbook

- 3 primary textbooks (links on syllabus):
  - Modern Dive (free online)
  - “Quantitative Social Science: An Introduction in tidyverse” by Kosuke Imai (not free)
  - Introduction to Modern Statistics (free online)
- We'll move back and forth.
- Sometimes same material in two/three different books. Choose which helps most!

# Assignments

- Roughly weekly homeworks throughout semester
  - Posted on Thursday morning, due following Wednesday.
  - Dates on syllabus
  - Lowest score dropped.
- Two take-home “exams” which are just HWs done by yourself.
- Final project: a data essay
  - Find data, pose a research question, answer it using data.
  - Submitted as a public GitHub repository and website
  - First item in your public data portfolio

# Tutorials

- Getting practice with R can be overwhelming, so we'll introduce new skills through online tutorials.
- Guided practice on R, helping to introduce new concepts.
  - Low stakes/stress: graded simply on completion.
  - Due on Monday nights
- Lecture/HW won't be the first time you're trying some code!

# Ed discussion board

The screenshot shows the ed platform's interface. On the left, there's a sidebar with a 'New Thread' button, course categories like COURSES (GOV 50 selected), and sections like Drafts and Scheduled. The main area has a search bar and a 'Welcome' thread in the GOV 50 course. The right side features a 'New Post' interface with tabs for Question, Post, and Announcement. It includes a title input, category selection (General, Lectures, Sections, Problem Sets, Assignments, Social), and a rich text editor with a code snippet example. At the bottom, there are checkboxes for Pinned, Private, Anonymous, Anonymous Comments, and Megathread, along with a Draft saved message and a Post button.

ed GOV 50 – Ed Discussion

New Thread

COURSES

GOV 50

Gov JLR

Drafts 1

Scheduled

CATEGORIES

- General
- Lectures
- Sections
- Problem Sets
- Assignments
- Social

Search

Welcome!

General Matt Blackwell 2w

Cancel

New Post

Question Post Announcement

Title

Category General Lectures Sections Problem Sets Assignments Social

Code

I can use **markdown** and insert code easily:

```
mean(some_data)
```

Pinned Keep at top of thread list

Private Visible to you and staff only

Anonymous Hide your name from students

Anonymous Comments Allow anonymous comments

Megathread Resolvable comments

Draft saved

Post

# Grades

- Grade breakdown as follows:
  - R tutorials (10% of final grade)
  - Homeworks (40% of final grade)
  - Exams (30% of final grade)
  - Final project (20% of final grade)
- Final grade is curved
- **Bump-up:** we bump up grades of students close to the cutoff who make valuable contributions to the course.

# Study Halls

- Study Halls: a place to work on Gov 50 and get help.
  - Will happen weekly, exact number of hours will depend on enrollment.
  - Peer tutors with experience in statistics and R will be on hand to help you if you get stuck or have question.
  - Best to come in groups and work together, grab a tutor when stuck.
- Bottom line: **we want you to succeed in this class!**

# What should you do today?

- Try to get everything set up on your computer (Problem Set 0)
- Start Tutorial 1 on basics of R and data visualization
  - Can be done on the web before installing R on your computer.
- Respond to sign-up requests for GovCodes and section times.
- **Tell your friends:** data science is more fun with friends along for the journey.

# Gov 50: 2. R, RStudio, and Rmarkdown

Matthew Blackwell

Harvard University

# Roadmap

1. Working in Plain Text
2. Let's take a touR
3. Using Rmarkdown
4. Getting R bearings
5. Our first visualizations

# 1/ Working in Plain Text

# The two computer revolutions



## **The frontier of computing**

- Touch-based interfaces
- Single app at a time
- Little multitasking between apps
- Hides the file system



## **Where statistical computing lives**

- Windows and pointers
- Multi-tasking, multiple windows
- Works heavily with the file system
- Underneath it's UNIX and the command line

# Plain-text tools for data analysis

## The Plain Person's Guide

~/>\_

## to Plain Text Social Science

Kieran Healy

- Often free, open-sourced, and powerful.
- Large, friendly communities around them.
- Tons of resources
- But... far from the touch-based paradigm of modern computing
- So why use them?

The process of data  
science is intrinsically  
messy

# Office vs engineering model of computing

What's real in the project? How are changes managed?

## In the Office model

- Formatted documents are real.
- Intermediate outputs copy/pasted into documents.
- Changes are tracked inside files.
- Final output is the file you are working on (e.g., Word file or maybe converted to a PDF).

## In the Engineering model

- Plain-text files are real.
- Intermediate outputs are produced via code, often inside documents.
- Changes are tracked outside files.
- Final outputs are assembled programmatically and converted to desired output format.

# Pros and cons to each approach

- Office model:
  - Everyone knows Word, Excel, Google Docs.
  - “Track changes” is powerful and easy.
  - Wait, how did I make this figure?
  - Which version of my code made this table?
  - `Blackwell_report_final_submitted_edits_FINAL_v2.docx`
- Engineering model:
  - Plain text is universally portable.
  - Push button, recreate analysis.
  - Why won’t R just do what I want!
  - Version control is a pain.
  - Object of type 'closure' is not subsettable

We'll tend toward the Engineering model because it's better suited to keep the mess in check.

**2/** Let's take a touR

# R versus RStudio

```
R version 4.2.1 (2022-06-23) -- "Funny-Looking Kid"  
Copyright (C) 2022 The R Foundation for Statistical Computing  
Platform: aarch64-apple-darwin20 (64-bit)  
  
R is free software and comes with ABSOLUTELY NO WARRANTY.  
You are welcome to redistribute it under certain conditions.  
Type 'licence()' or 'licence()' for distribution details.  
  
Natural language support but running in an English locale  
  
R is a collaborative project with many contributors.  
Type 'contributors()' for more information and  
'citation()' on how to cite R or R packages in publications.  
  
Type 'demo()' for some demos, 'help()' for on-line help, or  
'help.start()' for an HTML browser interface to help.  
Type 'q()' to quit R.  
? |
```



RStudio interface showing a project named "cars-project".

**Project View:**

- File: cars-project.Rmd
- Knit on Save
- Knit
- Run
- Environment
- History
- Connections
- Tutorial
- Import Dataset (158 MiB)
- Global Environment

The Global Environment pane shows "Environment is empty".

**File Explorer:**

- New Folder
- New Blank File
- Delete
- Rename
- More

Home > Dropbox > workland > tmp > cars-project

Name	Size	Modified
cars-project.Rproj	205 B	Sep 5, 2022, 9:57 PM
data		
cars-project.Rmd	845 B	Sep 5, 2022, 9:58 PM
figures		

**Console:**

```
R 4.2.1 - ~/Dropbox/workland/tmp/cars-project/
> 5 + 10
[1] 15
> library(tidyverse)
-- Attaching packages --
→ ggplot2 3.3.6     → purrr  0.3.4
→ tibble  3.1.8     → dplyr   1.0.10
→ tidyverse 1.3.2    → stringr 1.4.1
→ readr   2.1.2     → forcats 0.5.2
-- Conflicts --
* dplyr::filter() masks stats::filter()
* dplyr::lag()   masks stats::lag()
```

cars-project - RStudio

cars-project

cars-project.Rmd x Knit on Save Knit Run Outline

Source Visual

Environment History Connections Tutorial Import Dataset 158 MiB List Global Environment

Environment is empty

Write notes, paper in Rmarkdown

```
1 - ---
2 #title: "Car Project"
3 author: "Matthew Blackwell"
4 date: "2022-09-06"
5 output: pdf_document
6 -
7
8 ````{r setup, include=FALSE}
9 knitr::opts_chunk$set(echo = TRUE)
10 -
11
12 ## R Markdown
13
14 This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.
15
16 When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:
17
18 ````{r cars}
19 summary(cars)
20 -
21 Car Project
```

R Markdown

Files Plots Packages Help Viewer Presentation

New Folder New Blank File Delete Rename More

Home > Dropbox > workland > tmp > cars-project

Name	Size	Modified
cars-project.Rproj	205 B	Sep 5, 2022, 9:57 PM
data		
cars-project.Rmd	845 B	Sep 5, 2022, 9:58 PM
figures		

Console Background jobs

R 4.2.1 - ~/Dropbox/workland/tmp/cars-project/

```
> 5 + 10
[1] 15
> library(tidyverse)
-- Attaching packages -- tidyverse 1.3.2 --
✓ ggplot2 3.3.6   ✓ purrr  0.3.4
✓ tibble  3.1.8   ✓ dplyr   1.0.10
✓ tidyr   1.2.0    ✓ stringr 1.4.1
✓ readr   2.1.2    ✓ forcats 0.5.2
-- Conflicts -- tidyverse_conflicts() --
#> dplyr::filter() masks stats::filter()
#> dplyr::lag()   masks stats::lag()
```

cars-project - RStudio

cars-project.Rmd x Knit on Save Knit Run Addins Environment History Connections Tutorial Import Dataset 158 MiB Global Environment List Environment is empty

```
1 - ---  
2 title: "Car Project"  
3 author: "Matthew Blackwell"  
4 date: "2022-09-06"  
5 output: pdf_document  
6 - ---  
7  
8 ```{r setup, include=FALSE}  
9 knitr::opts_chunk$set(echo = TRUE)  
10 - ---  
11  
12 ## R Markdown  
13  
14 This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.  
15  
16 When you click the Knit button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:  
17  
18 ```{r cars}  
19 summary(cars)  
20 ```  
2:1 □ Car Project R Markdown
```

Files Plots Packages Help Viewer Presentation New Folder New Blank File Delete Rename More Home > Dropbox > workland > tmp > cars-project Name Size Modified ..

- cars-project.Rproj 205 B Sep 5, 2022, 9:57 PM
- data
- cars-project.Rmd 845 B Sep 5, 2022, 9:58 PM
- figures

Console Background Jobs

R 4.2.1 - ~/Dropbox/workland/tmp/cars-project/

```
> 5 + 10  
[1] 15  
> library(tidyverse)  
-- Attaching packages -- tidyverse 1.3.2 --  
✓ ggplot2 3.3.6 ✓ purrr 0.3.4  
✓ tibble 3.1.8 ✓ dplyr 1.0.10  
✓ tidyv 1.2.0 ✓ stringr 1.4.1  
✓ readr 2.1.2 ✓ forcats 0.5.2  
-- Conflicts -- tidyverse_conflicts() --  
# dplyr::filter() masks stats::filter()  
# dplyr::lag() masks stats::lag()  
>  
>  
>  
>  
>  
>
```

Console: run code, send code to here, inspect output

Environment History Connections Tutorial

Import Dataset 158 MB

Global Environment

Environment is empty

Files Plots Packages Help Viewer Presentation

New Folder New Blank File Delete Rename More

Home > Dropbox > workland > tmp > cars-project

Name	Size	Modified
cars-project.Rproj	205 B	Sep 5, 2022, 9:57 PM
data		
cars-project.Rmd	845 B	Sep 5, 2022, 9:58 PM
figures		

Project files, plots, and help

**cars-project - RStudio**

**cars-project.Rmd** x Knit on Save Knit Run Addins Outline Source Visual

```
1 - ---  
2 title: "Car Project"  
3 author: "Matthew Blackwell"  
4 date: "2022-09-06"  
5 output: pdf_document  
6 - ---  
7  
8 ```{r setup, include=FALSE}  
9 knitr:::opts_chunk$set(echo = TRUE)  
10 - ---  
11  
12 ## R Markdown  
13  
14 This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.  
15  
16 When you click the Knit button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:  
17  
18 ```{r cars}  
19 summary(cars)  
20 ```  
2:1 □ Car Project R Markdown
```

Console Background Jobs

```
R 4.2.1 - ~/Dropbox/workland/tmp/cars-project/  
> 5 + 10  
[1] 15  
> library(tidyverse)  
-- Attaching packages -- tidyverse 1.3.2 --  
✓ ggplot2 3.3.6 ✓ purrr 0.3.4  
✓ tibble 3.1.8 ✓ dplyr 1.0.10  
✓ tidyv 1.2.0 ✓ stringr 1.4.1  
✓ readr 2.1.2 ✓ forcats 0.5.2  
-- Conflicts -- tidyverse_conflicts() --  
* dplyr::filter() masks stats::filter()  
* dplyr::lag() masks stats::lag()  
>  
>  
>  
>  
>
```

Environment History Connections Tutorial Import Dataset 158 MiB Global Environment List

Environment is empty

Interacting with R objects, working with git, running local tutorials

Files Plots Packages Help Viewer Presentation

New Folder New Blank File Delete Rename More

Home > Dropbox > workland > tmp > cars-project

Name	Size	Modified
cars-project.Rproj	205 B	Sep 5, 2022, 9:57 PM
data		
cars-project.Rmd	845 B	Sep 5, 2022, 9:58 PM
figures		

# 3/ Using Rmarkdown

# The acts of coding

```
library(ggplot2)  
ggplot(mtcars, aes(x = wt, y = mpg)) +  
  geom_point()
```

Figure: 1. Writing code

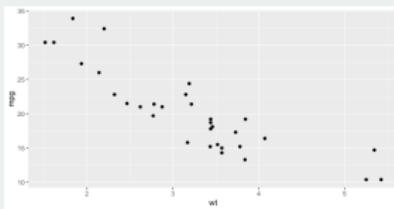


Figure: 2. Looking at output

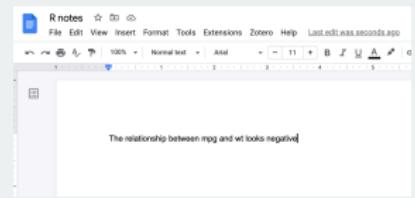


Figure: 3. Taking notes

How to do all of these efficiently?

# Rmarkdown files to the rescue

```
notes.Rmd x
Source Visual
1 # ...
2 title: "Car Project"
3 output: pdf_document
4 date: "2022-09-04"
5 ...
6
7 ```{r setup, include=FALSE}
8 knitr::opts_chunk$set(echo = TRUE)
9 library(ggplot2)
10 ...
11
12
13
14 Now I will produce a scatterplot of car weight against
15 mileage per gallon:
16 ...
17 ```{r}
18 ggplot(mtcars, aes(x = wt, y = mpg)) +
19   geom_point()
20 ...
21 As we can see, this relationship is negative.
```

Figure: Rmarkdown file

Keep code and notes  
together in plain text

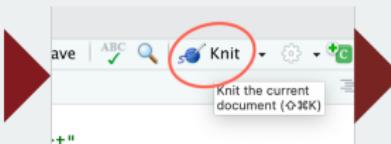


Figure: Knit in R

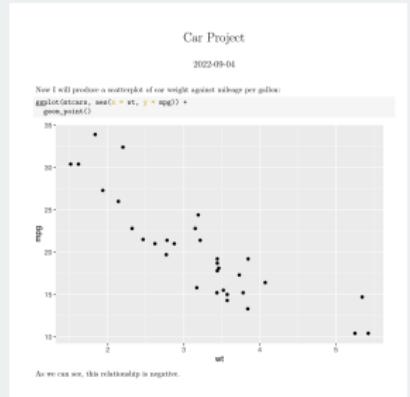


Figure: PDF output

Produce nice-looking  
outputs in different  
formats

# Markdown: formatting in plain text

Non-code text in Rmd files is plain text with formatting instructions

## syntax

```
Plain text
End a line with two spaces to start a new paragraph.
*italics* and _italics_
**bolds** and __bold__
superscript^2
~~strikethrough~~
[link](www.rstudio.com)

# Header 1

## Header 2

### Header 3

#### Header 4

##### Header 5

###### Header 6

endash: --
emdash: ---
ellipsis: ...
inline equation: $A = \pi * r^2$
image: 

horizontal rule (or slide break):

***

> block quote

* unordered list
* item 2
  + sub-item 1
  + sub-item 2

1. ordered list
2. item 2
  + sub-item 1
  + sub-item 2
```

## becomes

Plain text  
End a line with two spaces to start a new paragraph.  
*italics* and *italics*  
**bold** and **bold**  
superscript<sup>2</sup>  
~~strikethrough~~  
[link](#)

## Header 1

## Header 2

## Header 3

### Header 4

#### Header 5

##### Header 6

endash: --  
emdash: ---  
ellipsis: ...  
inline equation:  $A = \pi * r^2$   
image: ![]



horizontal rule (or slide break):

## block quote

- unordered list
  - item 2
    - sub-item 1
    - sub-item 2
1. ordered list
  2. item 2
    - sub-item 1
    - sub-item 2

```
---
```

```
title: "Car Project"
author: "Matthew Blackwell"
date: "2022-09-06"
output: pdf_document
```

```
--
```

```
```{r setup, include=FALSE}
knitr::opts_chunk$set(echo = TRUE)
```
```

```
## R Markdown
```

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <<http://rmarkdown.rstudio.com>>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
```{r cars}
summary(cars)
```
```

```
## Including Plots
```

You can also embed plots, for example:

```
```{r pressure, echo=FALSE}
plot(pressure)
```
```

Header contains metadata and sets options about the whole document

## Code Chunk



Plain text with markdown formatting



Can "play" chunks interactively



Chunks can have names and options

Code chunks replaced with output when Knitted



# Remember what's real

Options

General

Code

Console

Appearance

Pane Layout

Packages

R Markdown

Python

Sweave

Spelling

Git/SVN

Publishing

Terminal

Accessibility

Basic   Graphics   Advanced

**R Sessions**

Default working directory (when not in a project):

Restore most recently opened project at startup

Restore previously open source documents at startup

**Workspace**

Restore .RData into workspace at startup

Save workspace to .RData on exit:

**History**

Always save history (even when not saving .RData)

Remove duplicate entries in history

**Other**

Wrap around when navigating to previous/next tab

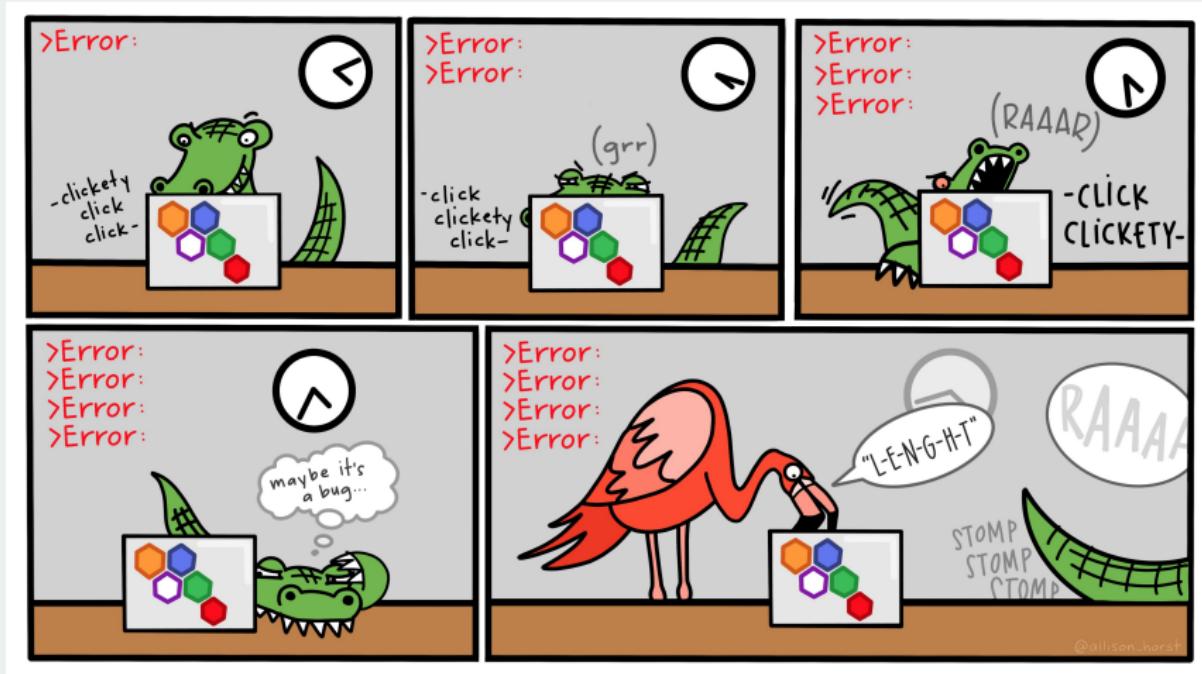
Automatically notify me of updates to RStudio

Send automated crash reports to RStudio

# 4| Getting R bearings

**Try to type your code by  
hand**

# Typing speeds up the try-fail cycle



Physically typing the code is best way to familiarize yourself with R and the try-fail-try-fail-try-succeed cycle

# What R looks like

Code that you can type and run:

```
## Any R code that begins with the # character is a comment
## Comments are ignored by R

my_numbers <- c(4, 8, 15, 16, 23, 42) # Anything after # is also a comment
```

Output from code prefixed by ## by convention:

```
my_numbers
```

```
## [1] 4 8 15 16 23 42
```

Output also has a counter in brackets when over one line:

```
letters
```

```
##  [1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l"
## [13] "m" "n" "o" "p" "q" "r" "s" "t" "u" "v" "w" "x"
## [25] "y" "z"
```

# Everything in R has a name

```
my_numbers # just created this
```

```
## [1] 4 8 15 16 23 42
```

```
letters # this is built into R
```

```
## [1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l"
```

```
## [13] "m" "n" "o" "p" "q" "r" "s" "t" "u" "v" "w" "x"
```

```
## [25] "y" "z"
```

```
pi # also built in
```

```
## [1] 3.14
```

Some names are forbidden (NA, TRUE, FALSE, etc) or strongly not recommended (c, mean, table)

# We do things in R with functions

Functions take in objects, perform actions, and return outputs:

```
mean(x = my_numbers)
```

```
## [1] 18
```

- `x` is the argument name,
- `my_numbers` is what we're passing to the that argument

If you omit the argument name, R will assume the default order:

```
mean(my_numbers)
```

```
## [1] 18
```

# Getting help with R

How do we know the default argument order? Look to help files:

```
help(mean)  
?mean # shorter
```

- Sometimes inscrutable, so look elsewhere:
  - Google, StackOverflow, Twitter, RStudio Community.
  - Ask on Ed or on class Slack.
  - Come to section, office hours, study hall.
- Get help **early** before becoming too frustrated!
  - Easy to overlook small issues like missing commas, etc.

# Functions live in packages

Packages are bundles of functions written by other users that we can use.

Install packages using `install.packages()` to have them on your machine:

```
install.packages("ggplot2")
```

Load them into your R session with `library()`:

```
library(ggplot2)
```

Now we can use any function provided by ggplot2.

# Functions live in packages

We can also use the `mypackage::` prefix to access package functions without loading:

```
knitr::kable(head(mtcars))
```

|                   | mpg  | cyl | disp | hp  | drat | wt   | qsec | vs | am | gear | carb |
|-------------------|------|-----|------|-----|------|------|------|----|----|------|------|
| Mazda RX4         | 21.0 | 6   | 160  | 110 | 3.90 | 2.62 | 16.5 | 0  | 1  | 4    | 4    |
| Mazda RX4 Wag     | 21.0 | 6   | 160  | 110 | 3.90 | 2.88 | 17.0 | 0  | 1  | 4    | 4    |
| Datsun 710        | 22.8 | 4   | 108  | 93  | 3.85 | 2.32 | 18.6 | 1  | 1  | 4    | 1    |
| Hornet 4 Drive    | 21.4 | 6   | 258  | 110 | 3.08 | 3.21 | 19.4 | 1  | 0  | 3    | 1    |
| Hornet Sportabout | 18.7 | 8   | 360  | 175 | 3.15 | 3.44 | 17.0 | 0  | 0  | 3    | 2    |
| Valiant           | 18.1 | 6   | 225  | 105 | 2.76 | 3.46 | 20.2 | 1  | 0  | 3    | 1    |

# 5/ Our first visualizations

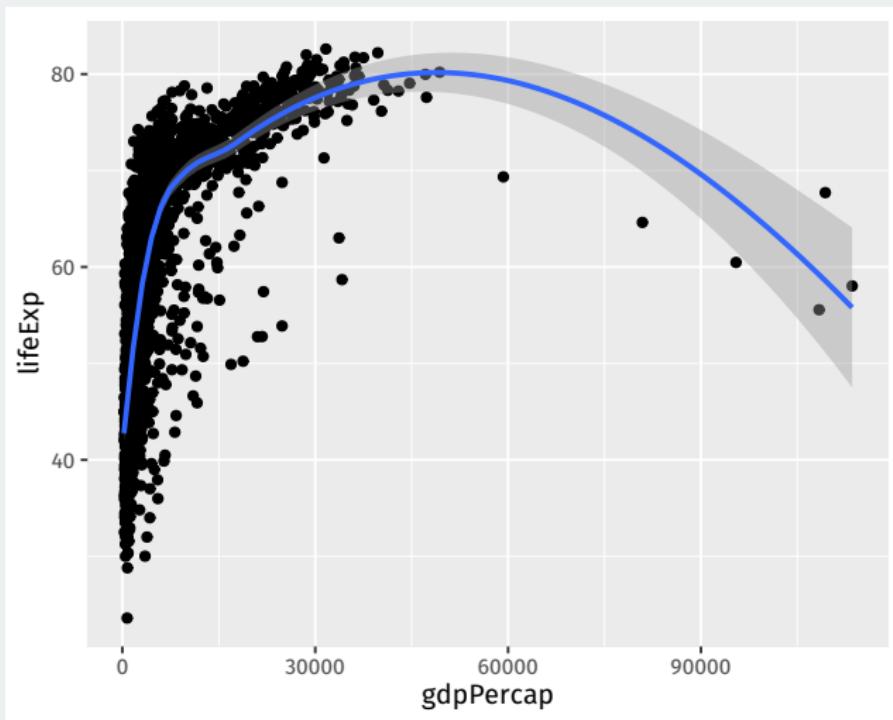
# Gapminder data

```
library(gapminder)
gapminder

## # A tibble: 1,704 x 6
##   country     continent year lifeExp      pop gdpPe~1
##   <fct>       <fct>    <int>   <dbl>     <int>   <dbl>
## 1 Afghanistan Asia     1952     28.8  8425333    779.
## 2 Afghanistan Asia     1957     30.3  9240934    821.
## 3 Afghanistan Asia     1962     32.0 10267083    853.
## 4 Afghanistan Asia     1967     34.0 11537966    836.
## 5 Afghanistan Asia     1972     36.1 13079460    740.
## 6 Afghanistan Asia     1977     38.4 14880372    786.
## 7 Afghanistan Asia     1982     39.9 12881816    978.
## 8 Afghanistan Asia     1987     40.8 13867957    852.
## 9 Afghanistan Asia     1992     41.7 16317921    649.
## 10 Afghanistan Asia    1997     41.8 22227415    635.
## # ... with 1,694 more rows, and abbreviated variable
## #   name 1: gdpPerCap
```

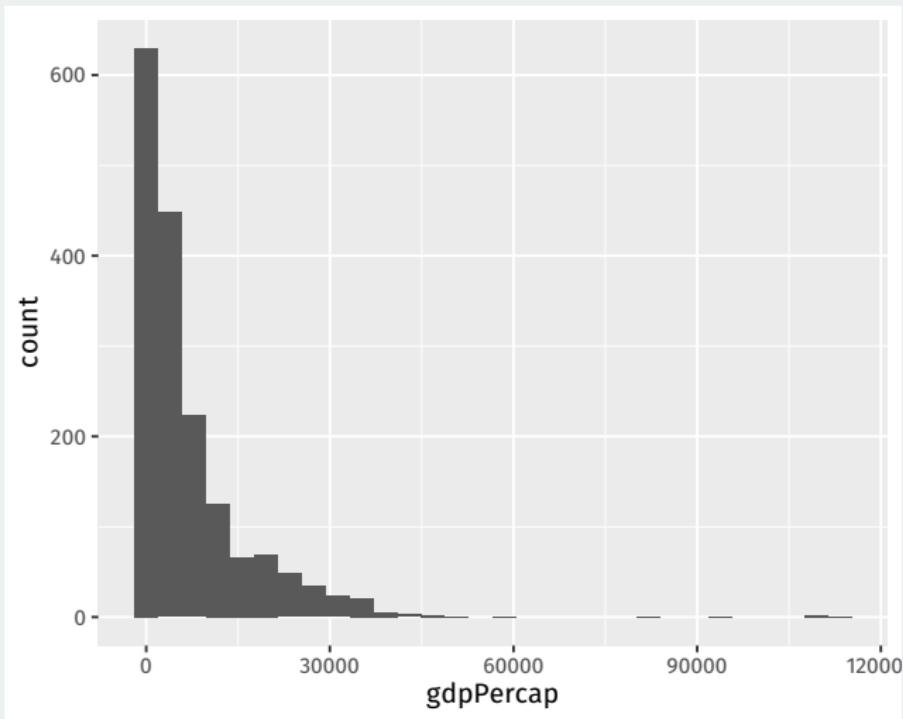
# Plotting life expectancy over time

```
ggplot(gapminder, mapping = aes(x = gdpPercap, y = lifeExp)) +  
  geom_point() + geom_smooth(method = "loess")
```



# A histogram of GDP per capita

```
ggplot(gapminder, mapping = aes(x = gdpPercap)) +  
  geom_histogram()
```



# Gov 50: 3. Data Visualization

Matthew Blackwell

Harvard University

# Roadmap

1. Building plots by layers
2. Histograms and boxplots
3. Grouped data

# 1/ Building plots by layers

# Midwest data

```
midwest
```

```
## # A tibble: 437 x 28
##   PID county state area popto~1 popde~2 popwh~3 popbl~4 popam~5
##   <int> <chr>  <chr> <dbl>  <int>   <dbl>  <int>   <int>   <int>
## 1 561 ADAMS   IL    0.052  66090  1271.  63917  1702    98
## 2 562 ALEXANDER IL    0.014  10626   759    7054  3496    19
## 3 563 BOND    IL    0.022  14991  681.   14477  429     35
## 4 564 BOONE   IL    0.017  30806  1812.  29344  127     46
## 5 565 BROWN   IL    0.018  5836   324.   5264   547     14
## 6 566 BUREAU  IL    0.05   35688  714.   35157  50      65
## 7 567 CALHOUN IL    0.017  5322   313.   5298    1      8
## 8 568 CARROLL IL    0.027  16805  622.   16519  111     30
## 9 569 CASS    IL    0.024  13437  560.   13384  16      8
## 10 570 CHAMPAIGN IL   0.058  173025  2983.  146506  16559   331
## # ... with 427 more rows, 19 more variables: popasian <int>,
## #   popother <int>, percwhite <dbl>, percblack <dbl>,
## #   percamerindan <dbl>, percasiain <dbl>, percother <dbl>,
## #   popadults <int>, perchsd <dbl>, percollege <dbl>, percprof <dbl>,
## #   poppovertyknown <int>, percpovertyknown <dbl>,
## #   percbelowpoverty <dbl>, percchildbelowpovert <dbl>,
## #   percadultpoverty <dbl>, percelderlypoverty <dbl>, ...
```

# Building up a graph in pieces

Create ggplot object and direct it to the correct data:

```
p <- ggplot(data = midwest)
```

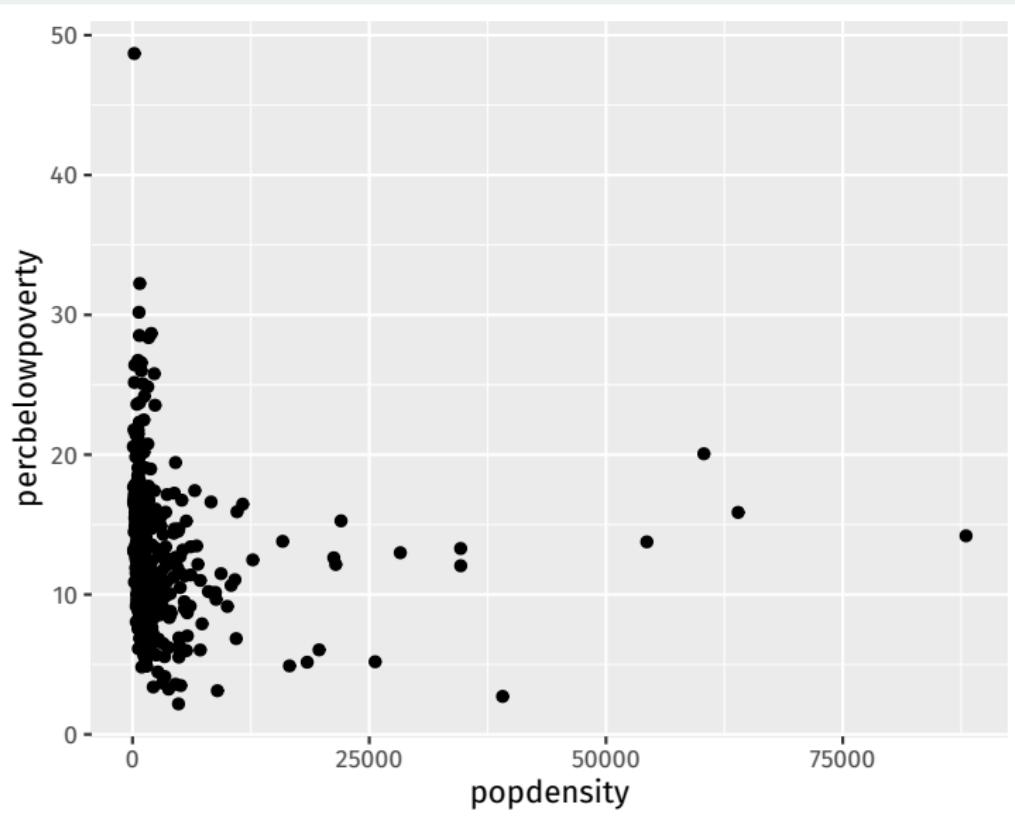
**Mapping:** tell ggplot what visual aesthetics correspond to which variables

```
p <- ggplot(data = midwest,  
             mapping = aes(x = popdensity,  
                            y = percbelowpoverty))
```

Other aesthetic mappings: color, shape, size, etc.

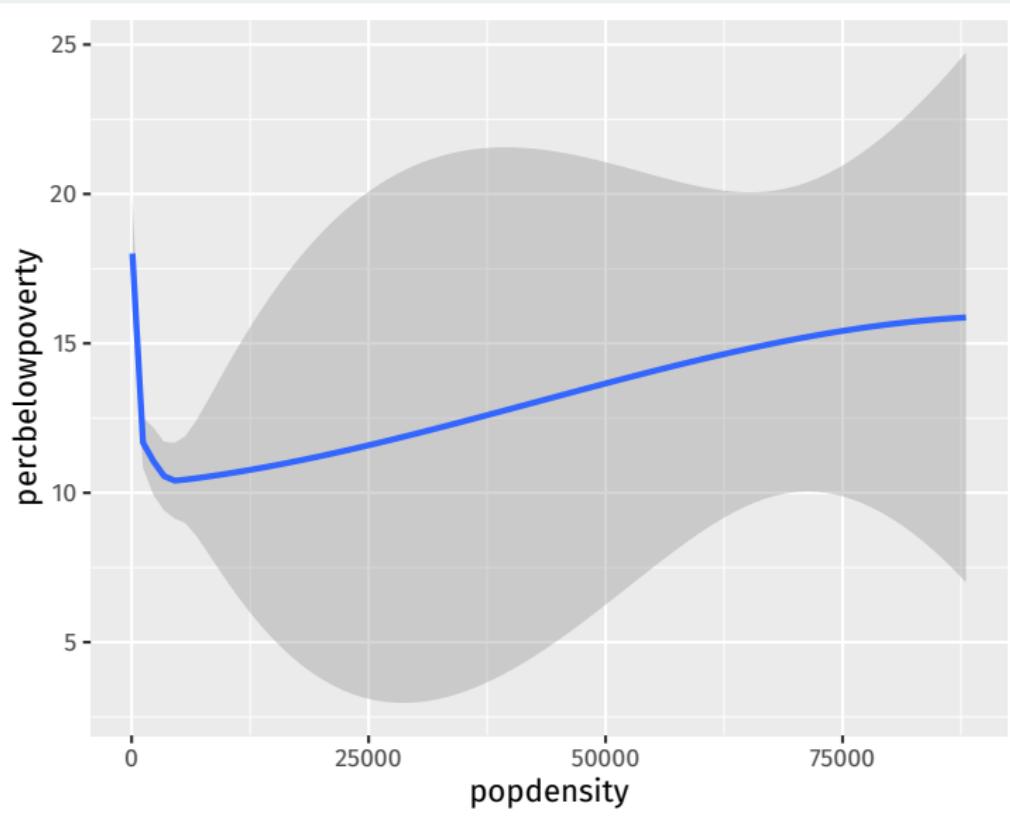
# Adding a geom layer

```
ggplot(data = midwest,
        mapping = aes(x = popdensity,
                      y = percbelowpoverty)) +
  geom_point()
```



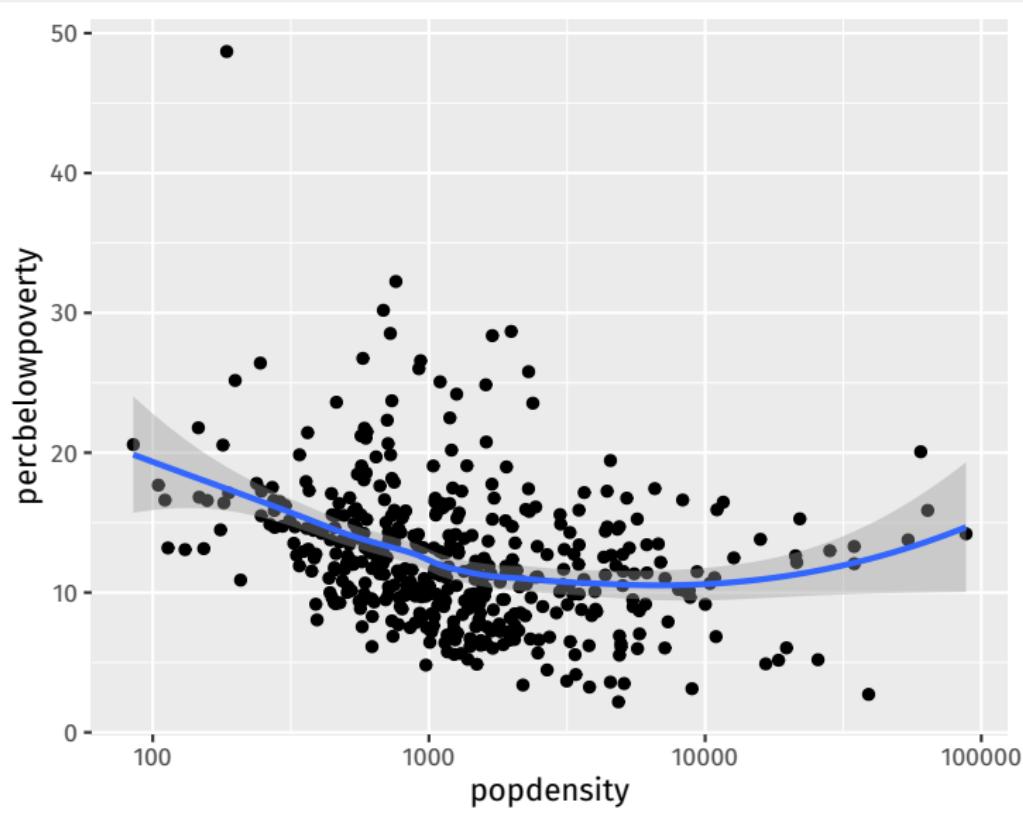
# Trying a new geom

```
ggplot(data = midwest,
        mapping = aes(x = popdensity,
                      y = percbelowpoverty)) +
  geom_smooth()
```



# Layering geoms is additive

```
ggplot(data = midwest,
        mapping = aes(x = popdensity,
                      y = percbelowpoverty)) +
  geom_point() +
  geom_smooth() +
  scale_x_log10()
```

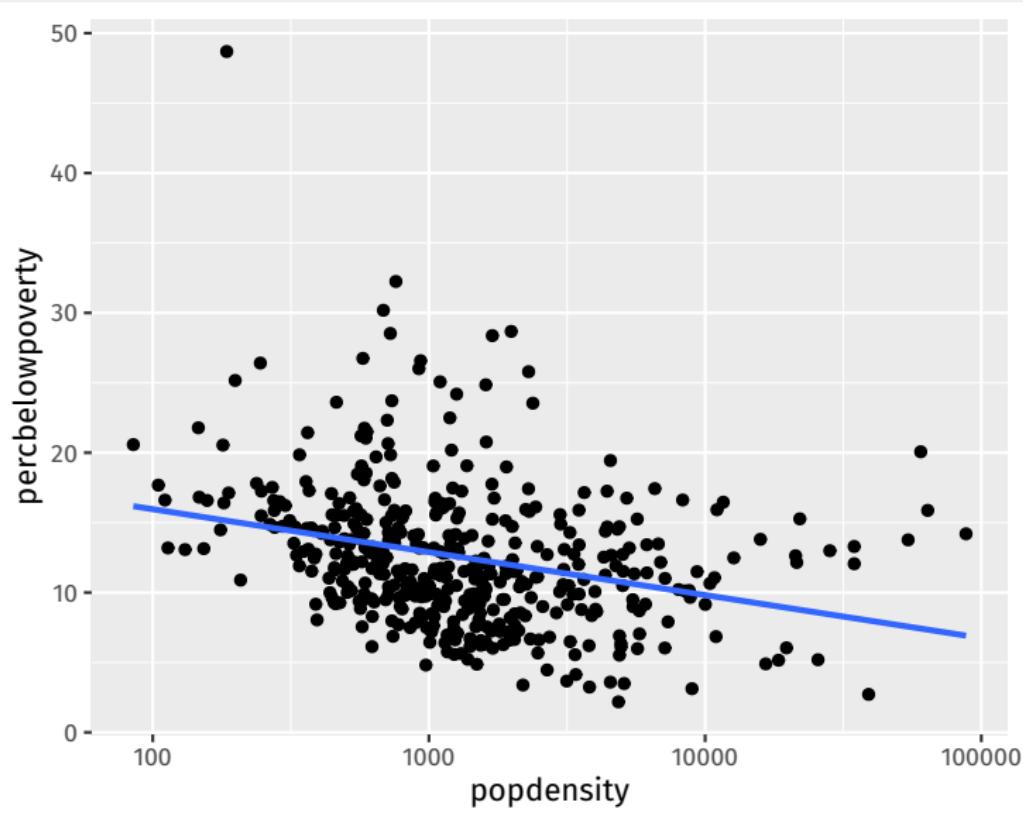


# Geoms are functions

Geoms can take arguments:

```
ggplot(data = midwest,
       mapping = aes(x = popdensity,
                     y = percbelowpoverty)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) +
  scale_x_log10()
```

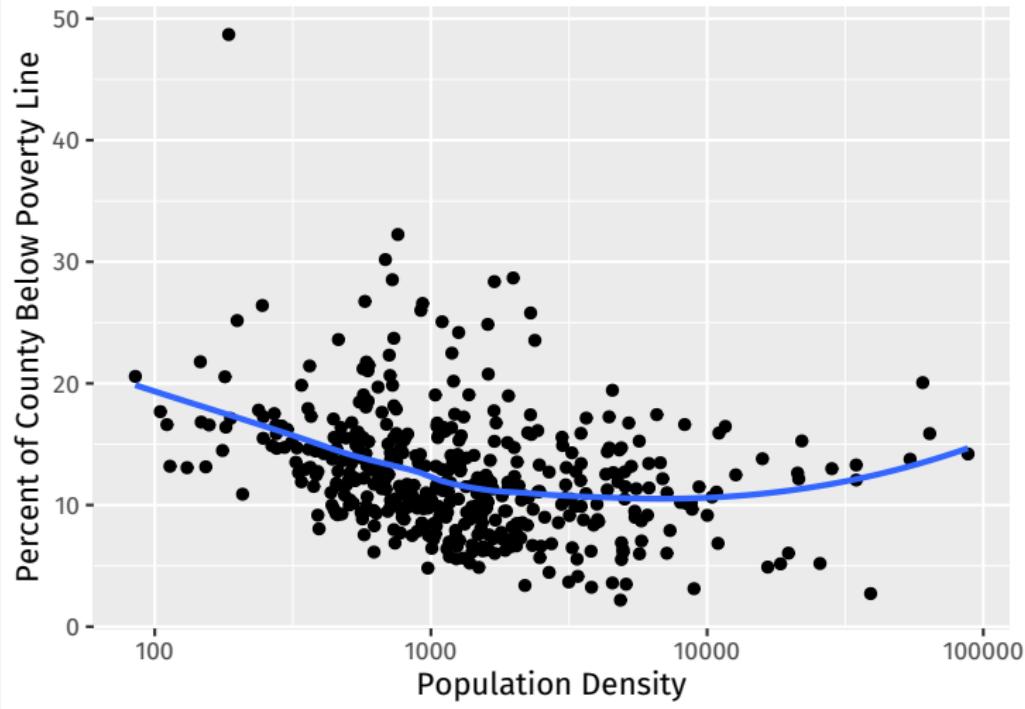
Tells geom\_smooth to do a linear fit with no error region



# Adding informative labels

```
ggplot(data = midwest,
        mapping = aes(x = popdensity,
                      y = percbelowpoverty)) +
  geom_point() +
  geom_smooth(method = "loess", se = FALSE) +
  scale_x_log10() +
  labs(x = "Population Density",
       y = "Percent of County Below Poverty Line",
       title = "Poverty and Population Density",
       subtitle = "Among Counties in the Midwest",
       source = "US Census, 2000")
```

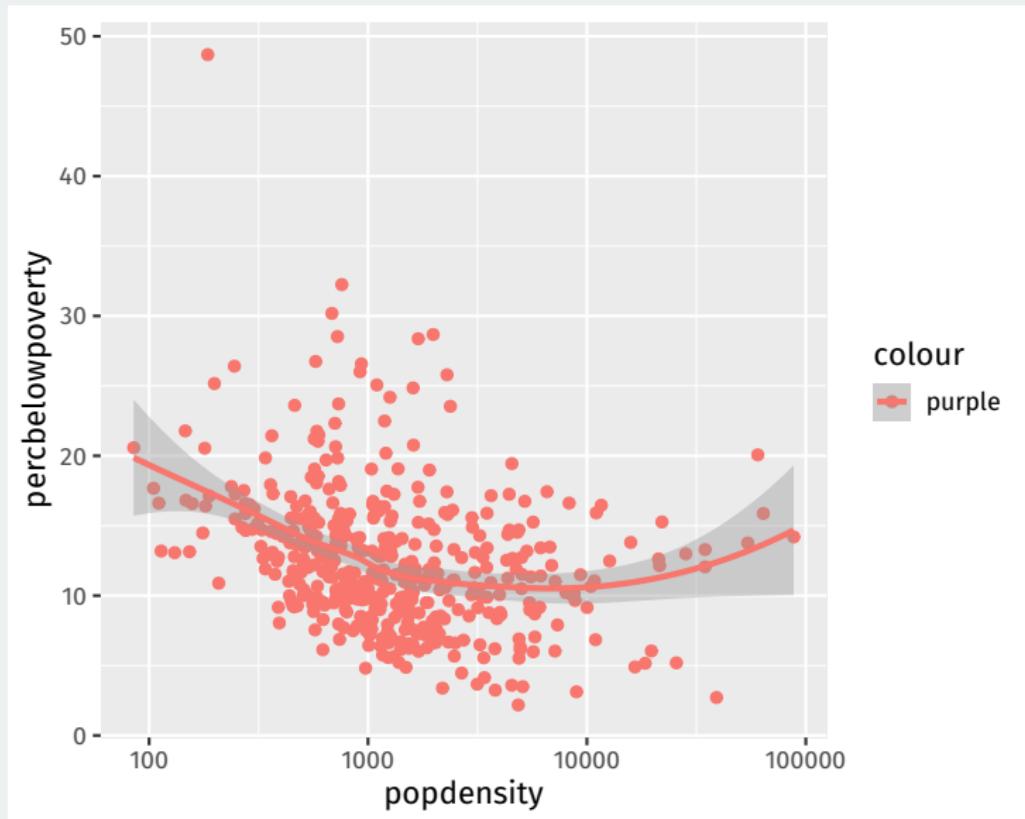
## Poverty and Population Density Among Counties in the Midwest



# Mapping vs setting aesthetics

```
ggplot(data = midwest,  
       mapping = aes(x = popdensity,  
                      y = percbelowpoverty,  
                      color = "purple")) +  
  geom_point() +  
  geom_smooth() +  
  scale_x_log10()
```

# Wait what?



# Mapping always refers to variables

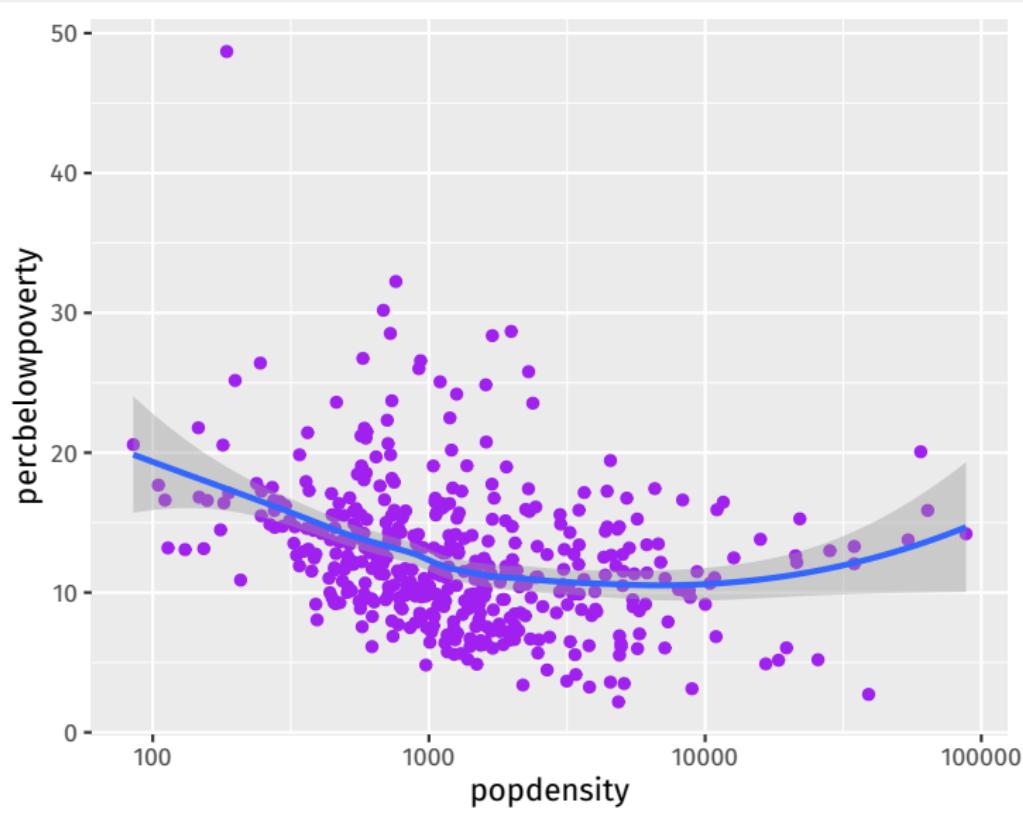
If passed a value other than a variable name, ggplot will implicitly create a variable with that value (in this case "purple" that is constant)

```
ggplot(data = midwest,
        mapping = aes(x = popdensity,
                      y = percbelowpoverty,
                      color = "purple")) +
  geom_point() +
  geom_smooth() +
  scale_x_log10()
```

# Setting aesthetics

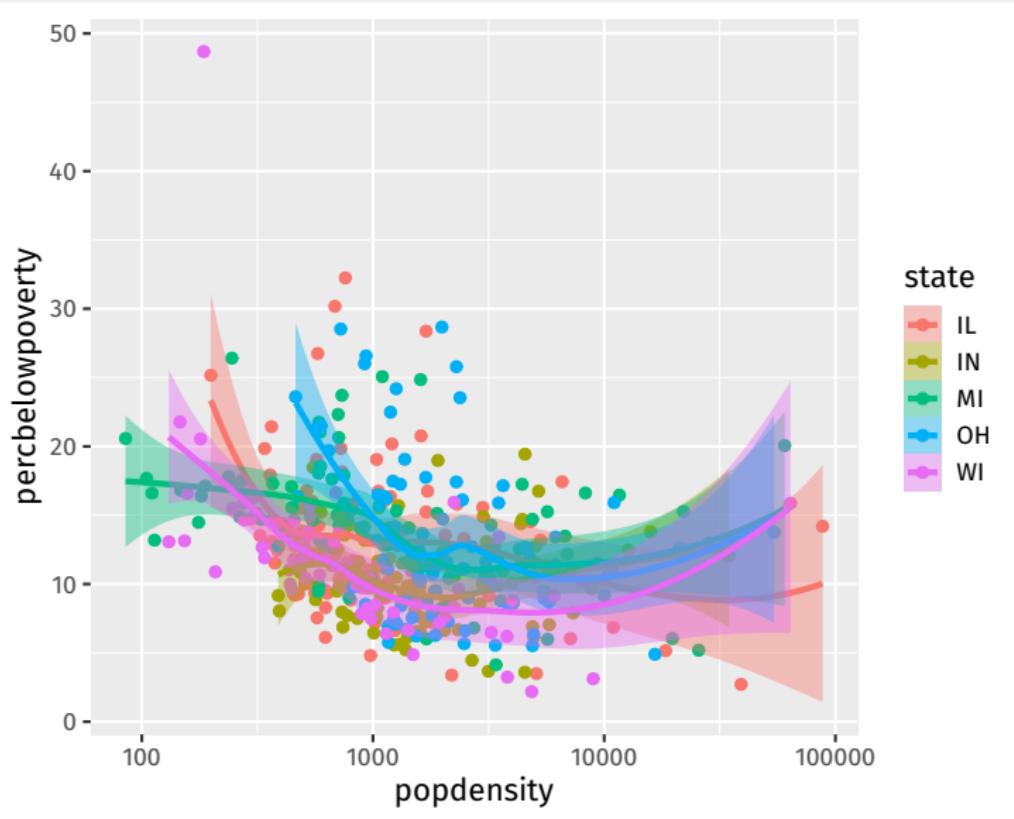
Set the color outside the `mapping = aes()` format.

```
ggplot(data = midwest,
        mapping = aes(x = popdensity,
                      y = percbelowpoverty)) +
  geom_point(color = "purple") +
  geom_smooth() +
  scale_x_log10()
```



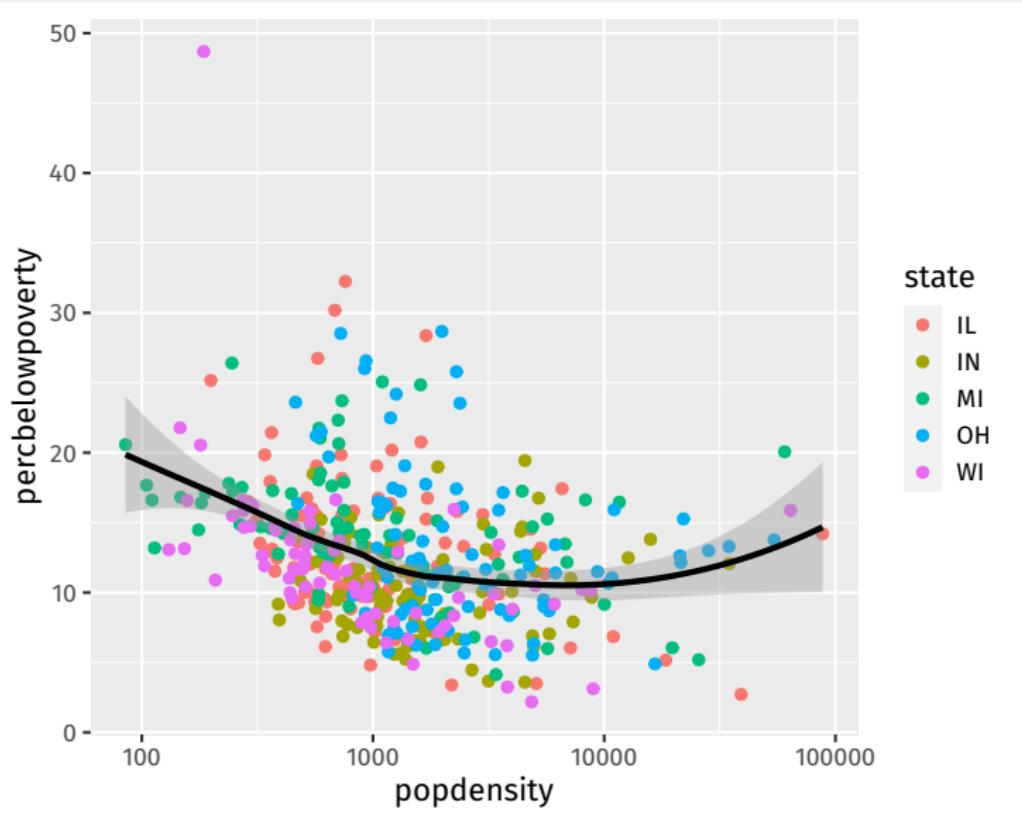
# Mapping more aesthetics

```
ggplot(data = midwest,
        mapping = aes(x = popdensity,
                      y = percbelowpoverty,
                      color = state,
                      fill = state)) +
  geom_point() +
  geom_smooth() +
  scale_x_log10()
```



# Mappings can be done on a per geom basis

```
ggplot(data = midwest,
       mapping = aes(x = popdensity,
                     y = percbelowpoverty)) +
  geom_point(mapping = aes(color = state)) +
  geom_smooth(color = "black") +
  scale_x_log10()
```



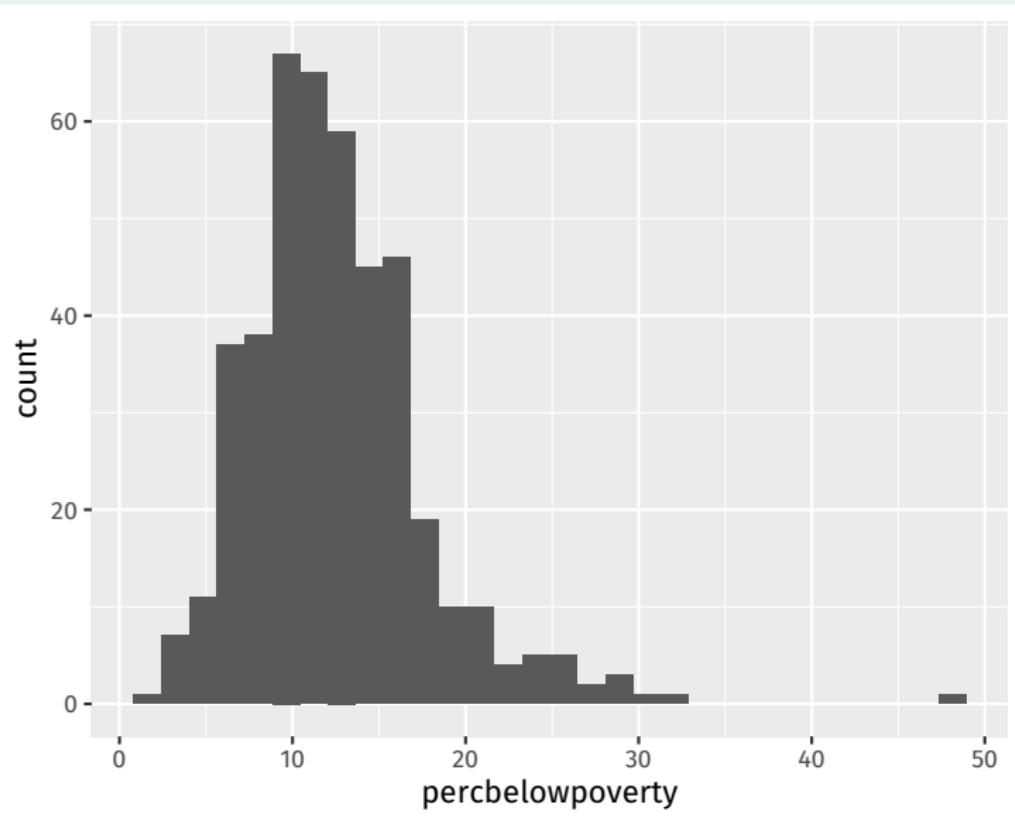
# **2/** Histograms and boxplots

# Histograms

**Histograms** show where there are more or fewer observations of a numeric variable.

```
ggplot(data = midwest,  
       mapping = aes(x = percbelowpoverty)) +  
       geom_histogram()
```

Split up range of variable into bins, count how many are in each bin.  
y aesthetic calculated automatically.

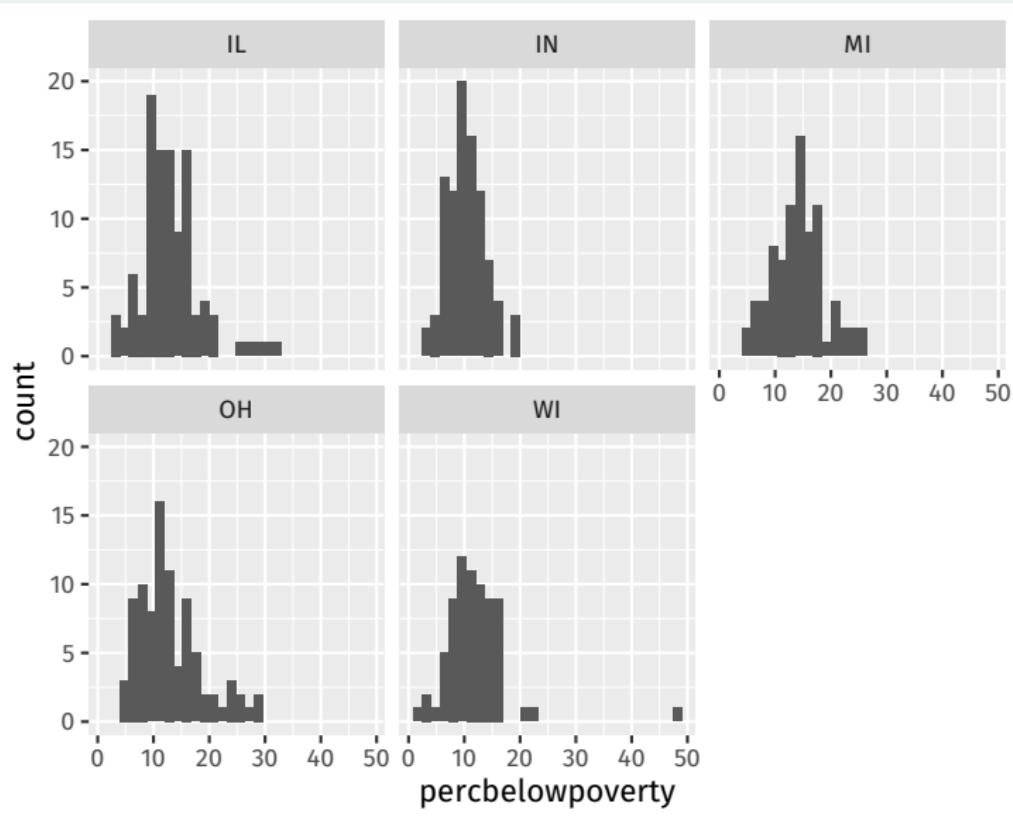


# Creating small multiples with facets

**Small multiples:** a series of similar graphs with the same scale/axes to help with comparing different partitions of a dataset.

```
ggplot(data = midwest,  
       mapping = aes(x = percbelowpoverty)) +  
  geom_histogram() +  
  facet_wrap(~ state)
```

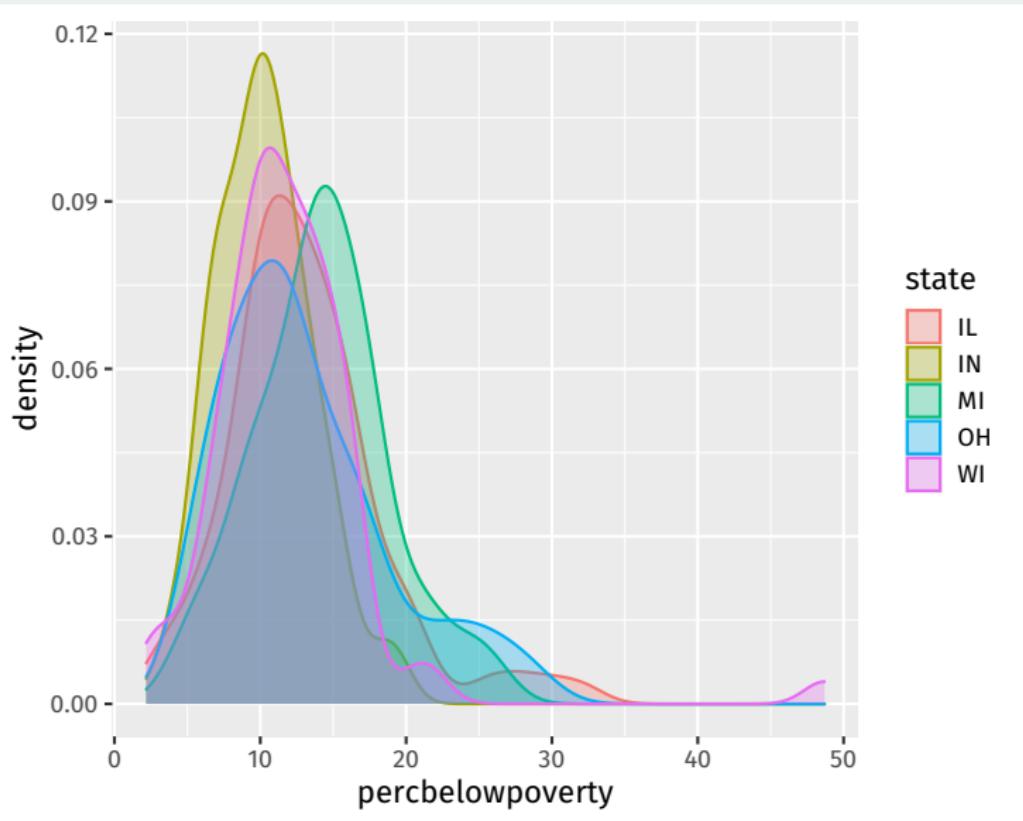
We'll see more of the `~` variable syntax (called a formula).



# Density as alternative to histograms

A **kernel density** plot is a smoothed version of a histogram and slightly easier to overlay.

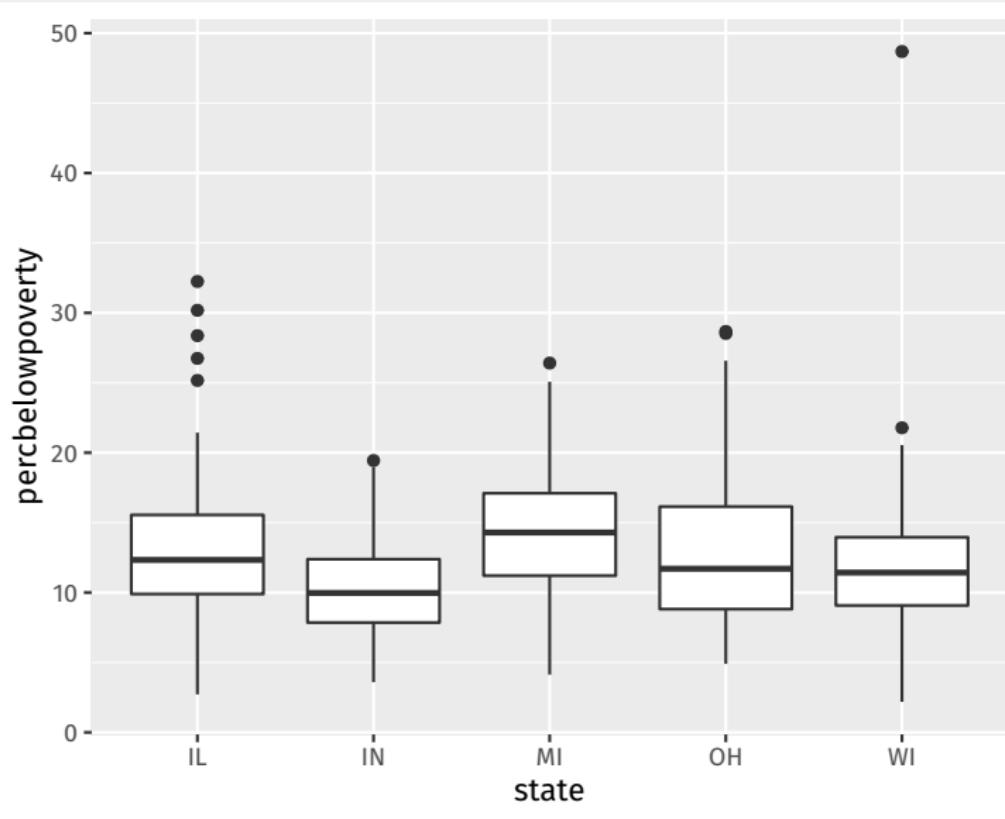
```
ggplot(data = midwest,  
       mapping = aes(x = percbelowpoverty,  
                      fill = state, color = state)) +  
  geom_density(alpha = 0.3)
```



# Boxplots

Boxplots are another way to compare distributions across discrete groups.

```
ggplot(data = midwest,
       mapping = aes(x = state,
                     y = percbelowpoverty)) +
  geom_boxplot()
```



# Boxplots in R

- “Box” represents middle 50% of the data.
  - 25% of the data above the box, 25% below
  - Width of the box is called the inter quartile range (IQR)
- Horizontal line in the box is the median
  - 50% of the data above the median, 50% below
- “Whiskers” represents either:
  - $1.5 \times \text{IQR}$  or max/min of the data, whichever is smaller.
  - Points beyond whiskers are outliers.

# **3/** Grouped data

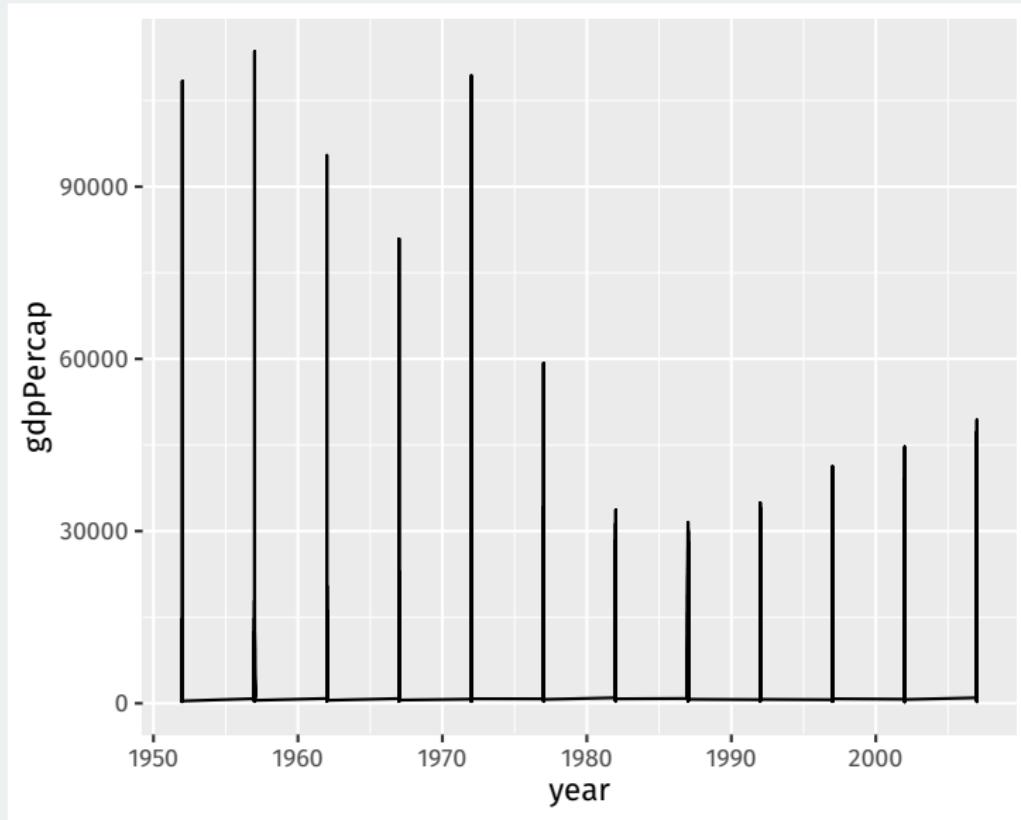
# Back to the gapminder data

```
glimpse(gapminder)

## # Rows: 1,704
## # Columns: 6
## # $ country <fct> "Afghanistan", "Afghanistan", "Afghanistan", "Afgh~
## # $ continent <fct> Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, As~
## # $ year <int> 1952, 1957, 1962, 1967, 1972, 1977, 1982, 1987, 19~
## # $ lifeExp <dbl> 28.8, 30.3, 32.0, 34.0, 36.1, 38.4, 39.9, 40.8, 41~
## # $ pop <int> 8425333, 9240934, 10267083, 11537966, 13079460, 14~
## # $ gdpPercap <dbl> 779, 821, 853, 836, 740, 786, 978, 852, 649, 635, ~
```

# Let's plot the trend in income

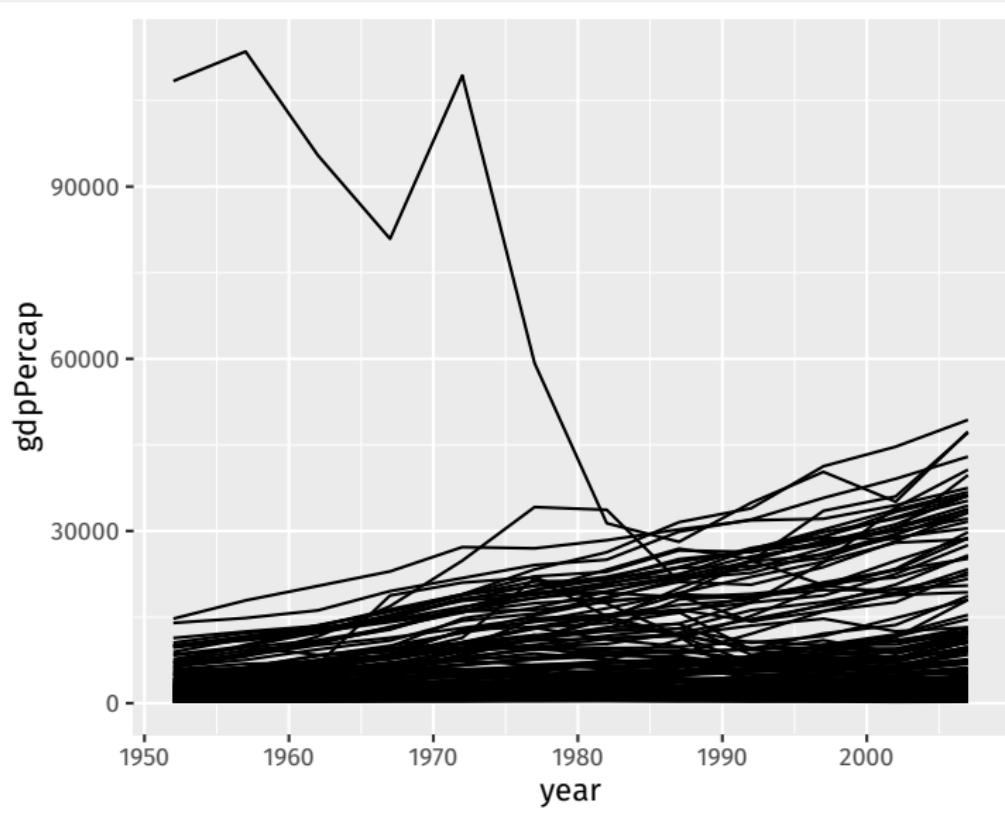
```
ggplot(data = gapminder,  
       mapping = aes(x = year,  
                      y = gdpPercap)) +  
  geom_line()
```



geom\_line connects points from different countries in the same year.

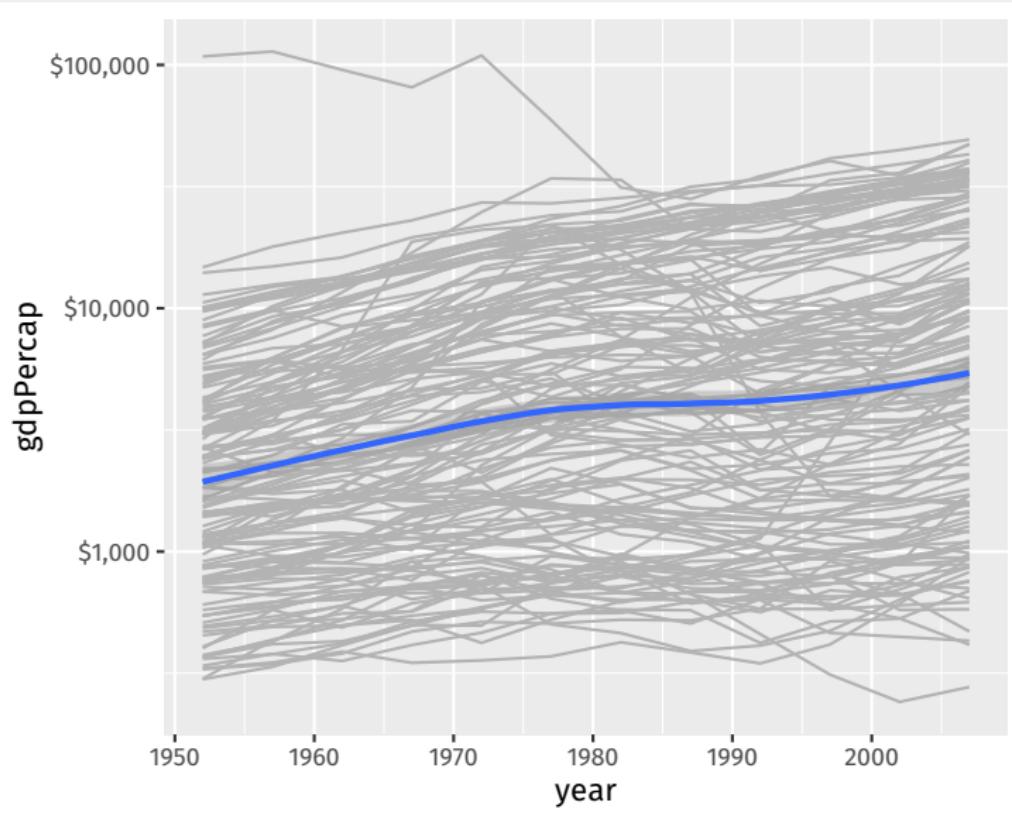
# Tell geom\_line how to group the lines

```
ggplot(data = gapminder,  
       mapping = aes(x = year,  
                      y = gdpPercap)) +  
  geom_line(mapping = aes(group = country))
```



# Scales

```
ggplot(data = gapminder,
       mapping = aes(x = year,
                     y = gdpPercap)) +
  geom_line(mapping = aes(group = country), color = "grey70") +
  geom_smooth(method = "loess") +
  scale_y_log10(labels = scales::dollar)
```



# Gov 50: 4. Data Wrangling

Matthew Blackwell

Harvard University

# Roadmap

1. Data Wrangling
2. Operating on rows
3. Operating on columns
4. Operating on groups

# 1/ Data Wrangling

# Why?

# data.frames vs tibbles

- The standard R object for datasets is the `data.frame`
  - Each column is a vector of the same length.
  - Columns can be different types
- Access columns with `$`: `mydata$myvariable`

```
mtcars$mpg
```

```
## [1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8
## [12] 16.4 17.3 15.2 10.4 10.4 14.7 32.4 30.4 33.9 21.5 15.5
## [23] 15.2 13.3 19.2 27.3 26.0 30.4 15.8 19.7 15.0 21.4
```

# Problems with data frames

mtcars

```
##          mpg cyl  disp  hp drat    wt  qsec vs am
## Mazda RX4      21.0   6 160.0 110 3.90 2.62 16.5  0  1
## Mazda RX4 Wag  21.0   6 160.0 110 3.90 2.88 17.0  0  1
## Datsun 710     22.8   4 108.0  93 3.85 2.32 18.6  1  1
## Hornet 4 Drive 21.4   6 258.0 110 3.08 3.21 19.4  1  0
## Hornet Sportabout 18.7   8 360.0 175 3.15 3.44 17.0  0  0
## Valiant        18.1   6 225.0 105 2.76 3.46 20.2  1  0
## Duster 360     14.3   8 360.0 245 3.21 3.57 15.8  0  0
## Merc 240D       24.4   4 146.7  62 3.69 3.19 20.0  1  0
## Merc 230        22.8   4 140.8  95 3.92 3.15 22.9  1  0
## Merc 280        19.2   6 167.6 123 3.92 3.44 18.3  1  0
## Merc 280C       17.8   6 167.6 123 3.92 3.44 18.9  1  0
## Merc 450SE      16.4   8 275.8 180 3.07 4.07 17.4  0  0
## Merc 450SL      17.3   8 275.8 180 3.07 3.73 17.6  0  0
## Merc 450SLC     15.2   8 275.8 180 3.07 3.78 18.0  0  0
## Cadillac Fleetwood 10.4   8 472.0 205 2.93 5.25 18.0  0  0
## Lincoln Continental 10.4   8 460.0 215 3.00 5.42 17.8  0  0
## Chrysler Imperial 14.7   8 440.0 230 3.23 5.34 17.4  0  0
## Fiat 128         32.4   4  78.7  66 4.08 2.20 19.5  1  1
## Honda Civic      30.4   4  75.7  52 4.93 1.61 18.5  1  1
## Toyota Corolla   22.8   /  71.1  65 4.23 1.83 19.9  1  1
```

# tibbles: a tidyverse alternative

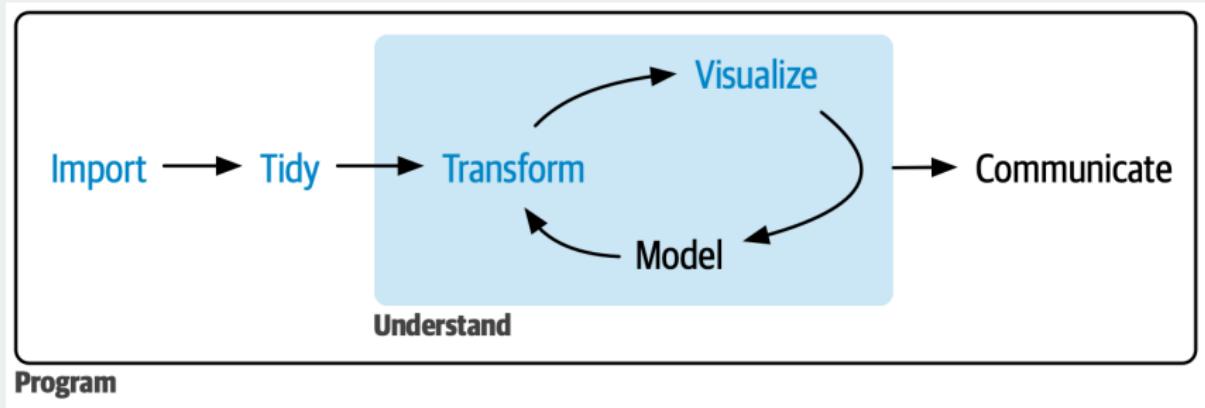
midwest

```
## # A tibble: 437 x 28 rows x columns
##   PID county    state  area poptotal popdensity
##   <int> <chr>     <chr> <dbl>   <int>      <dbl>
## 1 561 ADAMS     IL    0.052   66090     1271.
## 2 562 ALEXANDER IL    0.014   10626      759
## 3 563 BOND       IL    0.022   14991      681.
## 4 564 BOONE      IL    0.017   30806     1812.
## 5 565 BROWN      IL    0.018   5836       324.
## 6 566 BUREAU     IL    0.05    35688     714.
## 7 567 CALHOUN    IL    0.017   5322       313.
## 8 568 CARROLL    IL    0.027   16805     622.
## 9 569 CASS        IL    0.024   13437     560.
## 10 570 CHAMPAIGN IL    0.058   173025    2983.
## # ... with 427 more rows, and 22 more variables:
## #   popwhite <int>, popblack <int>,
## #   popamerindian <int>, popasian <int>,
## #   popother <int>, percwhite <dbl>, percblack <dbl>,
## #   percamerindan <dbl>, percasian <dbl>,
```

column types

abridged output

# Transform-Visualize-Model cycle



# dplyr: a package for data transformation



- All `dplyr` functions:

- Take a dataset as their first argument
- Manipulate the dataset in some way
- Returns the manipulated dataset

# pipe

Nested calls can be hard to read (have to read inside out):

```
f(g(h(r(x))))
```

The pipe `|>` allows us to move output between functions (`|>` = “and then”):

```
x |>  
  r() |>  
  h() |>  
  g() |>  
  h()
```

The piped output goes to the first argument by default.

# Local news data

- How does station ownership affect local news coverage?
- Martin and McCrain (2019) use data on local news at TV stations before and after a large acquisition by a conglomerate.

---

| Variable          | Description   |
|-------------------|---|
| callsign          | Callsign of the station                             |
| affiliation       | Network affiliation of the station                  |
| date              | Airdate of news                                     |
| weekday           | Day of the week of airdate                          |
| ideology          | Measure of news slant (bigger is more conservative) |
| national_politics | Avg proportion of segments on national politics     |
| local_politics    | Avg proportion of segments on national politics     |
| sinclair2017      | Station acquired by Sinclair group in Sept 2017     |
| post              | Date is before/after acquisition (0/1)              |

---

```
library(gov50data)
data(news)
news

## # A tibble: 3,137 x 10
##   callsign affiliation date      weekday ideology nation~1
##   <chr>     <chr>     <date>    <ord>      <dbl>      <dbl>
## 1 KRBC      NBC       2017-06-05 Mon        NA 0.0286
## 2 KTAB      CBS       2017-06-05 Mon        NA 0.0286
## 3 KXVA      FOX       2017-06-05 Mon        NA 0.0393
## 4 KPAX      CBS       2017-06-06 Tue        NA 0.00357
## 5 KTAB      CBS       2017-06-06 Tue        NA 0.0945
## 6 KECI      NBC       2017-06-07 Wed       0.0655 0.225
## 7 KPAX      CBS       2017-06-07 Wed       0.0853 0.283
## 8 KRBC      NBC       2017-06-07 Wed       0.0183 0.130
## 9 KTAB      CBS       2017-06-07 Wed       0.0850 0.0901
## 10 KTMF     ABC       2017-06-07 Wed      0.0842 0.152
## # ... with 3,127 more rows, 4 more variables:
## #   local_politics <dbl>, sinclair2017 <dbl>, post <dbl>,
## #   month <ord>, and abbreviated variable name
## #   1: national_politics
```

# 2/ Operating on rows

# filter()

filter() selects rows that satisfy the argument you pass it:

**dplyr:: filter()** KEEP ROWS THAT  
satisfy  
*your CONDITIONS*

keep rows from... this data... ONLY IF... type is "otter" AND site is "bay"  
filter(df, type == "otter" & site == "bay")

A cartoon illustration featuring a smiling orange otter on the left, pointing towards a map of a coastal area labeled "BAY". To the right of the map is a purple seal, a green crab, and a purple starfish. A small copyright notice at the bottom right reads "© Allison Horst".

| type  | food    | site    |
|-------|---------|---------|
| otter | urchin  | bay     |
| Shark | seal    | channel |
| otter | abalone | bay     |
| otter | crab    | wharf   |

```
news |>  
  filter(weekday == "Tue")
```

```
## # A tibble: 626 x 10  
##   callsign affiliation date      weekday ideology nation~1  
##   <chr>     <chr>    <date>    <ord>      <dbl>       <dbl>  
## 1 KPAX      CBS      2017-06-06 Tue        NA 0.00357  
## 2 KTAB      CBS      2017-06-06 Tue        NA 0.0945  
## 3 KAEF      ABC      2017-06-13 Tue       0.0242 0.180  
## 4 KBVU      FOX      2017-06-13 Tue       0.00894 0.186  
## 5 KBZK      CBS      2017-06-13 Tue       0.129 0.306  
## 6 KCVU      FOX      2017-06-13 Tue       0.114 0.124  
## 7 KECI      NBC      2017-06-13 Tue       0.115 0.283  
## 8 KHSL      CBS      2017-06-13 Tue       0.0821 0.274  
## 9 KNVN      NBC      2017-06-13 Tue       0.120 0.261  
## 10 KPAX     CBS      2017-06-13 Tue       0.0984 0.208  
## # ... with 616 more rows, 4 more variables:  
## #   local_politics <dbl>, sinclair2017 <dbl>, post <dbl>,  
## #   month <ord>, and abbreviated variable name  
## #   1: national_politics
```

# Multiple conditions means “and”

```
news |>  
  filter(weekday == "Tue",  
         affiliation == "FOX")
```

```
## # A tibble: 130 x 10  
##   callsign affiliation date     weekday ideology nation~1  
##   <chr>      <chr>    <date>    <ord>      <dbl>       <dbl>  
## 1 KBVU        FOX     2017-06-13 Tue     0.00894     0.186  
## 2 KCVU        FOX     2017-06-13 Tue     0.114       0.124  
## 3 WEMT        FOX     2017-06-13 Tue     0.235       0.149  
## 4 WYDO        FOX     2017-06-13 Tue     0.0949      0.182  
## 5 KBVU        FOX     2017-06-20 Tue     NA          0.0229  
## 6 KCVU        FOX     2017-06-20 Tue     NA          0.0170  
## 7 KXVA        FOX     2017-06-20 Tue     NA          0.0203  
## 8 WEMT        FOX     2017-06-20 Tue     0.268       0.134  
## 9 WYDO        FOX     2017-06-20 Tue     0.0590      0.155  
## 10 KBVU       FOX     2017-06-27 Tue     NA          0.0601  
## # ... with 120 more rows, 4 more variables:  
## #   local_politics <dbl>, sinclair2017 <dbl>, post <dbl>,  
## #   month <ord>, and abbreviated variable name  
## #   1: national_politics
```

# logicals

- Comparing two values/vectors:
  - $>/>=$ : greater than/greater than or equal to
  - $</<=$ : less than/less than or equal to
  - $==/!=$ : equal to/not equal to
- Combining multiple logical statements:
  - $\&$ : and
  - $|$ : or

# Common gotcha!

```
news |>  
  filter(weekday = "Tue")
```

```
## Error in `filter()`:  
## ! We detected a named input.  
## i This usually means that you've used `=` instead of `==`.  
## i Did you mean `weekday == "Tue"`?
```

```
news |>  
  filter(affiliation == "FOX" | affiliation == "ABC")
```

```
## # A tibble: 1,525 x 10  
##   callsign affiliation date      weekday  ideology natio~1  
##   <chr>     <chr>    <date>    <ord>      <dbl>    <dbl>  
## 1 KXVA      FOX      2017-06-05 Mon       NA        0.0393  
## 2 KTMF      ABC      2017-06-07 Wed       0.0842    0.152  
## 3 KTXS      ABC      2017-06-07 Wed      -0.000488  0.0925  
## 4 KXVA      FOX      2017-06-07 Wed       NA        0.00718  
## 5 KAEF      ABC      2017-06-08 Thu       0.0426    0.213  
## 6 KBVU      FOX      2017-06-08 Thu      -0.0860    0.169  
## 7 KTMF      ABC      2017-06-08 Thu       0.0433    0.179  
## 8 KTXS      ABC      2017-06-08 Thu       0.0627    0.158  
## 9 KXVA      FOX      2017-06-08 Thu       NA        0.0124  
## 10 WCTI     ABC      2017-06-08 Thu      0.139     0.225  
## # ... with 1,515 more rows, 4 more variables:  
## #   local_politics <dbl>, sinclair2017 <dbl>, post <dbl>,  
## #   month <ord>, and abbreviated variable name  
## #   1: national_politics
```

```
news |>  
  filter(ideology < 0 & weekday == "Tue")
```

```
## # A tibble: 66 x 10  
##   callsign affiliation date      weekday ideology nation~1  
##   <chr>     <chr>    <date>    <ord>      <dbl>       <dbl>  
## 1 KAEF      ABC      2017-06-27 Tue     -0.0117     0.162  
## 2 KECI      NBC      2017-06-27 Tue     -0.00362    0.177  
## 3 KHSL      CBS      2017-06-27 Tue     -0.0735     0.170  
## 4 KNVN      NBC      2017-06-27 Tue     -0.0175     0.180  
## 5 KPAX      CBS      2017-06-27 Tue     -0.134      0.219  
## 6 KTXS      ABC      2017-06-27 Tue     -0.0307     0.129  
## 7 WCTI      ABC      2017-06-27 Tue     -0.0308     0.187  
## 8 WITN      NBC      2017-06-27 Tue     -0.0233     0.155  
## 9 WJHL      CBS      2017-06-27 Tue     -0.00388    0.166  
## 10 WNCT     CBS      2017-06-27 Tue     -0.130      0.181  
## # ... with 56 more rows, 4 more variables:  
## #   local_politics <dbl>, sinclair2017 <dbl>, post <dbl>,  
## #   month <ord>, and abbreviated variable name  
## #   1: national_politics
```

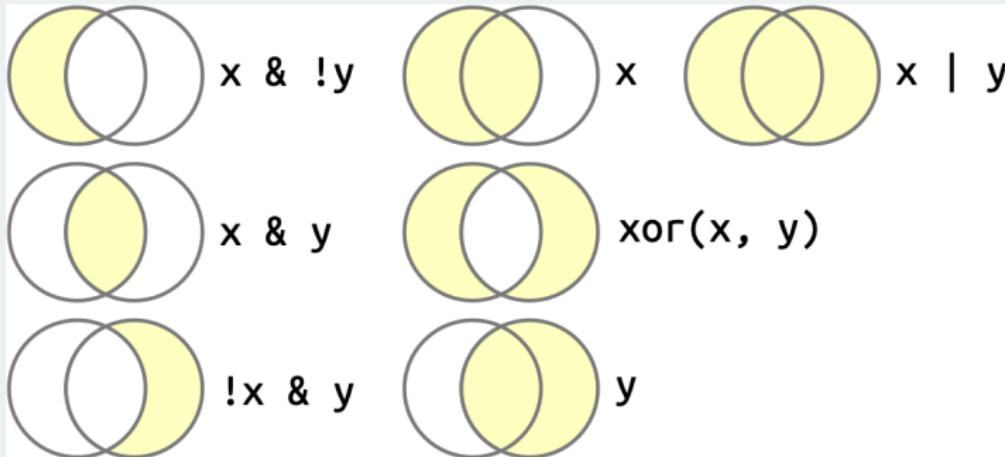
# Combining %in%

When combining | and ==, useful to use %in%:

```
news |>  
  filter(weekday %in% c("Mon", "Fri"))
```

```
## # A tibble: 1,253 x 10  
##   callsign affiliation date      weekday ideology nation~1  
##   <chr>     <chr>    <date>    <ord>     <dbl>       <dbl>  
## 1 KRBC      NBC      2017-06-05 Mon        NA       0.0286  
## 2 KTAB      CBS      2017-06-05 Mon        NA       0.0286  
## 3 KXVA      FOX      2017-06-05 Mon        NA       0.0393  
## 4 KAEF      ABC      2017-06-09 Fri       0.0870      0.153  
## 5 KBVU      FOX      2017-06-09 Fri       NA       0.0553  
## 6 KECI      NBC      2017-06-09 Fri       0.115      0.216  
## 7 KPAX      CBS      2017-06-09 Fri       0.0882      0.315  
## 8 KRBC      NBC      2017-06-09 Fri       0.0929      0.152  
## 9 KTAB      CBS      2017-06-09 Fri       0.0588      0.0711  
## 10 KTMF     ABC      2017-06-09 Fri       NA       0.0495  
## # ... with 1,243 more rows, 4 more variables:  
## #   local_politics <dbl>, sinclair2017 <dbl>, post <dbl>,  
## #   month <ord>, and abbreviated variable name  
## #   1: national_politics
```

# Complicated logicals



# arrange()

arrange( ) will reorder the rows based on the values of the columns.

With multiple arguments, sort by first argument, then second, then third...

# Arrange by callsign then date

```
news |>
  arrange(callsign, date)

## # A tibble: 3,137 x 10
##   callsign affiliation date      weekday ideology nation~1
##   <chr>     <chr>    <date>    <ord>      <dbl>      <dbl>
## 1 KAEF      ABC      2017-06-08 Thu       0.0426    0.213
## 2 KAEF      ABC      2017-06-09 Fri       0.0870    0.153
## 3 KAEF      ABC      2017-06-12 Mon      0.0135    0.149
## 4 KAEF      ABC      2017-06-13 Tue      0.0242    0.180
## 5 KAEF      ABC      2017-06-14 Wed      0.123     0.182
## 6 KAEF      ABC      2017-06-15 Thu      0.0778    0.114
## 7 KAEF      ABC      2017-06-16 Fri      NA        0.109
## 8 KAEF      ABC      2017-06-19 Mon     0.778     0.0823
## 9 KAEF      ABC      2017-06-20 Tue     0.115     0.131
## 10 KAEF     ABC      2017-06-21 Wed     -0.315    0.130
## # ... with 3,127 more rows, 4 more variables:
## #   local_politics <dbl>, sinclair2017 <dbl>, post <dbl>,
## #   month <ord>, and abbreviated variable name
## #   1: national_politics
```

# Which station-dates were the most liberal?

```
news |>
  arrange(ideology)

## # A tibble: 3,137 x 10
##   callsign affiliation date      weekday ideology nation~1
##   <chr>     <chr>    <date>    <ord>      <dbl>      <dbl>
## 1 KRBC      NBC      2017-10-19 Thu       -0.674     0.0731
## 2 WJHL      CBS      2017-12-08 Fri       -0.673     0.0364
## 3 KRBC      NBC      2017-10-18 Wed       -0.586     0.0470
## 4 KCVU      FOX      2017-06-22 Thu       -0.414     0.158
## 5 KRBC      NBC      2017-12-11 Mon      -0.365     0.0674
## 6 KAEF      ABC      2017-06-21 Wed      -0.315     0.130
## 7 KTMF      ABC      2017-12-01 Fri      -0.303     0.179
## 8 KWYB      ABC      2017-12-01 Fri      -0.303     0.160
## 9 KTVM      NBC      2017-09-01 Fri      -0.302     0.0507
## 10 KNVN     NBC      2017-12-08 Fri      -0.299     0.121
## # ... with 3,127 more rows, 4 more variables:
## #   local_politics <dbl>, sinclair2017 <dbl>, post <dbl>,
## #   month <ord>, and abbreviated variable name
## #   1: national_politics
```

# Which station-dates were the most conservative?

Use `desc()` to reverse the order:

```
news |>
  arrange(desc(ideology))

## # A tibble: 3,137 x 10
##   callsign affiliation date      weekday ideology nation~1
##   <chr>     <chr>    <date>    <ord>      <dbl>     <dbl>
## 1 KAEF      ABC      2017-06-19 Mon       0.778    0.0823
## 2 WYDO      FOX      2017-07-19 Wed       0.580    0.126
## 3 KRCR      ABC      2017-10-03 Tue       0.566    0.123
## 4 KAEF      ABC      2017-10-18 Wed       0.496    0.0892
## 5 KBVU      FOX      2017-11-16 Thu       0.491    0.159
## 6 KTMF      ABC      2017-11-06 Mon      0.455    0.138
## 7 KAEF      ABC      2017-06-29 Thu       0.447    0.126
## 8 KPAX      CBS      2017-11-23 Thu       0.437    0.125
## 9 KTAB      CBS      2017-11-16 Thu       0.427    0.0631
## 10 KCVU     FOX      2017-07-06 Thu      0.406    0.154
## # ... with 3,127 more rows, 4 more variables:
## #   local_politics <dbl>, sinclair2017 <dbl>, post <dbl>,
## #   month <ord>, and abbreviated variable name
## #   1: national_politics
```

# 3/ Operating on columns

# `select()`:

`select()` selects columns via their names.

# Selecting based on names

```
news |>  
  select(callsign, date, ideology)
```

```
## # A tibble: 3,137 x 3  
##   callsign     date   ideology  
##   <chr>       <date>    <dbl>  
## 1 KRBC        2017-06-05    NA  
## 2 KTAB        2017-06-05    NA  
## 3 KXVA        2017-06-05    NA  
## 4 KPAX        2017-06-06    NA  
## 5 KTAB        2017-06-06    NA  
## 6 KECI        2017-06-07  0.0655  
## 7 KPAX        2017-06-07  0.0853  
## 8 KRBC        2017-06-07  0.0183  
## 9 KTAB        2017-06-07  0.0850  
## 10 KTMF       2017-06-07  0.0842  
## # ... with 3,127 more rows
```

# Selecting based on a range of variables

```
news |>
  select(callsign:ideology)

## # A tibble: 3,137 x 5
##   callsign affiliation date     weekday ideology
##   <chr>     <chr>    <date>    <ord>      <dbl>
## 1 KRBC      NBC      2017-06-05 Mon       NA
## 2 KTAB      CBS      2017-06-05 Mon       NA
## 3 KXVA      FOX      2017-06-05 Mon       NA
## 4 KPAX      CBS      2017-06-06 Tue       NA
## 5 KTAB      CBS      2017-06-06 Tue       NA
## 6 KECI      NBC      2017-06-07 Wed      0.0655
## 7 KPAX      CBS      2017-06-07 Wed      0.0853
## 8 KRBC      NBC      2017-06-07 Wed      0.0183
## 9 KTAB      CBS      2017-06-07 Wed      0.0850
## 10 KTMF     ABC      2017-06-07 Wed     0.0842
## # ... with 3,127 more rows
```

# Selecting all not in a range

```
news |>
  select(!callsign, ideology)

## # A tibble: 3,137 x 5
##   national_politics local_politics sinclair2017 post month
##       <dbl>           <dbl>           <dbl> <dbl> <ord>
## 1      0.0286          0.0190          0     0 Jun
## 2      0.0286          0.0190          0     0 Jun
## 3      0.0393          0.0262          0     0 Jun
## 4      0.00357          0.194           0     0 Jun
## 5      0.0945          0.109           0     0 Jun
## 6      0.225            0.148           1     0 Jun
## 7      0.283            0.123           0     0 Jun
## 8      0.130            0.189           0     0 Jun
## 9      0.0901           0.138           0     0 Jun
## 10     0.152            0.129           0     0 Jun
## # ... with 3,127 more rows
```

# Selecting all numeric columns

```
news |>  
  select(where(is.numeric))  
  
## # A tibble: 3,137 x 5  
##   ideology national_politics local_politics sinclai~1  post  
##   <dbl>          <dbl>          <dbl>        <dbl> <dbl>  
## 1 NA            0.0286        0.0190        0     0  
## 2 NA            0.0286        0.0190        0     0  
## 3 NA            0.0393        0.0262        0     0  
## 4 NA            0.00357       0.194         0     0  
## 5 NA            0.0945        0.109         0     0  
## 6 0.0655        0.225         0.148         1     0  
## 7 0.0853        0.283         0.123         0     0  
## 8 0.0183        0.130         0.189         0     0  
## 9 0.0850        0.0901        0.138         0     0  
## 10 0.0842       0.152         0.129         0     0  
## # ... with 3,127 more rows, and abbreviated variable name  
## #   1: sinclair2017
```

# Combining multiple selections

```
news |>  
  select(callsign:weekday, ends_with("politics"))
```

```
## # A tibble: 3,137 x 6  
##   callsign affiliation date     weekday nationa~1 local~2  
##   <chr>     <chr>    <date>    <ord>      <dbl>      <dbl>  
## 1 KRBC      NBC      2017-06-05 Mon       0.0286    0.0190  
## 2 KTAB      CBS      2017-06-05 Mon       0.0286    0.0190  
## 3 KXVA      FOX      2017-06-05 Mon       0.0393    0.0262  
## 4 KPAX      CBS      2017-06-06 Tue       0.00357   0.194  
## 5 KTAB      CBS      2017-06-06 Tue       0.0945    0.109  
## 6 KECI      NBC      2017-06-07 Wed      0.225     0.148  
## 7 KPAX      CBS      2017-06-07 Wed      0.283     0.123  
## 8 KRBC      NBC      2017-06-07 Wed      0.130     0.189  
## 9 KTAB      CBS      2017-06-07 Wed      0.0901   0.138  
## 10 KTMF     ABC      2017-06-07 Wed     0.152     0.129  
## # ... with 3,127 more rows, and abbreviated variable names  
## #   1: national_politics, 2: local_politics
```

# rename()

`rename(new_name = old_name)` renames the `old_name` variable to `new_name`

```
news |>  
  rename(call_sign = callsign)
```

```
## # A tibble: 3,137 x 10  
##   call_sign affiliation date      weekday ideology natio~1  
##   <chr>     <chr>       <date>     <ord>      <dbl>    <dbl>  
## 1 KRBC      NBC        2017-06-05 Mon        NA     0.0286  
## 2 KTAB      CBS        2017-06-05 Mon        NA     0.0286  
## 3 KXVA      FOX        2017-06-05 Mon        NA     0.0393  
## 4 KPAX      CBS        2017-06-06 Tue        NA     0.00357  
## 5 KTAB      CBS        2017-06-06 Tue        NA     0.0945  
## 6 KECI      NBC        2017-06-07 Wed       0.0655  0.225  
## 7 KPAX      CBS        2017-06-07 Wed       0.0853  0.283  
## 8 KRBC      NBC        2017-06-07 Wed       0.0183  0.130  
## 9 KTAB      CBS        2017-06-07 Wed       0.0850  0.0901  
## 10 KTMF     ABC        2017-06-07 Wed      0.0842  0.152  
## # ... with 3,127 more rows, 4 more variables:  
## #   local_politics <dbl>, sinclair2017 <dbl>, post <dbl>,  
## #   month <ord>, and abbreviated variable name  
## #   1: national_politics
```

# mutate()

`mutate(new_var = fun(old_vars))` adds new columns that are functions of existing columns.

```

news |>
  mutate(
    national_local_diff = national_politics - local_politics,
    national_politics_perc = national_politics * 100
  ) |>
  select(callsign, date, national_politics, local_politics,
         national_local_diff, national_politics_perc)

## # A tibble: 3,137 x 6
##   callsign date      national_politics local_politics national_local_diff national_politics_perc
##   <chr>     <date>            <dbl>          <dbl>             <dbl>                  <dbl>
## 1 KRBC     2017-06-05       0.0286        0.0190           0.00952                2.86
## 2 KTAB     2017-06-05       0.0286        0.0190           0.00952                2.86
## 3 KXVA     2017-06-05       0.0393        0.0262           0.0131                 3.93
## 4 KPAX     2017-06-06       0.00357       0.194            -0.191                 0.357
## 5 KTAB     2017-06-06       0.0945        0.109            -0.0145                9.45
## 6 KECI     2017-06-07       0.225          0.148            0.0761                 22.5
## 7 KPAX     2017-06-07       0.283          0.123            0.160                  28.3
## 8 KRBC     2017-06-07       0.130          0.189            -0.0589                13.0
## 9 KTAB     2017-06-07       0.0901        0.138            -0.0476                9.01
## 10 KTMF    2017-06-07      0.0152         0.129            0.0229                 15.2
## # ... with 3,127 more rows

```

# if\_else()

`if_else(test_condition, yes, no)` allows us to create a vector that depends on a logical

New vector gets `yes` expression when `test_condition` is `TRUE`, `no` otherwise

```
news |>  
  mutate(Ownership = if_else(sinclair2017 == 1,  
                            "Acquired by Sinclair",  
                            "Not Acquired")) |>  
  select(callsign, affiliation, date, Ownership)
```

```
## # A tibble: 3,137 x 4  
##   callsign affiliation date      Ownership  
##   <chr>     <chr>     <date>    <chr>  
## 1 KRBC      NBC       2017-06-05 Not Acquired  
## 2 KTAB      CBS       2017-06-05 Not Acquired  
## 3 KXVA      FOX       2017-06-05 Not Acquired  
## 4 KPAX      CBS       2017-06-06 Not Acquired  
## 5 KTAB      CBS       2017-06-06 Not Acquired  
## 6 KECI      NBC       2017-06-07 Acquired by Sinclair  
## 7 KPAX      CBS       2017-06-07 Not Acquired  
## 8 KRBC      NBC       2017-06-07 Not Acquired  
## 9 KTAB      CBS       2017-06-07 Not Acquired  
## 10 KTMF     ABC       2017-06-07 Not Acquired  
## # ... with 3,127 more rows
```

# 4| Operating on groups

## group\_by()

`group_by(var)` divides the data into groups based on the `var` variable.

Doesn't change data yet, but subsequent operations will by var.

```
news |>
  group_by(month)

## # A tibble: 3,137 x 10
## # Groups:   month [7]
##   callsign affil~1 date      weekday ideol~2 natio~3 local~4 sincl~5
##   <chr>     <chr>    <date>    <ord>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 KRBC      NBC      2017-06-05 Mon       NA     0.0286   0.0190     0
## 2 KTAB      CBS      2017-06-05 Mon       NA     0.0286   0.0190     0
## 3 KXVA      FOX      2017-06-05 Mon       NA     0.0393   0.0262     0
## 4 KPAX      CBS      2017-06-06 Tue       NA     0.00357  0.194     0
## 5 KTAB      CBS      2017-06-06 Tue       NA     0.0945   0.109     0
## 6 KECI      NBC      2017-06-07 Wed      0.0655  0.225     0.148    1
## 7 KPAX      CBS      2017-06-07 Wed      0.0853  0.283     0.123     0
## 8 KRBC      NBC      2017-06-07 Wed      0.0183  0.130     0.189     0
## 9 KTAB      CBS      2017-06-07 Wed      0.0850  0.0901   0.138     0
## 10 KTMF     ABC      2017-06-07 Wed     0.0842  0.152     0.129     0
## # ... with 3,127 more rows, 2 more variables: post <dbl>,
## #   month <ord>, and abbreviated variable names 1: affiliation,
## #   2: ideology, 3: national_politics, 4: local_politics,
## #   5: sinclair2017
```

# summarize()

`summarize(sum_var = fun(curr_var))` calculates summaries of variables by groups.

# Ideological slant by weekday

```
news |>
  group_by(month) |>
  summarize(
    slant_mean = mean(ideology, na.rm = TRUE)
  )
```

```
## # A tibble: 7 x 2
##   month slant_mean
##   <ord>     <dbl>
## 1 Jun      0.0786
## 2 Jul      0.103 
## 3 Aug      0.105 
## 4 Sep      0.0751
## 5 Oct      0.0862
## 6 Nov      0.0972
## 7 Dec      0.0774
```

# Summaries by ownership and pre/post

```
news |>
  group_by(sinclair2017, post) |>
  summarize(
    slant_mean = mean(ideology, na.rm = TRUE),
    national_mean = mean(national_politics, na.rm = TRUE)
  )

## # A tibble: 4 x 4
## # Groups:   sinclair2017 [2]
##   sinclair2017  post slant_mean national_mean
##   <dbl> <dbl>     <dbl>        <dbl>
## 1 0       0      0.100       0.118
## 2 0       1      0.0768      0.107
## 3 1       0      0.0936      0.124
## 4 1       1      0.0938      0.144
```

# Summarize across types of variables

`across()` will apply a summary function across many variables

```
news |>
  group_by(sinclair2017, post) |>
  summarize(
    across(where(is.numeric), mean, na.rm = TRUE),
  )
```

```
## # A tibble: 4 x 5
## # Groups:   sinclair2017 [2]
##   sinclair2017 post ideology national_politics local_politics
##   <dbl> <dbl>     <dbl>           <dbl>           <dbl>
## 1 0     0     0.100          0.118          0.158
## 2 0     1     0.0768         0.107          0.150
## 3 1     0     0.0936         0.124          0.170
## 4 1     1     0.0938         0.144          0.147
```

# Gov 50: 5. Data Wrangling and Barplots

Matthew Blackwell

Harvard University

# Roadmap

1. Operating on rows
2. Operating on columns
3. Operating on groups
4. Creating barplots

# Local news data

- How does station ownership affect local news coverage?
- Martin and McCrain (2019) use data on local news at TV stations before and after a large acquisition by a conglomerate.

---

| Variable          | Description   |
|-------------------|---|
| callsign          | Callsign of the station                             |
| affiliation       | Network affiliation of the station                  |
| date              | Airdate of news                                     |
| weekday           | Day of the week of airdate                          |
| ideology          | Measure of news slant (bigger is more conservative) |
| national_politics | Avg proportion of segments on national politics     |
| local_politics    | Avg proportion of segments on national politics     |
| sinclair2017      | Station acquired by Sinclair group in Sept 2017     |
| post              | Date is before/after acquisition (0/1)              |

---

```
library(gov50data)
data(news)
news

## # A tibble: 3,137 x 10
##   callsign affil~1 date      weekday ideol~2 natio~3 local~4 sincl~5
##   <chr>     <chr>    <date>    <ord>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 KRBC      NBC     2017-06-05 Mon       NA     0.0286   0.0190     0
## 2 KTAB      CBS     2017-06-05 Mon       NA     0.0286   0.0190     0
## 3 KXVA      FOX     2017-06-05 Mon       NA     0.0393   0.0262     0
## 4 KPAX      CBS     2017-06-06 Tue       NA     0.00357  0.194     0
## 5 KTAB      CBS     2017-06-06 Tue       NA     0.0945   0.109     0
## 6 KECI      NBC     2017-06-07 Wed      0.0655  0.225     0.148    1
## 7 KPAX      CBS     2017-06-07 Wed      0.0853  0.283     0.123     0
## 8 KRBC      NBC     2017-06-07 Wed      0.0183  0.130     0.189     0
## 9 KTAB      CBS     2017-06-07 Wed      0.0850  0.0901   0.138     0
## 10 KTMF     ABC     2017-06-07 Wed     0.0842  0.152     0.129     0
## # ... with 3,127 more rows, 2 more variables: post <dbl>,
## #   month <ord>, and abbreviated variable names 1: affiliation,
## #   2: ideology, 3: national_politics, 4: local_politics,
## #   5: sinclair2017
```

# 1/ Operating on rows

# slice()

slice() can give you a specific set of rows:

```
## first and third row  
news |>  
  slice(1, 3)
```

```
## # A tibble: 2 x 10  
##   callsign affili~1 date      weekday ideol~2 natio~3 local~4 sincl~5  
##   <chr>     <chr>    <date>    <ord>     <dbl>    <dbl>    <dbl>    <dbl>  
## 1 KRBC      NBC      2017-06-05 Mon        NA  0.0286  0.0190      0  
## 2 KXVA      FOX      2017-06-05 Mon        NA  0.0393  0.0262      0  
## # ... with 2 more variables: post <dbl>, month <ord>, and abbreviated  
## #   variable names 1: affiliation, 2: ideology, 3: national_politics,  
## #   4: local_politics, 5: sinclair2017
```

You can ask for a range of rows with `start:stop` syntax:

```
## first three rows
news |>
  slice(1:3)

## # A tibble: 3 x 10
##   callsign affili~1 date      weekday ideol~2 natio~3 local~4 sincl~5
##   <chr>     <chr>    <date>    <ord>     <dbl>    <dbl>    <dbl>    <dbl>
## 1 KRBC      NBC      2017-06-05 Mon        NA  0.0286  0.0190     0
## 2 KTAB      CBS      2017-06-05 Mon        NA  0.0286  0.0190     0
## 3 KXVA      FOX      2017-06-05 Mon        NA  0.0393  0.0262     0
## # ... with 2 more variables: post <dbl>, month <ord>, and abbreviated
## #   variable names 1: affiliation, 2: ideology, 3: national_politics,
## #   4: local_politics, 5: sinclair2017
```

# slice\_max()

slice\_max(var, n = 5) will return the top 5 observations on column var

```
news |>  
  slice_max(ideology, n = 5)
```

```
## # A tibble: 5 x 10  
##   callsign affili~1 date      weekday ideol~2 natio~3 local~4 sincl~5  
##   <chr>     <chr>    <date>    <ord>      <dbl>    <dbl>    <dbl>    <dbl>  
## 1 KAEF      ABC      2017-06-19 Mon       0.778   0.0823   0.179    1  
## 2 WYDO      FOX      2017-07-19 Wed       0.580   0.126    0.121    1  
## 3 KRCR      ABC      2017-10-03 Tue       0.566   0.123    0.192    1  
## 4 KAEF      ABC      2017-10-18 Wed       0.496   0.0892   0.217    1  
## 5 KBVU      FOX      2017-11-16 Thu       0.491   0.159    0.184    1  
## # ... with 2 more variables: post <dbl>, month <ord>, and abbreviated  
## #   variable names 1: affiliation, 2: ideology, 3: national_politics,  
## #   4: local_politics, 5: sinclair2017
```

# slice\_min()

slice\_min(var, n = 5) will return the bottom 5 observations on column var

```
news |>  
  slice_min(ideology, n = 5)
```

```
## # A tibble: 5 x 10  
##   callsign affili~1 date      weekday ideol~2 natio~3 local~4 sincl~5  
##   <chr>     <chr>    <date>    <ord>     <dbl>    <dbl>    <dbl>    <dbl>  
## 1 KRBC      NBC      2017-10-19 Thu     -0.674   0.0731   0.161     0  
## 2 WJHL      CBS      2017-12-08 Fri     -0.673   0.0364   0.206     0  
## 3 KRBC      NBC      2017-10-18 Wed    -0.586   0.0470   0.135     0  
## 4 KCVU      FOX      2017-06-22 Thu    -0.414   0.158    0.172     1  
## 5 KRBC      NBC      2017-12-11 Mon    -0.365   0.0674   0.163     0  
## # ... with 2 more variables: post <dbl>, month <ord>, and abbreviated  
## #   variable names 1: affiliation, 2: ideology, 3: national_politics,  
## #   4: local_politics, 5: sinclair2017
```

# 2/ Operating on columns

# rename()

`rename(new_name = old_name)` renames the `old_name` variable to `new_name`

```
news |>
  rename(call_sign = callsign)

## # A tibble: 3,137 x 10
##   call_s~1 affil~2 date      weekday ideol~3 natio~4 local~5 sincl~6
##   <chr>     <chr>    <date>    <ord>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 KRBC      NBC      2017-06-05 Mon       NA     0.0286   0.0190    0
## 2 KTAB      CBS      2017-06-05 Mon       NA     0.0286   0.0190    0
## 3 KXVA      FOX      2017-06-05 Mon       NA     0.0393   0.0262    0
## 4 KPAX      CBS      2017-06-06 Tue       NA     0.00357  0.194     0
## 5 KTAB      CBS      2017-06-06 Tue       NA     0.0945   0.109     0
## 6 KECI      NBC      2017-06-07 Wed      0.0655  0.225    0.148     1
## 7 KPAX      CBS      2017-06-07 Wed      0.0853  0.283    0.123     0
## 8 KRBC      NBC      2017-06-07 Wed      0.0183  0.130    0.189     0
## 9 KTAB      CBS      2017-06-07 Wed      0.0850  0.0901   0.138     0
## 10 KTMF     ABC      2017-06-07 Wed     0.0842  0.152    0.129     0
## # ... with 3,127 more rows, 2 more variables: post <dbl>,
## #   month <ord>, and abbreviated variable names 1: call_sign,
## #   2: affiliation, 3: ideology, 4: national_politics,
## #   5: local_politics, 6: sinclair2017
```

# mutate()

`mutate(new_var = fun(old_vars))` adds new columns that are functions of existing columns.

```

news |>
  mutate(
    national_local_diff = national_politics - local_politics,
    national_politics_perc = national_politics * 100
  ) |>
  select(callsign, date, national_politics, local_politics,
         national_local_diff, national_politics_perc)

## # A tibble: 3,137 x 6
##   callsign date      national_politics local_politics national_local_diff national_politics_perc
##   <chr>     <date>            <dbl>          <dbl>             <dbl>                  <dbl>
## 1 KRBC     2017-06-05       0.0286        0.0190           0.00952                2.86
## 2 KTAB     2017-06-05       0.0286        0.0190           0.00952                2.86
## 3 KXVA     2017-06-05       0.0393        0.0262           0.0131                 3.93
## 4 KPAX     2017-06-06       0.00357       0.194            -0.191                 0.357
## 5 KTAB     2017-06-06       0.0945        0.109            -0.0145                9.45
## 6 KECI     2017-06-07       0.225          0.148            0.0761                 22.5
## 7 KPAX     2017-06-07       0.283          0.123            0.160                  28.3
## 8 KRBC     2017-06-07       0.130          0.189            -0.0589                13.0
## 9 KTAB     2017-06-07       0.0901        0.138            -0.0476                9.01
## 10 KTMF    2017-06-07      0.152          0.129            0.0229                 15.2
## # ... with 3,127 more rows

```

# if\_else()

`if_else(test_condition, yes, no)` allows us to create a vector that depends on a logical

New vector gets `yes` expression when `test_condition` is `TRUE`, `no` otherwise

```
news |>  
  mutate(Ownership = if_else(sinclair2017 == 1,  
                            "Acquired by Sinclair",  
                            "Not Acquired")) |>  
  select(callsign, affiliation, date, Ownership)
```

```
## # A tibble: 3,137 x 4  
##   callsign affiliation date      Ownership  
##   <chr>     <chr>     <date>    <chr>  
## 1 KRBC      NBC       2017-06-05 Not Acquired  
## 2 KTAB      CBS       2017-06-05 Not Acquired  
## 3 KXVA      FOX       2017-06-05 Not Acquired  
## 4 KPAX      CBS       2017-06-06 Not Acquired  
## 5 KTAB      CBS       2017-06-06 Not Acquired  
## 6 KECI      NBC       2017-06-07 Acquired by Sinclair  
## 7 KPAX      CBS       2017-06-07 Not Acquired  
## 8 KRBC      NBC       2017-06-07 Not Acquired  
## 9 KTAB      CBS       2017-06-07 Not Acquired  
## 10 KTMF     ABC       2017-06-07 Not Acquired  
## # ... with 3,127 more rows
```

# **3/** Operating on groups

# group\_by()

`group_by(var)` divides the data into groups based on the `var` variable.

Doesn't change data yet, but subsequent operations will by var.

```
news |>
  group_by(month)

## # A tibble: 3,137 x 10
## # Groups:   month [7]
##   callsign affil~1 date      weekday ideol~2 natio~3 local~4 sincl~5
##   <chr>     <chr>    <date>    <ord>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 KRBC      NBC      2017-06-05 Mon       NA     0.0286   0.0190     0
## 2 KTAB      CBS      2017-06-05 Mon       NA     0.0286   0.0190     0
## 3 KXVA      FOX      2017-06-05 Mon       NA     0.0393   0.0262     0
## 4 KPAX      CBS      2017-06-06 Tue       NA     0.00357  0.194     0
## 5 KTAB      CBS      2017-06-06 Tue       NA     0.0945   0.109     0
## 6 KECI      NBC      2017-06-07 Wed      0.0655  0.225     0.148    1
## 7 KPAX      CBS      2017-06-07 Wed      0.0853  0.283     0.123     0
## 8 KRBC      NBC      2017-06-07 Wed      0.0183  0.130     0.189     0
## 9 KTAB      CBS      2017-06-07 Wed      0.0850  0.0901   0.138     0
## 10 KTMF     ABC      2017-06-07 Wed     0.0842  0.152     0.129     0
## # ... with 3,127 more rows, 2 more variables: post <dbl>,
## #   month <ord>, and abbreviated variable names 1: affiliation,
## #   2: ideology, 3: national_politics, 4: local_politics,
## #   5: sinclair2017
```

# summarize()

`summarize(sum_var = fun(curr_var))` calculates summaries of variables by groups.

# Ideological slant by weekday

```
news |>
  group_by(month) |>
  summarize(
    slant_mean = mean(ideology, na.rm = TRUE)
  )
```

```
## # A tibble: 7 x 2
##   month slant_mean
##   <ord>     <dbl>
## 1 Jun      0.0786
## 2 Jul      0.103 
## 3 Aug      0.105 
## 4 Sep      0.0751
## 5 Oct      0.0862
## 6 Nov      0.0972
## 7 Dec      0.0774
```

# Summaries by ownership and pre/post

```
news |>
  group_by(sinclair2017, post) |>
  summarize(
    slant_mean = mean(ideology, na.rm = TRUE),
    national_mean = mean(national_politics, na.rm = TRUE)
  )

## # A tibble: 4 x 4
## # Groups:   sinclair2017 [2]
##   sinclair2017  post slant_mean national_mean
##   <dbl> <dbl>     <dbl>        <dbl>
## 1 0       0      0.100       0.118
## 2 0       1      0.0768      0.107
## 3 1       0      0.0936      0.124
## 4 1       1      0.0938      0.144
```

# Summarize across types of variables

across( ) will apply a summary function across many variables

```
news |>
  group_by(sinclair2017, post) |>
  summarize(
    across(where(is.numeric), mean, na.rm = TRUE),
  )
```

```
## # A tibble: 4 x 5
## # Groups:   sinclair2017 [2]
##   sinclair2017 post ideology national_politics local_politics
##   <dbl> <dbl>     <dbl>           <dbl>           <dbl>
## 1 0     0     0.100          0.118          0.158
## 2 0     1     0.0768         0.107          0.150
## 3 1     0     0.0936         0.124          0.170
## 4 1     1     0.0938         0.144          0.147
```

# kable( ) to produce nice tables

```
news |>
  group_by(month) |>
  summarize(
    slant_mean = mean(ideology, na.rm = TRUE)
  ) |>
  knitr::kable()
```

| month | slant_mean |
|-------|------------|
| Jun   | 0.079      |
| Jul   | 0.103      |
| Aug   | 0.105      |
| Sep   | 0.075      |
| Oct   | 0.086      |
| Nov   | 0.097      |
| Dec   | 0.077      |

# Giving nicer column names

```
news |>
  group_by(month) |>
  summarize(
    slant_mean = mean(ideology, na.rm = TRUE)
  ) |>
  knitr::kable(col.names = c("Month", "Avg. Slant"))
```

| Month | Avg. Slant |
|-------|------------|
| Jun   | 0.079      |
| Jul   | 0.103      |
| Aug   | 0.105      |
| Sep   | 0.075      |
| Oct   | 0.086      |
| Nov   | 0.097      |
| Dec   | 0.077      |

# Producing a table of counts of a categorical variable

```
news |>  
  group_by(affiliation) |>  
  summarize(n = n())
```

```
## # A tibble: 4 x 2  
##   affiliation     n  
##   <chr>       <int>  
## 1 ABC           863  
## 2 CBS           807  
## 3 FOX           662  
## 4 NBC           805
```

# Helper function count()

count() does the same thing:

```
news |>  
  count(affiliation)
```

```
## # A tibble: 4 x 2  
##   affiliation     n  
##   <chr>        <int>  
## 1 ABC            863  
## 2 CBS            807  
## 3 FOX            662  
## 4 NBC            805
```

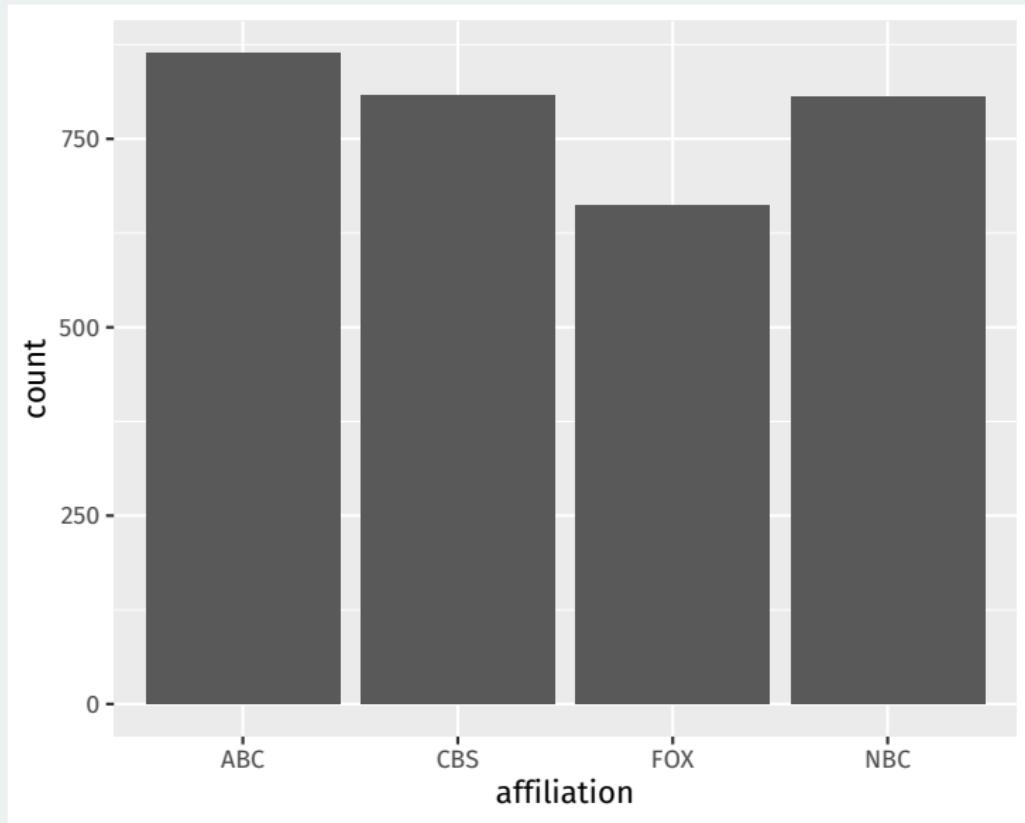
# 4| Creating barplots

# Combining our skills

Let's combine our tools to produce a bar plot with `geom_bar()`

By default, bar plots take a single variable and show the number of observations in each category.

```
ggplot(news, mapping = aes(x = affiliation)) +  
  geom_bar()
```



# Barplots of non-counts

Barplots can represent a lot beyond counts, including variables in our dataset or group summaries.

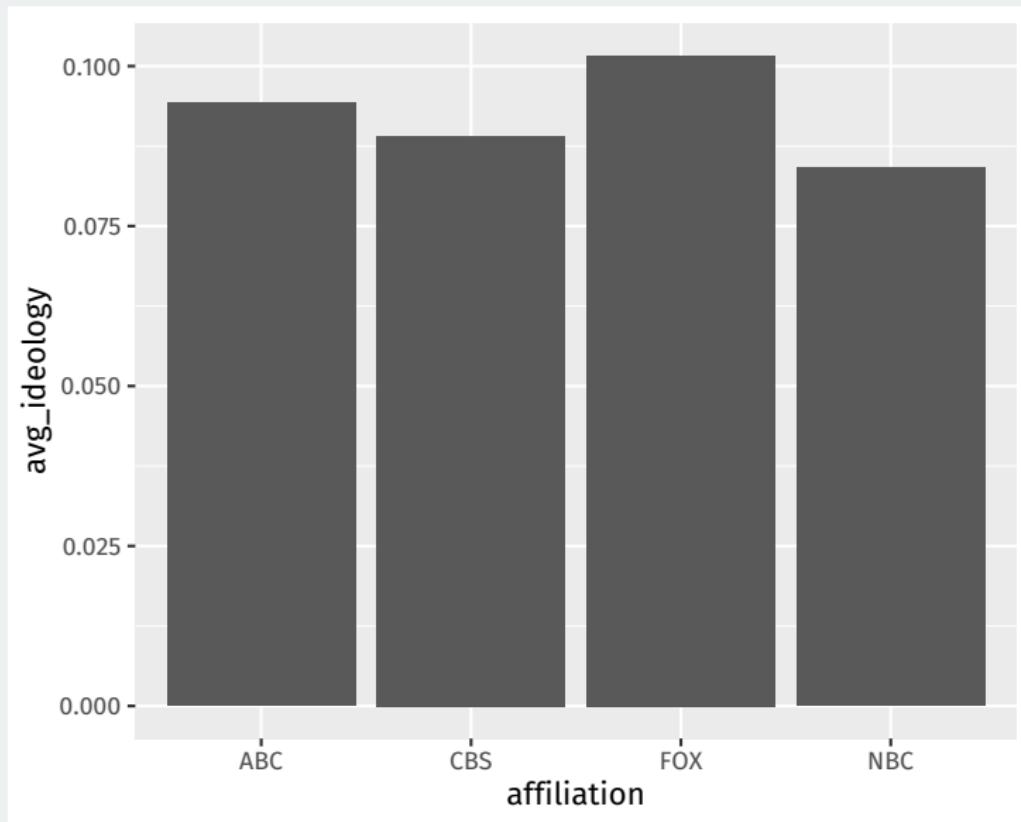
When the height of the bar is another variable in our data and not just a count, we set the x and y aesthetics and use `geom_col()` instead of `geom_bar()`.

Let's create a group summary:

```
aff_ideology_means <- news |>
  group_by(affiliation) |>
  summarize(avg_ideology = mean(ideology, na.rm = TRUE))
aff_ideology_means
```

```
## # A tibble: 4 x 2
##   affiliation avg_ideology
##   <chr>          <dbl>
## 1 ABC            0.0943
## 2 CBS            0.0891
## 3 FOX            0.102 
## 4 NBC            0.0841
```

```
ggplot(aff_ideology_means, aes(x = affiliation, y = avg_ideology)) +
  geom_col()
```

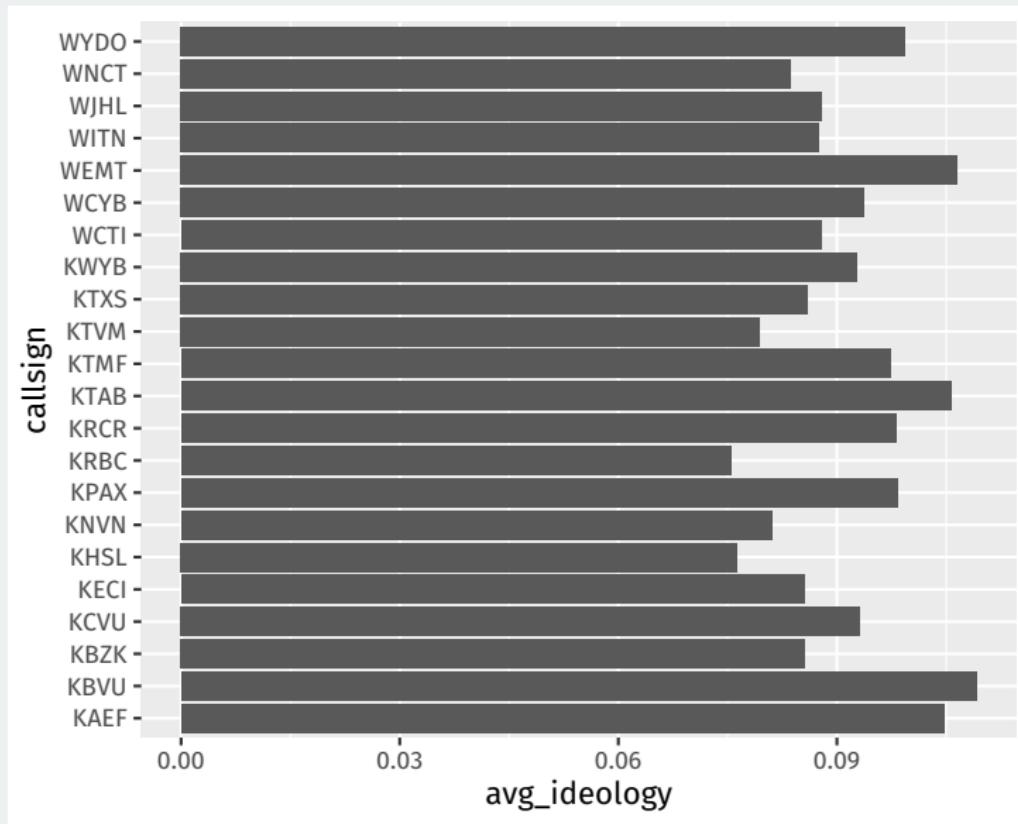


# A more complicated example

Let's create a barplot that shows the top 10 stations in terms of slant. First, let's get the data:

```
station_ideology <- news |>  
  group_by(callsign, affiliation) |>  
  summarize(avg_ideology = mean(ideology, na.rm = TRUE)) |>  
  slice_max(avg_ideology, n = 20)
```

```
ggplot(station_ideology, aes(x = avg_ideology, y = callsign)) +  
  geom_col()
```

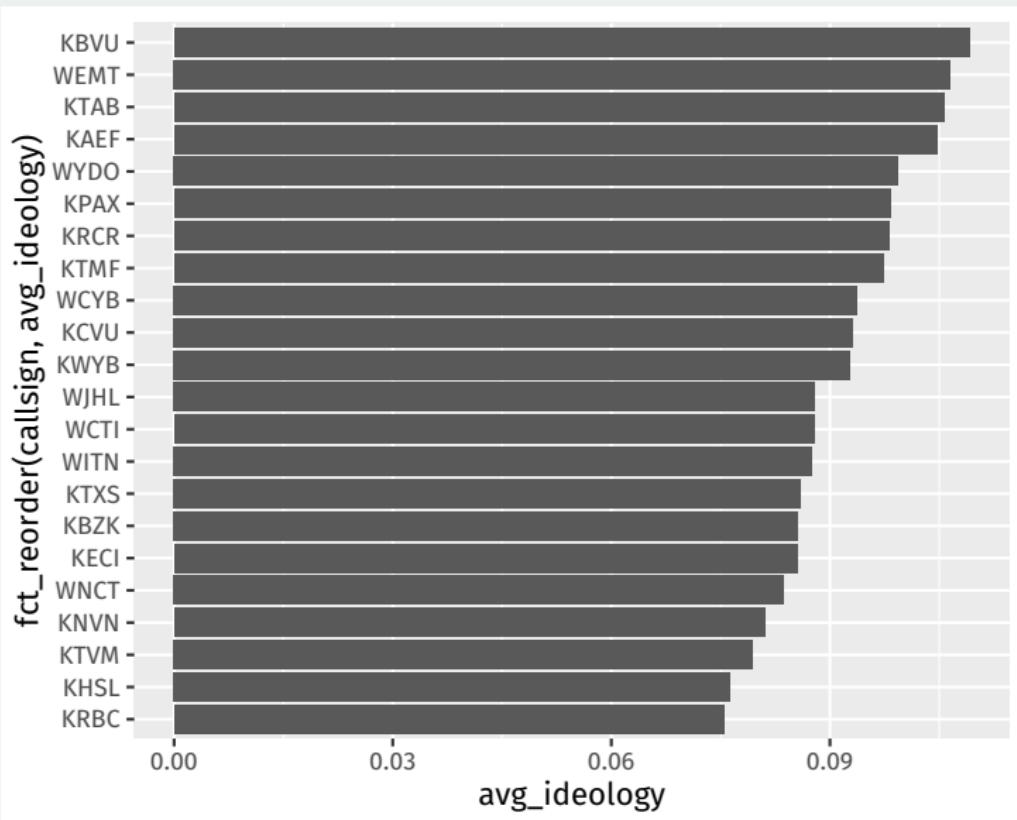


# How do we reorder the stations?

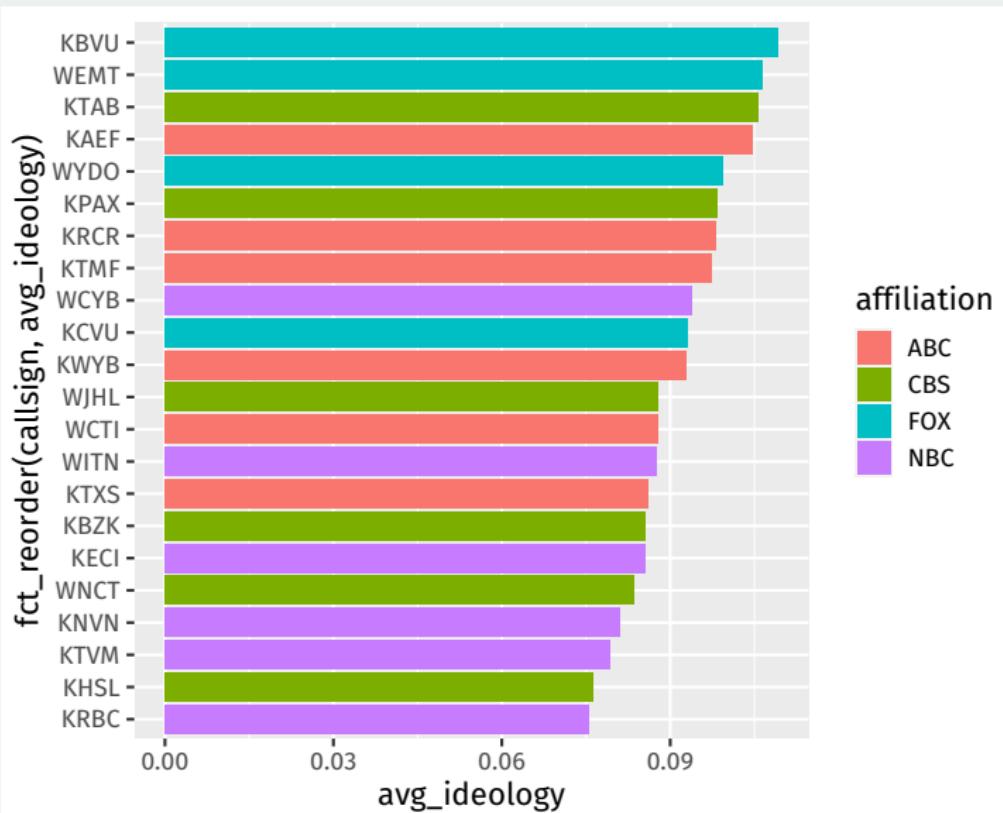
We would like to order the stations by ideology.

`fct_reorder(group, order_var)` function (loaded with tidyverse) will reorder the groups by the order bar (low to high). Easiest to put this in the mapping.

```
ggplot(station_ideology,  
       mapping = aes(x = avg_ideology,  
                      y = fct_reorder(callsign, avg_ideology))) +  
  geom_col()
```



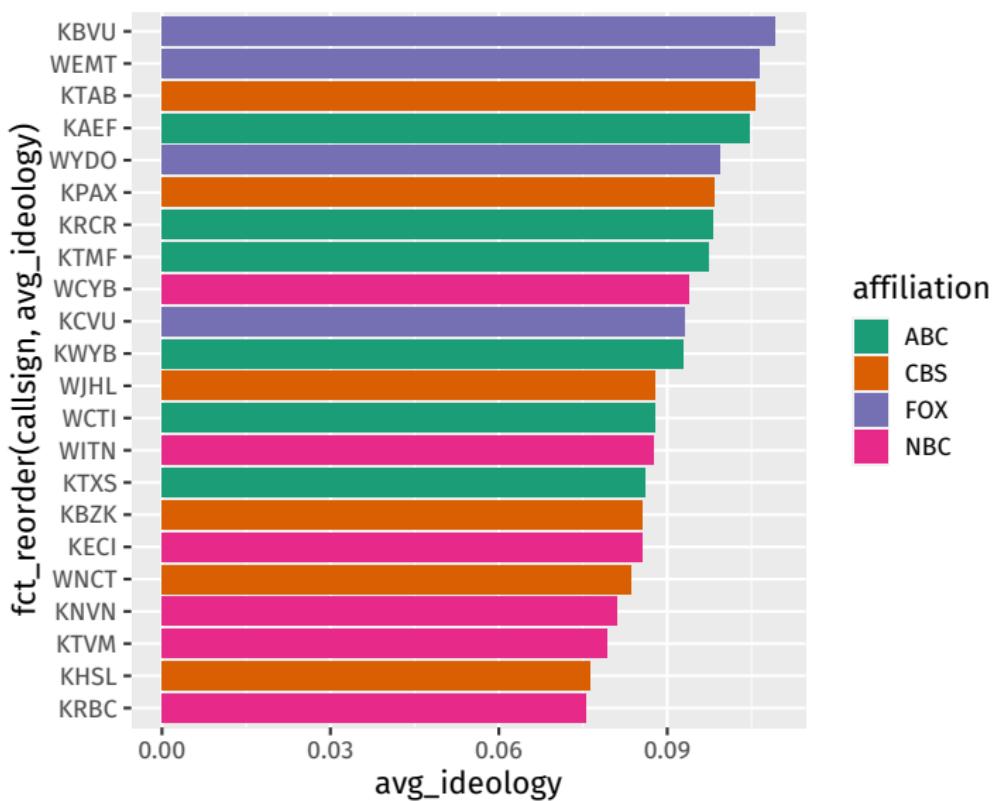
```
ggplot(station_ideology,  
       mapping = aes(x = avg_ideology,  
                      y = fct_reorder(callsign, avg_ideology))) +  
  geom_col(mapping = aes(fill = affiliation))
```



# Setting the color palette

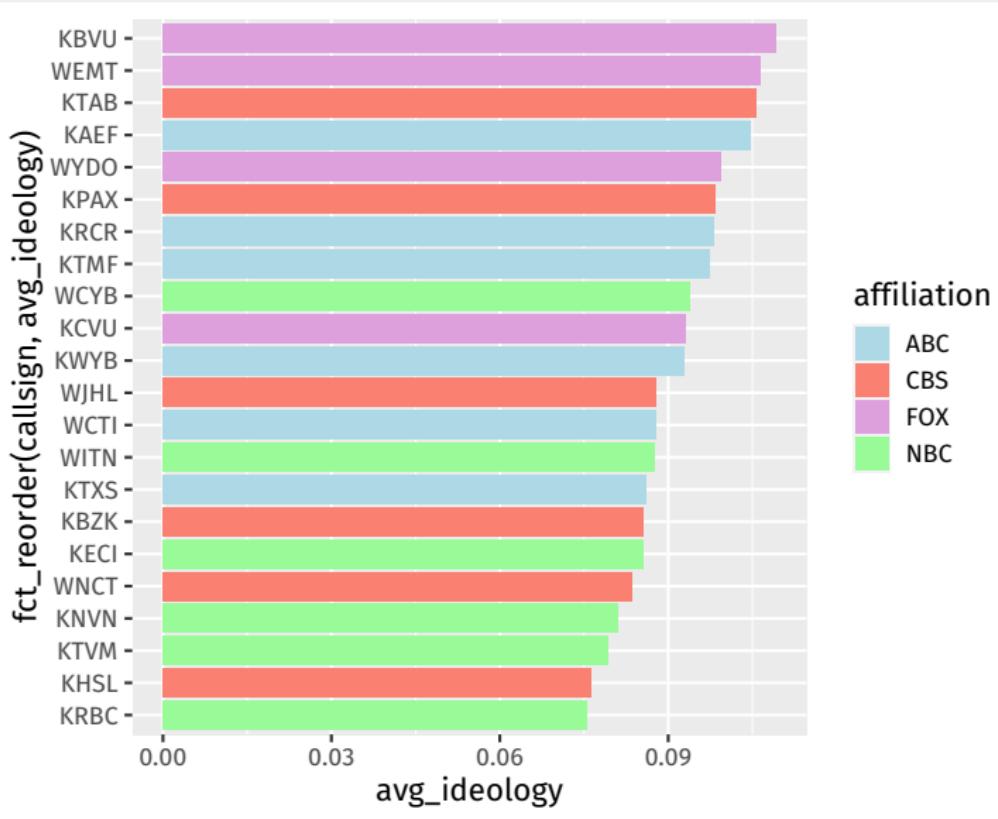
We can use color palettes from a project called ColorBrewer

```
ggplot(station_ideology,  
       mapping = aes(x = avg_ideology,  
                      y = fct_reorder(callsign, avg_ideology))) +  
  geom_col(mapping = aes(fill = affiliation)) +  
  scale_fill_brewer(palette = "Dark2")
```



# Manually setting the color palette

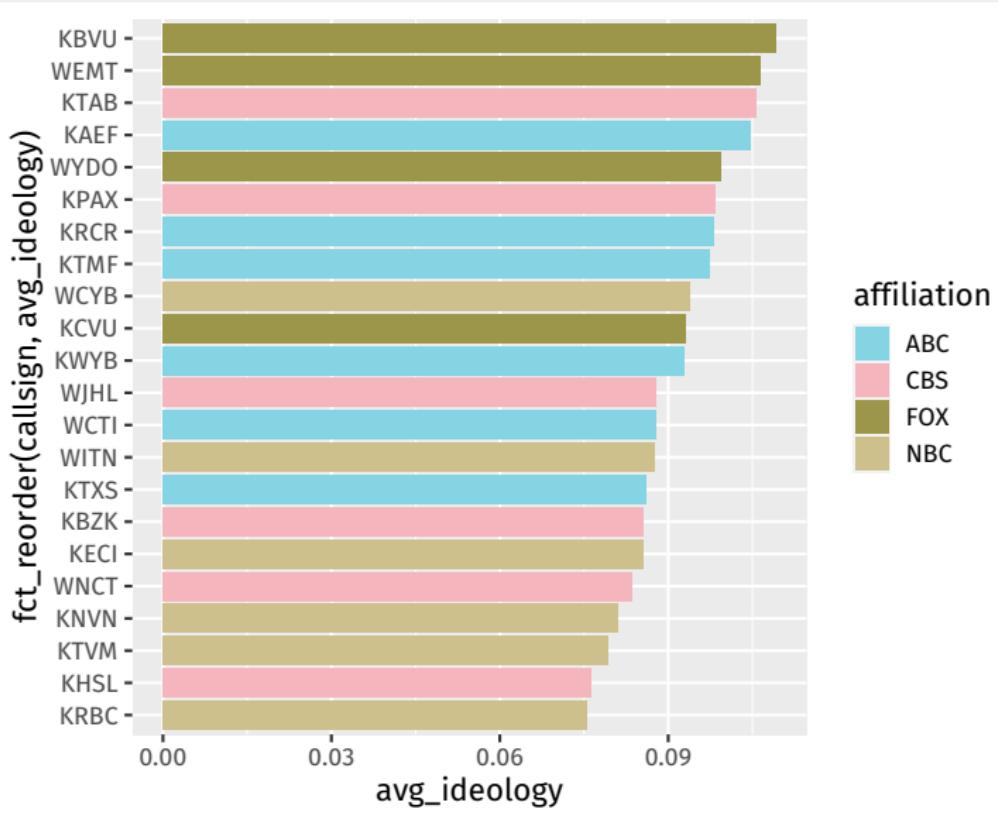
```
ggplot(station_ideology,
       mapping = aes(x = avg_ideology,
                     y = fct_reorder(callsign, avg_ideology))) +
  geom_col(mapping = aes(fill = affiliation)) +
  scale_fill_manual(values = c(ABC = "lightblue",
                               CBS = "salmon",
                               FOX = "plum",
                               NBC = "palegreen"))
```



# Fun with colors

Other packages provide more palettes:

```
library(wesanderson)
ggplot(station_ideology,
       mapping = aes(x = avg_ideology,
                      y = fct_reorder(callsign, avg_ideology))) +
  geom_col(mapping = aes(fill = affiliation)) +
  scale_fill_manual(values = wes_palette("Moonrise3"))
```



# Gov 50: 6. Causality

Matthew Blackwell

Harvard University

# Roadmap

1. What is causality?
2. Randomized experiments
3. Calculating effects

# **1/** What is causality?



Two roads diverged in a yellow wood,  
And sorry I could not travel both  
And be one traveler, long I stood  
And looked down one as far as I could  
To where it bent in the undergrowth;

# What is a causal effect?

factual

vs.

counterfactual

- Does increasing the minimum wage increase the unemployment rate?
  - Unemployment rate went up after the minimum wage increased
  - Would it have gone up if the minimum wage increase not occurred?
- Does having girls affect a judge's rulings in court?
  - A judge with a daughter gave a pro-choice ruling.
  - Would they have done that if had a son instead?
- **Fundamental problem of causal inference:**
  - Can never observe counterfactuals, must be inferred.

# Political canvassing study



POLITICAL SCIENCE

## Durably reducing transphobia: A field experiment on door-to-door canvassing

David Brockman<sup>1\*</sup> and Joshua Kalla<sup>2</sup>

Existing research depicts intergroup prejudices as deeply ingrained, requiring intense intervention to lastingly reduce. Here, we show that a single approximately 10-minute conversation encouraging actively taking the perspective of others can markedly reduce prejudice for at least 3 months. We illustrate this potential with a door-to-door canvassing intervention in South Florida targeting antitransgender prejudice. Despite declines in homophobic transphobia remains pervasive. For the intervention, 56 canvassers went door to door encouraging active perspective-taking with 501 voters at voters' doorsteps. A randomized trial found that these conversations substantially reduced transphobia, with decreases greater than Americans' average decrease in homophobia from 1998 to 2012. These effects persisted for 3 months, and both transgender and nontransgender canvassers were effective. The intervention also increased support for a nondiscrimination law, even after exposing voters to counterarguments.

- Can canvassers change minds about topics like transgender rights?
- Experimental setting:
  - Randomly assign canvassers to have a conversation about transgender right or a conversation about recycling.
  - Trans rights conversations focused on “perspective taking”
- Outcome of interest: support for trans rights policies.

# A tale of two respondents

|              | Conversation Script | Support for Nondiscrimination Law |
|--------------|---------------------|-----------------------------------|
| Respondent 1 | Recycling           | No                                |
| Respondent 2 | Trans rights        | Yes                               |

Did the second respondent support the law **because** of the perspective-taking conversation?

# Translating into math

Useful to have **compact** notation for referring to **treatment variable**:

$$T_i = \begin{cases} 1 & \text{if respondent } i \text{ had trans rights conversation} \\ 0 & \text{if respondent } i \text{ had recycling conversation} \end{cases}$$

Similar notation for the **outcome variable**:

$$Y_i = \begin{cases} 1 & \text{if respondent } i \text{ supports trans nondiscrimination laws} \\ 0 & \text{if respondent } i \text{ doesn't support nondiscrimination laws} \end{cases}$$

*i* is a placeholder to refer to a generic unit/respondent:  $Y_{42}$  is the outcome for the 42nd unit.

# A tale of two respondents (redux)

|              | Conversation Script | Support for Nondiscrimination Law |
|--------------|---------------------|-----------------------------------|
| Respondent 1 | Recycling           | No                                |
| Respondent 2 | Trans rights        | Yes                               |

becomes...

| $i$          | $T_i$ | $Y_i$ |
|--------------|-------|-------|
| Respondent 1 | 0     | 0     |
| Respondent 2 | 1     | 1     |

# Causal effects & counterfactuals

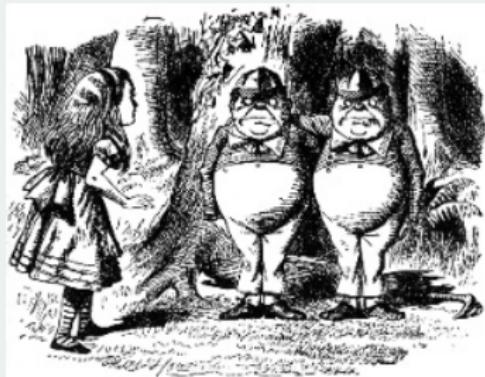
- What does “ $T_i$  causes  $Y_i$ ” mean?  $\rightsquigarrow$  **counterfactuals**, “what if”
- Would respondent change their support based on the conversation?
- Two **potential outcomes**:
  - $Y_i(1)$ : would respondent  $i$  support ND laws if they had trans rights script?
  - $Y_i(0)$ : would respondent  $i$  support ND laws if they had recycling script?
- **Causal effect**:  $Y_i(1) - Y_i(0)$ 
  - $Y_i(1) - Y_i(0) = 0 \rightsquigarrow$  script has no effect on policy views
  - $Y_i(1) - Y_i(0) = -1 \rightsquigarrow$  trans rights script lower support for laws
  - $Y_i(1) - Y_i(0) = +1 \rightsquigarrow$  trans rights script increases support for laws

# Potential outcomes

| $i$          | $T_i$ | $Y_i$ | $Y_i(1)$ | $Y_i(0)$ |
|--------------|-------|-------|----------|----------|
| Respondent 1 | 0     | 0     | ???      | 0        |
| Respondent 2 | 1     | 1     | 1        | ???      |

- **Fundamental problem of causal inference:**
  - We only observe one of the two potential outcomes.
  - Observe  $Y_i = Y_i(1)$  if  $T_i = 1$  or  $Y_i = Y_i(0)$  if  $T_i = 0$
- To infer causal effect, we need to infer the missing counterfactuals!

# How can we figure out counterfactuals?



- Find a similar unit!  $\rightsquigarrow$  **matching**
  - Mill's method of difference
- Does respondent support law because of the trans rights script?
  - $\rightsquigarrow$  find a identical respondent who got the recycling script.
- NJ increased the minimum wage. Causal effect on unemployment?
  - $\rightsquigarrow$  find a state similar to NJ that didn't increase minimum wage.

# Imperfect matches



- The problem: imperfect matches!
- Say we match  $i$  (treated) and  $j$  (control)
- **Selection Bias:**  $Y_i(1) \neq Y_j(1)$
- Those who take treatment may be different than those who take control.
- How can we correct for that?

# **2/** Randomized experiments

# Match groups not individuals



- **Randomized control trial:** each unit's treatment assignment is determined by chance.
  - Flip a coin; draw red and blue chips from a hat; etc
- Randomization ensures **balance** between treatment and control group.
  - Treatment and control group are identical **on average**
  - Similar on both observable and unobservable characteristics.

# A little more notation

- We will often refer to the **sample size** (number of units) as  $n$ .
- We often have  $n$  measurements of some variable:  $(Y_1, Y_2, \dots, Y_n)$
- How many in our sample support nondiscrimination laws?

$$Y_1 + Y_2 + Y_3 + \cdots + Y_n$$

- Notation is a bit clunky, so we often use the **Sigma notation**:

$$\sum_{i=1}^n Y_i = Y_1 + Y_2 + Y_3 + \cdots + Y_n$$

- $\sum_{i=1}^n$  means sum each value from  $Y_1$  to  $Y_n$

# Averages

- The **sample average** or **sample mean** is simply the sum of all values divided by the number of values.
- Sigma notation allows us to write this in a compact way:

$$\bar{Y} = \frac{1}{n} \sum_{i=1}^n Y_i$$

- Suppose we surveyed 6 people and 3 supported nondiscrim. laws:

$$\bar{Y} = \frac{1}{6} (1 + 1 + 1 + 0 + 0 + 0) = 0.5$$

# Quantity of interest

- We want to estimate the average causal effects over all units:

$$\begin{aligned}\text{Sample Average Treatment Effect (SATE)} &= \frac{1}{n} \sum_{i=1}^n \{Y_i(1) - Y_i(0)\} \\ &= \frac{1}{n} \sum_{i=1}^n Y_i(1) - \frac{1}{n} \sum_{i=1}^n Y_i(0)\end{aligned}$$

- Why can't we just calculate this quantity directly?
- What we can estimate instead:

$$\text{Difference in means} = \bar{Y}_{\text{treated}} - \bar{Y}_{\text{control}}$$

- $\bar{Y}_{\text{treated}}$ : sample average outcome for treated group
- $\bar{Y}_{\text{control}}$ : sample average outcome for control group
- When will the difference-in-means is a good estimate of the SATE?

# Why randomization works

- Under an RCT, treatment and control groups are random samples.
- Average in the treatment group will be similar to average if all treated:

$$\bar{Y}_{\text{treated}} \approx \frac{1}{n} \sum_{i=1}^n Y_i(1)$$

- Average in the control group will be similar to average if all untreated:

$$\bar{Y}_{\text{control}} \approx \frac{1}{n} \sum_{i=1}^n Y_i(0)$$

- Implies difference-in-means should be close to SATE:

$$\bar{Y}_{\text{treated}} - \bar{Y}_{\text{control}} \approx \frac{1}{n} \sum_{i=1}^n Y_i(1) - \frac{1}{n} \sum_{i=1}^n Y_i(0) = \frac{1}{n} \sum_{i=1}^n \{Y_i(1) - Y_i(0)\} = \text{SATE}$$

# Some potential problems with RCTs

- **Placebo effects:**
  - Respondents will be affected by any intervention, even if they shouldn't have any effect.
  - Reason to have control group be recycling script
- **Hawthorne effects:**
  - Respondents act differently just knowing that they are under study.

# Balance checking

- Can we determine if randomization “worked”?
- If it did, we shouldn’t see large differences between treatment and control group on **pretreatment variable**.
  - Pretreatment variable are those that are unaffected by treatment.
- We can check in the actual data for some pretreatment variable  $X$ 
  - $\bar{X}_{\text{treated}}$ : average value of variable for treated group.
  - $\bar{X}_{\text{control}}$ : average value of variable for control group.
  - Under randomization,  $\bar{X}_{\text{treated}} - \bar{X}_{\text{control}} \approx 0$

# Multiple treatments

- Instead of 1 treatment, we might have multiple **treatment arms**:
  - Control condition
  - Treatment A
  - Treatment B
  - Treatment C, etc
- In this case, we will look at multiple comparisons:
  - $\bar{Y}_{\text{treated, A}} - \bar{Y}_{\text{control}}$
  - $\bar{Y}_{\text{treated, B}} - \bar{Y}_{\text{control}}$
  - $\bar{Y}_{\text{treated, A}} - \bar{Y}_{\text{treated, B}}$
- If treatment arms are randomly assigned, these differences will be good estimators for each causal contrast.

# **3/** Calculating effects

# Transphobia study data

```
## reinstall gov50data if necessary  
library(gov50data)
```

---

| Variable Name   | Description   |
|-----------------|---|
| age             | Age of the R in years                                   |
| female          | 1=R marked “Female” on voter reg., 0 otherwise          |
| voted_gen_14    | 1 if R voted in the 2014 general election               |
| vote_gen_12     | 1 if R voted in the 2012 general election               |
| treat_ind       | 1 if R assigned to trans rights script, 0 for recycling |
| racename        | name of racial identity indicated on voter file         |
| democrat        | 1 if R is a registered Democrat                         |
| nondiscrim_pre  | 1 if R supports nondiscrim. law at baseline             |
| nondiscrim_post | 1 if R supports nondiscrim. law after 3 months          |

---

# Peak at the data

```
trans
```

```
## # A tibble: 565 x 9
##       age female voted_gen_14 voted~1 treat~2 racen~3 democ~4
##     <dbl>   <dbl>      <dbl>    <dbl>    <dbl> <chr>    <dbl>
## 1     29     0          0        1        0 Africa~     1
## 2     59     1          1        0        1 Africa~     1
## 3     35     1          1        1        1 Africa~     1
## 4     63     1          1        1        1 Africa~     1
## 5     65     0          1        1        1 Africa~     0
## 6     51     1          1        1        0 Caucas~     0
## 7     26     1          1        1        0 Africa~     1
## 8     62     1          1        1        1 Africa~     1
## 9     37     0          1        1        0 Caucas~     0
## 10    51     1          1        1        0 Caucas~     0
## # ... with 555 more rows, 2 more variables:
## #   nondiscrim_pre <dbl>, nondiscrim_post <dbl>, and
## #   abbreviated variable names 1: voted_gen_12,
## #   2: treat_ind, 3: racename, 4: democrat
```

# Calculate the average outcomes in each group

```
treat_mean <- trans |>
  filter(treat_ind == 1) |>
  summarize(nondiscrim_mean = mean(nondiscrim_post))
treat_mean
```

```
## # A tibble: 1 x 1
##   nondiscrim_mean
##             <dbl>
## 1           0.687
```

```
control_mean <- trans |>
  filter(treat_ind == 0) |>
  summarize(nondiscrim_mean = mean(nondiscrim_post))
control_mean
```

```
## # A tibble: 1 x 1
##   nondiscrim_mean
##             <dbl>
## 1           0.648
```

# Calculating the difference in means

```
treat_mean - control_mean
```

```
##     nondiscrim_mean  
## 1             0.039
```

We'll see more ways to do this throughout the semester.

# Checking balance on numeric covariates

We can use `group_by` to see how the mean of covariates varies by group:

```
trans |>  
  group_by(treat_ind) |>  
  summarize(age_mean = mean(age))
```

```
## # A tibble: 2 x 2  
##   treat_ind age_mean  
##       <dbl>     <dbl>  
## 1         0     48.2  
## 2         1     48.3
```

# Checking balance on categorical covariates

Or we can group by treatment and a categorical control:

```
trans |>  
  group_by(treat_ind, racename) |>  
  summarize(n = n())
```

```
## # A tibble: 9 x 3  
## # Groups:   treat_ind [2]  
##   treat_ind racename      n  
##       <dbl> <chr>     <int>  
## 1       0 African American    58  
## 2       0 Asian             2  
## 3       0 Caucasian        77  
## 4       0 Hispanic          150  
## 5       1 African American    68  
## 6       1 Asian              4  
## 7       1 Caucasian         75  
## 8       1 Hispanic           130  
## 9       1 Native American     1
```

Hard to read!

# pivot\_wider

`pivot_wider()` takes data from a single column and moves it into multiple columns based on a grouping variable:

```
trans |>  
  group_by(treat_ind, racename) |>  
  summarize(n = n()) |>  
  pivot_wider(  
    names_from = treat_ind,  
    values_from = n  
  )
```

`names_from` tells us what variable will map onto the columns

`values_from` tell us what values should be mapped into those columns

```
trans |>
  group_by(treat_ind, racename) |>
  summarize(n = n()) |>
  pivot_wider(
    names_from = treat_ind,
    values_from = n
  )
```

```
## # A tibble: 5 x 3
##   racename      `0`     `1`
##   <chr>        <int> <int>
## 1 African American    58     68
## 2 Asian            2       4
## 3 Caucasian        77     75
## 4 Hispanic         150    130
## 5 Native American    NA      1
```

# Calculating diff-in-means by group

```
trans |>
  mutate(
    treat_ind = if_else(treat_ind == 1, "Treated", "Control"),
    party = if_else(democrat == 1, "Democrat", "Non-Democrat")
  ) |>
  group_by(treat_ind, party) |>
  summarize(nondiscrim_mean = mean(nondiscrim_post)) |>
  pivot_wider(
    names_from = treat_ind,
    values_from = nondiscrim_mean
  ) |>
  mutate(
    diff_in_means = Treated - Control
  )
```

```
## # A tibble: 2 x 4
##   party      Control Treated diff_in_means
##   <chr>     <dbl>    <dbl>        <dbl>
## 1 Democrat   0.704    0.754       0.0498
## 2 Non-Democrat 0.605    0.628       0.0234
```

```
## # A tibble: 2 x 4
##   party      Control Treated diff_in_means
##   <chr>     <dbl>    <dbl>        <dbl>
## 1 Democrat   0.704    0.754       0.0498
## 2 Non-Democrat  0.605    0.628       0.0234
```

# Gov 50: 7. Observational Studies

Matthew Blackwell

Harvard University

# Roadmap

1. Calculating effects
2. Observational Studies

# 1/ Calculating effects

# Transphobia study data

```
## reinstall gov50data if necessary  
library(gov50data)
```

---

| Variable Name   | Description   |
|-----------------|---|
| age             | Age of the R in years                                   |
| female          | 1=R marked “Female” on voter reg., 0 otherwise          |
| voted_gen_14    | 1 if R voted in the 2014 general election               |
| vote_gen_12     | 1 if R voted in the 2012 general election               |
| treat_ind       | 1 if R assigned to trans rights script, 0 for recycling |
| racename        | name of racial identity indicated on voter file         |
| democrat        | 1 if R is a registered Democrat                         |
| nondiscrim_pre  | 1 if R supports nondiscrim. law at baseline             |
| nondiscrim_post | 1 if R supports nondiscrim. law after 3 months          |

---

# Peak at the data

```
trans
```

```
## # A tibble: 565 x 9
##       age female voted_gen_14 voted~1 treat~2 racen~3 democ~4
##     <dbl>   <dbl>      <dbl>    <dbl>    <dbl> <chr>    <dbl>
## 1     29     0          0        1        0 Africa~     1
## 2     59     1          1        0        1 Africa~     1
## 3     35     1          1        1        1 Africa~     1
## 4     63     1          1        1        1 Africa~     1
## 5     65     0          1        1        1 Africa~     0
## 6     51     1          1        1        0 Caucas~     0
## 7     26     1          1        1        0 Africa~     1
## 8     62     1          1        1        1 Africa~     1
## 9     37     0          1        1        0 Caucas~     0
## 10    51     1          1        1        0 Caucas~     0
## # ... with 555 more rows, 2 more variables:
## #   nondiscrim_pre <dbl>, nondiscrim_post <dbl>, and
## #   abbreviated variable names 1: voted_gen_12,
## #   2: treat_ind, 3: racename, 4: democrat
```

# Calculate the average outcomes in each group

```
treat_mean <- trans |>
  filter(treat_ind == 1) |>
  summarize(nondiscrim_mean = mean(nondiscrim_post))
treat_mean
```

```
## # A tibble: 1 x 1
##   nondiscrim_mean
##             <dbl>
## 1           0.687
```

```
control_mean <- trans |>
  filter(treat_ind == 0) |>
  summarize(nondiscrim_mean = mean(nondiscrim_post))
control_mean
```

```
## # A tibble: 1 x 1
##   nondiscrim_mean
##             <dbl>
## 1           0.648
```

# Calculating the difference in means

```
treat_mean - control_mean
```

```
##     nondiscrim_mean  
## 1             0.039
```

We'll see more ways to do this throughout the semester.

# Checking balance on numeric covariates

We can use `group_by` to see how the mean of covariates varies by group:

```
trans |>  
  group_by(treat_ind) |>  
  summarize(age_mean = mean(age))
```

```
## # A tibble: 2 x 2  
##   treat_ind age_mean  
##       <dbl>     <dbl>  
## 1         0     48.2  
## 2         1     48.3
```

# Checking balance on categorical covariates

Or we can group by treatment and a categorical control:

```
trans |>  
  group_by(treat_ind, racename) |>  
  summarize(n = n())
```

```
## # A tibble: 9 x 3  
## # Groups:   treat_ind [2]  
##   treat_ind racename      n  
##       <dbl> <chr>     <int>  
## 1       0 African American    58  
## 2       0 Asian             2  
## 3       0 Caucasian        77  
## 4       0 Hispanic          150  
## 5       1 African American    68  
## 6       1 Asian              4  
## 7       1 Caucasian         75  
## 8       1 Hispanic           130  
## 9       1 Native American     1
```

Hard to read!

# pivot\_wider

`pivot_wider()` takes data from a single column and moves it into multiple columns based on a grouping variable:

```
trans |>  
  group_by(treat_ind, racename) |>  
  summarize(n = n()) |>  
  pivot_wider(  
    names_from = treat_ind,  
    values_from = n  
  )
```

`names_from` tells us what variable will map onto the columns

`values_from` tell us what values should be mapped into those columns

```
trans |>  
  group_by(treat_ind, racename) |>  
  summarize(n = n()) |>  
  pivot_wider(  
    names_from = treat_ind,  
    values_from = n  
  )
```

```
## # A tibble: 5 x 3  
##   racename      `0`     `1`  
##   <chr>        <int> <int>  
## 1 African American     58     68  
## 2 Asian             2       4  
## 3 Caucasian         77     75  
## 4 Hispanic          150    130  
## 5 Native American    NA      1
```

# Calculating diff-in-means by group

```
trans |>
  mutate(
    treat_ind = if_else(treat_ind == 1, "Treated", "Control"),
    party = if_else(democrat == 1, "Democrat", "Non-Democrat")
  ) |>
  group_by(treat_ind, party) |>
  summarize(nondiscrim_mean = mean(nondiscrim_post)) |>
  pivot_wider(
    names_from = treat_ind,
    values_from = nondiscrim_mean
  ) |>
  mutate(
    diff_in_means = Treated - Control
  )
```

```
## # A tibble: 2 x 4
##   party      Control Treated diff_in_means
##   <chr>     <dbl>    <dbl>        <dbl>
## 1 Democrat    0.704    0.754       0.0498
## 2 Non-Democrat 0.605    0.628       0.0234
```

```
## # A tibble: 2 x 4
##   party      Control Treated diff_in_means
##   <chr>     <dbl>    <dbl>        <dbl>
## 1 Democrat   0.704    0.754       0.0498
## 2 Non-Democrat  0.605    0.628       0.0234
```

# 2/ Observational Studies

# Do newspaper endorsements matter?



- Can newspaper endorsements change voters' minds?
- Why not compare vote choice of readers of different papers?
  - Problem: readers choose papers based on their previous beliefs.
  - Liberals ↗ New York Times, conservatives ↗ Wall Street Journal.
- Study for today: British newspapers switching their endorsements.
  - Some newspapers endorsing Tories in 1992 switched to Labour in 1997.
  - **Treated group:** readers of Tory → Labour papers.
  - **Control group:** readers of papers who didn't switch.

# Data

| Name          | Description   |
|---------------|---|
| to_labour     | Read a newspaper that switched endorsement to Labour between 1992 and 1997 (1=Yes, 0=No)? |
| vote_lab_92   | Did respondent vote for Labour in 1992 election (1=Yes, 0=No)?                            |
| vote_lab_97   | Did respondent vote for Labour in 1997 election (1=Yes, 0=No)?                            |
| age           | Age of respondent   |
| male          | Does the respondent identify as Male (1=Yes, 0=No)?                                       |
| parent_labour | Did the respondent's parents vote for Labour (1=Yes, 0=No)?                               |
| work_class    | Does the respondent identify as working class (1=Yes, 0=No)?                              |

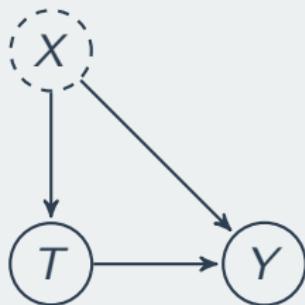
```
library(tidyverse)
library(gov50data)
newspapers

## # A tibble: 1,593 x 7
##   to_labour vote_lab_92 vote_~1 age     male paren~2 work_~3
##   <dbl>      <dbl>    <dbl> <dbl> <dbl>    <dbl>    <dbl>
## 1 0          1        1 33 0       1       1
## 2 0          1        0 51 0       1       0
## 3 0          0        0 46 0       1       1
## 4 0          1        1 45 1       1       1
## 5 0          1        1 29 0       1       1
## 6 0          1        1 47 1       1       1
## 7 0          1        1 34 1       0       1
## 8 0          1        1 31 0       1       1
## 9 0          1        1 24 1       1       1
## 10 1         1        1 48 0       1       1
## # ... with 1,583 more rows, and abbreviated variable names
## #   1: vote_lab_97, 2: parent_labour, 3: work_class
```

# Observational studies

- Example of an **observational study**:
  - We as researchers observe a naturally assigned treatment
  - Very common: often can't randomize for ethical/logistical reasons.
- **Internal validity**: are the causal assumption satisfied? Can we interpret this as a causal effect?
  - RCTs usually have higher internal validity.
  - Observational studies less so because treatment and control groups may differ in ways that are hard to measure
- **External validity**: can the conclusions/estimated effects be generalized beyond this study?
  - RCTs weaker here because often very expensive to conduct on representative samples.
  - Observational studies often have larger/more representative samples that improve external validity.

# Confounding



- **Confounder:** pre-treatment variable affecting treatment & the outcome.
  - Leftists ( $X$ ) more likely to read newspapers switching to Labour ( $T$ ).
  - Leftists ( $X$ ) also more likely to vote for Labour ( $Y$ ).
- **Confounding bias** in the estimated SATE due to these differences
  - $\bar{Y}_{\text{control}}$  not a good proxy for  $\frac{1}{n} \sum_{i=1}^n Y_i(0)$  in treated group.
  - one type: **selection bias** from self-selection into treatment

# Research designs

- How can we find a good comparison group?
- Depends on the data we have available.
- Three general types of observational study **research designs:**
  1. **Cross-sectional design:** compare outcomes treated and control units at one point in time.
  2. **Before-and-after design:** compare outcomes before and after a unit has been treated, but need over-time data on treated group.
  3. **Difference-in-differences design:** use before/after information for the treated and control group; need over-time on treated & control group.

# Cross-sectional design

- Compare treatment and control groups after treatment happens.
  - Readers of switching papers vs readers of non-switching papers in 1997.
- Treatment & control groups assumed identical on average as in RCT.
  - Sometimes called **unconfoundedness** or **as-if randomized**.
- Cross-section comparison estimate:

$$\bar{Y}_{\text{treated}}^{\text{after}} - \bar{Y}_{\text{control}}^{\text{after}}$$

- Could there be confounders?

# Cross-sectional design in R

```
switched <- newspapers |>
  filter(to_labour == 1) |>
  summarize(mean(vote_lab_97))

no_change <- newspapers %>%
  filter(to_labour == 0) |>
  summarize(mean(vote_lab_97))

switched - no_change

##   mean(vote_lab_97)
## 1          0.14
```

# Statistical control

- **Statistical control:** adjust for confounders using statistical procedures.
  - Can help to reduce confounding bias.
- One type of statistical control: **subclassification**
  - Compare treated and control groups within levels of a confounder.
  - Remaining effect can't be due to the confounder.
- Threat to inference: we can only control for observed variables ↪ threat of **unmeasured confounding**

# Statistical control in R

```
newspapers %>%
  group_by(parent_labour, to_labour) %>%
  summarize(avg_vote = mean(vote_lab_97)) %>%
  pivot_wider(
    names_from = to_labour,
    values_from = avg_vote
  ) %>%
  mutate(diff_by_parent = `1` - `0`)
```

```
## # A tibble: 2 x 4
## # Groups:   parent_labour [2]
##   parent_labour    `0`    `1` diff_by_parent
##   <dbl> <dbl> <dbl>        <dbl>
## 1 0     0.279 0.434      0.155
## 2 1     0.597 0.698      0.101
```

# Before-and-after comparison

- Compare readers of party-switching newspapers before & after switch.
- Advantage: all person-specific features held fixed
  - comparing within a person over time.
- Before-and-after estimate:

$$\bar{Y}_{\text{treated}}^{\text{after}} - \bar{Y}_{\text{treated}}^{\text{before}}$$

- Threat to inference: **time-varying confounders**
  - Time trend: Labour just did better overall in 1997 compared to 1992.

# Before and after in R

```
newspapers |>
  mutate(
    vote_change = vote_lab_97 - vote_lab_92
  ) |>
  summarize(avg_change = mean(vote_change))
```

```
## # A tibble: 1 x 1
##   avg_change
##       <dbl>
## 1     0.119
```

# Differences in differences

- Key idea: use the before-and-after difference of **control group** to infer what would have happened to **treatment group** without treatment.
- DiD estimate:

$$\underbrace{\left( \bar{Y}_{\text{treated}}^{\text{after}} - \bar{Y}_{\text{treated}}^{\text{before}} \right)}_{\text{trend in treated group}} - \underbrace{\left( \bar{Y}_{\text{control}}^{\text{after}} - \bar{Y}_{\text{control}}^{\text{before}} \right)}_{\text{trend in control group}}$$

- Change in treated group above and beyond the change in control group.
- **Parallel time trend assumption**
  - Changes in vote of readers of non-switching papers roughly the same as changes that readers of switching papers would have been if they read non-switching papers.
  - Threat to inference: non-parallel trends.

# Difference-in-differences in R

```
newspapers |>
  mutate(
    vote_change = vote_lab_97 - vote_lab_92,
    to_labour = if_else(to_labour == 1, "switched", "unswitched")
  ) |>
  group_by(to_labour) |>
  summarize(avg_change = mean(vote_change)) |>
  pivot_wider(
    names_from = to_labour,
    values_from = avg_change
  ) |>
  mutate(DID = switched - unswitched)
```

```
## # A tibble: 1 x 3
##   switched unswitched     DID
##       <dbl>      <dbl>  <dbl>
## 1     0.190      0.110  0.0796
```

# Summarizing approaches

## 1. Cross-sectional comparison

- Compare treated units with control units after treatment
- Assumption: treated and controls units are comparable
- Possible confounding

## 2. Before-and-after comparison

- Compare the same units before and after treatment
- Assumption: no time-varying confounding

## 3. Differences-in-differences

- Assumption: parallel trends assumptions
- Under this assumption, it accounts for unit-specific and time-varying confounding.
- All rely on assumptions that can't be verified to handle confounding.
- RCTs handle confounding by design.

# Causality understanding check

I USED TO THINK  
CORRELATION IMPLIED  
CAUSATION.



THEN I TOOK A  
STATISTICS CLASS.  
NOW I DON'T.



SOUNDS LIKE THE  
CLASS HELPED.  
WELL, MAYBE.



# Gov 50: 8. Summarizing Data

Matthew Blackwell

Harvard University

# Roadmap

1. Descriptive Statistics
2. Missing data
3. Proportion tables

# 1/ Descriptive Statistics

# Lots of data

```
library(tidyverse)
library(gapminder)
gapminder

## # A tibble: 1,704 x 6
##   country   continent year lifeExp      pop gdpPercap
##   <fct>     <fct>    <int>   <dbl>    <int>     <dbl>
## 1 Afghanistan Asia     1952    28.8  8425333    779.
## 2 Afghanistan Asia     1957    30.3  9240934    821.
## 3 Afghanistan Asia     1962    32.0  10267083   853.
## 4 Afghanistan Asia     1967    34.0  11537966   836.
## 5 Afghanistan Asia     1972    36.1  13079460   740.
## 6 Afghanistan Asia     1977    38.4  14880372   786.
## 7 Afghanistan Asia     1982    39.9  12881816   978.
## 8 Afghanistan Asia     1987    40.8  13867957   852.
## 9 Afghanistan Asia     1992    41.7  16317921   649.
## 10 Afghanistan Asia    1997    41.8  22227415   635.
## # ... with 1,694 more rows
```

# Lots and lots of data

```
head(gapminder$gdpPerCap, n = 200)
```

```
## [1] 779 821 853 836 740 786 978 852 649
## [10] 635 727 975 1601 1942 2313 2760 3313 3533
## [19] 3631 3739 2497 3193 4604 5937 2449 3014 2551
## [28] 3247 4183 4910 5745 5681 5023 4797 5288 6223
## [37] 3521 3828 4269 5523 5473 3009 2757 2430 2628
## [46] 2277 2773 4797 5911 6857 7133 8053 9443 10079
## [55] 8998 9140 9308 10967 8798 12779 10040 10950 12217
## [64] 14526 16789 18334 19477 21889 23425 26998 30688 34435
## [73] 6137 8843 10751 12835 16662 19749 21597 23688 27042
## [82] 29096 32418 36126 9867 11636 12753 14805 18269 19340
## [91] 19211 18524 19036 20292 23404 29796 684 662 686
## [100] 721 630 660 677 752 838 973 1136 1391
## [109] 8343 9715 10991 13149 16672 19118 20980 22526 25576
## [118] 27561 30486 33693 1063 960 949 1036 1086 1029
## [127] 1278 1226 1191 1233 1373 1441 2677 2128 2181
## [136] 2587 2980 3548 3157 2754 2962 3326 3413 3822
## [145] 974 1354 1710 2172 2860 3528 4127 4314 2547
## [154] 4766 6019 7446 851 918 984 1215 2264 3215
## [163] 4551 6206 7954 8647 11004 12570 2109 2487 3337
## [172] 3430 4986 6660 7031 7807 6950 7958 8131 9066
```

# How to summarize data

- How should we summarize the wages data? Many possibilities!
  - Up to now: focus on **averages** or means of variables.
- Two salient features of a variable that we want to know:
  - **Central tendency:** where is the middle/typical/average value.
  - **Spread** around the center: are all values to the center or spread out?

# Center of the data

- “Center” of the data: typical/average value.
- **Mean:** sum of the values divided by the number of observations

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

- **Median:**

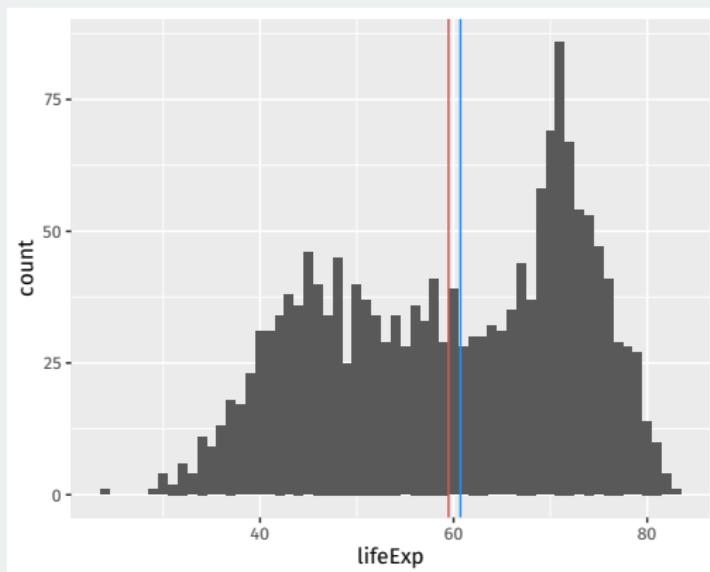
$$\text{median} = \begin{cases} \text{middle value} & \text{if number of entries is odd} \\ \frac{\text{sum of two middle values}}{2} & \text{if number of entries is even} \end{cases}$$

- In R: `mean( )` and `median( )`.

# Mean vs median

- Median more robust to **outliers**:
  - Example 1: data = {0, 1, 2, 3, 5}. Mean? Median?
  - Example 2: data = {0, 1, 2, 3, 100}. Mean? Median?
- What does Mark Zuckerberg do to the mean vs median income?

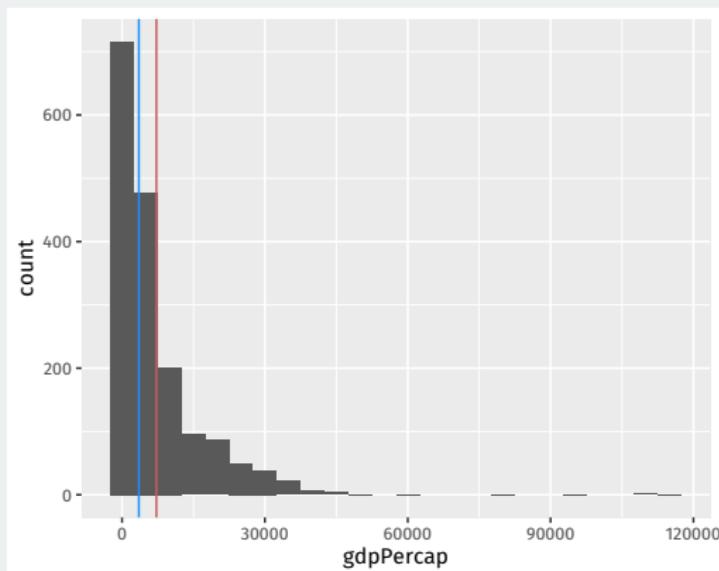
```
ggplot(gapminder, aes(x = lifeExp)) +  
  geom_histogram(binwidth = 1) +  
  geom_vline(aes(xintercept = mean(lifeExp)), color = "indianred") +  
  geom_vline(aes(xintercept = median(lifeExp)), color = "dodgerblue")
```



```
summary(gapminder$lifeExp)
```

|    | Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. |
|----|------|---------|--------|------|---------|------|
| ## | 23.6 | 48.2    | 60.7   | 59.5 | 70.8    | 82.6 |

```
ggplot(gapminder, aes(x = gdpPercap)) +  
  geom_histogram(binwidth = 5000) +  
  geom_vline(aes(xintercept = mean(gdpPercap)), color = "indianred") +  
  geom_vline(aes(xintercept = median(gdpPercap)), color = "dodgerblue")
```

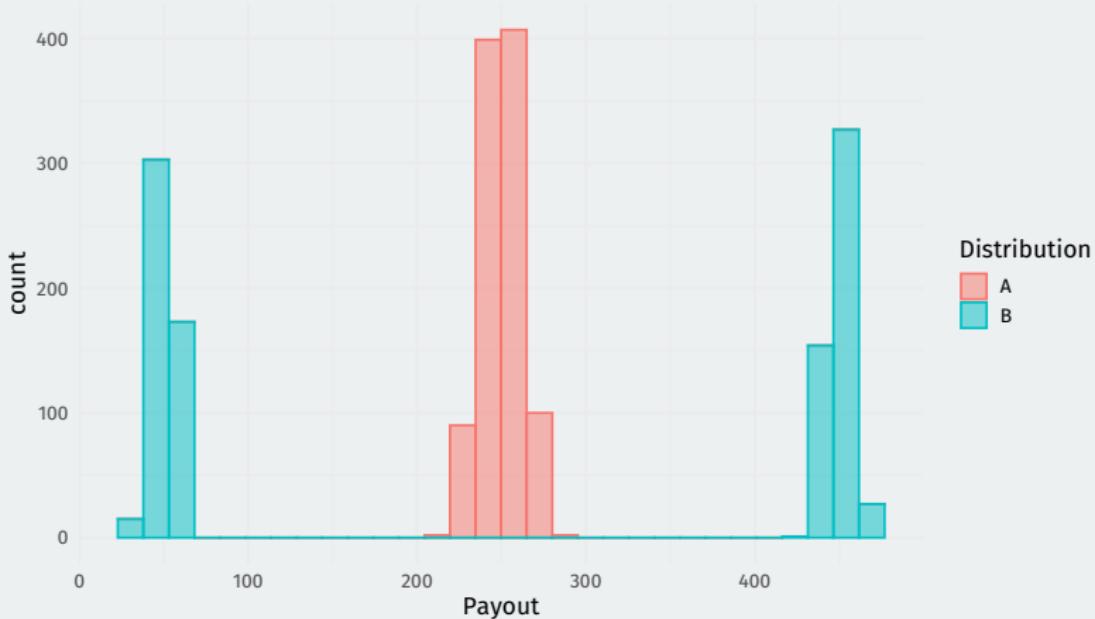


```
summary(gapminder$gdpPercap)
```

|    | Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max.   |
|----|------|---------|--------|------|---------|--------|
| ## | 241  | 1202    | 3532   | 7215 | 9325    | 113523 |

# Which distribution would you prefer?

Lottery where we randomly draw one value from A or B:



They have the same mean, so why do we care about the difference? **Spread!!**

# Spread of the data

- Are the values of the variable close to the center?
- **Range:**  $[\min(X), \max(X)]$
- **Quantile** (quartile, percentile, etc): divide data into equal sized groups.
  - 25th percentile = lower quartile (25% of the data below this value)
  - 50th percentile = median (50% of the data below this value)
  - 75th percentile = upper quartile (75% of the data below this value)
- **Interquartile range (IQR):** a measure of variability
  - How spread out is the middle half of the data?
  - Is most of the data really close to the median or are the values spread out?
- **R function:** `range()`, `summary()`, `IQR()`

# Standard deviation

- **Standard deviation:** On average, how far away are data points from the mean?

$$\text{standard deviation} = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2}$$

- Steps:
  1. Subtract each data point by the mean.
  2. Square each resulting difference.
  3. Take the sum of these values
  4. Divide by  $n - 1$  (or  $n$ , doesn't matter much)
  5. Take the square root.
- **Variance** = standard deviation<sup>2</sup>
- Why not just take the average deviations from mean without squaring?

# 2/ Missing data

# Missing data

- **Nonresponse:** respondent can't or won't answer question.
  - Sensitive questions ↪ **social desirability bias**
  - Some countries lack official statistics like unemployment.
  - Leads to missing data.
- Missing data in R: a special value NA
- Have already seen how to use `na.rm = TRUE`

# CCES data

```
library(gov50data)
cces_2020

## # A tibble: 51,551 x 6
##   gender race  educ          pid3  turno~1 pres_~2
##   <fct>  <fct> <fct>        <fct>    <dbl> <fct>
## 1 Male   White 2-year      Republ~     1 Donald~
## 2 Female White Post-grad  Democr~     NA <NA>
## 3 Female White 4-year     Indepe~     1 Joe Bi~
## 4 Female White 4-year     Democr~     1 Joe Bi~
## 5 Male   White 4-year     Indepe~     1 Other
## 6 Male   White Some college Republ~     1 Donald~
## 7 Male   Black Some college Not su~     NA <NA>
## 8 Female White Some college Indepe~     1 Donald~
## 9 Female White High school graduate Republ~     1 Donald~
## 10 Female White 4-year    Democr~     1 Joe Bi~
## # ... with 51,541 more rows, and abbreviated variable names
## #   1: turnout_self, 2: pres_vote
```

# drop\_na() to remove rows with missing values

```
cces_2020 |>
  drop_na()

## # A tibble: 45,651 x 6
##   gender race  educ      pid3  turno~1 pres_~2
##   <fct>  <fct> <fct>    <fct>    <dbl> <fct>
## 1 Male    White 2-year  Republ~     1 Donald~
## 2 Female  White 4-year  Indepe~     1 Joe Bi~
## 3 Female  White 4-year  Democr~     1 Joe Bi~
## 4 Male    White 4-year  Indepe~     1 Other
## 5 Male    White Some college Republ~     1 Donald~
## 6 Female  White Some college Indepe~     1 Donald~
## 7 Female  White High school graduate Republ~     1 Donald~
## 8 Female  White 4-year   Democr~     1 Joe Bi~
## 9 Female  White 4-year   Democr~     1 Joe Bi~
## 10 Female White 4-year  Democr~     1 Joe Bi~
## # ... with 45,641 more rows, and abbreviated variable names
## #   1: turnout_self, 2: pres_vote
```

# Drop rows based on certain variables

```
cces_2020 |>  
dim_desc()
```

```
## [1] "[51,551 x 6]"
```

```
cces_2020 |>  
drop_na() |>  
dim_desc()
```

```
## [1] "[45,651 x 6]"
```

```
cces_2020 |>  
drop_na(turnout_self) |>  
dim_desc()
```

```
## [1] "[48,462 x 6]"
```

# Available-case vs complete-case analysis

**Available-case analysis:** use the data you have for that variable:

```
cces_2020 |>  
  summarize(mean(turnout_self, na.rm = TRUE)) |>  
  pull()
```

```
## [1] 0.942
```

**Complete-case analysis:** only use units that have data on all variables

```
cces_2020 |>  
  drop_na() |>  
  summarize(mean(turnout_self)) |>  
  pull()
```

```
## [1] 0.999
```

(also called **listwise deletion**)

# is.na( ) to detect missingness

Trying to detect missingness with == doesn't work:

```
c(5, 6, NA, 0) == NA
```

```
## [1] NA NA NA NA
```

Use is.na() instead:

```
is.na(c(5, 6, NA, 0))
```

```
## [1] FALSE FALSE TRUE FALSE
```

Can use sum( ) or mean( ) on this to get number/proportion missing:

```
sum(is.na(c(5, 6, NA, 0)))
```

```
## [1] 1
```

# Nonresponse bias

Nonresponse can create bias if lower turnout  $\Rightarrow$  more non-response:

```
cces_2020 |>  
  group_by(pid3) |>  
  summarize(  
    mean_turnout = mean(turnout_self, na.rm = TRUE),  
    missing_turnout = mean(is.na(turnout_self))  
)
```

```
## # A tibble: 5 x 3  
##   pid3      mean_turnout missing_turnout  
##   <fct>        <dbl>          <dbl>  
## 1 Democrat     0.963         0.0280  
## 2 Republican   0.953         0.0403  
## 3 Independent  0.924         0.0718  
## 4 Other        0.957         0.0709  
## 5 Not sure     0.630         0.431
```

# 3/ Proportion tables

# Review of getting counts

First, let's review how to get counts:

```
cces_2020 |>  
  group_by(pres_vote) |>  
  summarize(n = n())
```

```
## # A tibble: 7 x 2  
##   pres_vote                 n  
##   <fct>                  <int>  
## 1 Joe Biden (Democrat)    26188  
## 2 Donald J. Trump (Republican) 17702  
## 3 Other                   1458  
## 4 I did not vote in this race  100  
## 5 I did not vote            13  
## 6 Not sure                 190  
## 7 <NA>                      5900
```

# First attempt to create proportions

```
cces_2020 |>  
  group_by(pres_vote) |>  
  summarize(prop = n() / sum(n()))
```

```
## # A tibble: 7 x 2  
##   pres_vote          prop  
##   <fct>            <dbl>  
## 1 Joe Biden (Democrat)     1  
## 2 Donald J. Trump (Republican) 1  
## 3 Other                 1  
## 4 I did not vote in this race 1  
## 5 I did not vote           1  
## 6 Not sure                1  
## 7 <NA>                   1
```

Inside `summarize()` all operations are done within groups!

# Mutate after summarizing

```
cces_2020 |>
  group_by(pres_vote) |>
  summarize(n = n()) |>
  mutate(prop = n / sum(n))

## # A tibble: 7 x 3
##   pres_vote                 n     prop
##   <fct>              <int>    <dbl>
## 1 Joe Biden (Democrat) 26188  0.508
## 2 Donald J. Trump (Republican) 17702  0.343
## 3 Other                  1458  0.0283
## 4 I did not vote in this race 100  0.00194
## 5 I did not vote          13  0.000252
## 6 Not sure                190  0.00369
## 7 <NA>                   5900  0.114
```

Grouping is silently dropped after summarize()

# Multiple grouping variables

What happens with multiple grouping variables

```
cces_2020 |>  
  filter(pres_vote %in% c("Joe Biden (Democrat)",  
                           "Donald J. Trump (Republican)")) |>  
  group_by(pid3, pres_vote) |>  
  summarize(n = n()) |>  
  mutate(prop = n / sum(n))
```

```
## # A tibble: 10 x 4
## # Groups: pid3 [5]
##   pid3      pres_vote          n    prop
##   <fct>     <fct>        <int>  <dbl>
## 1 Democrat  Joe Biden (Democrat) 17649  0.968
## 2 Democrat  Donald J. Trump (Republican) 581  0.0319
## 3 Republican Joe Biden (Democrat) 856  0.0712
## 4 Republican Donald J. Trump (Republican) 11164  0.929
## 5 Independent Joe Biden (Democrat) 6601  0.571
## 6 Independent Donald J. Trump (Republican) 4951  0.429
## 7 Other     Joe Biden (Democrat) 735  0.487
## 8 Other     Donald J. Trump (Republican) 774  0.513
## 9 Not sure  Joe Biden (Democrat) 347  0.599
## 10 Not sure Donald J. Trump (Republican) 232  0.401
```

With multiple grouping variables, summarize( ) drops the last one.

# Dropping all groups

If we want the proportion of all rows, need to drop all groups.

```
cces_2020 |>
  filter(pres_vote %in% c("Joe Biden (Democrat)",
                           "Donald J. Trump (Republican)")) |>
  group_by(pid3, pres_vote) |>
  summarize(n = n(), .groups = "drop") |>
  mutate(prop = n / sum(n))
```

```
## # A tibble: 10 x 4
##   pid3     pres_vote          n     prop
##   <fct>    <fct>      <int>    <dbl>
## 1 Democrat Joe Biden (Democrat) 17649  0.402
## 2 Democrat Donald J. Trump (Republican) 581  0.0132
## 3 Republican Joe Biden (Democrat) 856  0.0195
## 4 Republican Donald J. Trump (Republican) 11164  0.254
## 5 Independent Joe Biden (Democrat) 6601  0.150
## 6 Independent Donald J. Trump (Republican) 4951  0.113
## 7 Other     Joe Biden (Democrat) 735  0.0167
## 8 Other     Donald J. Trump (Republican) 774  0.0176
## 9 Not sure  Joe Biden (Democrat) 347  0.00791
## 10 Not sure Donald J. Trump (Republican) 232  0.00529
```

# Gov 50: 9. Survey Sampling

Matthew Blackwell

Harvard University

# Roadmap

1. Proportion tables
2. Measurement

# 1/ Proportion tables

# CCES Data

```
library(gov50data)
cces_2020

## # A tibble: 51,551 x 6
##   gender race  educ          pid3  turno~1 pres_~2
##   <fct>  <fct> <fct>        <fct>    <dbl> <fct>
## 1 Male   White 2-year      Republ~     1 Donald~
## 2 Female White Post-grad  Democr~     NA <NA>
## 3 Female White 4-year     Indepe~     1 Joe Bi~
## 4 Female White 4-year     Democr~     1 Joe Bi~
## 5 Male   White 4-year     Indepe~     1 Other
## 6 Male   White Some college Republ~     1 Donald~
## 7 Male   Black Some college Not su~     NA <NA>
## 8 Female White Some college Indepe~     1 Donald~
## 9 Female White High school graduate Republ~     1 Donald~
## 10 Female White 4-year    Democr~     1 Joe Bi~
## # ... with 51,541 more rows, and abbreviated variable names
## #   1: turnout_self, 2: pres_vote
```

# Mutate after summarizing

```
cces_2020 |>  
  group_by(pres_vote) |>  
  summarize(n = n()) |>  
  mutate(prop = n / sum(n))
```

```
## # A tibble: 7 x 3  
##   pres_vote                 n     prop  
##   <fct>                  <int>    <dbl>  
## 1 Joe Biden (Democrat)    26188  0.508  
## 2 Donald J. Trump (Republican) 17702  0.343  
## 3 Other                   1458  0.0283  
## 4 I did not vote in this race  100  0.00194  
## 5 I did not vote            13  0.000252  
## 6 Not sure                 190  0.00369  
## 7 <NA>                      5900  0.114
```

# Another approach

```
cces_2020 |>  
  group_by(pres_vote) |>  
  summarize(prop = n() / nrow(cces_2020))
```

```
## # A tibble: 7 x 2  
##   pres_vote                      prop  
##   <fct>                          <dbl>  
## 1 Joe Biden (Democrat)          0.508  
## 2 Donald J. Trump (Republican)  0.343  
## 3 Other                          0.0283  
## 4 I did not vote in this race  0.00194  
## 5 I did not vote                0.000252  
## 6 Not sure                       0.00369  
## 7 <NA>                           0.114
```

Doesn't work if you have filtered the data in any way during the pipe

# Multiple grouping variables

What happens with multiple grouping variables

```
vote_by_party <- cces_2020 |>
  filter(pres_vote %in% c("Joe Biden (Democrat)",
                          "Donald J. Trump (Republican)")) |>
  mutate(pres_vote = if_else(pres_vote == "Joe Biden (Democrat)",
                            "Biden", "Trump")) |>
  group_by(pid3, pres_vote) |>
  summarize(n = n()) |>
  mutate(prop = n / sum(n)) |>
  select(-n)

vote_by_party
```

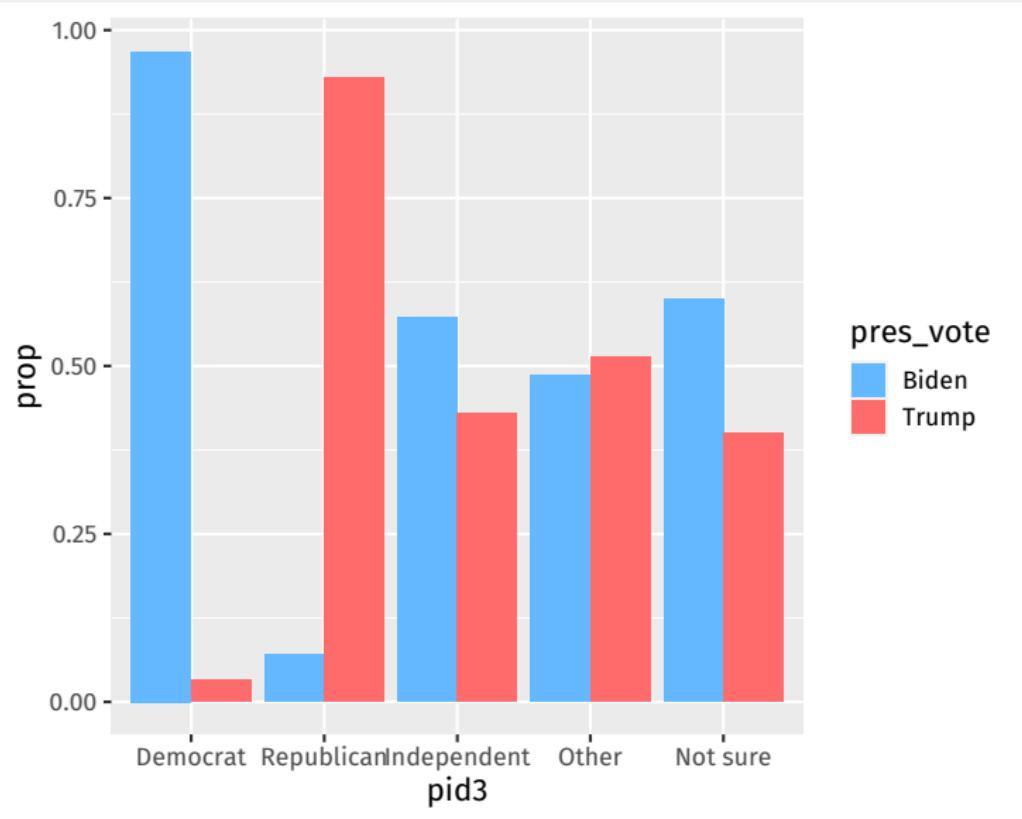
```
## # A tibble: 10 x 3
## # Groups: pid3 [5]
##   pid3      pres_vote    prop
##   <fct>      <chr>     <dbl>
## 1 Democrat    Biden    0.968
## 2 Democrat    Trump    0.0319
## 3 Republican Biden    0.0712
## 4 Republican Trump    0.929
## 5 Independent Biden    0.571
## 6 Independent Trump    0.429
## 7 Other        Biden    0.487
## 8 Other        Trump    0.513
## 9 Not sure    Biden    0.599
## 10 Not sure   Trump    0.401
```

With multiple grouping variables, summarize( ) drops the last one.

# Visualizing the cross-tab

We can visualize this using the `fill` aesthetic and `position="dodge"`:

```
ggplot(vote_by_party,  
       aes(x = pid3, y = prop, fill = pres_vote)) +  
  geom_col(position = "dodge") +  
  scale_fill_manual(values = c(Biden = "steelblue1", Trump = "indianred1"))
```



# Pivoting to create cross-tab

```
cces_2020 |>
  filter(pres_vote %in% c("Joe Biden (Democrat)",
                           "Donald J. Trump (Republican)")) |>
  mutate(pres_vote = if_else(pres_vote == "Joe Biden (Democrat)",
                             "Biden", "Trump")) |>
  group_by(pid3, pres_vote) |>
  summarize(n = n()) |>
  mutate(prop = n / sum(n)) |>
  select(-n) |>
  pivot_wider(
    names_from = pid3,
    values_from = prop
  )
```

```
## # A tibble: 2 x 6
##   pres_vote Democrat Republican Independent Other `Not sure`
##   <chr>      <dbl>       <dbl>        <dbl> <dbl>       <dbl>
## 1 Biden      0.968       0.0712       0.571 0.487      0.599
## 2 Trump      0.0319      0.929        0.429 0.513      0.401
```

# What if we want row proportions?

Switch the grouping variables to switch denominator:

```
cces_2020 |>
  filter(pres_vote %in% c("Joe Biden (Democrat)",
                           "Donald J. Trump (Republican)")) |>
  mutate(pres_vote = if_else(pres_vote == "Joe Biden (Democrat)",
                             "Biden", "Trump")) |>
  group_by(pres_vote, pid3) |>
  summarize(n = n()) |>
  mutate(prop = n / sum(n)) |>
  select(-n) |>
  pivot_wider(
    names_from = pid3,
    values_from = prop
  )
```

```
## # A tibble: 2 x 6
## # Groups: pres_vote [2]
##   pres_vote Democrat Republican Independent Other Not sur~1
##   <chr>      <dbl>      <dbl>      <dbl>    <dbl>    <dbl>
## 1 Biden       0.674     0.0327     0.252  0.0281   0.0133
## 2 Trump       0.0328     0.631      0.280  0.0437   0.0131
## # ... with abbreviated variable name 1: `Not sure`
```

# Proportion of all observations

If we want the proportion of all rows, drop all groups

```
cces_2020 |>
  filter(pres_vote %in% c("Joe Biden (Democrat)",
                           "Donald J. Trump (Republican)")) |>
  mutate(pres_vote = if_else(pres_vote == "Joe Biden (Democrat)",
                             "Biden", "Trump")) |>
  group_by(pid3, pres_vote) |>
  summarize(n = n(), .groups = "drop") |>
  mutate(prop = n / sum(n)) |>
  select(-n) |>
  pivot_wider(
    names_from = pid3,
    values_from = prop
  )
```

```
## # A tibble: 2 x 6
##   pres_vote Democrat Republican Independent Other Not sur~1
##   <chr>      <dbl>       <dbl>       <dbl>    <dbl>    <dbl>
## 1 Biden      0.402       0.0195     0.150  0.0167  0.00791
## 2 Trump      0.0132      0.254      0.113  0.0176  0.00529
## # ... with abbreviated variable name 1: `Not sure`
```

# 2/ Measurement

# Where does data come from?

- Social science is about developing and testing **causal theories**:
  - Does minimum wage change levels of employment?
  - Does outgroup contact influence views on immigration?
- Theories are made up of **concepts**:
  - Minimum wage, level of employment, outgroup contact, views on immigration.
  - We took these for granted when talking about causality.
- Need **operational definition** to concretely measure these concepts

# Concepts vary in how observable they are

Kinds of measurement arranged by how direct we can measure them:



## Observable in the world

- Minimum wage laws
- Sensor measurements
- Election results

## Observable by survey

- Age of a person
- Employment status
- Presidential approval

## Not directly observable

- A person's ideology
- Levels of democracy
- Extent of gerrymandering

# Example

- Concept: presidential approval.
- Conceptual definition:
  - Extent to which US adults support the actions and policies of the current US president.
- Operational definition:
  - “On a scale from 1 to 5, where 1 is least supportive and 5 is more supportive, how much would you say you support the job that Joe Biden is doing as president?”

# Measurement error

**Table 1**

Response to citizenship question across two-waves of CCES panel.

| Response in 2010 | Response in 2012 | Number of respondents | Percentage |
|------------------|------------------|-----------------------|------------|
| Citizen          | Citizen          | 18,737                | 99.25      |
| Citizen          | Non-Citizen      | 20                    | 0.11       |
| Non-Citizen      | Citizen          | 36                    | 0.19       |
| Non-Citizen      | Non-Citizen      | 85                    | 0.45       |

- **Measurement error:** chance variation in our measurements.
  - individual measurement = exact value + chance error
  - chance errors tend to cancel out when we take averages.
  - why? often data entry errors or faulty memories.

# Bias

VZW Wi-Fi 18:23 33% gop.com

Official Presidential Job Performance Poll

1. How would you rate President Trump's job performance so far?

Great  
 Good  
 Okay  
 Other

2. (Optional) Please explain why you selected your response.

[Large empty rectangular box for writing]

- **Bias:** systematic errors for all units in the same direction.
- individual measurement = exact value + bias + chance error.
- “What did you eat yesterday?”  
~~ underreporting

# 1936 Literary Digest Poll

## The Literary Digest

NEW YORK

OCTOBER 31, 1936

### *Topics of the day*

**LANDON, 1,293,669; ROOSEVELT, 972,897**

Final Returns in The Digest's Poll of Ten Million Voters

Well, the great battle of the ballots in the Poll of ten million voters, scattered throughout the forty-eight States of the

litan National Committee purchased The LITERARY DIGEST? And all types and varieties, including: "Have the Jews purchased

returned and let the people of the Nation draw their conclusions as to our accuracy. So far, we have been right in every Poll. Will we be right in the current Poll? That, as Mrs. Roosevelt said concerning the President's reelection, is in the 'lap of the gods.'

"We never make any claims before election but we respectfully refer you to the

*opinion of one of the most noted citizens*

- Literary Digest predicted elections using mail-in polls.
- Source of addresses: automobile registrations, phone books, etc.
- In 1936, sent out 10 million ballots, over 2.3 million returned.
- George Gallup used only 50,000 respondents.

# Poll fail



|                 | FDR's Vote Share |
|-----------------|------------------|
| Literary Digest | 43%              |
| George Gallup   | 56%              |
| Actual Outcome  | 62%              |

- **Selection bias:** ballots skewed toward the wealthy (with cars, phones)
  - Only 1 in 4 households had a phone in 1936.
- **Nonresponse bias:** respondents differ from nonrespondents.
  - ↗ when selection procedure is biased, adding more units won't help!

# 1948 Election



# The Polling Disaster

|          | Truman | Dewey | Thurmond | Wallace |
|----------|--------|-------|----------|---------|
| Crossley | 45     | 50    | 2        | 3       |
| Gallup   | 44     | 50    | 2        | 4       |
| Roper    | 38     | 53    | 5        | 4       |
| Actual   | 50     | 45    | 3        | 2       |

- **Quota sampling:** fixed quota of certain respondents for each interviewer
  - If black women make up 5% of the population, stop interviewing them once they make up 5% of your sample.
- Sample resembles the population on these characteristics
- Potential unobserved confounding ↵ **selection bias**
- Republicans easier to find within quotas (phones, listed addresses)

# Sample surveys

- **Probability sampling** to ensure representativeness
  - Definition: every unit in the population has a known, non-zero probability of being selected into sample.
- **Simple random sampling:** every unit has an **equal** selection probability.
- Random digit dialing:
  - Take a particular area code + exchange: 617-495-XXXX.
  - Randomly choose each digit in XXXX to call a particular phone.
  - Every phone in America has an equal chance of being included in sample.

# Sampling lingo

- **Target population:** set of people we want to learn about.
  - Ex: people who will vote in the next election.
- **Sampling frame:** list of people from which we will actually sample.
  - Frame bias: list of registered voters (frame) might include nonvoters!
- **Sample:** set of people contacted.
- **Respondents:** subset of sample that actually responds to the survey.
  - Unit non-response: sample  $\neq$  respondents.
  - Not everyone picks up their phone.
- **Completed items:** subset of questions that respondents answer.
  - Item non-response: refusing to disclose their vote preference.

# Difficulties of sampling

- Problems of telephone survey
  - Cell phones (double counting for the wealthy)
  - Caller ID screening (unit non-response)
  - Response rates down to 9%!
- An alternative: Internet surveys
  - Opt-in panels, respondent-driven sampling ↽ **non-probability sampling**
  - Cheaper, but non-representative
  - Digital divide: rich vs. poor, young vs. old
  - Correct for potential sampling bias via statistical methods.

# Gov 50: 10. Summarizing Bivariate Relationships

Matthew Blackwell

Harvard University

# Roadmap

1. Z-scores and standardization
2. Correlation
3. Writing our own functions

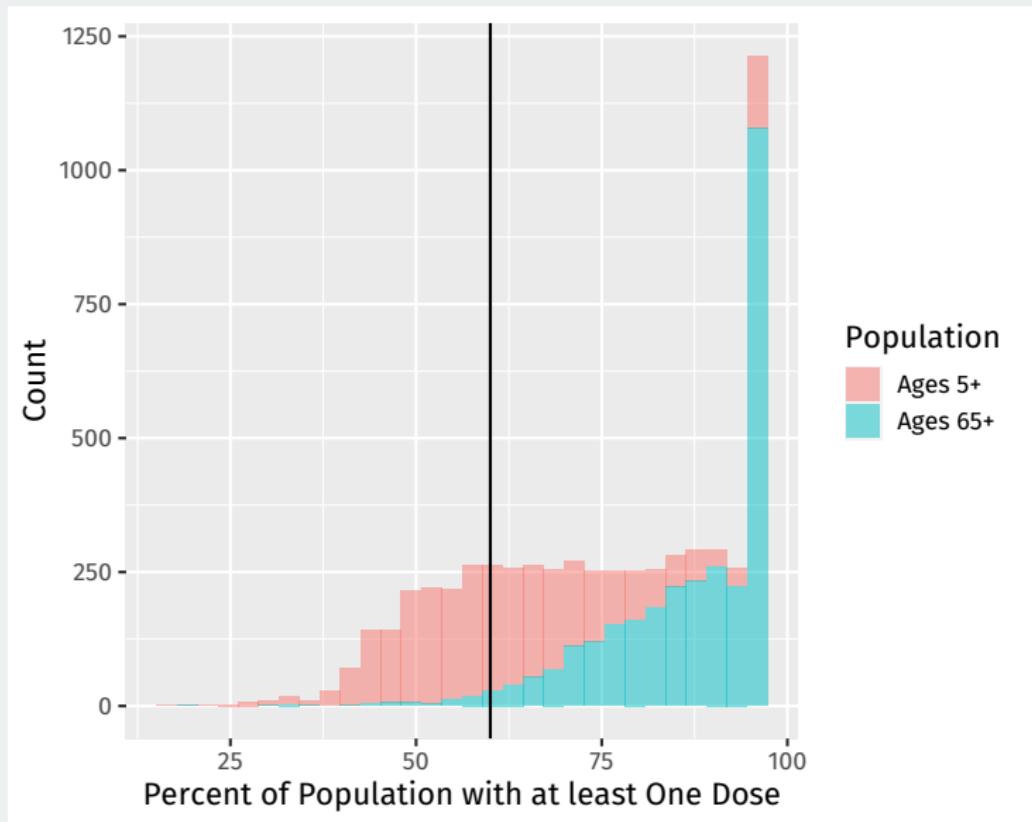
# **1/** Z-scores and standardization

# COVID vaccination rates and votes

```
library(tidyverse)
library(gov50data)
covid_votes

## # A tibble: 3,114 x 8
##   fips    county      state one_d~1 one_d~2 boost~3 dem_p~4
##   <chr>   <chr>       <chr>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 26039 Crawford Cou~ MI        55.7     77.3    31.2    43.8
## 2 40015 Caddo County OK        83.3     95.0    30.3    46.4
## 3 17007 Boone County IL        71.1     94.5    35.1    41.8
## 4 12055 Highlands Co~ FL       68.9     93.7    24.7    40.3
## 5 34029 Ocean County NJ        71.0     95.0    32.1    47.2
## 6 01067 Henry County AL       58.5     85.5    18.2    40.1
## 7 27037 Dakota County MN       81.0     95.0    49.5    46.9
## 8 27115 Pine County MN       56.5     85.0    31.7    47.0
## 9 51750 Radford city VA       41.5     73.8    1.79    46.4
## 10 22009 Avoyelles Pa~ LA      59.7     80.1    21.9    45.7
## # ... with 3,104 more rows, 1 more variable:
## #   dem_pct_2020 <dbl>, and abbreviated variable names
## #   1: one_dose_5plus_pct, 2: one_dose_65plus_pct,
## #   3: booster_5plus_pct, 4: dem_pct_2000
```

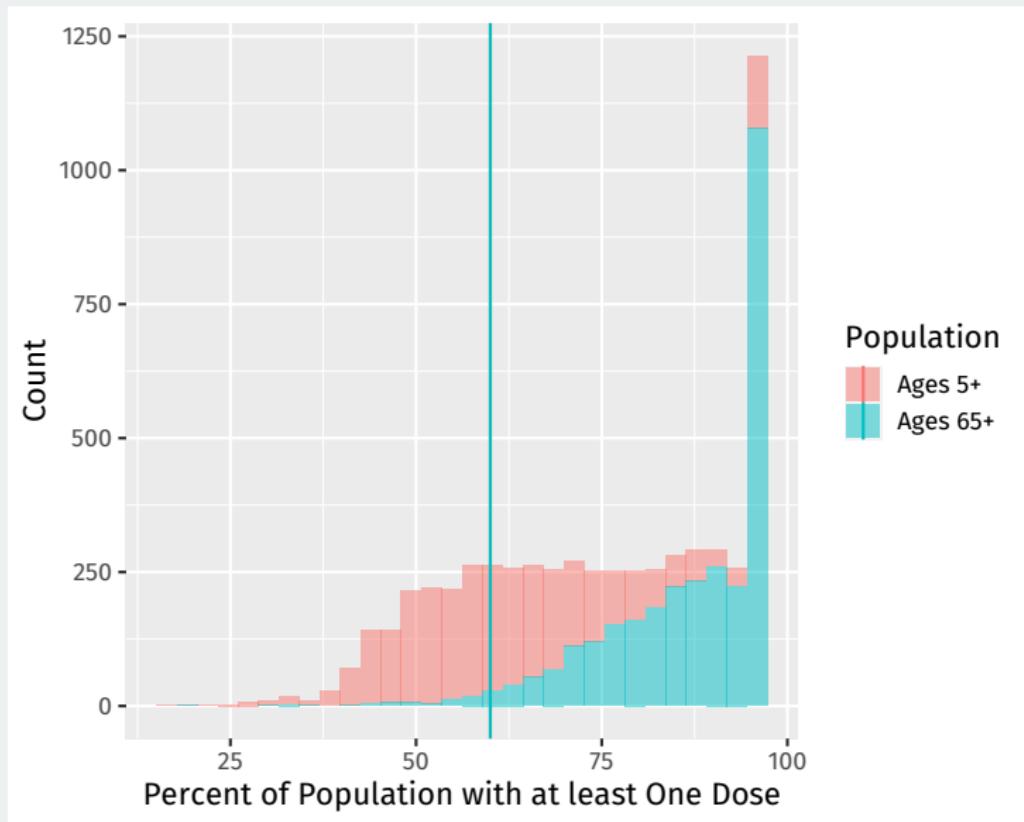
# Is 60% vaccinated a lot?



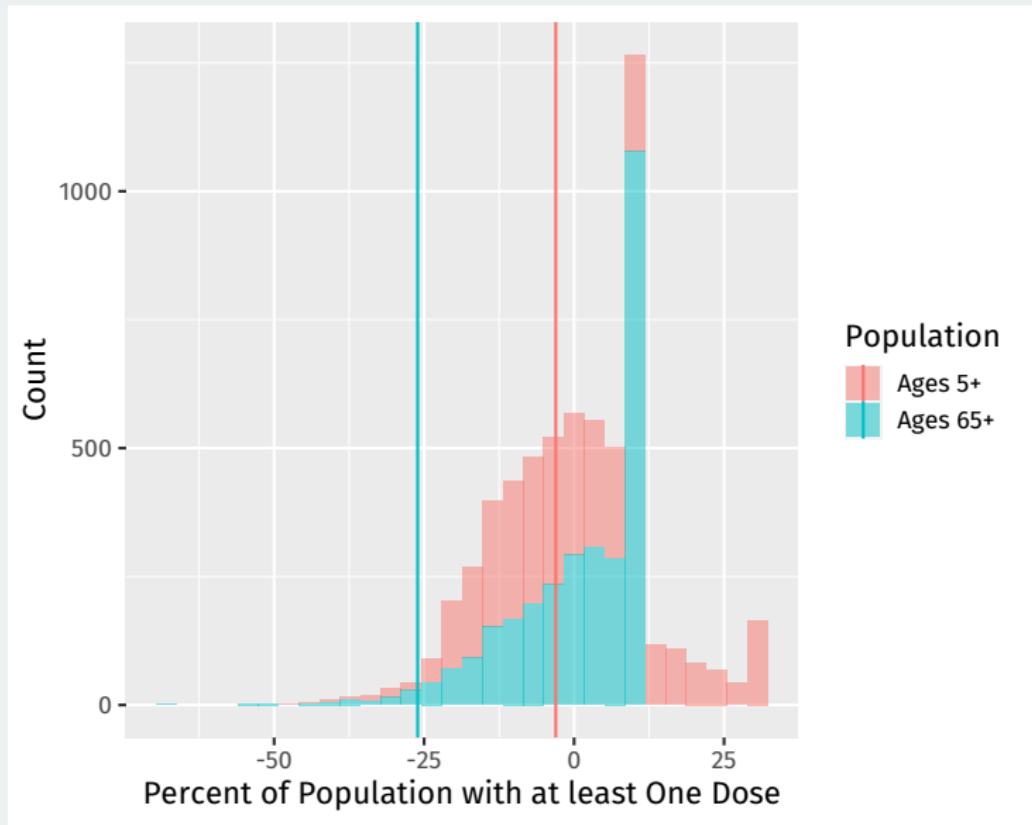
# How large is large?

- How large 60% vaccinated is depends on the distribution!
  - Clear to see from the histogram
  - Middling for the 5+ group, but very low for the 65+ group.
- Can we transform the values of our variables to be **common units**?
- Yes, with two transformations:
  - **Centering**: subtract the mean of the variable from each value.
  - **Scaling**: dividing deviations from the mean by the standard deviation.

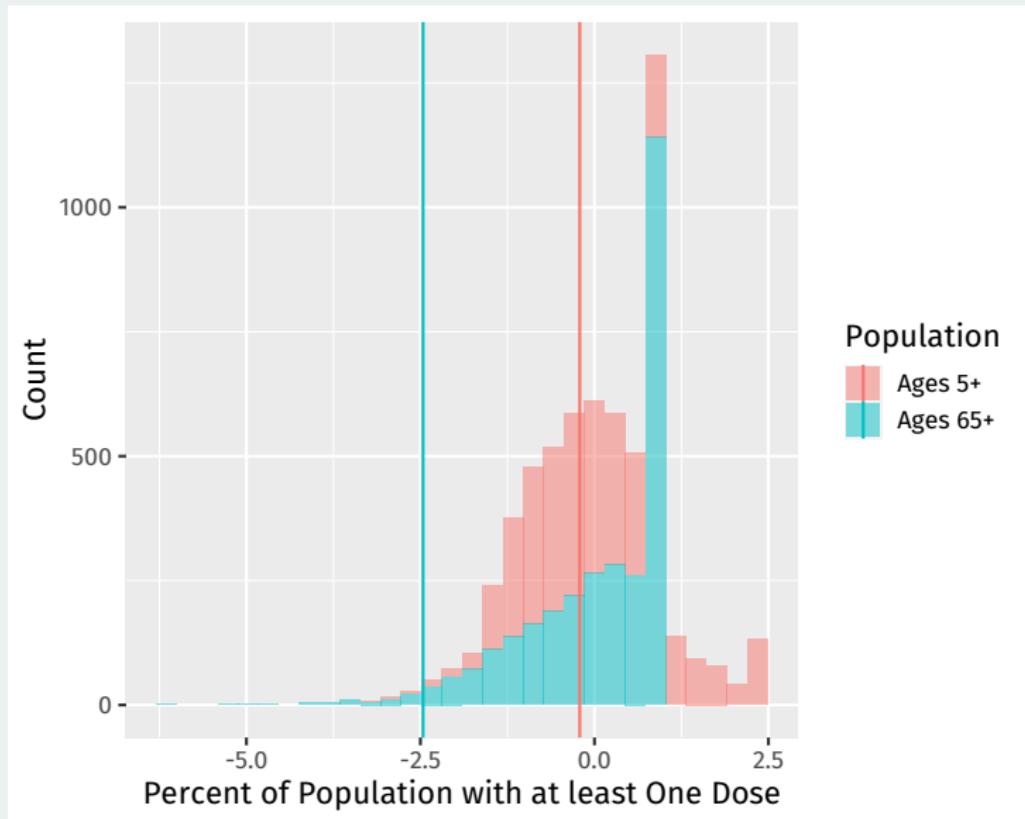
# Original distributions



# Centered distributions



# Centered and scaled distributions



# Z-scores

- Centering tells us immediately if a value is above or below the mean.
- Scaling tells us how many standard deviations away from the mean it is.
- Combine them with the **z-score** transformation:

$$\text{z-score of } x_i = \frac{x_i - \text{mean of } x}{\text{standard deviation of } x}$$

- Useful heuristic: data more than 3 SDs away from mean are rare.

# z-score example

```
covid_votes |>  
  mutate(one_dose_centered = one_dose_5plus_pct -  
         mean(one_dose_5plus_pct, na.rm = TRUE)) |>  
  select(fips:state, one_dose_5plus_pct, one_dose_centered)  
  
## # A tibble: 3,114 x 5  
##   fips    county      state one_dose_5plus_pct one_dose_centered  
##   <chr>   <chr>       <chr>          <dbl>            <dbl>  
## 1 26039 Crawford County MI             55.7            -7.35  
## 2 40015 Caddo County   OK            83.3             20.2  
## 3 17007 Boone County   IL            71.1             8.05  
## 4 12055 Highlands County FL            68.9             5.85  
## 5 34029 Ocean County  NJ            71              7.95  
## 6 01067 Henry County   AL            58.5            -4.55  
## 7 27037 Dakota County  MN            81              17.9  
## 8 27115 Pine County   MN            56.5            -6.55  
## 9 51750 Radford city  VA            41.5            -21.6  
## 10 22009 Avoyelles Parish LA           59.7            -3.35  
## # ... with 3,104 more rows, and abbreviated variable name  
## #   1: one_dose_centered
```

# z-score example

```
covid_votes |>
  mutate(
    one_dose_z =
      (one_dose_5plus_pct - mean(one_dose_5plus_pct, na.rm = TRUE)) /
      sd(one_dose_5plus_pct, na.rm = TRUE)) |>
  select(fips:state, one_dose_5plus_pct, one_dose_z)
```

```
## # A tibble: 3,114 x 5
##   fips   county       state one_dose_5plus_pct one_dose_z
##   <chr> <chr>        <chr>          <dbl>      <dbl>
## 1 26039 Crawford County MI            55.7     -0.508
## 2 40015 Caddo County   OK            83.3      1.40 
## 3 17007 Boone County  IL            71.1      0.556
## 4 12055 Highlands County FL            68.9      0.404
## 5 34029 Ocean County  NJ            71         0.549
## 6 01067 Henry County  AL            58.5     -0.314
## 7 27037 Dakota County MN            81         1.24 
## 8 27115 Pine County   MN            56.5     -0.452
## 9 51750 Radford city  VA            41.5     -1.49 
## 10 22009 Avoyelles Parish LA           59.7     -0.231
## # ... with 3,104 more rows, and abbreviated variable name
## #   1: one_dose_z
```

# 2/ Correlation

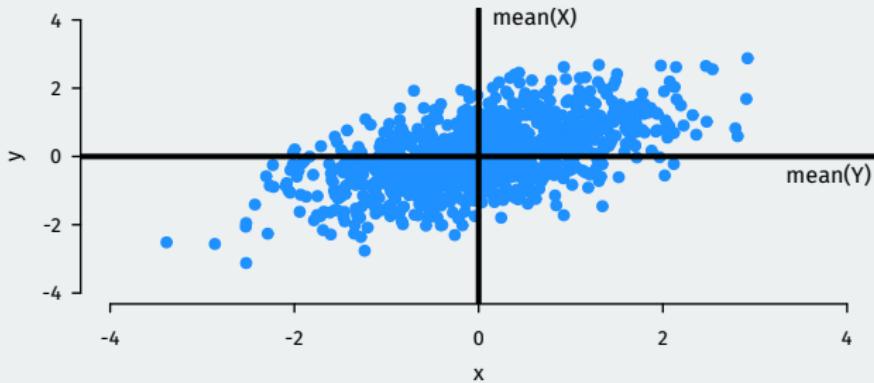
# Correlation

- How do variables move together on average?
- When  $x_i$  is big, what is  $y_i$  likely to be?
  - Positive correlation: when  $x_i$  is big,  $y_i$  is also big
  - Negative correlation: when  $x_i$  is big,  $y_i$  is small
  - High magnitude of correlation: data cluster tightly around a line.
- The technical definition of the **correlation coefficient**:

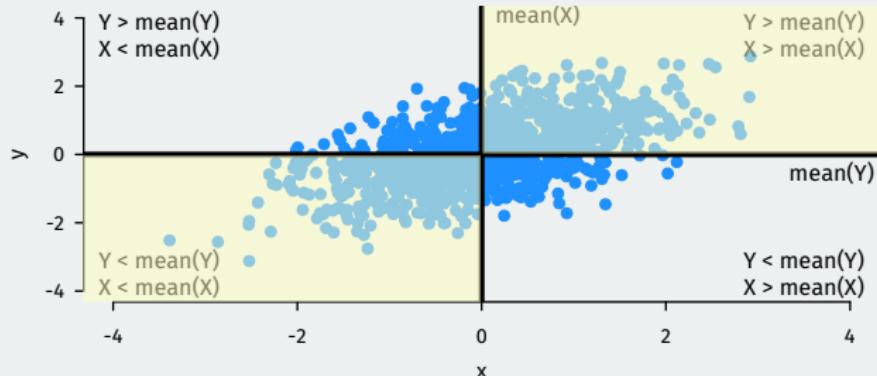
$$\frac{1}{n-1} \sum_{i=1}^n [(\text{z-score for } x_i) \times (\text{z-score for } y_i)]$$

- Interpretation:
  - Correlation is between -1 and 1
  - Correlation of 0 means no linear association.
  - Positive correlations  $\rightsquigarrow$  positive associations.
  - Negative correlations  $\rightsquigarrow$  negative associations.
  - Closer to -1 or 1 means stronger association.

# Correlation intuition

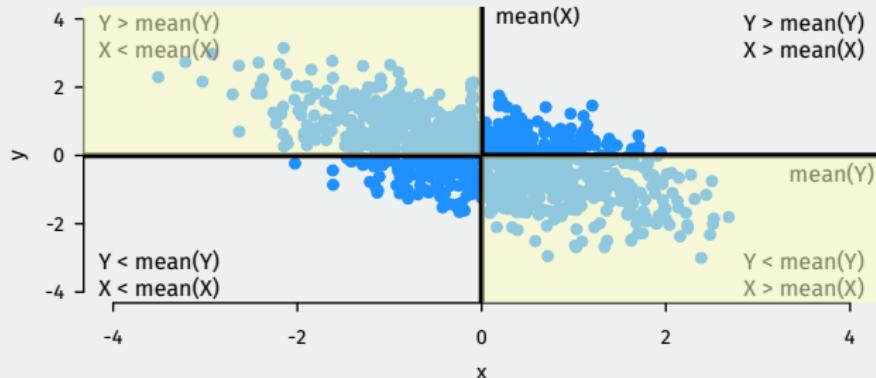


# Correlation intuition



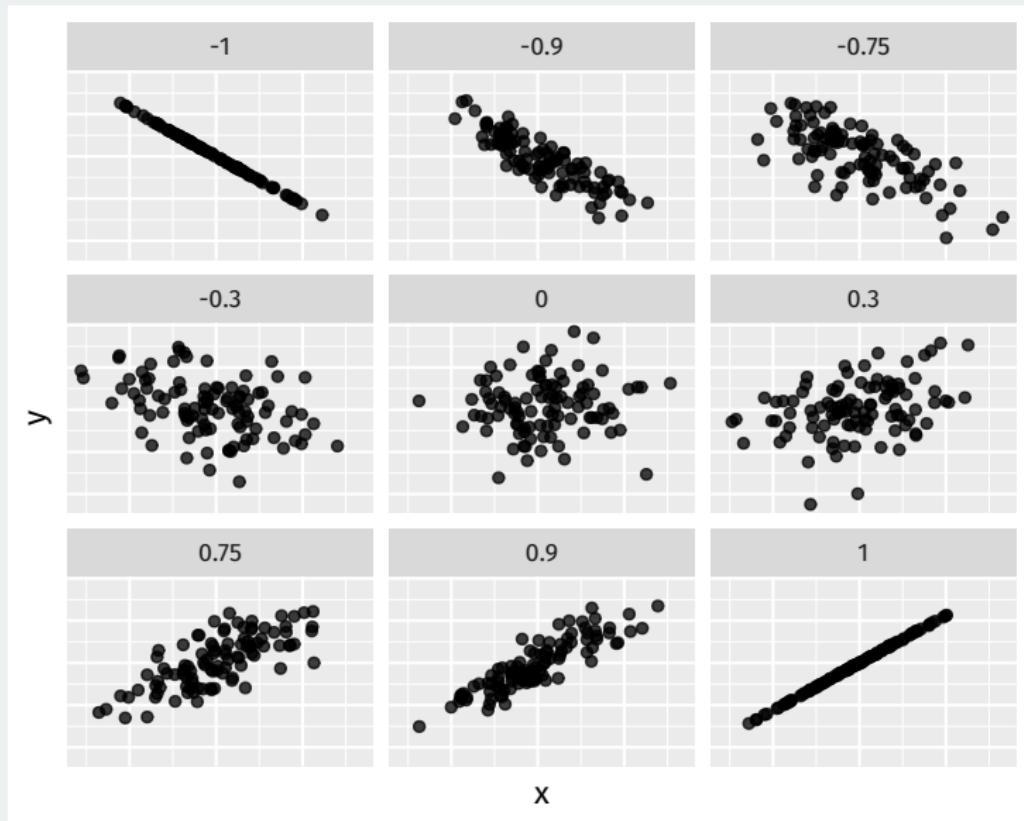
- Large values of  $X$  tend to occur with large values of  $Y$ :
  - $(\text{z-score for } x_i) \times (\text{z-score for } y_i) = (\text{pos. num.}) \times (\text{pos. num.}) = +$
- Small values of  $X$  tend to occur with small values of  $Y$ :
  - $(\text{z-score for } x_i) \times (\text{z-score for } y_i) = (\text{neg. num.}) \times (\text{neg. num.}) = +$
- If these dominate  $\rightsquigarrow$  positive correlation.

# Correlation intuition



- Large values of  $X$  tend to occur with small values of  $Y$ :
  - $(\text{z-score for } x_i) \times (\text{z-score for } y_i) = (\text{pos. num.}) \times (\text{neg. num.}) = -$
- Small values of  $X$  tend to occur with large values of  $Y$ :
  - $(\text{z-score for } x_i) \times (\text{z-score for } y_i) = (\text{neg. num.}) \times (\text{pos. num.}) = -$
- If these dominate  $\rightsquigarrow$  negative correlation.

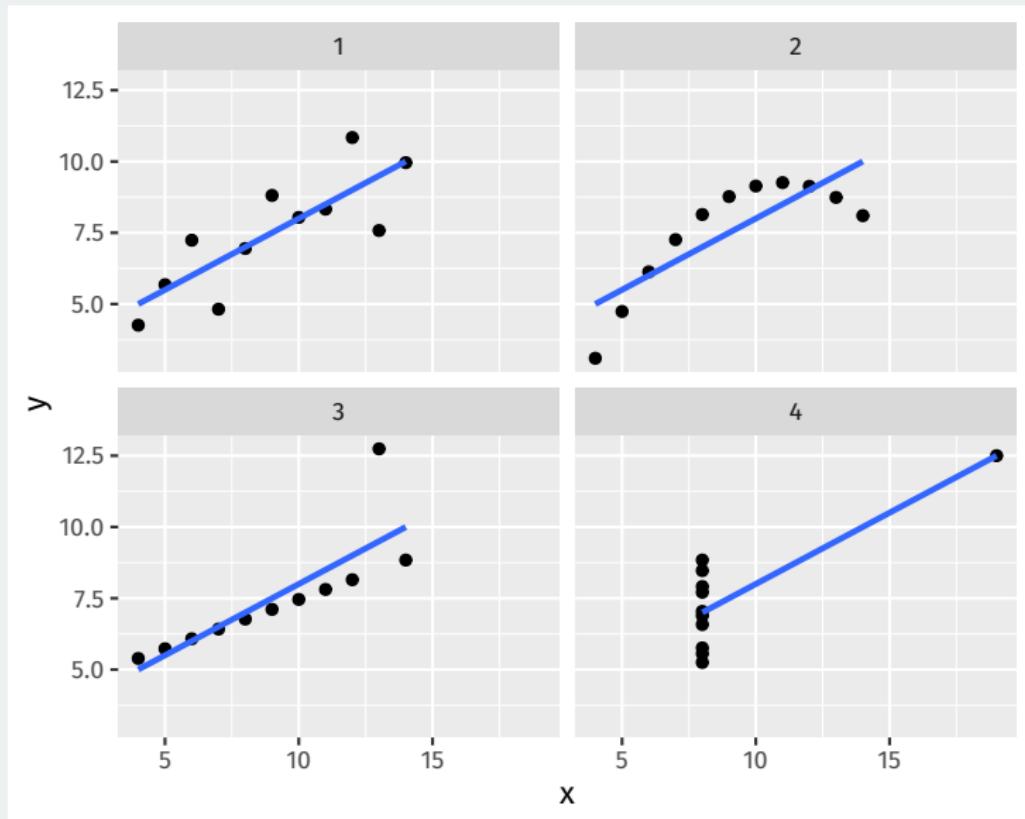
# Correlation examples



# Properties of correlation coefficient

- Correlation measures **linear** association.
- Order doesn't matter:  $\text{cor}(x, y) = \text{cor}(y, x)$
- Not affected by changes of scale:
  - $\text{cor}(x, y) = \text{cor}(ax+b, cy+d)$
  - Celsius vs. Fahrenheit; dollars vs. pesos; cm vs. in.

# All 4 relationships have 0.816 correlation



# Correlation in R

Use the `cor()` function:

```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2020)
```

```
## [1] NA
```

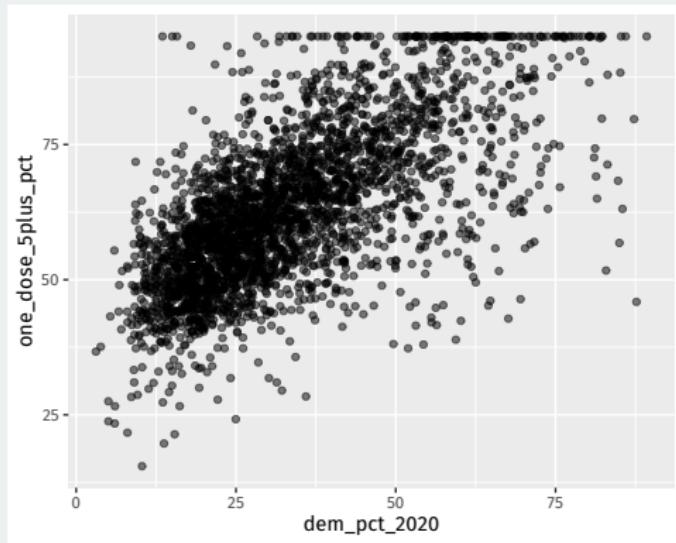
Missing values: set the `use = "pairwise"` → available case analysis

```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2020,  
    use = "pairwise")
```

```
## [1] 0.666
```

# Comparing correlations

```
covid_votes |>  
  ggplot(aes(x = dem_pct_2020, y = one_dose_5plus_pct)) +  
  geom_point(alpha = 0.5)
```

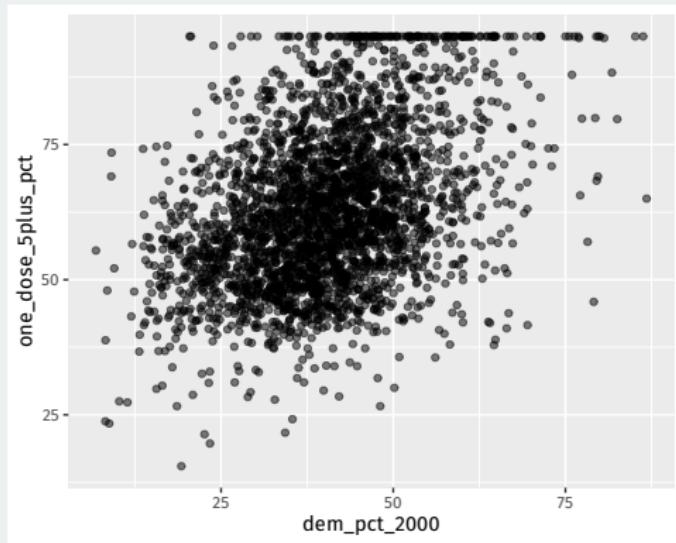


```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2020,  
  use = "pairwise")
```

```
## [1] 0.666
```

# Comparing correlations

```
covid_votes |>  
  ggplot(aes(x = dem_pct_2000, y = one_dose_5plus_pct)) +  
  geom_point(alpha = 0.5)
```

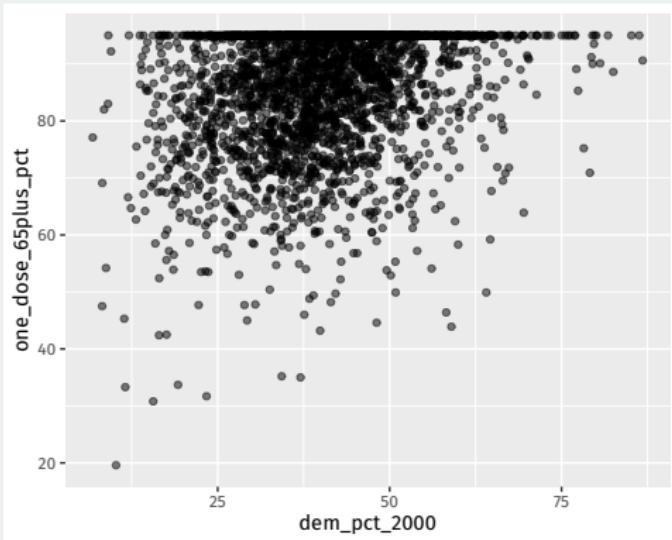


```
cor(covid_votes$one_dose_5plus_pct, covid_votes$dem_pct_2000,  
    use = "pairwise")
```

```
## [1] 0.394
```

# Comparing correlations

```
covid_votes |>  
  ggplot(aes(x = dem_pct_2000, y = one_dose_65plus_pct)) +  
  geom_point(alpha = 0.5)
```



```
cor(covid_votes$one_dose_65plus_pct, covid_votes$dem_pct_2000,  
  use = "pairwise")
```

```
## [1] 0.263
```

# **3/** Writing our own functions

# Why write functions?

Copy-pasting code tedious and prone to failure:

```
covid_votes |>  
  mutate(  
    one_dose_5p_z =  
      (one_dose_5plus_pct - mean(one_dose_5plus_pct, na.rm = TRUE)) /  
      sd(one_dose_5plus_pct, na.rm = TRUE),  
    one_dose_65p_z =  
      (one_dose_65plus_pct - mean(one_dose_65plus_pct, na.rm = TRUE)) /  
      sd(one_dose_65plus_pct, na.rm = TRUE),  
    booster_z =  
      (booster_5plus_pct - mean(booster_5plus_pct, na.rm = TRUE)) /  
      sd(booster_5plus_pct, na.rm = TRUE),  
    dem_pct_2000_z =  
      (dem_pct_2000 - mean(dem_pct_2000, na.rm = TRUE)) /  
      sd(dem_pct_2000, na.rm = TRUE),  
    dem_pct_2020_z =  
      (dem_pct_2020 - mean(dem_pct_2020, na.rm = TRUE)) /  
      sd(dem_pct_2020, na.rm = TRUE)  
)
```

# Writing a new function

Notice that all of the mutations follow the same template:

```
([] - mean([], na.rm = TRUE)) / sd([], na.rm = TRUE)
```

Only one thing varies: the column of data, represented with []

# Components of a function

We create functions like so:

```
name <- function(arguments) {  
  body  
}
```

Three components:

1. **Name:** the name of the function that we'll use to call it. Maybe `z_score?`
2. **Arguments:** things that we want to vary across calls of our function. We'll use `x`.
3. **Body:** the code that the function performs.

# Our first function

Convert our template to a function:

```
z_score <- function(x) {  
  (x - mean(x, na.rm = TRUE)) / sd(x, na.rm = TRUE)  
}
```

Check that it seems to work:

```
z_score(c(1,2, 3, 4, 5))
```

```
## [1] -1.265 -0.632  0.000  0.632  1.265
```

# Cleaning up our code

```
covid_votes |>  
  mutate(  
    one_dose_5p_z = z_score(one_dose_5plus_pct),  
    one_dose_65p_z = z_score(one_dose_65plus_pct),  
    booster_z = z_score(booster_5plus_pct),  
    dem_pct_2000_z = z_score(dem_pct_2000),  
    dem_pct_2020_z = z_score(dem_pct_2020)  
  )
```

# across() function

If we want to replace our variables with z-scores, we can use the `across()` function to perform many mutations at once:

```
covid_votes |>  
  mutate(across(one_dose_5plus_pct:dem_pct_2020, z_score))
```

```
## # A tibble: 3,114 x 8  
##   fips    county      state one_d~1 one_d~2 boost~3 dem_p~4  
##   <chr>   <chr>       <chr>    <dbl>    <dbl>    <dbl>    <dbl>  
## 1 26039 Crawford Cou~ MI     -0.508   -0.829   0.531   0.340  
## 2 40015 Caddo County OK     1.40     0.843   0.439   0.556  
## 3 17007 Boone County IL     0.556   0.795   0.927   0.163  
## 4 12055 Highlands Co~ FL     0.404   0.720   -0.135   0.0402  
## 5 34029 Ocean County NJ     0.549   0.843   0.623   0.624  
## 6 01067 Henry County AL     -0.314   -0.0545  -0.799   0.0255  
## 7 27037 Dakota County MN     1.24     0.843   2.40    0.598  
## 8 27115 Pine County MN     -0.452   -0.102   0.577   0.612  
## 9 51750 Radford city VA     -1.49    -1.16   -2.47   0.556  
## 10 22009 Avoyelles Pa~ LA    -0.231   -0.564  -0.424   0.501  
## # ... with 3,104 more rows, 1 more variable:  
## #   dem_pct_2020 <dbl>, and abbreviated variable names  
## #   1: one_dose_5plus_pct, 2: one_dose_65plus_pct,
```

# Alternative approach

We could also target all the numeric variables:

```
covid_votes |>  
  mutate(across(where(is.numeric), z_score))
```

```
## # A tibble: 3,114 x 8  
##   fips county      state one_d~1 one_d~2 boost~3 dem_p~4  
##   <chr> <chr>       <chr>    <dbl>    <dbl>    <dbl>    <dbl>  
## 1 26039 Crawford Cou~ MI     -0.508   -0.829    0.531   0.340  
## 2 40015 Caddo County OK      1.40     0.843    0.439   0.556  
## 3 17007 Boone County IL      0.556   0.795    0.927   0.163  
## 4 12055 Highlands Co~ FL     0.404   0.720    -0.135   0.0402  
## 5 34029 Ocean County NJ      0.549   0.843    0.623   0.624  
## 6 01067 Henry County AL     -0.314   -0.0545   -0.799   0.0255  
## 7 27037 Dakota County MN     1.24     0.843    2.40    0.598  
## 8 27115 Pine County MN     -0.452   -0.102    0.577   0.612  
## 9 51750 Radford city VA     -1.49    -1.16    -2.47   0.556  
## 10 22009 Avoyelles Pa~ LA    -0.231   -0.564   -0.424   0.501  
## # ... with 3,104 more rows, 1 more variable:  
## #   dem_pct_2020 <dbl>, and abbreviated variable names  
## #   1: one_dose_5plus_pct, 2: one_dose_65plus_pct,  
## #   3: booster_5plus_pct, 4: dem_pct_2000
```

# Alternative approach

We could also target only the first dose variables:

```
covid_votes |>  
  mutate(across(starts_with("one_dose"), z_score))
```

```
## # A tibble: 3,114 x 8  
##   fips    county      state one_d~1 one_d~2 boost~3 dem_p~4  
##   <chr>   <chr>       <chr>    <dbl>    <dbl>    <dbl>    <dbl>  
## 1 26039  Crawford Cou~ MI     -0.508   -0.829    31.2    43.8  
## 2 40015  Caddo County OK      1.40     0.843    30.3    46.4  
## 3 17007  Boone County IL      0.556    0.795    35.1    41.8  
## 4 12055  Highlands Co~ FL     0.404    0.720    24.7    40.3  
## 5 34029  Ocean County NJ      0.549    0.843    32.1    47.2  
## 6 01067  Henry County AL     -0.314   -0.0545   18.2    40.1  
## 7 27037  Dakota County MN     1.24     0.843    49.5    46.9  
## 8 27115  Pine County MN     -0.452   -0.102    31.7    47.0  
## 9 51750  Radford city VA     -1.49    -1.16     1.79    46.4  
## 10 22009 Avoyelles Pa~ LA     -0.231   -0.564    21.9    45.7  
## # ... with 3,104 more rows, 1 more variable:  
## #   dem_pct_2020 <dbl>, and abbreviated variable names  
## #   1: one_dose_5plus_pct, 2: one_dose_65plus_pct,  
## #   3: booster_5plus_pct, 4: dem_pct_2000
```

# Adding arguments to our function

What if we want to be able to control `na.rm` in the calls to `mean()` and `sd()` in our `z_score` function? Add an argument!

```
z_score2 <- function(x, na.rm = FALSE) {  
  (x - mean(x, na.rm = na.rm)) / sd(x, na.rm = na.rm)  
}
```

```
head(z_score2(covid_votes$one_dose_5plus_pct))
```

```
## [1] NA NA NA NA NA NA
```

```
head(z_score2(covid_votes$one_dose_5plus_pct, na.rm = TRUE))
```

```
## [1] -0.508  1.398  0.556  0.404  0.549 -0.314
```

# Gov 50: 11. Tidying and Joining Data

Matthew Blackwell

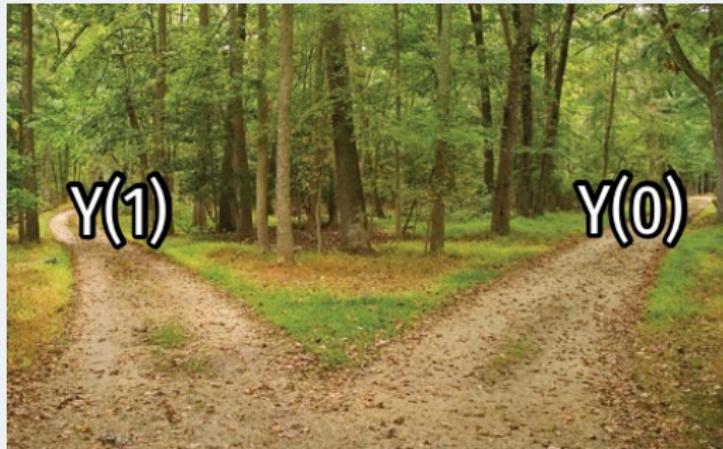
Harvard University

# Roadmap

1. Causality review
2. Pivoting data longer
3. Joining data sets

# 1/ Causality review

# Potential outcomes



Potential outcomes:

- $Y_i(1)$  is the value that the outcome would take if gave unit  $i$  **treatment** and changed nothing else about them.
- $Y_i(0)$  is the value that the outcome would take if gave unit  $i$  **no treatment** and changed nothing else about them.
- Not the **possible values** of the outcome

# COVID-19 vaccine trials



**Treatment:**  $T_i = 1$  if vaccinated,  $T_i = 0$  if not

**Outcome:**  $Y_i = 1$  if acquired COVID after 12 weeks,  $Y_i = 0$  if not

1. What are the potential outcomes  $Y_i(1)$  and  $Y_i(0)$ ?
2. Why not compare early volunteers for the vaccine to the overall population?

# 2/ Pivoting data longer

# Mortality data

```
library(tidyverse)
library(gov50data)
mortality
```

```
## # A tibble: 217 x 52
##   country      count~1 indic~2 `1972` `1973` `1974` `1975`
##   <chr>        <chr>    <chr>    <dbl>   <dbl>   <dbl>   <dbl>
## 1 Aruba        ABW     Mortal~    NA     NA     NA     NA
## 2 Afghanistan  AFG     Mortal~    291    285.   280.   274.
## 3 Angola       AGO     Mortal~    NA     NA     NA     NA
## 4 Albania      ALB     Mortal~    NA     NA     NA     NA
## 5 Andorra      AND     Mortal~    NA     NA     NA     NA
## 6 United Arab ~ ARE     Mortal~    80.1   72.6   65.7   59.4
## 7 Argentina    ARG     Mortal~    69.7   68.2   66.1   63.3
## 8 Armenia      ARM     Mortal~    NA     NA     NA     NA
## 9 American Sam~ ASM     Mortal~    NA     NA     NA     NA
## 10 Antigua and ~ ATG    Mortal~    26.9   25.1   23.5   22.1
## # ... with 207 more rows, 45 more variables: `1976` <dbl>,
## #   `1977` <dbl>, `1978` <dbl>, `1979` <dbl>, `1980` <dbl>,
## #   `1981` <dbl>, `1982` <dbl>, `1983` <dbl>, `1984` <dbl>,
## #   `1985` <dbl>, `1986` <dbl>, `1987` <dbl>, `1988` <dbl>,
## #   `1989` <dbl>, `1990` <dbl>, `1991` <dbl>, `1992` <dbl>,
```

# Pivoting longer

Mortality data in a “wide” format (years in columns).

We can convert this to country-year rows with `pivot_longer()`.

```
mydata |>  
  pivot_longer(  
    cols = <<variables to pivot>>,  
    names_to = <<new variable to put column names>>,  
    values_to = <<new variable to put column values>>  
  )
```

# Pivoting the mortality data

```
mortality |>
  select(-indicator) |>
  pivot_longer(
    cols = `1972`:`2020`,
    names_to = "year",
    values_to = "child_mortality"
  )
```

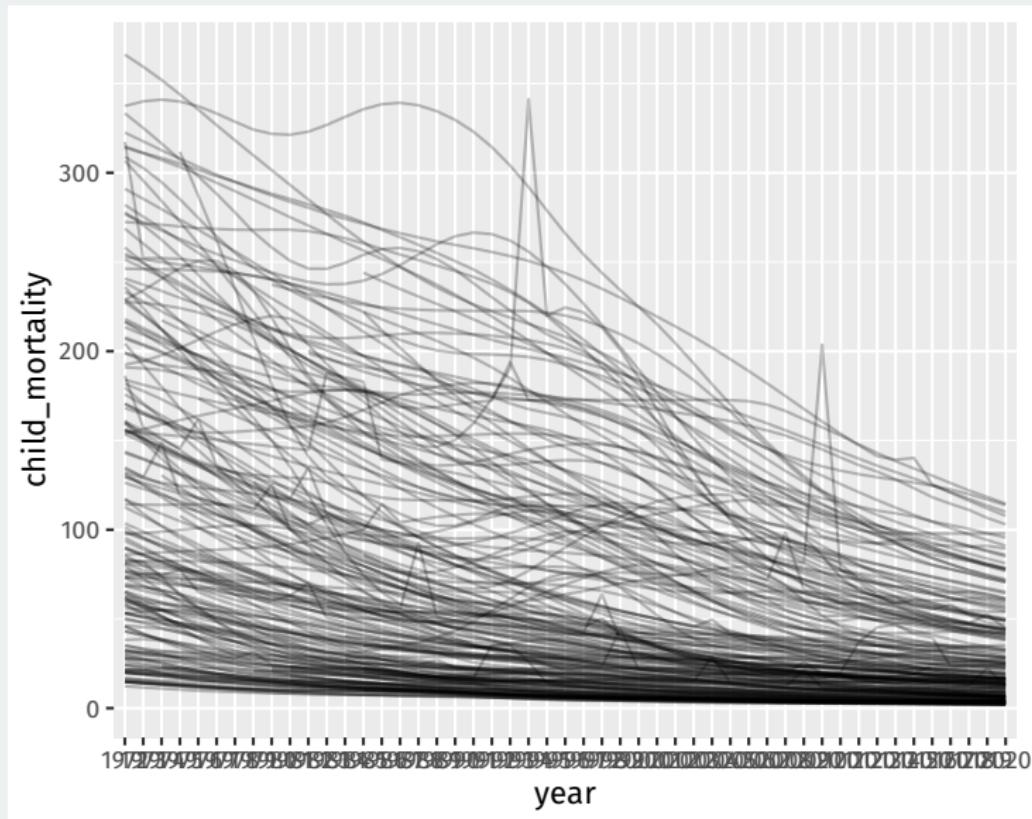
```
## # A tibble: 10,633 x 4
##   country country_code year child_mortality
##   <chr>    <chr>      <chr>        <dbl>
## 1 Aruba    ABW        1972         NA
## 2 Aruba    ABW        1973         NA
## 3 Aruba    ABW        1974         NA
## 4 Aruba    ABW        1975         NA
## 5 Aruba    ABW        1976         NA
## 6 Aruba    ABW        1977         NA
## 7 Aruba    ABW        1978         NA
## 8 Aruba    ABW        1979         NA
## 9 Aruba    ABW        1980         NA
## 10 Aruba   ABW       1981         NA
## # ... with 10,623 more rows
```

# Let's do line plots!

```
mortality |>
  select(-indicator) |>
  pivot_longer(
    cols = `1972`:`2020`,
    names_to = "year",
    values_to = "child_mortality"
  ) |>
  ggplot(mapping = aes(x = year, y = child_mortality, group = country)) +
  geom_line(alpha = 0.25)
```

+

# Hmm, what's going on?



# Making sure year is numeric

By default, pivoted column names are characters, but we can transform them:

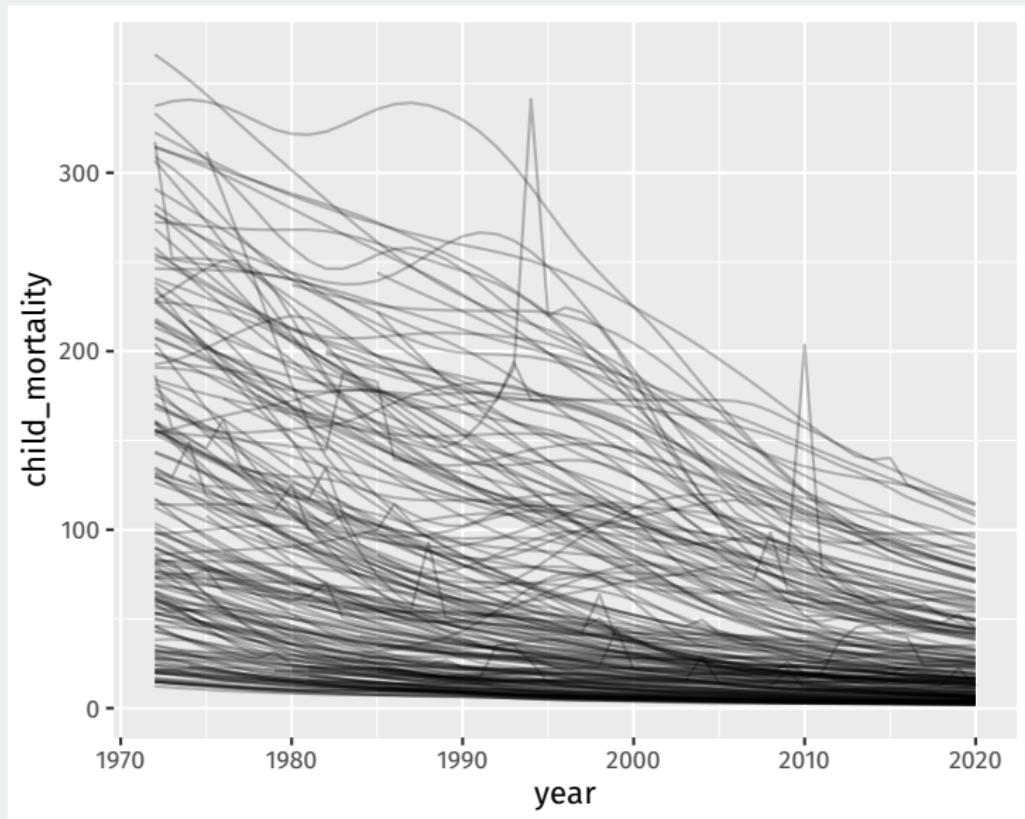
```
mortality_long <- mortality |>
  select(-indicator) |>
  pivot_longer(
    cols = `1972`:`2020`,
    names_to = "year",
    values_to = "child_mortality"
  ) |>
  mutate(year = as.integer(year))
mortality_long
```

```
## # A tibble: 10,633 x 4
##   country country_code  year child_mortality
##   <chr>     <chr>     <int>          <dbl>
## 1 Aruba      ABW      1972            NA
## 2 Aruba      ABW      1973            NA
## 3 Aruba      ABW      1974            NA
## 4 Aruba      ABW      1975            NA
## 5 Aruba      ABW      1976            NA
## 6 Aruba      ABW      1977            NA
```

# Let's (re)do line plots!

```
mortality_long |>  
  ggplot(mapping = aes(x = year, y = child_mortality, group = country)) +  
    geom_line(alpha = 0.25)
```

# There we go



# Spotify data

spotify

```
## # A tibble: 490 x 54
##   Track ~1 Artist week1 week2 week3 week4 week5 week6 week7
##   <chr>    <chr>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 The Box  Roddy~     1     1     1     1     1     1     1
## 2 ROXANNE Arizo~     2     4     5     4     4     4     6
## 3 Yummy    Justi~     3     6    17    17    17    24    15
## 4 Circles   Post ~     4     7     9    10     7    10    11
## 5 BOP       DaBaby     5     5     7     5    11    12    18
## 6 Falling   Trevo~     6     8    10     7     6     8    10
## 7 Dance M~ Tones~     7    13    13    12    12    13    17
## 8 Bandit ~ Juice~     8    11    14    14    15    20    27
## 9 Futsal ~ Lil U~     9     9    19    21    24    32    40
## 10 everyth~ Billi~    10    17    28     9     8    11    14
## # ... with 480 more rows, 45 more variables: week8 <dbl>,
## #   week9 <dbl>, week10 <dbl>, week11 <dbl>, week12 <dbl>,
## #   week13 <dbl>, week14 <dbl>, week15 <dbl>, week16 <dbl>,
## #   week17 <dbl>, week18 <dbl>, week19 <dbl>, week20 <dbl>,
## #   week21 <dbl>, week22 <dbl>, week23 <dbl>, week24 <dbl>,
## #   week25 <dbl>, week26 <dbl>, week27 <dbl>, week28 <dbl>,
## #   week29 <dbl>, week30 <dbl>, week31 <dbl>, ...
```

# Pivoting not ideal

Last approach isn't ideal because of the week prefix:

```
spotify |>
  pivot_longer(
    cols = c(-`Track Name`, -Artist),
    names_to = "week_of_year",
    values_to = "rank"
  )
```

```
## # A tibble: 25,480 x 4
##   `Track Name` Artist     week_of_year   rank
##   <chr>        <chr>      <chr>       <dbl>
## 1 The Box      Roddy Ricch week1         1
## 2 The Box      Roddy Ricch week2         1
## 3 The Box      Roddy Ricch week3         1
## 4 The Box      Roddy Ricch week4         1
## 5 The Box      Roddy Ricch week5         1
## 6 The Box      Roddy Ricch week6         1
## 7 The Box      Roddy Ricch week7         1
## 8 The Box      Roddy Ricch week8         1
## 9 The Box      Roddy Ricch week9         1
## 10 The Box     Roddy Ricch week10        1
## # ... with 25,470 more rows
```

# Removing a column name prefix

When the data in the column name has a fixed prefix, we can use the `names_prefix` to remove it when moving the data to rows

```
spotify |>  
  pivot_longer(  
    cols = c(`Track Name`, -Artist),  
    names_to = "week_of_year",  
    values_to = "rank",  
    names_prefix = "week"  
  ) |>  
  mutate(  
    week_of_year = as.integer(week_of_year)  
  )
```

# Removing a column name prefix

```
## # A tibble: 25,480 x 4
##   `Track Name` Artist    week_of_year   rank
##   <chr>       <chr>           <int> <dbl>
## 1 The Box     Roddy Ricch     1      1
## 2 The Box     Roddy Ricch     2      1
## 3 The Box     Roddy Ricch     3      1
## 4 The Box     Roddy Ricch     4      1
## 5 The Box     Roddy Ricch     5      1
## 6 The Box     Roddy Ricch     6      1
## 7 The Box     Roddy Ricch     7      1
## 8 The Box     Roddy Ricch     8      1
## 9 The Box     Roddy Ricch     9      1
## 10 The Box    Roddy Ricch    10     1
## # ... with 25,470 more rows
```

# 3/ Joining data sets

# Gapminder data

```
library(gapminder)  
gapminder
```

```
## # A tibble: 1,704 x 6  
##   country   continent year lifeExp      pop gdpPercap  
##   <fct>     <fct>    <int>   <dbl>    <int>     <dbl>  
## 1 Afghanistan Asia     1952     28.8  8425333     779.  
## 2 Afghanistan Asia     1957     30.3  9240934     821.  
## 3 Afghanistan Asia     1962     32.0 10267083     853.  
## 4 Afghanistan Asia     1967     34.0 11537966     836.  
## 5 Afghanistan Asia     1972     36.1 13079460     740.  
## 6 Afghanistan Asia     1977     38.4 14880372     786.  
## 7 Afghanistan Asia     1982     39.9 12881816     978.  
## 8 Afghanistan Asia     1987     40.8 13867957     852.  
## 9 Afghanistan Asia     1992     41.7 16317921     649.  
## 10 Afghanistan Asia    1997     41.8 22227415     635.  
## # ... with 1,694 more rows
```

# Joining data sets

What if we want to add the child\_mortality variable to the gampinder data?

Just add the columns? Rows are not aligned properly!

```
gapminder |>  
  select(country, year) |>  
  head()
```

```
## # A tibble: 6 x 2  
##   country     year  
##   <fct>     <int>  
## 1 Afghanistan 1952  
## 2 Afghanistan 1957  
## 3 Afghanistan 1962  
## 4 Afghanistan 1967  
## 5 Afghanistan 1972  
## 6 Afghanistan 1977
```

```
mortality_long |>  
  select(country, year) |>  
  head()
```

```
## # A tibble: 6 x 2  
##   country     year  
##   <chr>      <int>  
## 1 Aruba      1972  
## 2 Aruba      1973  
## 3 Aruba      1974  
## 4 Aruba      1975  
## 5 Aruba      1976  
## 6 Aruba      1977
```

# Key variables

A **primary key** is a variable or set of variables that uniquely identifies rows in the data.

- {country, year} in the gapminder data

A **foreign key** is the corresponding variable(s) in another table.

- {country, year} in the mortality\_long data

If we align the two tables based on these variables, we can add variables from one to the other.

# Checking that the keys are unique

Things get weird if these keys are not unique. Let's check.

Checking primary key is unique:

```
gapminder |>  
  count(country, year) |>  
  filter(n > 1)
```

```
## # A tibble: 0 x 3
```

Checking foreign key:

```
mortality_long |>  
  count(country, year) |>  
  filter(n > 1)
```

```
## # A tibble: 0 x 3
```

# left\_join( ): add variables to primary table

left\_join() keeps all rows from the first argument/piped data:

```
gapminder |>
  left_join(mortality_long) |>
  select(country, year, lifeExp, pop, gdpPercap, child_mortality) |>
  head(n = 6)
```

```
## Joining, by = c("country", "year")

## # A tibble: 6 x 6
##   country     year lifeExp     pop gdpPercap child_morta~1
##   <chr>      <int>   <dbl>   <int>      <dbl>          <dbl>
## 1 Afghanistan 1952    28.8  8425333     779.          NA
## 2 Afghanistan 1957    30.3  9240934     821.          NA
## 3 Afghanistan 1962    32.0  10267083    853.          NA
## 4 Afghanistan 1967    34.0  11537966    836.          NA
## 5 Afghanistan 1972    36.1  13079460    740.         291
## 6 Afghanistan 1977    38.4  14880372    786.         262.
## # ... with abbreviated variable name 1: child_mortality
```

Rows in primary table not in foreign table: new values are NA.

# inner\_join(): add and filter

inner\_join() adds the variables from the foreign table and filters to rows in both tables:

```
gapminder |>
  inner_join(mortality_long) |>
  select(country, year, lifeExp, pop, gdpPercap, child_mortality) |>
  head(n = 6)
```

```
## Joining, by = c("country", "year")

## # A tibble: 6 x 6
##   country     year lifeExp     pop gdpPercap child_morta~1
##   <chr>       <int>   <dbl>   <int>      <dbl>        <dbl>
## 1 Afghanistan 1972    36.1 13079460      740.        291
## 2 Afghanistan 1977    38.4 14880372      786.        262.
## 3 Afghanistan 1982    39.9 12881816      978.        231.
## 4 Afghanistan 1987    40.8 13867957      852.        198.
## 5 Afghanistan 1992    41.7 16317921      649.        166.
## 6 Afghanistan 1997    41.8 22227415      635.        142.
## # ... with abbreviated variable name 1: child_mortality
```

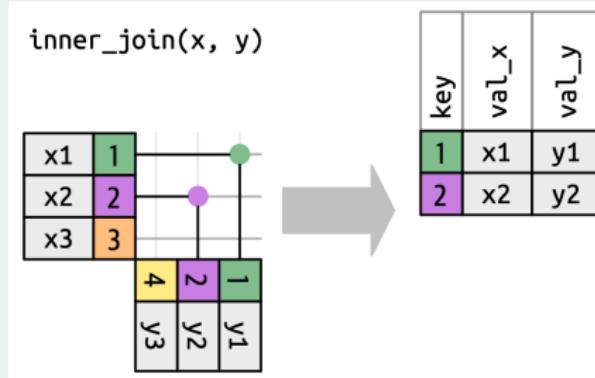
# How inner joins work

Two data sets:

| x | y  |
|---|----|
| 1 | x1 |
| 2 | x2 |
| 3 | x3 |

| y | y1 |
|---|----|
| 1 | y1 |
| 2 | y2 |
| 4 | y3 |

Find matching keys:

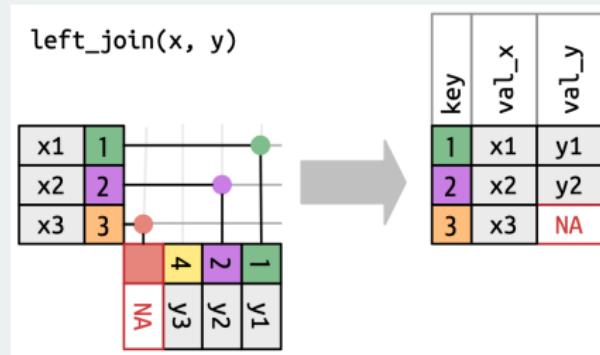


# How left joins work

Two data sets:

|   | x  | y  |
|---|----|----|
| 1 | x1 | y1 |
| 2 | x2 | y2 |
| 3 | x3 | y3 |

Keep all x keys:



# More complicated example

```
library(nycflights13)
flights2 <- flights |>
  select(year, time_hour, origin, dest, tailnum, carrier)
flights2

## # A tibble: 336,776 x 6
##       year   time_hour      origin   dest   tailnum carrier
##   <int> <dttm>      <chr>     <chr>   <chr>   <chr>
## 1  2013 2013-01-01 05:00:00 EWR      IAH      N14228  UA
## 2  2013 2013-01-01 05:00:00 LGA      IAH      N24211  UA
## 3  2013 2013-01-01 05:00:00 JFK      MIA      N619AA  AA
## 4  2013 2013-01-01 05:00:00 JFK      BQN      N804JB  B6
## 5  2013 2013-01-01 06:00:00 LGA      ATL      N668DN  DL
## 6  2013 2013-01-01 05:00:00 EWR      ORD      N39463  UA
## 7  2013 2013-01-01 06:00:00 EWR      FLL      N516JB  B6
## 8  2013 2013-01-01 06:00:00 LGA      IAD      N829AS  EV
## 9  2013 2013-01-01 06:00:00 JFK      MCO      N593JB  B6
## 10 2013 2013-01-01 06:00:00 LGA      ORD      N3ALAA  AA
## # ... with 336,766 more rows
```

# Planes data

```
planes2 <- planes |>
  select(tailnum, year, type, engine, seats)
planes2

## # A tibble: 3,322 x 5
##   tailnum    year type          engine  seats
##   <chr>     <int> <chr>        <chr>   <int>
## 1 N10156    2004 Fixed wing multi engine Turbo-fan 55
## 2 N102UW    1998 Fixed wing multi engine Turbo-fan 182
## 3 N103US    1999 Fixed wing multi engine Turbo-fan 182
## 4 N104UW    1999 Fixed wing multi engine Turbo-fan 182
## 5 N10575    2002 Fixed wing multi engine Turbo-fan 55
## 6 N105UW    1999 Fixed wing multi engine Turbo-fan 182
## 7 N107US    1999 Fixed wing multi engine Turbo-fan 182
## 8 N108UW    1999 Fixed wing multi engine Turbo-fan 182
## 9 N109UW    1999 Fixed wing multi engine Turbo-fan 182
## 10 N110UW   1999 Fixed wing multi engine Turbo-fan 182
## # ... with 3,312 more rows
```

year here is manufacture year.

# What happens with naive join?

```
flights2 |>  
  left_join(planes2)  
  
## Joining, by = c("year", "tailnum")  
  
## # A tibble: 336,776 x 9  
##   year time_hour      origin dest tailnum carrier type engine  
##   <int> <dttm>      <chr>  <chr> <chr>  <chr>  <chr> <chr>  
## 1 2013 2013-01-01 05:00:00 EWR     IAH    N14228 UA     <NA> <NA>  
## 2 2013 2013-01-01 05:00:00 LGA     IAH    N24211 UA     <NA> <NA>  
## 3 2013 2013-01-01 05:00:00 JFK     MIA    N619AA AA     <NA> <NA>  
## 4 2013 2013-01-01 05:00:00 JFK     BQN    N804JB B6     <NA> <NA>  
## 5 2013 2013-01-01 06:00:00 LGA     ATL    N668DN DL     <NA> <NA>  
## 6 2013 2013-01-01 05:00:00 EWR     ORD    N39463 UA     <NA> <NA>  
## 7 2013 2013-01-01 06:00:00 EWR     FLL    N516JB B6     <NA> <NA>  
## 8 2013 2013-01-01 06:00:00 LGA     IAD    N829AS EV     <NA> <NA>  
## 9 2013 2013-01-01 06:00:00 JFK     MCO    N593JB B6     <NA> <NA>  
## 10 2013 2013-01-01 06:00:00 LGA    ORD    N3ALAA AA     <NA> <NA>  
## # ... with 336,766 more rows, and 1 more variable: seats <int>
```

# Specify the joining variables

```
flights2 |>
  left_join(planes2, by = "tailnum")  
  
## # A tibble: 336,776 x 10  
##   year.x time_hour          origin dest tailnum carrier year.y  
##   <int> <dttm>           <chr>  <chr> <chr>   <chr>    <int>  
## 1 2013 2013-01-01 05:00:00 EWR     IAH     N14228  UA      1999  
## 2 2013 2013-01-01 05:00:00 LGA     IAH     N24211  UA      1998  
## 3 2013 2013-01-01 05:00:00 JFK     MIA     N619AA  AA      1990  
## 4 2013 2013-01-01 05:00:00 JFK     BQN     N804JB  B6      2012  
## 5 2013 2013-01-01 06:00:00 LGA     ATL     N668DN  DL      1991  
## 6 2013 2013-01-01 05:00:00 EWR     ORD     N39463  UA      2012  
## 7 2013 2013-01-01 06:00:00 EWR     FLL     N516JB  B6      2000  
## 8 2013 2013-01-01 06:00:00 LGA     IAD     N829AS  EV      1998  
## 9 2013 2013-01-01 06:00:00 JFK     MCO     N593JB  B6      2004  
## 10 2013 2013-01-01 06:00:00 LGA    ORD     N3ALAA  AA      NA  
## # ... with 336,766 more rows, and 3 more variables: type <chr>,  
## #   engine <chr>, seats <int>
```

# Change variables names

```
flights2 |>
  left_join(planes2 |> rename(manufacture_year = year))

## Joining, by = "tailnum"

## # A tibble: 336,776 x 10
##   year time_hour      origin dest tailnum carrier manufacture~1
##   <int> <dttm>       <chr>  <chr> <chr>  <chr>      <int>
## 1 2013 2013-01-01 05:00:00 EWR    IAH    N14228  UA        1999
## 2 2013 2013-01-01 05:00:00 LGA    IAH    N24211  UA        1998
## 3 2013 2013-01-01 05:00:00 JFK    MIA    N619AA  AA        1990
## 4 2013 2013-01-01 05:00:00 JFK    BQN    N804JB  B6        2012
## 5 2013 2013-01-01 06:00:00 LGA    ATL    N668DN  DL        1991
## 6 2013 2013-01-01 05:00:00 EWR    ORD    N39463  UA        2012
## 7 2013 2013-01-01 06:00:00 EWR    FLL    N516JB  B6        2000
## 8 2013 2013-01-01 06:00:00 LGA    IAD    N829AS  EV        1998
## 9 2013 2013-01-01 06:00:00 JFK    MCO    N593JB  B6        2004
## 10 2013 2013-01-01 06:00:00 LGA   ORD    N3ALAA  AA          NA
## # ... with 336,766 more rows, 3 more variables: type <chr>,
## #   engine <chr>, seats <int>, and abbreviated variable name
## #   1: manufacture_year
```

# Gov 50: 12. Prediction and Iteration

Matthew Blackwell

Harvard University

# Roadmap

1. Prediction
2. Loops
3. Evaluating the predictions
4. Time-series plot

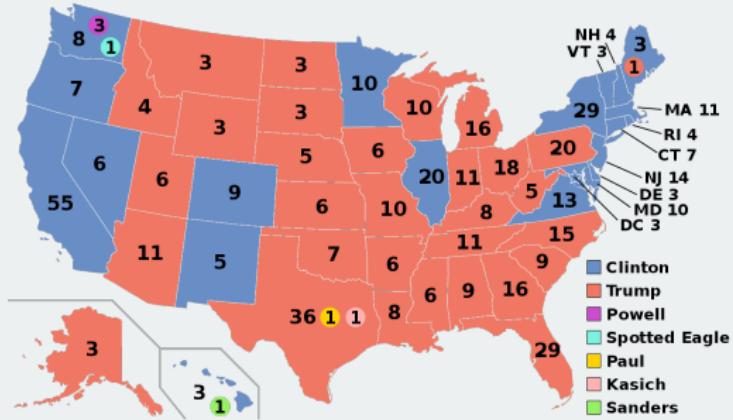
# 1/ Prediction

# 2016 US Presidential Election



- 2016 election popular vote:
  - Clinton: 65,853,516 (48.2%)
  - Trump: 62,984,825 (46.1%)
- Why did Trump win? **Electoral college**
  - Trump: 304, Clinton: 227
- Election determined by 77,744 votes (margins in WI, MI, and PA)
  - 0.056% of the electorate (~136 million)

# Predicting US Presidential Elections



- **Electoral college system**
  - Must win an absolute majority of 538 electoral votes
  - $538 = 435$  (House of Representatives) + 100 (Senators) + 3 (DC)
  - Must win at least 270 votes
  - nobody wins an absolute majority  $\rightsquigarrow$  House vote
- Must predict winner of each state

# Prediction strategy

- Predict state-level support for each candidate using polls
- Allocate electoral college votes of that state to its predicted winner
- Aggregate EC votes across states to determine the predicted winner
- Coding strategy:
  1. For each state, subset to polls within that state.
  2. Further subset the latest polls
  3. Average the latest polls to estimate support for each candidate
  4. Allocate the electoral votes to the candidate who has greatest support
  5. Repeat this for all states and aggregate the electoral votes
- Sounds like a lot of subsets, ugh...

# 2/ Loops

# A simple example

What if we wanted to know the number of unique values of each column of the cces\_2020 data?

```
library(gov50data)
cces_2020

## # A tibble: 51,551 x 6
##   gender race   educ                      pid3  turno~1 pres_~2
##   <fct>  <fct> <fct>                    <fct>    <dbl> <fct>
## 1 Male    White 2-year                   Republ~      1 Donald~
## 2 Female  White Post-grad               Democr~     NA <NA>
## 3 Female  White 4-year                  Indepe~      1 Joe Bi~
## 4 Female  White 4-year                  Democr~      1 Joe Bi~
## 5 Male    White 4-year                  Indepe~      1 Other
## 6 Male    White Some college            Republ~      1 Donald~
## 7 Male    Black  Some college           Not su~     NA <NA>
## 8 Female  White Some college           Indepe~      1 Donald~
## 9 Female  White High school graduate  Republ~      1 Donald~
## 10 Female White 4-year                 Democr~      1 Joe Bi~
## # ... with 51,541 more rows, and abbreviated variable names
## #   1: turnout_self, 2: pres_vote
```

# Manually changing values

```
length(unique(cces_2020$gender))
```

```
## [1] 2
```

```
length(unique(cces_2020$race))
```

```
## [1] 8
```

```
length(unique(cces_2020$educ))
```

```
## [1] 6
```

```
length(unique(cces_2020$pid3))
```

```
## [1] 5
```

```
length(unique(cces_2020$turnout_self))
```

```
## [1] 3
```

```
length(unique(cces_2020$pres_vote))
```

```
## [1] 7
```

# Subsetting with brackets

Note that we can also access variables with [ [] ]:

```
unique(cces_2020$gender)
```

```
## [1] Male   Female  
## Levels: Male Female skipped not asked
```

```
unique(cces_2020[[1]])
```

```
## [1] Male   Female  
## Levels: Male Female skipped not asked
```

```
unique(cces_2020$pid3)
```

```
## [1] Republican   Democrat     Independent Not sure  
## [5] Other  
## 7 Levels: Democrat Republican Independent ... not asked
```

```
unique(cces_2020[[4]])
```

```
## [1] Republican   Democrat     Independent Not sure  
## [5] Other  
## 7 Levels: Democrat Republican Independent ... not asked
```

# Manually changing values, alternative

```
length(unique(cces_2020[[1]]))
```

```
## [1] 2
```

```
length(unique(cces_2020[[2]]))
```

```
## [1] 8
```

```
length(unique(cces_2020[[3]]))
```

```
## [1] 6
```

```
length(unique(cces_2020[[4]]))
```

```
## [1] 5
```

```
length(unique(cces_2020[[5]]))
```

```
## [1] 3
```

```
length(unique(cces_2020[[6]]))
```

```
## [1] 7
```

# Recognizing the template

What if you had more values? Not efficient!

Recognize the template:

```
length(unique(cces_2020[["<>column number"]]))
```

Can we give R this template and a set of column numbers have it do our task repeatedly?

# Loops in R

**for loop** provide a way to execute these templates multiple times:

```
output <- rep(NA, times = ncol(cces_2020))      # 1. output
for (i in seq_along(cces_2020)) {                  # 2. sequence
  output[i] <- length(unique(cces_2020[[i]]))    # 3. body
}
output
```

```
## [1] 2 8 6 5 3 7
```

- Elements of a loop:
  1. output: vector to hold the
  2. i: placeholder name we'll use to swap values between iterations.
  3. seq\_along(cces\_2020): vector of values we want the placeholder to take.
  4. body: a set of expressions that will be repeatedly evaluated.
  5. {}: curly braces to define beginning and end of the loop.
- Indentation is important for readability of the code.

# 2020 polling prediction

Election data: pres20

| Name  | Description                                     |
|-------|---|
| state | abbreviated name of state                       |
| biden | Biden's vote share (percentage)                 |
| trump | Trump's vote share (percentage)                 |
| ev    | number of electoral college votes for the state |

Polling data polls20:

| Name        | Description                                      |
|-------------|--|
| state       | state in which poll was conducted                |
| end_date    | end date the period when poll was conducted      |
| daysleft    | number of days between end date and election day |
| pollster    | name of organization conducting poll             |
| sample_size | name of organization conducting poll             |
| biden       | predicted support for Biden (percentage)         |
| trump       | predicted support for Trump (percentage)         |

## Some preprocessing

# Reminder of our goal

- Coding strategy:
  1. For each state, subset to polls within that state.
  2. Further subset the latest polls
  3. Average the latest polls to estimate support for each candidate
  4. Allocate the electoral votes to the candidate who has greatest support
  5. Repeat this for all states and aggregate the electoral votes

# Poll prediction for each state

```
poll_pred <- rep(NA, 51) # place holder

# get list of unique state names to iterate over
state_names <- sort(unique(polls20$state))

# add labels to holder
names(poll_pred) <- state_names

for (i in 1:51) {
  state_data <- subset(polls20, subset = (state == state_names[i]))

  latest <- state_data$days_left == min(state_data$days_left)

  poll_pred[i] <- mean(state_data$margin[latest])
}

head(poll_pred)

##      AK      AL      AR      AZ      CA      CO
## -9.00 -26.00 -23.00   4.25  26.00  11.00
```

# Tidyverse alternative version

```
poll_pred <- polls20 |>
  group_by(state) |>
  filter(days_left == min(days_left)) |>
  summarize(margin_pred = mean(margin))
poll_pred
```

```
## # A tibble: 51 x 2
##   state    margin_pred
##   <chr>     <dbl>
## 1 AK        -9
## 2 AL       -26
## 3 AR       -23
## 4 AZ        4.25
## 5 CA         26
## 6 CO         11
## 7 CT         22
## 8 DC         89
## 9 DE         22
## 10 FL        0.0800
## # ... with 41 more rows
```

# **3/** Evaluating the predictions

# Polling errors

**Prediction error** = actual outcome – predicted outcome

```
poll_pred <- poll_pred |>  
  left_join(pres20) |>  
  mutate(errors = margin - margin_pred)  
poll_pred
```

```
## # A tibble: 51 x 8  
##   state margin_pred     ev biden trump other  margin errors  
##   <chr>      <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl>    <dbl>  
## 1 AK        -9       3  42.8  52.8  0.732  -10.1   -1.06  
## 2 AL       -26       9  36.6  62.0  0.699  -25.5    0.538  
## 3 AR       -23       6  34.8  62.4  0.257  -27.6   -4.62  
## 4 AZ        4.25     11  49.4  49.1  0.263    0.309  -3.94  
## 5 CA        26       55  63.5  34.3  0.244   29.2    3.16  
## 6 CO        11       9  55.0  41.6  0.161   13.4    2.41  
## 7 CT        22       7  59.3  39.2  0.129   20.1   -1.93  
## 8 DC        89       3  92.1  5.40  0.491   86.8   -2.25  
## 9 DE        22       3  58.7  39.8  0.0780  19.0   -3.03  
## 10 FL       0.0800   29  47.9  51.2  0.0835  -3.36  -3.44  
## # ... with 41 more rows
```

# Assessing the prediction error

**Bias:** average prediction error

```
mean(poll_pred$errors)
```

```
## [1] -3.98
```

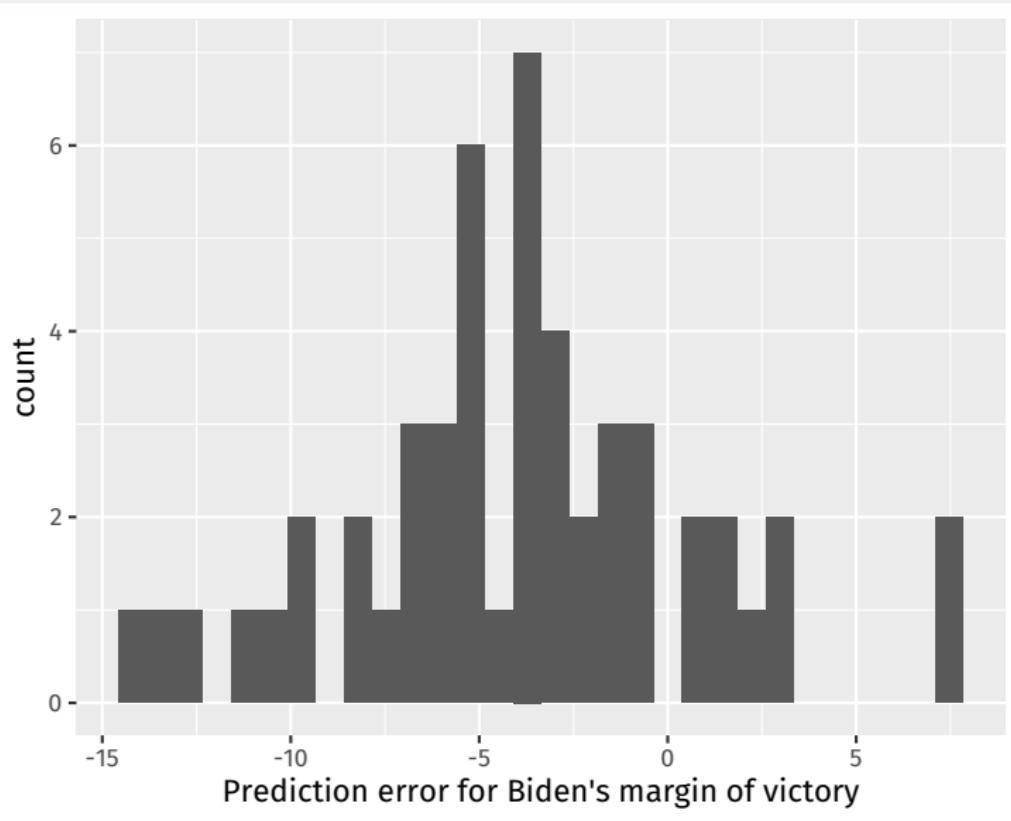
**Root mean-square error:** average magnitude of the prediction error

```
sqrt(mean(poll_pred$errors^2))
```

```
## [1] 6.07
```

# Histogram of the errors

```
ggplot(poll_pred, aes(x = errors)) +  
  geom_histogram() +  
  labs(  
    x = "Prediction error for Biden's margin of victory"  
  )
```

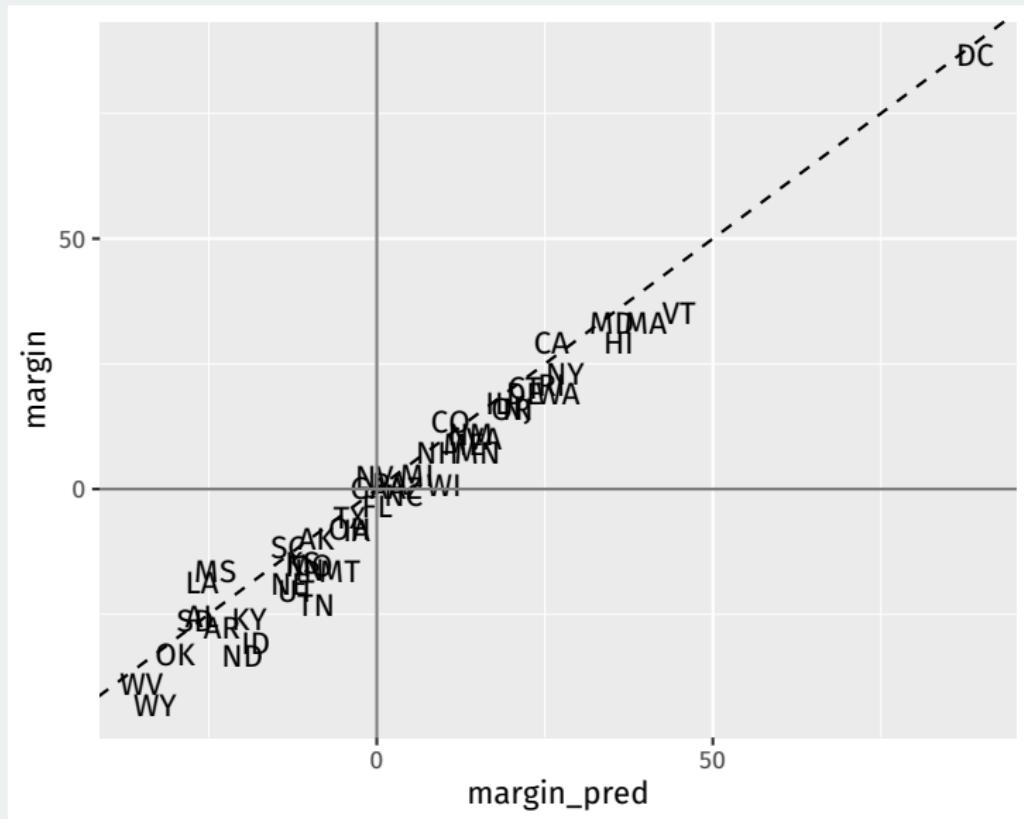


# Comparing polls to outcome

Sometimes we want plot text labels instead of point and we use `geom_text` and the `label` aesthetic:

```
## merge the actual results
ggplot(poll_pred, aes(x = margin_pred, y = margin)) +
  geom_text(aes(label = state)) +
  geom_abline(xintercept = 0, slope = 1, linetype = 2) +
  geom_hline(yintercept = 0, color = "grey50") +
  geom_vline(xintercept = 0, color = "grey50")
```

# Comparing polls to outcome



# Classification

Election prediction: need to predict winner in each state:

```
poll_pred |>  
  filter(margin > 0) |>  
  summarize(sum(ev)) |> pull()
```

```
## [1] 306
```

```
poll_pred |>  
  filter(margin_pred > 0) |>  
  summarize(sum(ev)) |> pull()
```

```
## [1] 328
```

- Prediction of binary outcome variable = **classification problem**
- Wrong prediction  $\rightsquigarrow$  misclassification
  1. **true positive**: predict Trump wins when he actually wins.
  2. **false positive**: predict Trump wins when he actually loses.
  3. **true negative**: predict Trump loses when he actually loses.
  4. **false negative**: predict Trump loses when he actually wins.
- Sometimes false negatives are more/less important: e.g., civil war.

# Classification based on polls

Accuracy: `sign()` returns 1 for a positive number, -1 for a negative number, and 0 for 0.

```
poll_pred |>
  summarize(prop_correct = mean(sign(margin_pred) == sign(margin))) |>
  pull()
```

```
## [1] 0.922
```

Which states did polls call wrong?

```
poll_pred |>
  filter(sign(margin_pred) != sign(margin))
```

```
## # A tibble: 4 x 8
##   state margin_pred    ev biden trump other margin errors
##   <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 FL        0.0800    29  47.9  51.2  0.0835 -3.36 -3.44
## 2 GA       -1.15     16  49.5  49.2  0.0759  0.236  1.39
## 3 NC        3.95     15  48.6  49.9  0.296  -1.35 -5.30
## 4 NV       -0.350     6  50.1  47.7  0.759   2.39  2.74
```

# 4| Time-series plot

# National polls

We often want to show a time series of the national-level polls to get a sense of the popular vote:

```
national_polls20
```

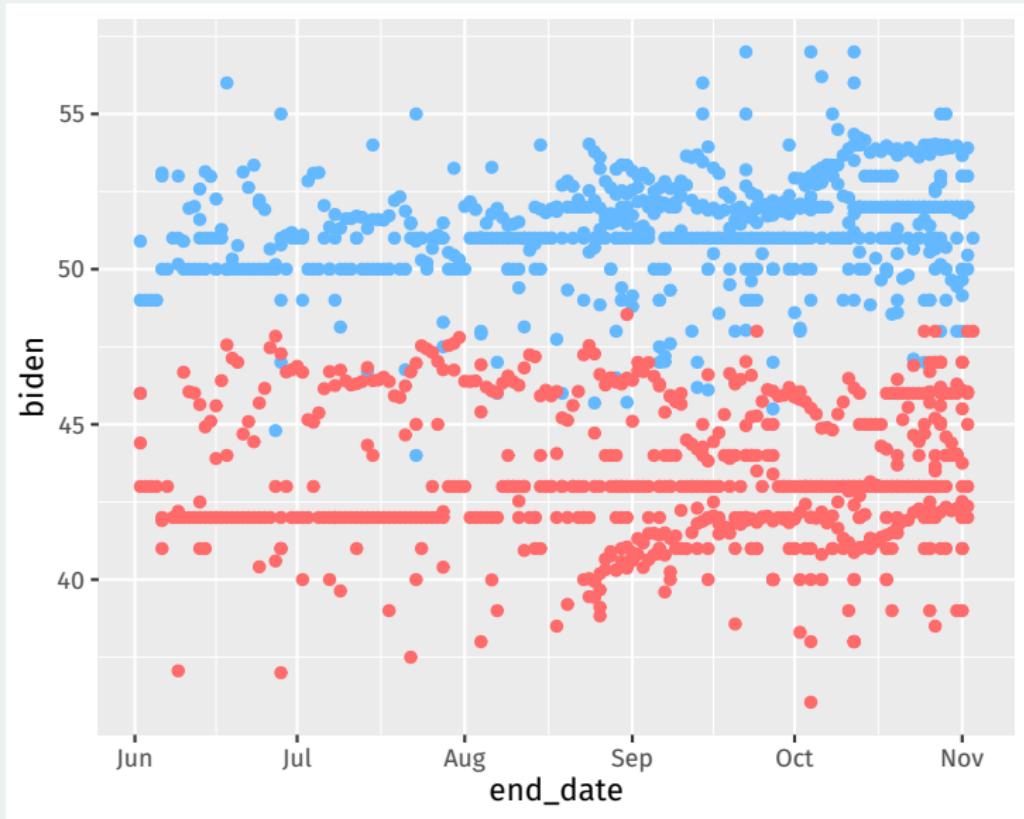
```
## # A tibble: 654 x 5
##   end_date   pollster      sampl~1 biden trump
##   <date>     <chr>        <dbl>  <dbl> <dbl>
## 1 2020-11-03 Lake Research    2400    51    48
## 2 2020-11-02 Research Co.   1025    50    42
## 3 2020-11-02 YouGov       1363    53    43
## 4 2020-11-02 Ipsos         914     52    45
## 5 2020-11-02 SurveyMonkey  28240   52    46
## 6 2020-11-02 HarrisX       2297    52    48
## 7 2020-11-02 TIPP          1212    50.4   46.0
## 8 2020-11-02 USC Dornsife  5423    53.9   42.4
## 9 2020-11-01 John Zogby Strategies/EMI~  1008    49.6   43.8
## 10 2020-11-01 Swayable      5174    51.8   46.1
## # ... with 644 more rows, and abbreviated variable name
## #   1: sample_size
```

# Plotting the raw results

```
national_polls20 |>  
  ggplot(aes(x = end_date)) +  
  geom_point(aes(y = biden), color = "steelblue1") +  
  geom_point(aes(y = trump), color = "indianred1")
```

# Plotting the raw results

Fairly messy:



# Clean the mess by taking moving averages

**Goal:** plot the average of polls in the last 7 days (very difficult with dplyr).

Loop over each day in the data and do:

1. Subset to all polls in the previous 7 days of that day.
2. Calculate the average of these polls for Biden and Trump.
3. Save the result as a 1-row tibble.

# Dates in R

You can get R to properly understand dates and do arithmetic with them:

```
head(national_polls20$end_date)
```

```
## [1] "2020-11-03" "2020-11-02" "2020-11-02" "2020-11-02"  
## [5] "2020-11-02" "2020-11-02"
```

```
head(national_polls20$end_date + 3)
```

```
## [1] "2020-11-06" "2020-11-05" "2020-11-05" "2020-11-05"  
## [5] "2020-11-05" "2020-11-05"
```

# Lubridate to create dates

We can convert a string to a date using the lubridate package:

```
"2020-11-03" + 3 ## R doesn't know this is a date yet!
```

```
## Error in "2020-11-03" + 3: non-numeric argument to binary operator  
lubridate::ymd("2020-11-03") + 3
```

```
## [1] "2020-11-06"  
lubridate::mdy("11/03/2020") + 3
```

```
## [1] "2020-11-06"
```

# Getting a vector of dates

Setup the vector of dates to cover:

```
election_day <- lubridate::ymd("2020-11-03")
all_dates <- seq(from = min(national_polls20$end_date) + 1,
                  to = election_day,
                  by = "days")
head(all_dates)
```

```
## [1] "2020-06-03" "2020-06-04" "2020-06-05" "2020-06-06"
## [5] "2020-06-07" "2020-06-08"
```

# Moving window loop

```
output <- vector("list", length = length(all_dates))

for (i in seq_along(all_dates)) {
  this_date <- all_dates[[i]]

  this_week <- national_polls20 |>
    filter(
      this_date - end_date >= 0,      # this_date is after end_date
      this_date - end_date < 7        # within a week
    )

  output[[i]] <- this_week |>
    summarize(
      date = this_date,
      biden = mean(biden, na.rm = TRUE),
      trump = mean(trump, na.rm = TRUE)
    )
}

output <- bind_rows(output)
```

# Result

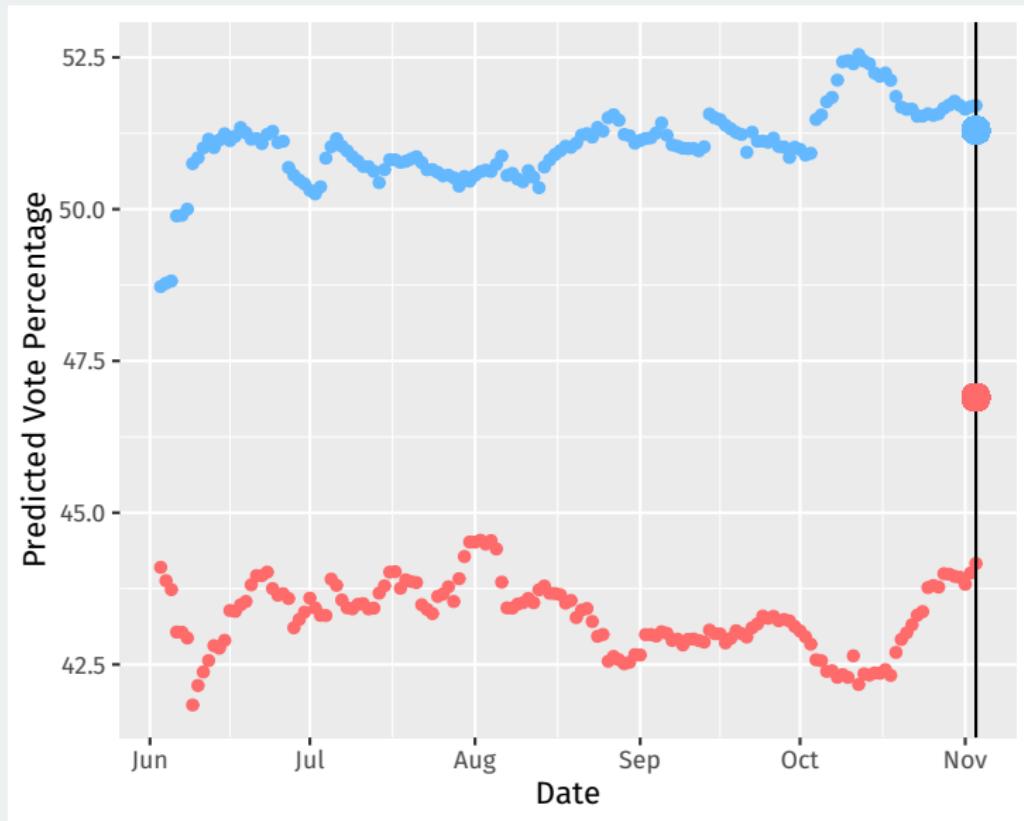
output

```
## # A tibble: 154 x 3
##   date     biden trump
##   <date>    <dbl> <dbl>
## 1 2020-06-03  48.7  44.1
## 2 2020-06-04  48.8  43.9
## 3 2020-06-05  48.8  43.7
## 4 2020-06-06  49.9  43.0
## 5 2020-06-07  49.9  43.0
## 6 2020-06-08  50.   42.9
## 7 2020-06-09  50.8  41.8
## 8 2020-06-10  50.8  42.2
## 9 2020-06-11  51.0  42.4
## 10 2020-06-12 51.2  42.6
## # ... with 144 more rows
```

# Let's plot

```
output |>
  ggplot(aes(x = date)) +
  geom_point(aes(y = biden), color = "steelblue1") +
  geom_point(aes(y = trump), color = "indianred1") +
  geom_vline(xintercept = election_day) +
  geom_point(aes(x = election_day, y = 51.3), color = "steelblue1", size = 1) +
  geom_point(aes(x = election_day, y = 46.9), color = "indianred1", size = 1) +
  labs(
    x = "Date",
    y = "Predicted Vote Percentage"
  )
```

# Let's plot



# Gov 50: 13. Regression

Matthew Blackwell

Harvard University

# Roadmap

1. Prediction
2. Modeling with a line
3. Linear regression in R

# 1/ Prediction

# Predicting my weight

Predicting weight with activity: health data

| Name            | Description                       |
|-----------------|-----------------------------------|
| date            | date of measurements              |
| active_calories | calories burned                   |
| steps           | number of steps taken (in 1,000s) |
| weight          | weight (lbs)                      |
| steps_lag       | steps on day before (in 1,000s)   |
| calories_lag    | calories burned on day before     |

# Predicting using bivariate relationship

- Goal: what's our best guess about  $Y_i$  if we know what  $X_i$  is?
  - what's our best guess about my weight this morning if I know how many steps I took yesterday?
- Terminology:
  - **Dependent/outcome variable:** what we want to predict (weight).
  - **Independent/explanatory variable:** what we're using to predict (steps).

# Weight data

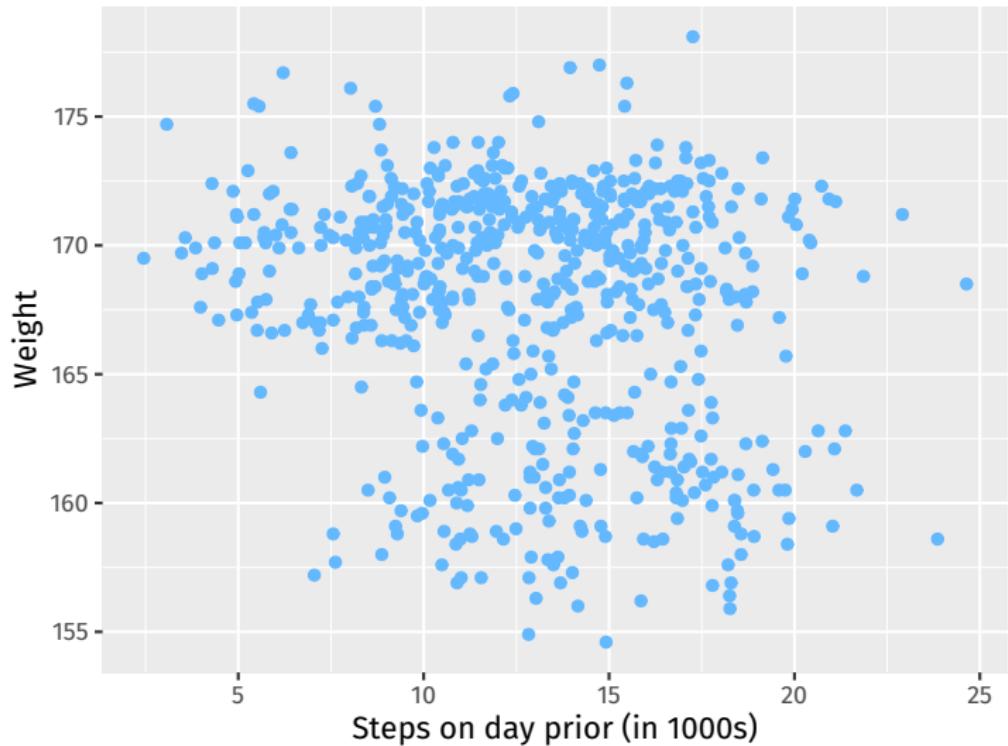
- Load the data:

```
library(gov50data)
health <- drop_na(health)
```

- Plot the data:

```
ggplot(health, aes(x = steps_lag, y = weight)) +
  geom_point(color = "steelblue1") +
  labs(
    x = "Steps on day prior (in 1000s)",
    y = "Weight",
    title = "Weight and Steps"
  )
```

## Weight and Steps



# Prediction one variable with another

- Prediction with access to just  $Y$ : average of the  $Y$  values.
- Prediction with another variable: for any value of  $X$ , what's the best guess about  $Y$ ?
  - Need a function  $y = f(x)$  that maps values of  $X$  into predictions.
  - **Machine learning:** fancy ways to determine  $f(x)$
- Example: what if did 5,000 steps today? What's my best guess about weight?

# Start with looking at a narrow strip of X

Let's find all values that round to 5,000 steps:

```
health |>  
  filter(round(steps_lag) == 5)
```

```
## # A tibble: 12 x 6  
##   date      active.calories steps weight steps_lag calor~1  
##   <date>            <dbl> <dbl>   <dbl>     <dbl>    <dbl>  
## 1 2015-09-08        1111.  15.2    169.     5.02    410.  
## 2 2015-12-12         728.  14.7    167.     5.36    259.  
## 3 2015-12-28         430.  8.94    170.     5.19    314  
## 4 2016-01-29         475.  8.26    171.     4.95    314.  
## 5 2016-02-14         264.  5.42    172.     4.86    297.  
## 6 2016-02-15         892.  13.1   171.     5.42    264.  
## 7 2016-05-02         627.  11.8   170.     5.04    283.  
## 8 2016-06-27         352.  7.21   169.     4.93    212.  
## 9 2016-07-22         766.  14.8   167.     4.96    251.  
## 10 2016-11-25        452.  9.4    173.     5.26    295  
## 11 2016-11-28        577.  11.8   171.     4.97    304.  
## 12 2016-12-30        621.  12.4   176.     5.42    371.  
## # ... with abbreviated variable name 1: calorie_lag
```

# Best guess about Y for this X

Best prediction about weight for a step count of roughly 5,000 is the average weight for observations around that value:

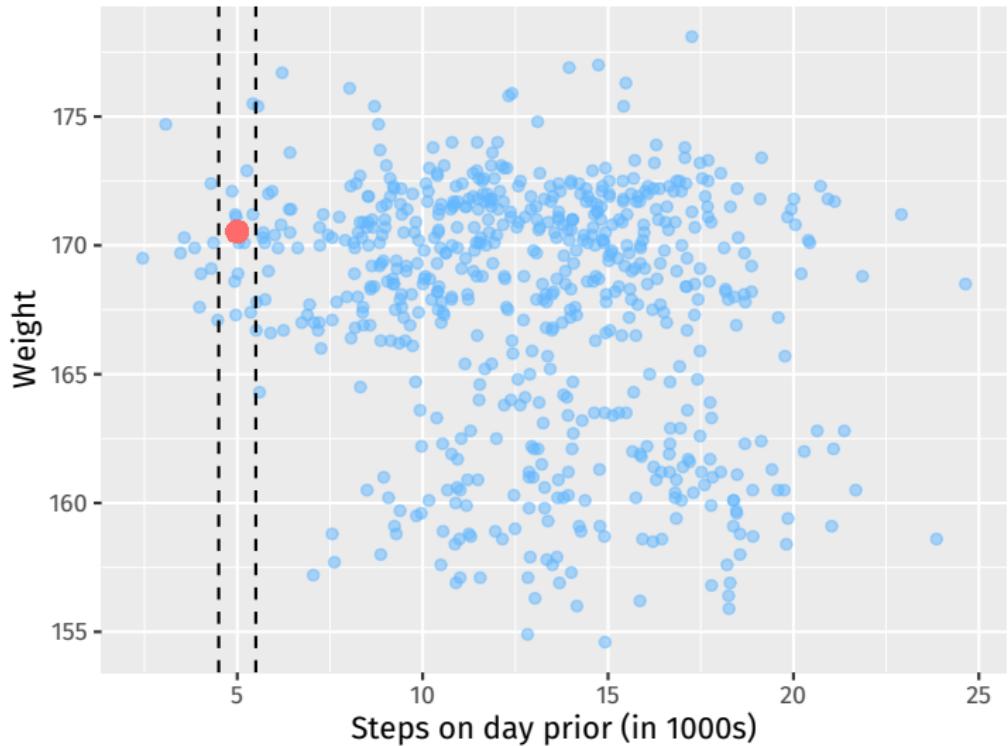
```
mean_wt_5k_steps <- health |>  
  filter(round(steps_lag) == 5) |>  
  summarize(mean(weight)) |>  
  pull()  
mean_wt_5k_steps
```

```
## [1] 171
```

# Plotting the best guess

```
ggplot(health, aes(x = steps_lag, y = weight)) +
  geom_point(color = "steelblue1", alpha = 0.5) +
  labs(
    x = "Steps on day prior (in 1000s)",
    y = "Weight",
    title = "Weight and Steps"
  ) +
  geom_vline(xintercept = c(4.5, 5.5), linetype = "dashed") +
  geom_point(aes(x = 5, y = mean_wt_5k_steps), color = "indianred1",
             size = 3)
```

## Weight and Steps

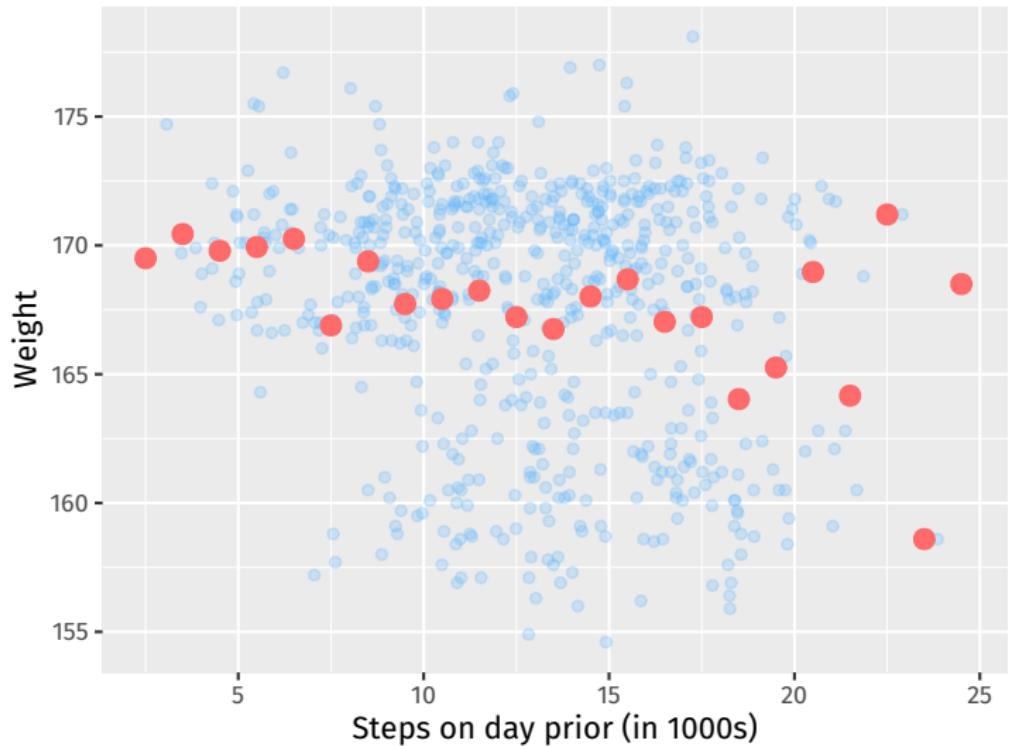


# Binned means

We can use a `stat_summary_bin()` to add these binned means all over the scatter plot:

```
ggplot(health, aes(x = steps_lag, y = weight)) +
  geom_point(color = "steelblue1", alpha = 0.25) +
  labs(
    x = "Steps on day prior (in 1000s)",
    y = "Weight",
    title = "Weight and Steps"
  ) +
  stat_summary_bin(fun = "mean", color = "indianred1", size = 3,
                  geom = "point", binwidth = 1)
```

## Weight and Steps

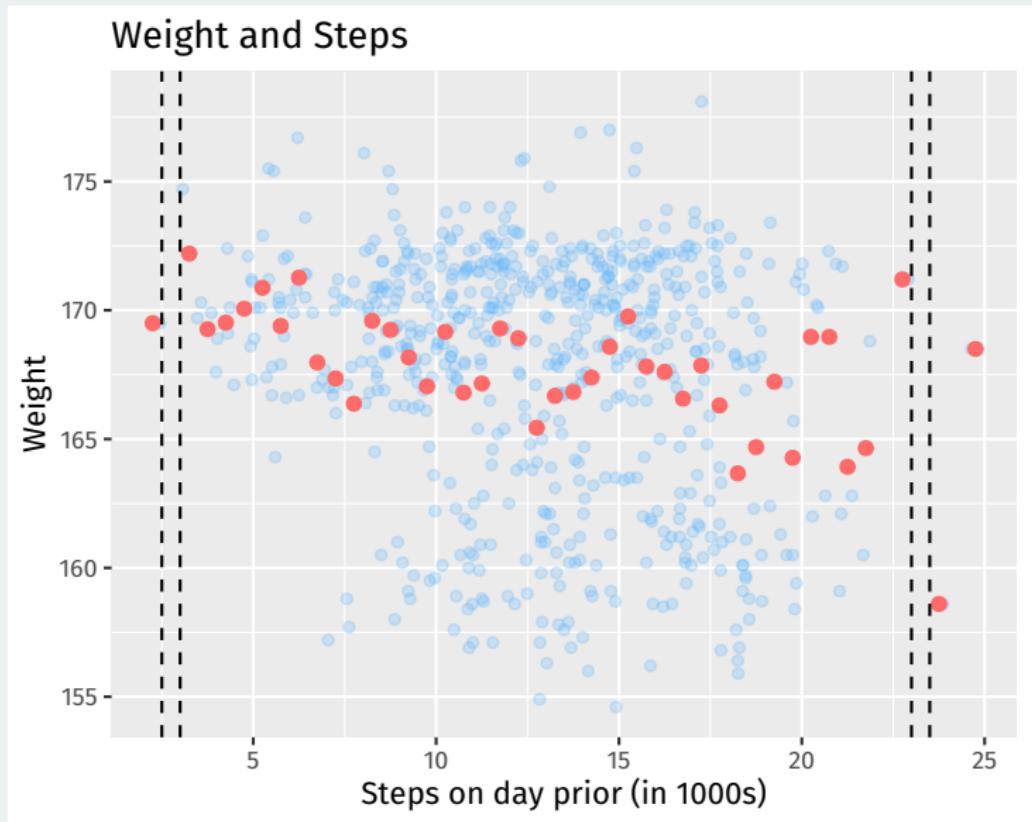


# Smaller bins

But what happens when we make the bins too small?

```
ggplot(health, aes(x = steps_lag, y = weight)) +  
  geom_point(color = "steelblue1", alpha = 0.25) +  
  labs(  
    x = "Steps on day prior (in 1000s)",  
    y = "Weight",  
    title = "Weight and Steps"  
  ) +  
  stat_summary_bin(fun = "mean", color = "indianred1", size = 2,  
                  geom = "point", binwidth = 0.5) +  
  geom_vline(xintercept = c(2.5, 3, 23, 23.5), linetype = "dashed")
```

Gaps and bumps:



# **2/** Modeling with a line

# Using a line to predict

- Can we smooth out these binned means and close gaps? **A model.**
- Simplest possible way to relate two variables: a line.

$$y = mx + b$$

- Problem: for any line we draw, not all the data is on the line.
  - Some points will be above the line, some below.
  - Need a way to account for **chance variation** away from the line.

# Linear regression model

- Model for the line of best fit:

$$Y_i = \underbrace{\alpha}_{\text{intercept}} + \underbrace{\beta}_{\text{slope}} \cdot X_i + \underbrace{\epsilon_i}_{\text{error term}}$$

- Coefficients/parameters** ( $\alpha, \beta$ ): true unknown intercept/slope of the line of best fit.
- Chance error**  $\epsilon_i$ : accounts for the fact that the line doesn't perfectly fit the data.
  - Each observation allowed to be off the regression line.
  - Chance errors are 0 on average.
- Useful fiction: this model represents the **data generating process**
  - George Box: "all models are wrong, some are useful"

# Interpreting the regression line

$$Y_i = \alpha + \beta \cdot X_i + \epsilon_i$$

- **Intercept**  $\alpha$ : average value of  $Y$  when  $X$  is 0
  - Average weight when I take 0 steps the day prior.
- **Slope**  $\beta$ : average change in  $Y$  when  $X$  increases by one unit.
  - Average decrease in weight for each additional 1,000 steps.

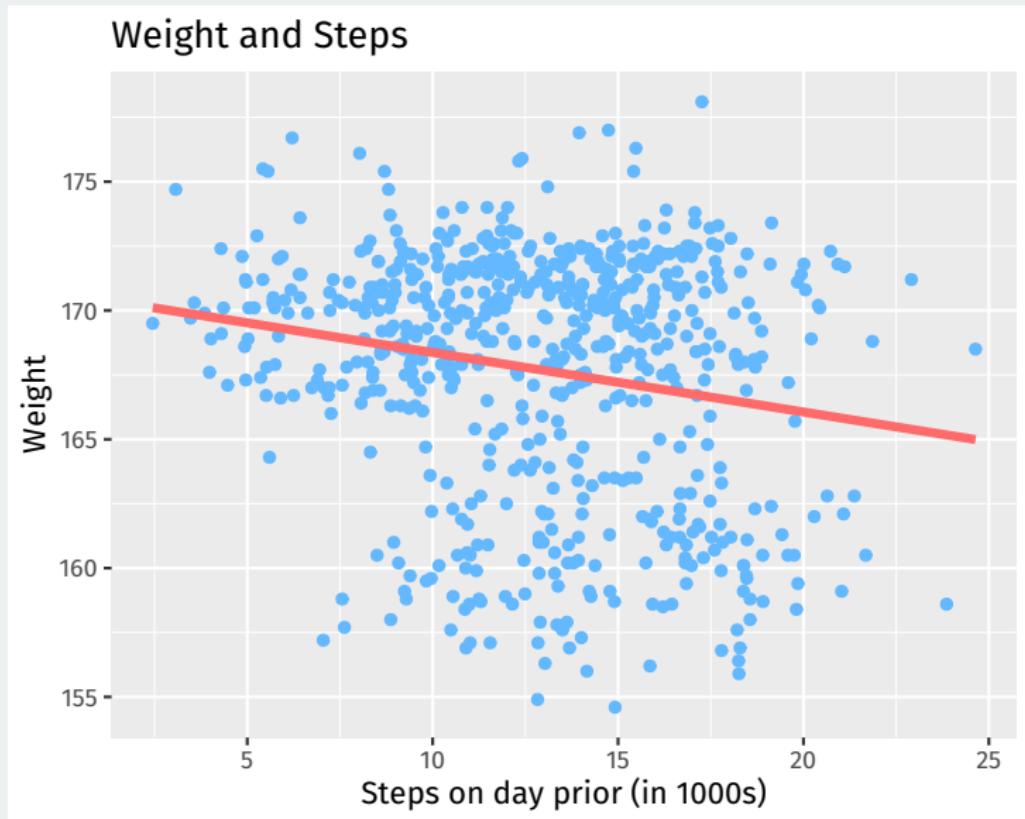
# Estimated coefficients

- Parameters:  $\alpha, \beta$ 
  - Unknown features of the **data-generating process**.
  - Chance error makes these impossible to observe directly.
- Estimates:  $\hat{\alpha}, \hat{\beta}$ 
  - An **estimate** is our best guess about some parameter.
- **Regression line:**  $\widehat{Y} = \hat{\alpha} + \hat{\beta} \cdot x$ 
  - Average value of  $Y$  when  $X$  is equal to  $x$ .
  - Represents the best guess or **predicted value** of the outcome at  $x$ .

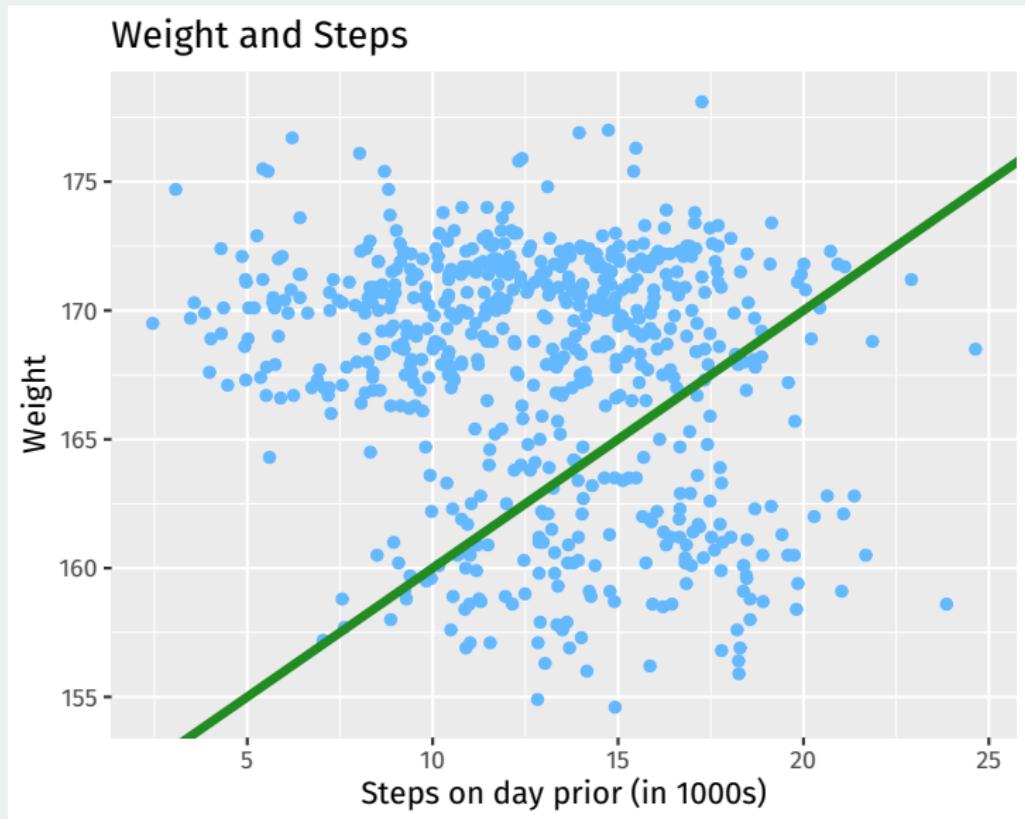
# Line of best fit

```
ggplot(health, aes(x = steps_lag, y = weight)) +
  geom_point(color = "steelblue1") +
  labs(
    x = "Steps on day prior (in 1000s)",
    y = "Weight",
    title = "Weight and Steps"
  ) +
  geom_smooth(method = "lm", se = FALSE, color = "indianred1", size = 1.5)
```

# Line of best fit



# Why not this line?



# Prediction error

Let's understand the **prediction error** for a line with intercept  $a$  and slope  $b$ .

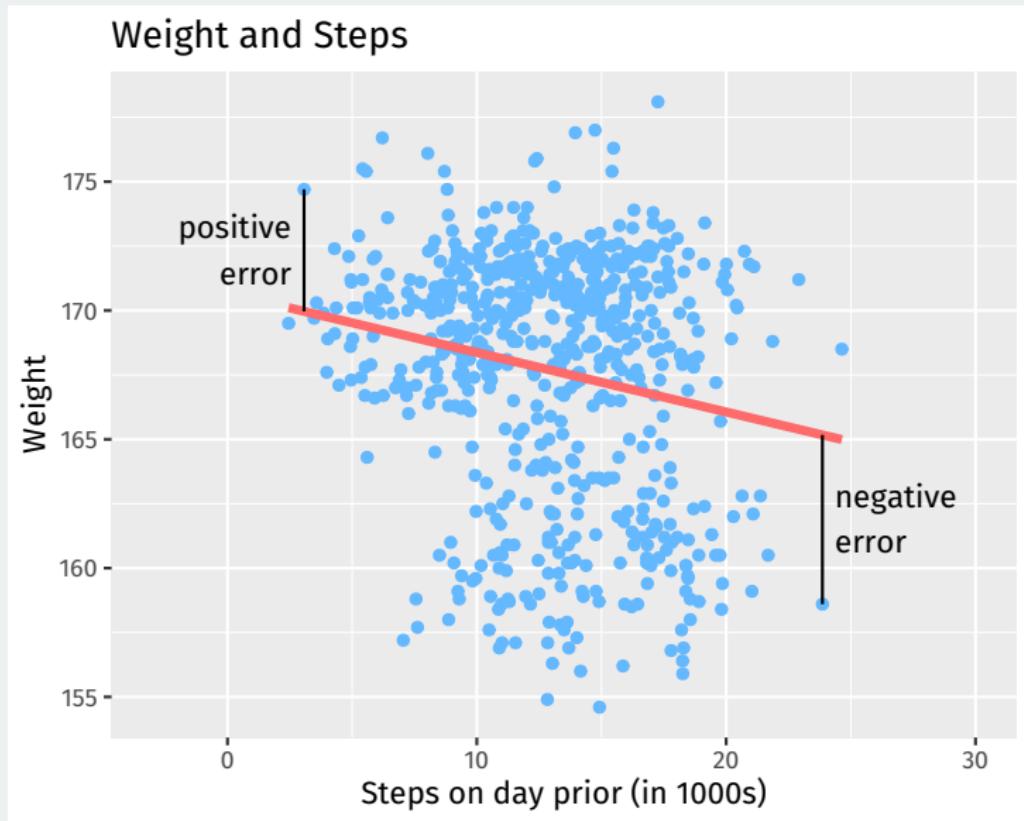
**Fitted/predicted value** for unit  $i$ :

$$a + b \cdot X_i$$

**Preditction error (residual):**

$$\text{error} = \text{actual} - \text{predicted} = Y_i - (a + b \cdot X_i)$$

# Prediction errors/residuals



# Least squares

- Get these estimates by the **least squares method**.
- Minimize the **sum of the squared residuals** (SSR):

$$\text{SSR} = \sum_{i=1}^n (\text{prediction error}_i)^2 = \sum_{i=1}^n (Y_i - a - b \cdot X_i)^2$$

- Finds the line that minimizes the magnitude of the prediction errors!

# 3/ Linear regression in R

# Linear regression in R

- R will calculate least squares line for a data set using `lm( )`.
  - Syntax: `lm(y ~ x, data = mydata)`
  - `y` is the name of the dependent variable
  - `x` is the name of the independent variable
  - `mydata` is the data.frame where they live

```
fit <- lm(weight ~ steps_lag, data = health)
fit

##
## Call:
## lm(formula = weight ~ steps_lag, data = health)
##
## Coefficients:
## (Intercept)    steps_lag
##           170.675        -0.231
```

# Coefficients

Use `coef()` to extract estimated coefficients:

```
coef(fit)
```

```
## (Intercept)  steps_lag  
##      170.675     -0.231
```

**Interpretation:** a 1-unit increase in  $X$  (1,000 steps) is associated with a decrease in the average weight of 0.231 pounds.

**Question:** what would this model predict about the change in average weight for a 10,000 step increase in steps?

# broom package

The `broom` package can provide nice summaries of the regression output.

`augment()` can show fitted values, residuals and other unit-level statistics:

```
library(broom)
augment(fit) |> head()
```

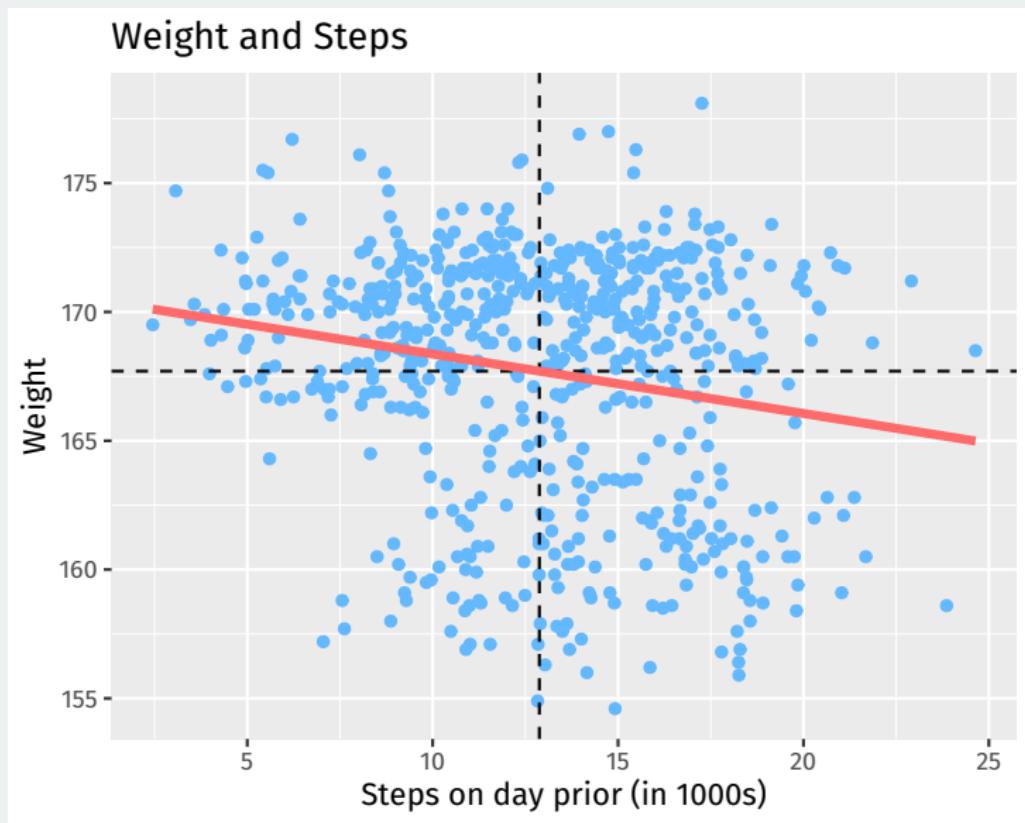
```
## # A tibble: 6 x 8
##   weight steps_lag .fitted   .resid     .hat .sigma   .cooksdf
##   <dbl>     <dbl>   <dbl>    <dbl>    <dbl>   <dbl>    <dbl>
## 1 169.      17.5    167.    2.46    0.00369  4.68  5.13e-4
## 2 168.      18.4    166.    1.57    0.00463  4.68  2.64e-4
## 3 167.      19.6    166.    1.05    0.00609  4.68  1.54e-4
## 4 168.      10.4    168.   -0.0750  0.00217  4.68  2.80e-7
## 5 168.      18.7    166.    1.44    0.00496  4.68  2.38e-4
## 6 166.      9.14    169.   -2.27    0.00296  4.68  3.49e-4
## # ... with 1 more variable: .std.resid <dbl>
```

# Properties of least squares

Least squares line always goes through  $(\bar{X}, \bar{Y})$ .

```
ggplot(health, aes(x = steps_lag, y = weight)) +  
  geom_point(color = "steelblue1") +  
  labs(  
    x = "Steps on day prior (in 1000s)",  
    y = "Weight",  
    title = "Weight and Steps"  
  ) +  
  geom_hline(yintercept = mean(health$weight), linetype = "dashed") +  
  geom_vline(xintercept = mean(health$steps_lag), linetype = "dashed") +  
  geom_smooth(method = "lm", se = FALSE, color = "indianred1", size = 1.5)
```

Least squares line always goes through  $(\bar{X}, \bar{Y})$ .



# Properties of least squares line

Estimated slope is related to correlation:

$$\hat{\beta} = (\text{correlation of } X \text{ and } Y) \times \frac{\text{SD of } Y}{\text{SD of } X}$$

Mean of residuals is always 0.

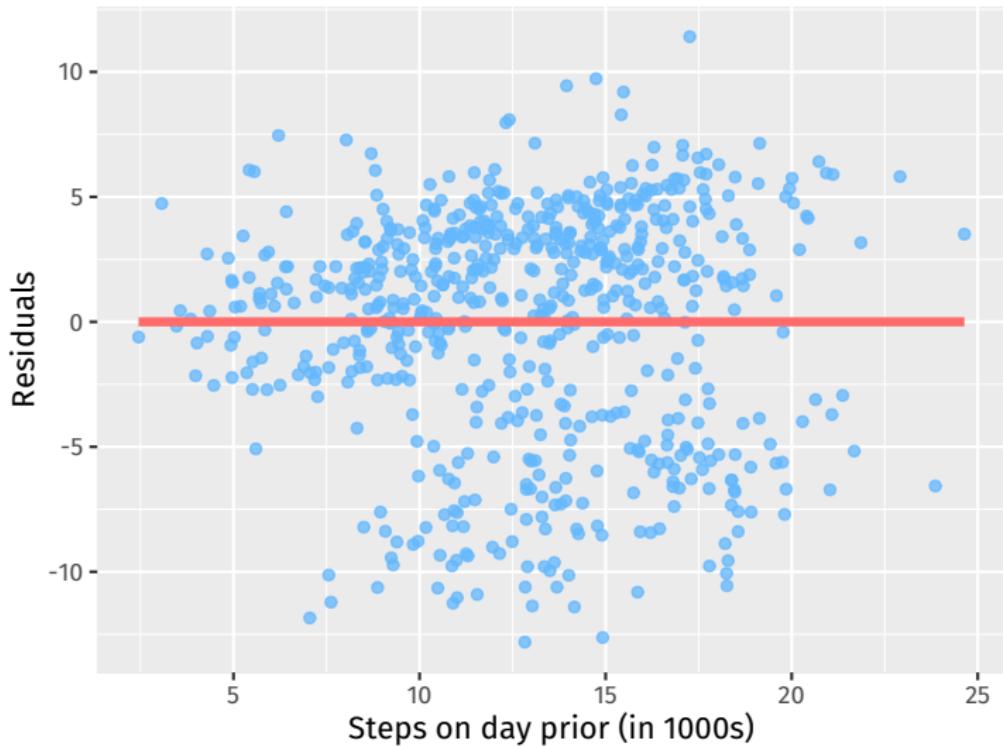
```
augment(fit) |>  
  summarize(mean(.resid))
```

```
## # A tibble: 1 x 1  
##   `mean(.resid)`  
##       <dbl>  
## 1      -1.21e-13
```

# Plotting the residuals

```
augment(fit) |>
  ggplot(aes(x = steps_lag, y = .resid)) +
  geom_point(color = "steelblue1", alpha = 0.75) +
  labs(
    x = "Steps on day prior (in 1000s)",
    y = "Residuals",
    title = "Residual plot"
  ) +
  geom_smooth(method = "lm", se = FALSE, color = "indianred1", size = 1.5)
```

## Residual plot

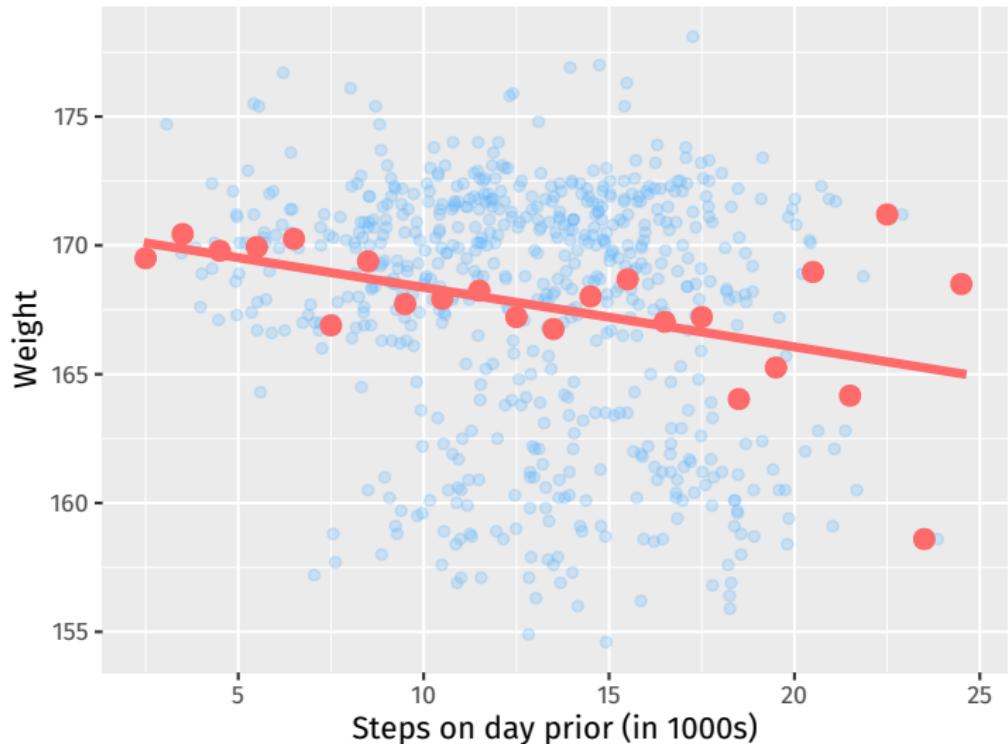


# Smoothed graph of averages

Another way to think of the regression line is a smoothed version of the binned means plot:

```
ggplot(health, aes(x = steps_lag, y = weight)) +  
  geom_point(color = "steelblue1", alpha = 0.25) +  
  labs(  
    x = "Steps on day prior (in 1000s)",  
    y = "Weight",  
    title = "Weight and Steps"  
  ) +  
  stat_summary_bin(fun = "mean", color = "indianred1", size = 3,  
                   geom = "point", binwidth = 1) +  
  geom_smooth(method = "lm", se = FALSE, color = "indianred1", size = 1.5)
```

## Weight and Steps



# Gov 50: 14. More Regression and Model Fit

Matthew Blackwell

Harvard University

# Roadmap

1. Model fit
2. Multiple regression

# 1/ Model fit

# Presidential popularity and the midterms

- Does popularity of the president or recent changes in the economy better predict midterm election outcomes?

| Name        | Description  |
|-------------|--|
| year        | midterm election year  |
| president   | name of president  |
| party       | Democrat or Republican   |
| approval    | Gallup approval rating at midterms                               |
| rdi_change  | % change in real disposable income over the year before midterms |
| seat_change | change in the number of House seats for the president's party    |

```
library(gov50data)
midterms
```

```
## # A tibble: 20 x 6
##   year president party approval seat_change rdi_change
##   <dbl> <chr>     <chr>      <dbl>        <dbl>       <dbl>
## 1 1946 Truman     D          33         -55        NA
## 2 1950 Truman     D          39         -29        8.2
## 3 1954 Eisenhower R          61          -4         1
## 4 1958 Eisenhower R          57         -47        1.1
## 5 1962 Kennedy    D          61          -4         5
## 6 1966 Johnson   D          44         -47        5.3
## 7 1970 Nixon     R          58          -8         6.6
## 8 1974 Ford      R          54         -43        6.4
## 9 1978 Carter    D          49         -11        7.7
## 10 1982 Reagan   R          42         -28        4.8
## 11 1986 Reagan   R          63          -5         5.1
## 12 1990 H.W. Bush R          58          -8         5.6
## 13 1994 Clinton  D          46         -53        3.9
## 14 1998 Clinton  D          66           5         5.6
## 15 2002 W. Bush   R          63           6         2.6
## 16 2006 W. Bush   R          38         -30        5.7
## 17 2010 Obama    D          45         -63        3.5
## 18 2014 Obama    D          40         -13        4.6
## 19 2018 Trump    R          38         -42        4.1
## 20 2022 Biden    D          42          NA       -0.003
```

# Fitting the approval model

```
fit.app <- lm(seat_change ~ approval, data = midterms)
fit.app

## 
## Call:
## lm(formula = seat_change ~ approval, data = midterms)
##
## Coefficients:
## (Intercept)      approval
##           -96.58          1.42
```

For a one-point increase in presidential approval, the predicted seat change increases by 1.42

# Fitting the income model

```
fit.rdi <- lm(seat_change ~ rdi_change, data = midterms)
fit.rdi

##
## Call:
## lm(formula = seat_change ~ rdi_change, data = midterms)
##
## Coefficients:
## (Intercept)    rdi_change
##           -29.41          1.21
```

For a one-point increase in the change in real disposable income, the predicted seat change increases by 1.21

# Comparing models



- How well do the models “fit the data”?
  - How well does the model predict the outcome variable in the data?

# Model fit

Model prediction error:

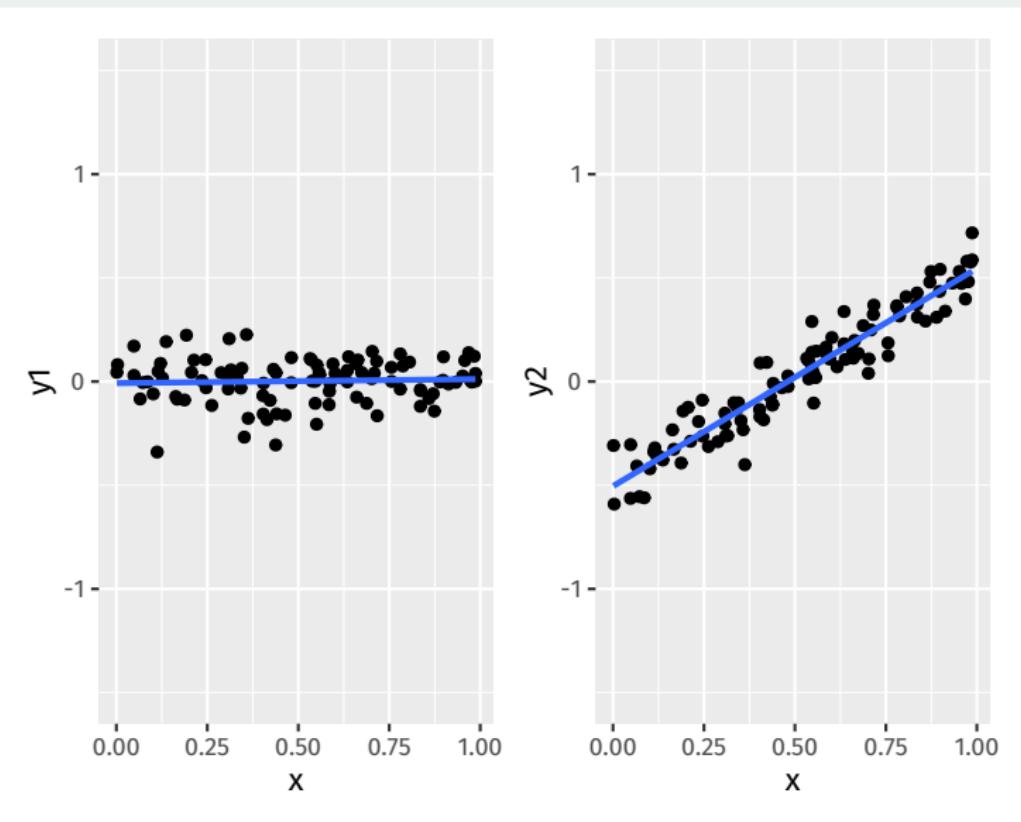
$$\text{prediction error} = \sum_{i=1}^n (\text{actual}_i - \text{predicted}_i)^2$$

Prediction error for regression: **Sum of squared residuals**

$$\text{SSR} = \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

Lower SSR is better, right?

These two regression lines have approximately the same SSR:



# Benchmarking model fit

Benchmarking our predictions using the **proportional reduction in error**:

$$\frac{\text{reduction in prediction error using model}}{\text{baseline prediction error}}$$

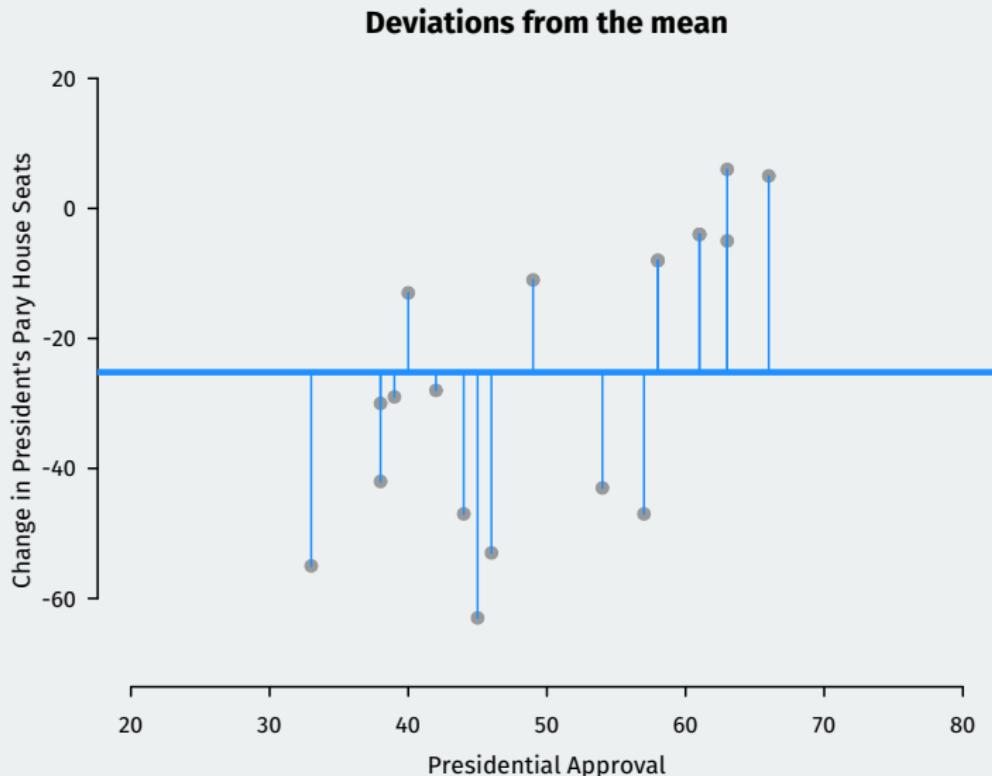
Baseline prediction error without a regression is using the mean of  $Y$  to predict. This is called the **Total sum of squares**:

$$\text{TSS} = \sum_{i=1}^n (Y_i - \bar{Y})^2$$

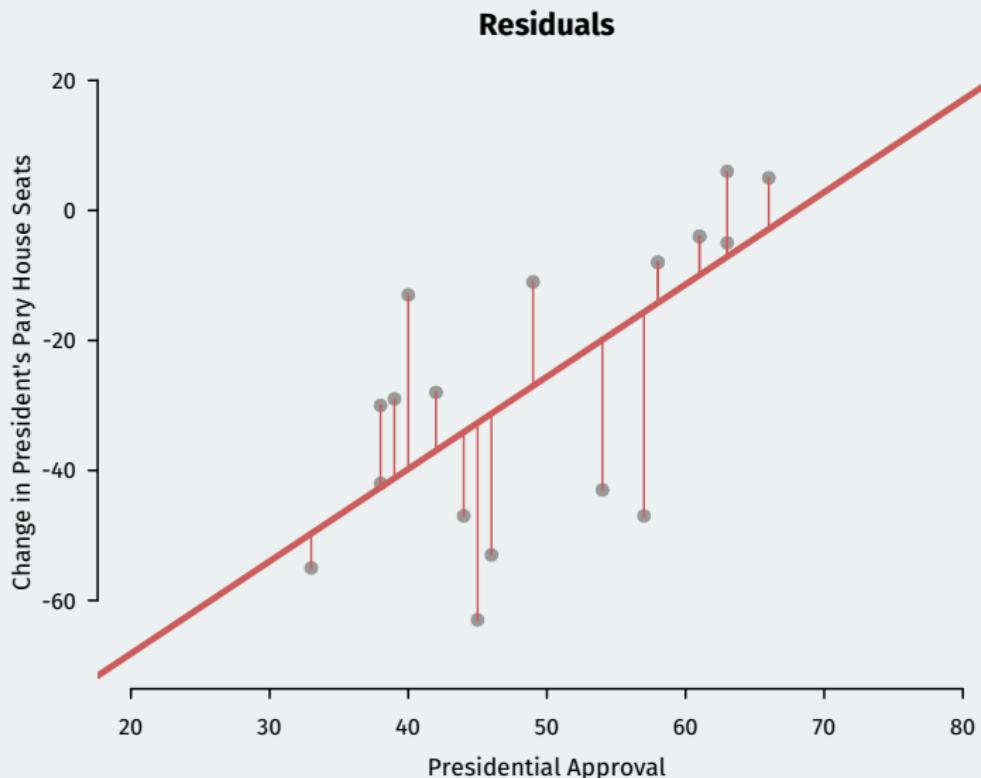
Leads to the **coefficient of determination**,  $R^2$ , one summary of LS model fit:

$$R^2 = \frac{TSS - SSR}{TSS} = \frac{\text{how much smaller LS prediction errors are vs mean prediction error using the mean}}{\text{prediction error using the mean}}$$

# Total SS vs SSR



# Total SS vs SSR



# Model fit in R

- To access  $R^2$  from the `lm()` output, use the `summary()` function:

```
fit.app.sum <- summary(fit.app)
fit.app.sum$r.squared
```

```
## [1] 0.45
```

- Compare to the fit using change in income:

```
fit.rdi.sum <- summary(fit.rdi)
fit.rdi.sum$r.squared
```

```
## [1] 0.012
```

- Which does a better job predicting midterm election outcomes?

# Accessing model fit via broom package

We can also access summary statistics like model fit using the `glance()` function from `broom`:

```
library(broom)
glance(fit.app)

## # A tibble: 1 x 12
##   r.squared adj.r~1 sigma stati~2 p.value    df logLik    AIC
##       <dbl>     <dbl> <dbl>    <dbl> <dbl> <dbl> <dbl>
## 1     0.450    0.418  16.9    13.9 0.00167     1   -79.6  165.
## # ... with 4 more variables: BIC <dbl>, deviance <dbl>,
## #   df.residual <int>, nobs <int>, and abbreviated variable
## #   names 1: adj.r.squared, 2: statistic
```

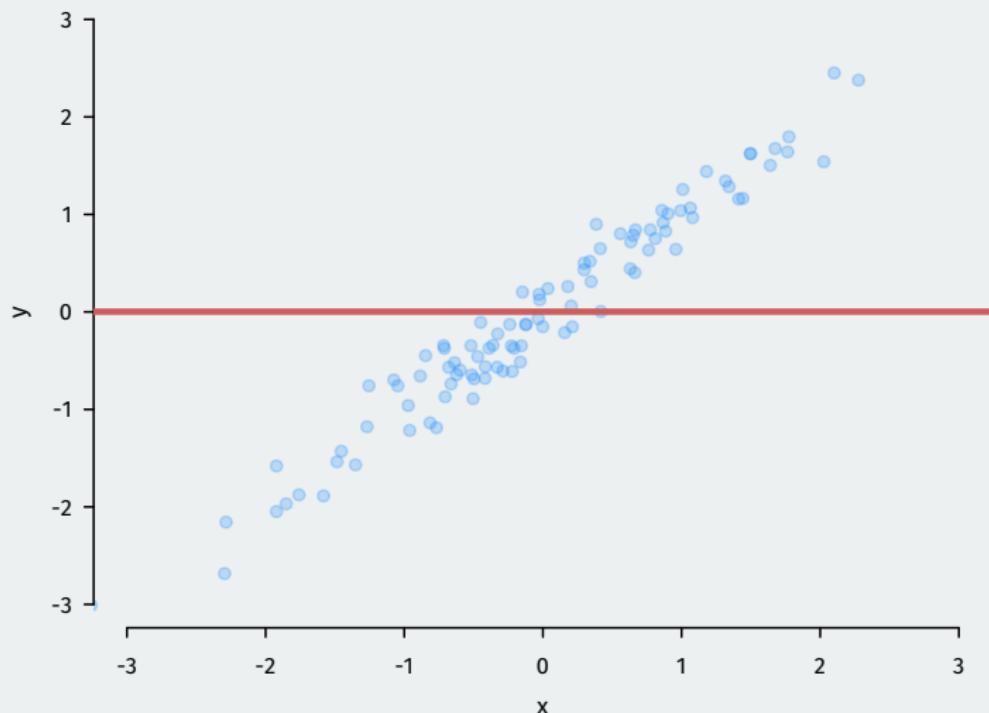
# Fake data, better fit

- Little hard to see what's happening in that example.
- Let's look at fake variables  $x$  and  $y$ :

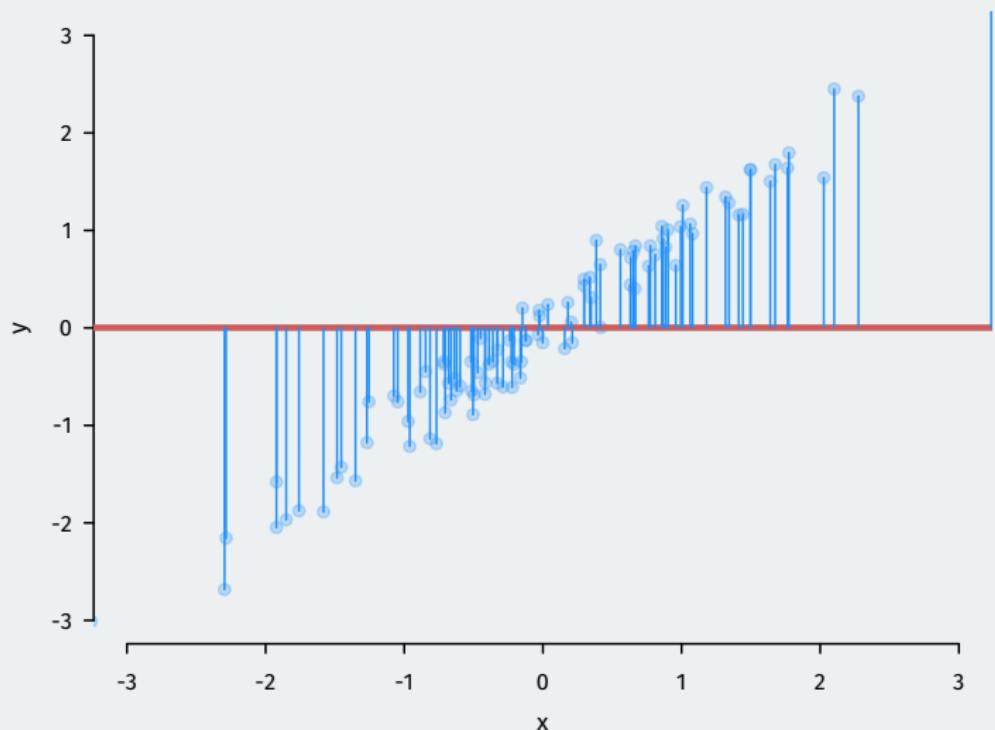
```
fit.x <- lm(y ~ x)
```

- Very good model fit:  $R^2 \approx 0.95$

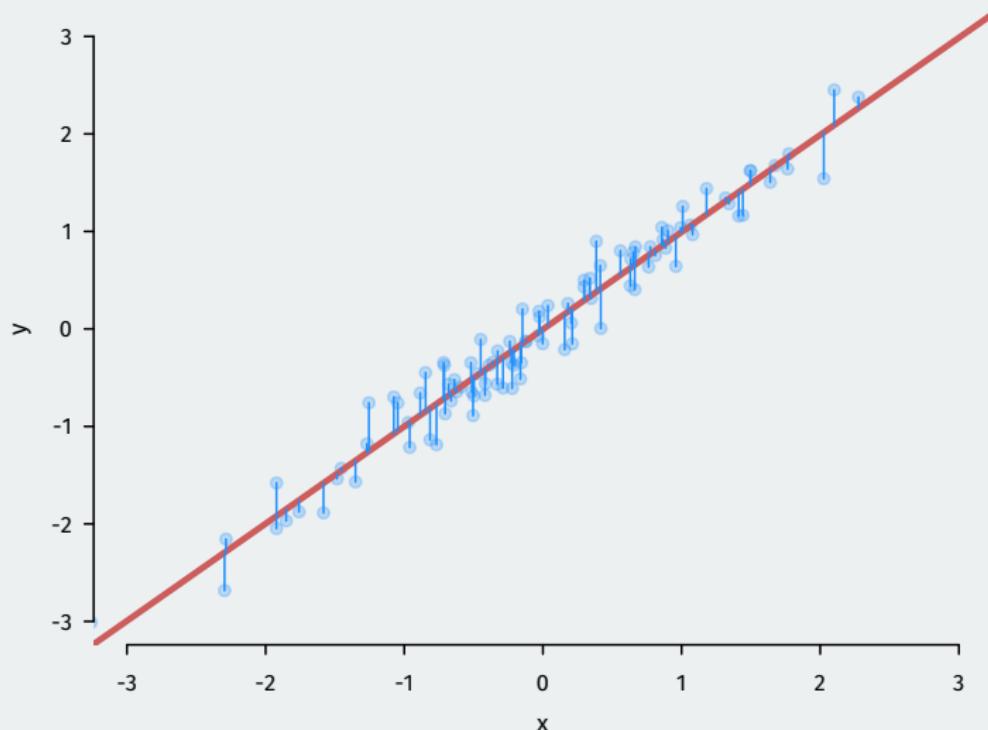
# Fake data, better fit



# Fake data, better fit

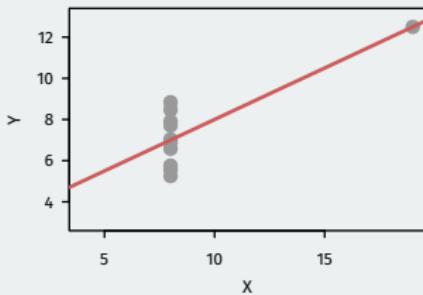
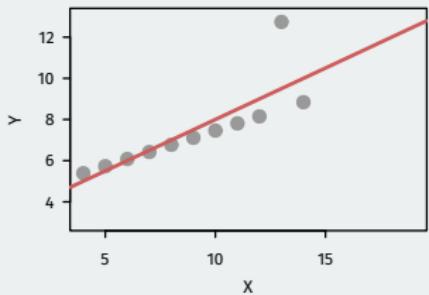
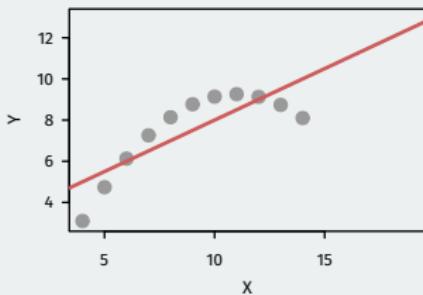
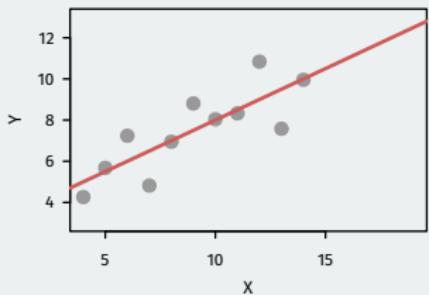


# Fake data, better fit



# Is R-squared useful?

- Can be very misleading. Each of these samples have the same  $R^2$  even though they are vastly different:



# Overfitting

- **In-sample fit:** how well your model predicts the data used to estimate it.
  - $R^2$  is a measure of in-sample fit.
- **Out-of-sample fit:** how well your model predicts new data.
- **Overfitting:** OLS optimizes in-sample fit; may do poorly out of sample.
  - Example: predicting winner of Democratic presidential primary with gender of the candidate.
  - Until 2016, gender was a **perfect** predictor of who wins the primary.
  - Prediction for 2016 based on this: Bernie Sanders as Dem. nominee.
  - Bad out-of-sample prediction due to overfitting!

# 2/ Multiple regression

# Multiple predictors

What if we want to predict  $Y$  as a function of many variables?

$$\text{seat\_change}_i = \alpha + \beta_1 \text{approval}_i + \beta_2 \text{rdi\_change}_i + \epsilon_i$$

Why?

- Better predictions (at least in-sample).
- Better interpretation as **ceteris paribus** relationships:
  - $\beta_1$  is the relationship between approval and seat\_change holding rdi\_change constant.
  - **Statistical control** in a cross-sectional study.

# Multiple regression in R

```
mult.fit <- lm(seat_change ~ approval + rdi_change,  
                data = midterms)  
  
mult.fit  
  
##  
## Call:  
## lm(formula = seat_change ~ approval + rdi_change, data = midterms)  
##  
## Coefficients:  
## (Intercept)      approval      rdi_change  
##       -117.23          1.53         3.22
```

- $\hat{\alpha} = -117.2$ : average seat change president has 0% approval and no change in income levels.
- $\hat{\beta}_1 = 1.53$ : average increase in seat change for additional percentage point of approval, **holding RDI change fixed**
- $\hat{\beta}_2 = 3.217$ : average increase in seat change for each additional percentage point increase of RDI, **holding approval fixed**

# Least squares with multiple regression

- How do we estimate the coefficients?
- The same exact way as before: minimize prediction error!
- Residuals (aka prediction error) with multiple predictors:

$$Y_i - \hat{Y}_i = \text{seat\_change}_i - \hat{\alpha} - \hat{\beta}_1 \text{approval}_i - \hat{\beta}_2 \text{rdi\_change}_i$$

- Find the coefficients that minimizes the **sum of the squared residuals**:

$$\text{SSR} = \sum_{i=1}^n \hat{\epsilon}_i^2 = (Y_i - \hat{\alpha} - \hat{\beta}_1 X_{i1} - \hat{\beta}_2 X_{i2})^2$$

# Model fit with multiple predictors

- $R^2$  mechanically increases when you add variables to the regression.
  - But this could be overfitting!!
- Solution: penalize regression models with more variables.
  - Occam's razor: **simpler models are preferred**
- Adjusted  $R^2$ : lowers regular  $R^2$  for each additional covariate.
  - If the added covariates doesn't help predict, adjusted  $R^2$  goes down!

# Comparing model fits

```
glance(fit.app) |>  
  select(r.squared, adj.r.squared, sigma)
```

```
## # A tibble: 1 x 3  
##   r.squared adj.r.squared sigma  
##       <dbl>         <dbl> <dbl>  
## 1     0.450         0.418 16.9
```

```
glance(mult.fit) |>  
  select(r.squared, adj.r.squared, sigma)
```

```
## # A tibble: 1 x 3  
##   r.squared adj.r.squared sigma  
##       <dbl>         <dbl> <dbl>  
## 1     0.468         0.397 16.7
```

# Gov 50: 15. Multiple Regression and Interpretation

Matthew Blackwell

Harvard University

# Roadmap

1. Multiple regression
2. Categorical independent variables

# 1/ Multiple regression

# Multiple predictors

What if we want to predict  $Y$  as a function of many variables?

$$\text{seat\_change}_i = \alpha + \beta_1 \text{approval}_i + \beta_2 \text{rdi\_change}_i + \epsilon_i$$

Why?

- Better predictions (at least in-sample).
- Better interpretation as **ceteris paribus** relationships:
  - $\beta_1$  is the relationship between approval and seat\_change holding rdi\_change constant.
  - **Statistical control** in a cross-sectional study.

# Multiple regression in R

```
mult.fit <- lm(seat_change ~ approval + rdi_change,  
                data = midterms)  
  
mult.fit  
  
##  
## Call:  
## lm(formula = seat_change ~ approval + rdi_change, data = midterms)  
##  
## Coefficients:  
## (Intercept)      approval      rdi_change  
##       -117.23           1.53          3.22
```

- $\hat{\alpha} = -117.2$ : average seat change president has 0% approval and no change in income levels.
- $\hat{\beta}_1 = 1.53$ : average increase in seat change for additional percentage point of approval, **holding RDI change fixed**
- $\hat{\beta}_2 = 3.217$ : average increase in seat change for each additional percentage point increase of RDI, **holding approval fixed**

# Least squares with multiple regression

- How do we estimate the coefficients?
- The same exact way as before: minimize prediction error!
- Residuals (aka prediction error) with multiple predictors:

$$Y_i - \hat{Y}_i = \text{seat\_change}_i - \hat{\alpha} - \hat{\beta}_1 \text{approval}_i - \hat{\beta}_2 \text{rdi\_change}_i$$

- Find the coefficients that minimizes the **sum of the squared residuals**:

$$\text{SSR} = \sum_{i=1}^n \hat{\epsilon}_i^2 = (Y_i - \hat{\alpha} - \hat{\beta}_1 X_{i1} - \hat{\beta}_2 X_{i2})^2$$

# Model fit with multiple predictors

- $R^2$  mechanically increases when you add variables to the regression.
  - But this could be overfitting!!
- Solution: penalize regression models with more variables.
  - Occam's razor: **simpler models are preferred**
- Adjusted  $R^2$ : lowers regular  $R^2$  for each additional covariate.
  - If the added covariates doesn't help predict, adjusted  $R^2$  goes down!

# Comparing model fits

```
library(broom)
fit.app <- lm(seat_change ~ approval, data = midterms)
glance(fit.app) |>
  select(r.squared, adj.r.squared, sigma)
```

```
## # A tibble: 1 x 3
##   r.squared adj.r.squared sigma
##       <dbl>         <dbl> <dbl>
## 1     0.450        0.418  16.9
```

```
glance(mult.fit) |>
  select(r.squared, adj.r.squared, sigma)
```

```
## # A tibble: 1 x 3
##   r.squared adj.r.squared sigma
##       <dbl>         <dbl> <dbl>
## 1     0.468        0.397  16.7
```

# Predicted values from R

We could plug in values into the equation, but R can do this for us. The `{modelr}` package gives some functions that allow us to predictions in a tidy way:

Let's use `add_predictions()` to predict the 2022 results

```
library(modelr)

midterms |>
  filter(year == 2022) |>
  add_predictions(mult.fit)

## # A tibble: 1 x 7
##   year president party approval seat_change rdi_cha~1 pred
##   <dbl> <chr>     <chr>      <dbl>        <dbl>      <dbl> <dbl>
## 1 2022 Biden      D          42          NA     -0.003 -53.2
## # ... with abbreviated variable name 1: rdi_change
```

# Predictions from several models

The `gather_predictions()` will return one row for each model passed to it with the prediction for that model:

```
midterms |>  
  filter(year == 2022) |>  
  gather_predictions(fit.app, mult.fit)
```

```
## # A tibble: 2 x 8  
##   model     year presi~1 party appro~2 seat_~3 rdi_c~4   pred  
##   <chr>     <dbl> <chr>    <chr>    <dbl>    <dbl>    <dbl>  
## 1 fit.app    2022 Biden     D        42       NA   -0.003 -36.9  
## 2 mult.fit   2022 Biden     D        42       NA   -0.003 -53.2  
## # ... with abbreviated variable names 1: president,  
## #   2: approval, 3: seat_change, 4: rdi_change
```

# Predictions from new data

What about predicted values not in data?

```
tibble(approval = c(50, 75), rdi_change = 0) |>  
  gather_predictions(fit.app, mult.fit)
```

```
## # A tibble: 4 x 4  
##   model     approval rdi_change   pred  
##   <chr>      <dbl>       <dbl>   <dbl>  
## 1 fit.app      50          0 -25.6  
## 2 fit.app      75          0  9.92  
## 3 mult.fit     50          0 -40.9  
## 4 mult.fit     75          0 -2.79
```

# Predictions from augment()

We can also get predicted values from the `augment()` function using the `newdata` argument:

```
newdata <- tibble(approval = c(50, 75), rdi_change = 0)

augment(mult.fit, newdata = newdata)
```

```
## # A tibble: 2 x 3
##   approval rdi_change .fitted
##       <dbl>      <dbl>     <dbl>
## 1        50          0    -40.9
## 2        75          0    -2.79
```

# **2|** Categorical independent variables

# Political effects of gov't programs



- *Progesa*: Mexican conditional cash transfer program (CCT) from ~2000
  - Welfare \$\$ given if kids enrolled in schools, get regular check-ups, etc.
- Do these programs have political effects?
  - Program had support from most parties.
  - Was implemented in a nonpartisan fashion.
  - Would the incumbent presidential party be rewarded?

# The data

- Randomized roll-out of the CCT program:
  - treatment: receive CCT 21 months before 2000 election
  - control: receive CCT 6 months before 2000 election
- Does having CCT longer mobilize voters for incumbent PRI party?

| Name      | Description   |
|-----------|---|
| treatment | early Progresa (1) or late Progresa (0)                         |
| pri2000s  | PRI votes in the 2000 election as a share of adults in precinct |
| t2000     | turnout in the 2000 election as share of adults in precinct     |

```
library(qss)
data("progesa", package = "qss")
cct <- as_tibble(progesa) |>
  select(treatment, pri2000s, t2000)
cct
```

```
## # A tibble: 417 x 3
##   treatment pri2000s t2000
##       <int>     <dbl>  <dbl>
## 1         1      40.8   55.8
## 2         1      22.4   31.2
## 3         1      38.9   47.0
## 4         1      31.2   45.0
## 5         0      76.9  100
## 6         0      23.9   37.4
## 7         1      47.3   64.9
## 8         1      21.4   58.1
## 9         1      56.5   71.3
## 10        1      36.6   51.2
## # ... with 407 more rows
```

# Difference in means estimates

Does CCT affect turnout?

```
cct |> group_by(treatment) |>
  summarize(t2000 = mean(t2000)) |>
  pivot_wider(names_from = treatment, values_from = t2000) |>
  mutate(ATE = `1` - `0`)
```

```
## # A tibble: 1 × 3
##   `0`   `1`   ATE
##   <dbl> <dbl> <dbl>
## 1 63.8  68.1  4.27
```

Does CCT affect PRI (incumbent) votes?

```
cct |> group_by(treatment) |>
  summarize(pri2000s = mean(pri2000s)) |>
  pivot_wider(names_from = treatment, values_from = pri2000s) |>
  mutate(ATE = `1` - `0`)
```

```
## # A tibble: 1 × 3
##   `0`   `1`   ATE
##   <dbl> <dbl> <dbl>
## 1 34.5  38.1  3.62
```

# Binary independent variables

$$Y_i = \alpha + \beta X_i + \varepsilon_i$$

- When independent variable  $X_i$  is **binary**:
  - Intercept  $\hat{\alpha}$  is the average outcome in the  $X = 0$  group.
  - Slope  $\hat{\beta}$  is the difference-in-means of  $Y$  between  $X = 1$  group and  $X = 0$  group.

$$\hat{\beta} = \bar{Y}_{\text{treated}} - \bar{Y}_{\text{control}}$$

- If there are other independent variables, this becomes the difference-in-means controlling for those covariates.

# Linear regression for experiments

- Under **randomization**, we can estimate the ATE with regression:

```
cct |> group_by(treatment) |>  
  summarize(pri2000s = mean(pri2000s)) |>  
  pivot_wider(names_from = treatment, values_from = pri2000s) |>  
  mutate(ATE = `1` - `0`)
```

```
## # A tibble: 1 x 3  
##       `0`     `1`   ATE  
##     <dbl> <dbl> <dbl>  
## 1    34.5   38.1  3.62
```

```
lm(pri2000s ~ treatment, data = cct) |> coef()
```

```
## (Intercept)  treatment  
##        34.49          3.62
```

# Categorical variables in regression

- We often have **categorical variables**:
  - Race/ethnicity: white, Black, Latino, Asian.
  - Partisanship: Democrat, Republican, Independent
- Strategy for including in a regression: create a **series of binary variables**

| Unit | Party       | Democrat | Republican | Independent |
|------|-------------|----------|------------|-------------|
| 1    | Democrat    | 1        | 0          | 0           |
| 2    | Democrat    | 1        | 0          | 0           |
| 3    | Independent | 0        | 0          | 1           |
| 4    | Republican  | 0        | 1          | 0           |
| ⋮    | ⋮           | ⋮        | ⋮          | ⋮           |

- Then include **all but one** of these binary variables:

$$\text{turnout}_i = \alpha + \beta_1 \text{Republican}_i + \beta_2 \text{Independent}_i + \varepsilon_i$$

# Interpreting categorical variables

$$\text{turnout}_i = \alpha + \beta_1 \text{Republican}_i + \beta_2 \text{Independent}_i + \varepsilon_i$$

- $\hat{\alpha}$ : average outcome in the **omitted group/baseline** (Democrats).
- $\hat{\beta}$  coefficients: average difference between each group and the baseline.
  - $\hat{\beta}_1$ : average difference in turnout between Republicans and Democrats
  - $\hat{\beta}_2$ : average difference in turnout between Independents and Democrats

# CCES data

```
library(gov50data)
cces_2020

## # A tibble: 51,551 x 6
##   gender race  educ          pid3  turno~1 pres_~2
##   <fct>  <fct> <fct>        <fct>    <dbl> <fct>
## 1 Male   White 2-year      Republ~     1 Donald~
## 2 Female White Post-grad  Democr~     NA <NA>
## 3 Female White 4-year     Indepe~     1 Joe Bi~
## 4 Female White 4-year     Democr~     1 Joe Bi~
## 5 Male   White 4-year     Indepe~     1 Other
## 6 Male   White Some college Republ~     1 Donald~
## 7 Male   Black Some college Not su~     NA <NA>
## 8 Female White Some college Indepe~     1 Donald~
## 9 Female White High school graduate Republ~     1 Donald~
## 10 Female White 4-year    Democr~     1 Joe Bi~
## # ... with 51,541 more rows, and abbreviated variable names
## #   1: turnout_self, 2: pres_vote
```

# Categorical variables in the CCES data

```
turnout_pred <- lm(turnout_self ~ pid3, data = cces_2020)
turnout_pred
```

```
##
## Call:
## lm(formula = turnout_self ~ pid3, data = cces_2020)
##
## Coefficients:
## (Intercept)    pid3Republican    pid3Independent
##          0.9635            -0.0103            -0.0394
## pid3Other      pid3Not sure
##         -0.0066            -0.3331
```

# What R does internally with factor variables in lm

```
cces_2020 |> drop_na(turnout_self, pid3) |> select(pid3) |> pull() |>  
head()
```

```
## [1] Republican  Independent Democrat    Independent  
## [5] Republican  Independent  
## 7 Levels: Democrat Republican Independent ... not asked
```

```
model.matrix(turnout_pred) |>  
head()
```

```
##   (Intercept) pid3Republican pid3Independent pid3Other  
## 1           1              1                0        0  
## 3           1              0                1        0  
## 4           1              0                0        0  
## 5           1              0                1        0  
## 6           1              1                0        0  
## 8           1              0                1        0  
##   pid3Not sure  
## 1           0  
## 3           0  
## 4           0  
## 5           0  
## 6           0
```

# Gov 50: 16. Sampling

Matthew Blackwell

Harvard University

# Roadmap

1. Sampling exercise
2. Sampling framework
3. Polls

# 1/ Sampling exercise

# Data on class years enrolled in Gov 50

```
library(gov50data)
class_years

## # A tibble: 122 x 1
##       year
##   <chr>
## 1 Senior
## 2 Junior
## 3 Sophomore
## 4 Junior
## 5 Graduate Year 2
## 6 Sophomore
## 7 Professional Year 2
## 8 First-Year
## 9 Sophomore
## 10 Junior
## # ... with 112 more rows
```

# What proportion of the class is first years?

```
class_years |>  
  count(year) |>  
  mutate(prop = n / nrow(class_years))
```

```
## # A tibble: 9 x 3  
##   year                 n     prop  
##   <chr>              <int>    <dbl>  
## 1 First-Year            25  0.205  
## 2 Graduate Year 1      2  0.0164  
## 3 Graduate Year 2      1  0.00820  
## 4 Junior                31  0.254  
## 5 Not Set               3  0.0246  
## 6 Professional Year 2   2  0.0164  
## 7 Senior                 14  0.115  
## 8 Sophomore              43  0.352  
## 9 Year 1, Semester 1     1  0.00820
```

# Let's take some samples!

We can use the `slice_sample()` function to take a random sample of rows of a tibble:

```
class_years |>  
  slice_sample(n = 5)
```

```
## # A tibble: 5 x 1  
##   year  
##   <chr>  
## 1 Sophomore  
## 2 Junior  
## 3 Junior  
## 4 Sophomore  
## 5 Sophomore
```

# Another sample

```
class_years |>  
  slice_sample(n = 5)
```

```
## # A tibble: 5 x 1  
##   year  
##   <chr>  
## 1 Junior  
## 2 Not Set  
## 3 First-Year  
## 4 First-Year  
## 5 Sophomore
```

# Sample proportion of first-years

```
class_years |>  
  slice_sample(n = 20) |>  
  summarize(fy_prop = mean(year == "First-Year"))
```

```
## # A tibble: 1 x 1  
##   fy_prop  
##     <dbl>  
##   1     0.15
```

# Repeated sampling

We sometimes want to draw multiple samples from a tibble. For this we can use `rep_slice_sample()` from the `infer` package:

```
library(infer)
class_years |>
  rep_slice_sample(n = 5, reps = 2)
```

```
## # A tibble: 10 x 2
## # Groups:   replicate [2]
##       replicate year
##       <int> <chr>
## 1 1     First-Year
## 2 1     Sophomore
## 3 1     First-Year
## 4 1     Sophomore
## 5 1     First-Year
## 6 2     Junior
## 7 2     First-Year
## 8 2     Sophomore
## 9 2     First-Year
## 10 2    Sophomore
```

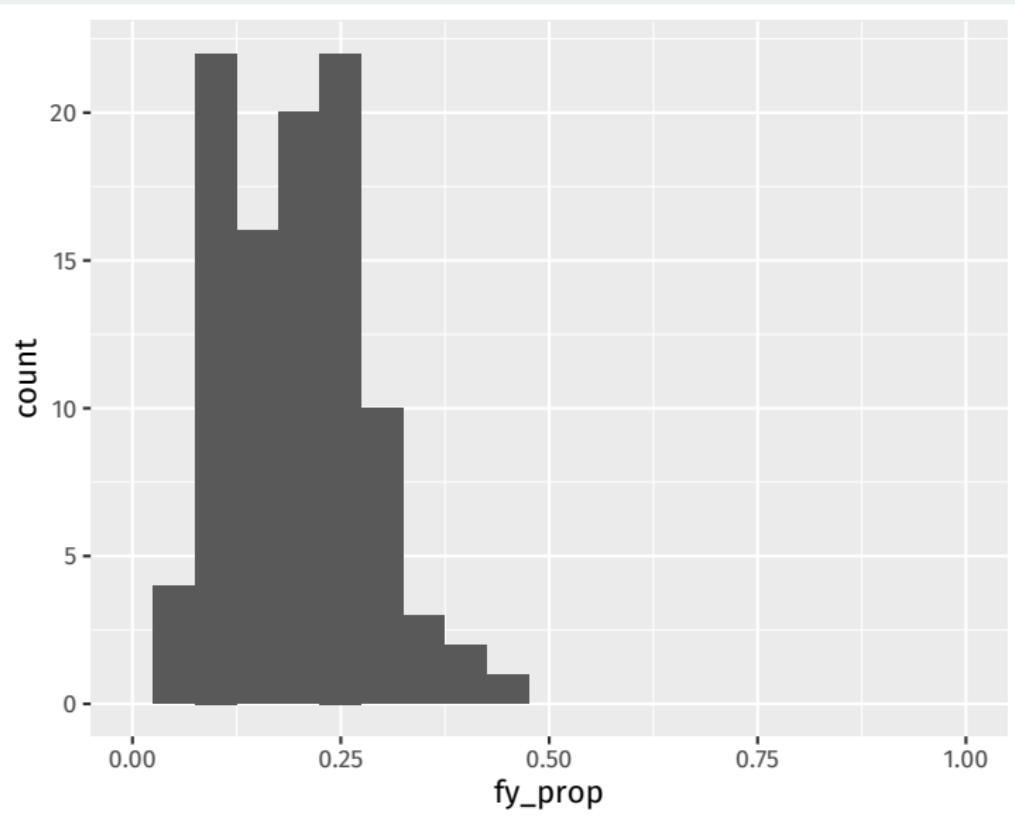
# Simulate many separate studies being done

```
samples_n20 <- class_years |>  
  rep_slice_sample(n = 20, reps = 100) |>  
  group_by(replicate) |>  
  summarize(fy_prop = mean(year == "First-Year"))  
samples_n20
```

```
## # A tibble: 100 x 2  
##   replicate fy_prop  
##       <int>    <dbl>  
## 1          1    0.25  
## 2          2    0.40  
## 3          3    0.30  
## 4          4    0.40  
## 5          5    0.20  
## 6          6    0.25  
## 7          7    0.10  
## 8          8    0.25  
## 9          9    0.35  
## 10        10    0.10  
## # ... with 90 more rows
```

# Distribution of these proportions

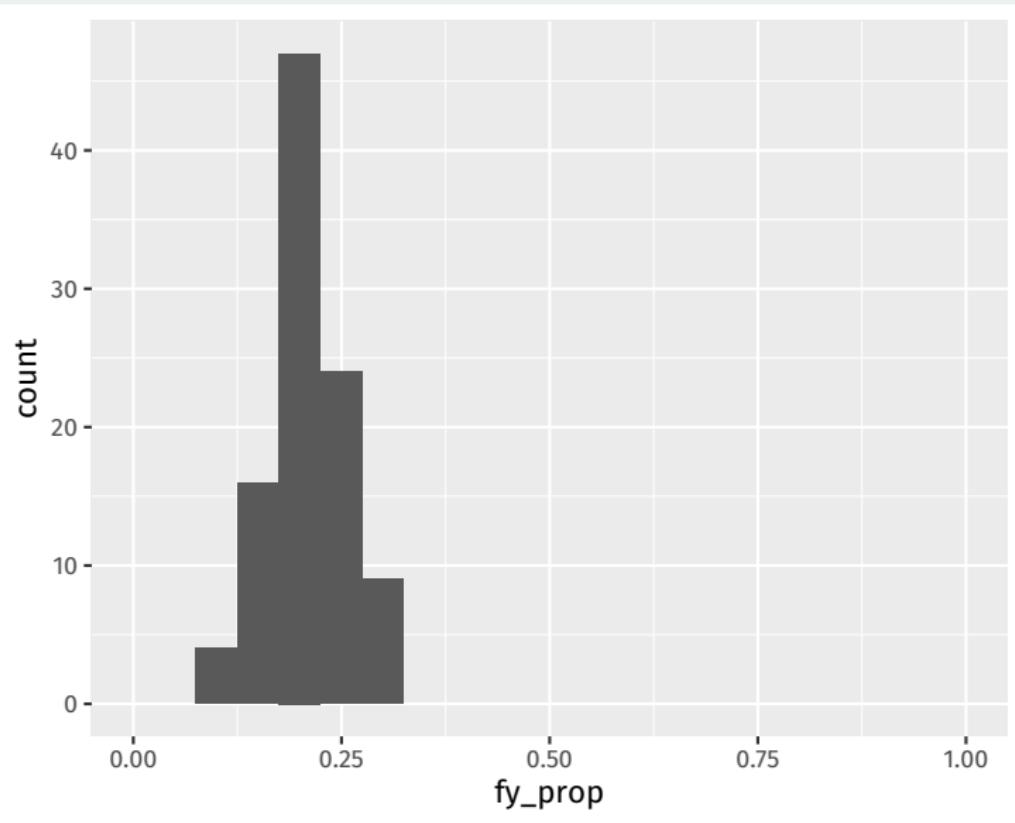
```
samples_n20 |>  
  ggplot(mapping = aes(x = fy_prop)) +  
  geom_histogram(binwidth=0.05) +  
  lims(x = c(0, 1))
```



# What if the sample sizes are bigger?

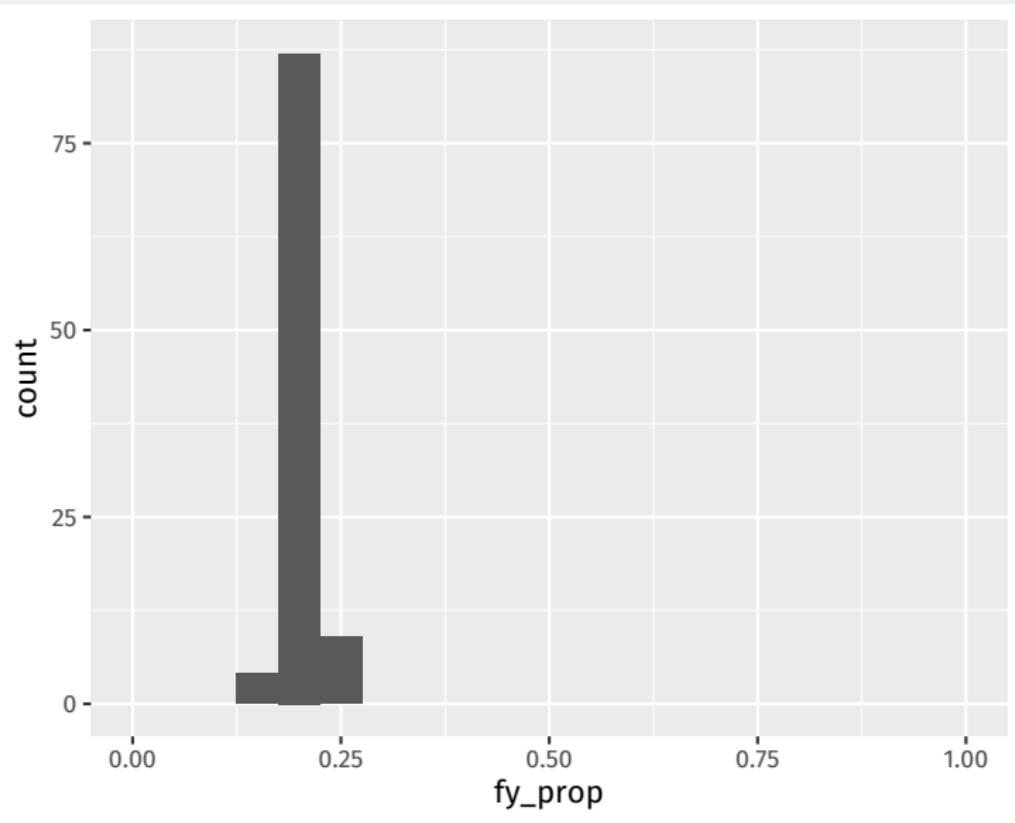
```
samples_n50 <- class_years |>  
  rep_slice_sample(n = 50, reps = 100) |>  
  group_by(replicate) |>  
  summarize(fy_prop = mean(year == "First-Year"))

samples_n50 |>  
  ggplot(mapping = aes(x = fy_prop)) +  
  geom_histogram(binwidth=0.05) +  
  lims(x = c(0, 1))
```



# What if the sample sizes are bigger?

```
samples_n100 <- class_years |>  
  rep_slice_sample(n = 100, reps = 100) |>  
  group_by(replicate) |>  
  summarize(fy_prop = mean(year == "First-Year"))  
  
samples_n100 |>  
  ggplot(mapping = aes(x = fy_prop)) +  
  geom_histogram(binwidth=0.05) +  
  lims(x = c(0, 1))
```



# Sample size and variability across samples

```
samples_n20 |>
  summarize(sd(fy_prop)) |> pull()

## [1] 0.0849

samples_n50 |>
  summarize(prop_sd = sd(fy_prop)) |> pull()

## [1] 0.0427

samples_n100 |>
  summarize(prop_sd = sd(fy_prop)) |> pull()

## [1] 0.0147
```

# 2/ Sampling framework

# Populations

**Population:** group of units/people we want to learn about.

**Population parameter:** some numerical summary of the population we would like to know. - population mean/proportion, population standard deviation.

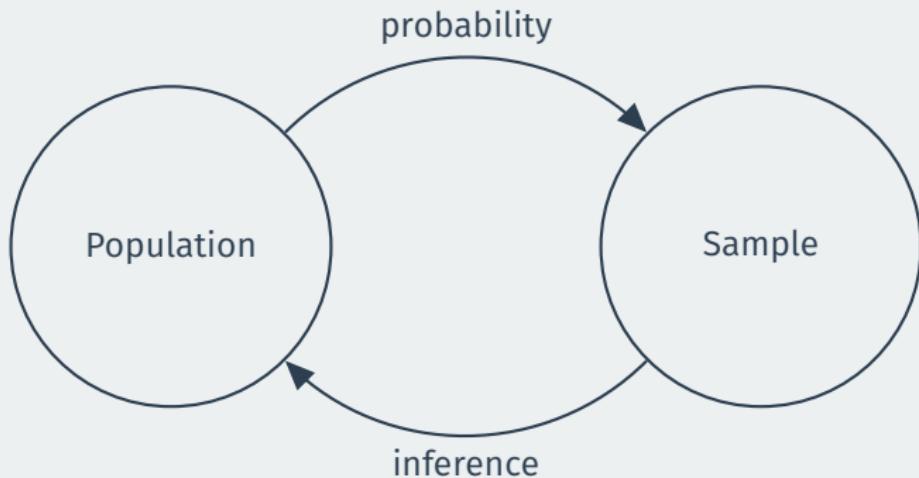
**Census:** complete recording of data on the entire population.

# Samples

**Sample:** subset of the population taken in some way (hopefully randomly).

**Estimator or sample statistic:** numerical summary of the sample that is our “best guess” for the unknown population parameter.

# Sampling framework



# Sampling at random

**Random sample:** units selected into sample from population with a non-zero probability.

**Simple random sample:** all units have the same probability of being selected into the sample.

# Our sampling exercise

- **Population:** all students enrolled in Gov 50.
- **Population parameter:** population proportion of first-years enrolled in Gov 50
  - Population proportions often denoted  $p$
- **Sample:** simple random sample of different sizes.
- **Sample statistic/estimator:** sample proportion of first-years
  - Estimators often denoted with a hat:  $\hat{p}$
  - We saw the  $\hat{p}$  varies with the random sample taken.

# Expected value

The **expected value** of a sample statistic,  $\mathbb{E}[\hat{p}]$ , is the average value of the statistic across repeated samples.

```
samples_n100 |>  
  summarize(mean(fy_prop)) |> pull()  
  
## [1] 0.205
```

The **expected value** of a sample proportion from a simple random sample is equal to the population proportion,  $\mathbb{E}[\hat{p}] = p$

# Standard error

The **standard error** is the standard deviation of the sample statistic across repeated samples.

```
samples_n100 |>  
  summarize(sd(fy_prop)) |> pull()
```

```
## [1] 0.0147
```

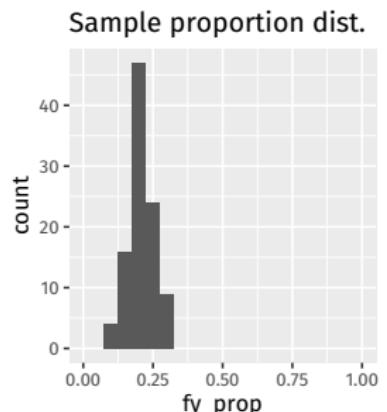
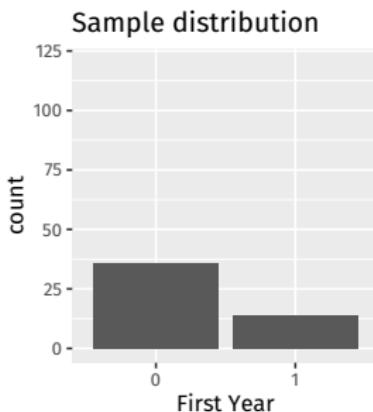
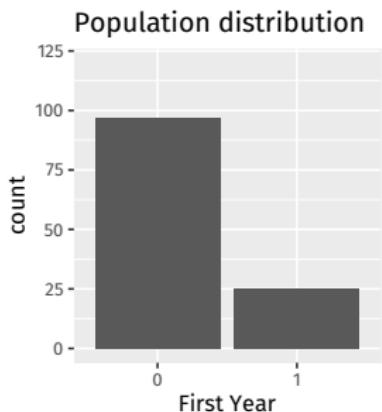
Tells us how far away, on average, the sample proportion will be from the population proportion.

# Standard error vs population standard deviation

The **standard error** is the SD of the statistic across repeated samples.

Should not be confused with the population standard deviation or sample standard deviation, both of which measure how far **units** are away from a mean.

# The three distributions



# 3/ Polls

# How popular is Joe Biden?

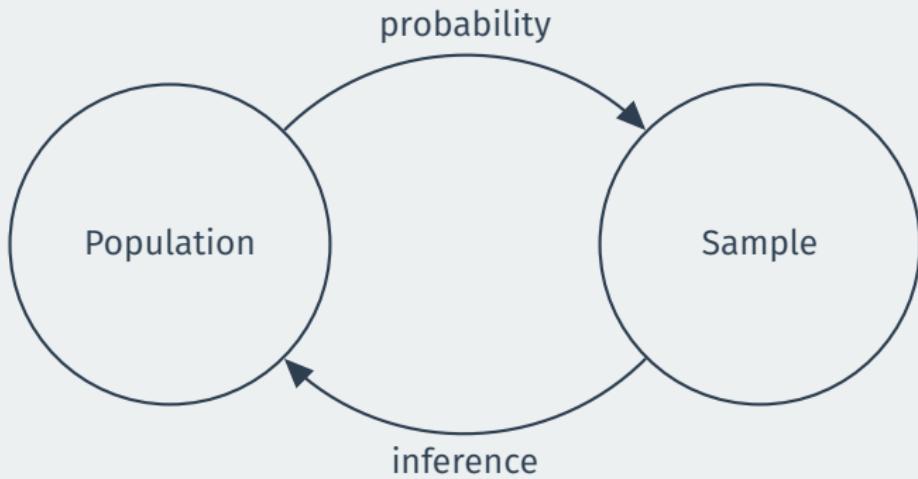


- What proportion of the public approves of Biden's job as president?
- Latest Gallup poll:
  - Sept 1st-16th
  - 812 adult Americans
  - Telephone interviews
  - Approve (42%), Disapprove (56%)

# Poll in our framework

- **Population:** adults 18+ living in 50 US states and DC.
- **Population parameter:** population proportion of all US adults that approve of Biden.
  - Census: not possible.
- **Sample:** random digit dialing phone numbers (cell and landline).
- **Point estimate:** sample proportion that approve of Biden

# Where are we going?



We only get 1 sample. Can we learn about the population from that sample?

# Gov 50: 17. Sampling Distributions

Matthew Blackwell

Harvard University

# Roadmap

1. Poll example
2. Random variables and probability distributions
3. Sampling distribution
4. Normal variables and the Central Limit Theorem

# 1/ Poll example

# How popular is Joe Biden?



- What proportion of the public approves of Biden's job as president?
- Latest Gallup poll:
  - Sept 1st-16th
  - 812 adult Americans
  - Telephone interviews
  - Approve (42%), Disapprove (56%)

# Poll in our framework

- **Population:** adults 18+ living in 50 US states and DC.
- **Population parameter:** population proportion of all US adults that approve of Biden.
  - Census: not possible.
- **Sample:** random digit dialing phone numbers (cell and landline).
- **Point estimate:** sample proportion that approve of Biden

# **2/** Random variables and probability distributions

# Random variables

**Random variables** are numerical summaries of chance processes:

$$X_i = \begin{cases} 1 & \text{if respondent } i \text{ supports Biden,} \\ 0 & \text{otherwise} \end{cases}$$

With a simple random sample, chance of  $X_i = 1$  is equal to the population proportion of people that support Biden.

# Types of random variables

- **Discrete:**  $X$  can take a finite (or countably infinite) number of values.
  - Number of heads in 5 coin flips
  - Sampled senator is a woman ( $X = 1$ ) or not ( $X = 0$ )
  - Number of battle deaths in a civil war
- **Continuous:**  $X$  can take any real value (usually within an interval).
  - GDP per capita (average income) in a country.
  - Share of population that approves of Biden.
  - Amount of time spent on a website.

# Probability distributions

**Probability distributions** tell us the chances of different values of a r.v. occurring

**Discrete variables:** like a frequency barplot for the population distribution.

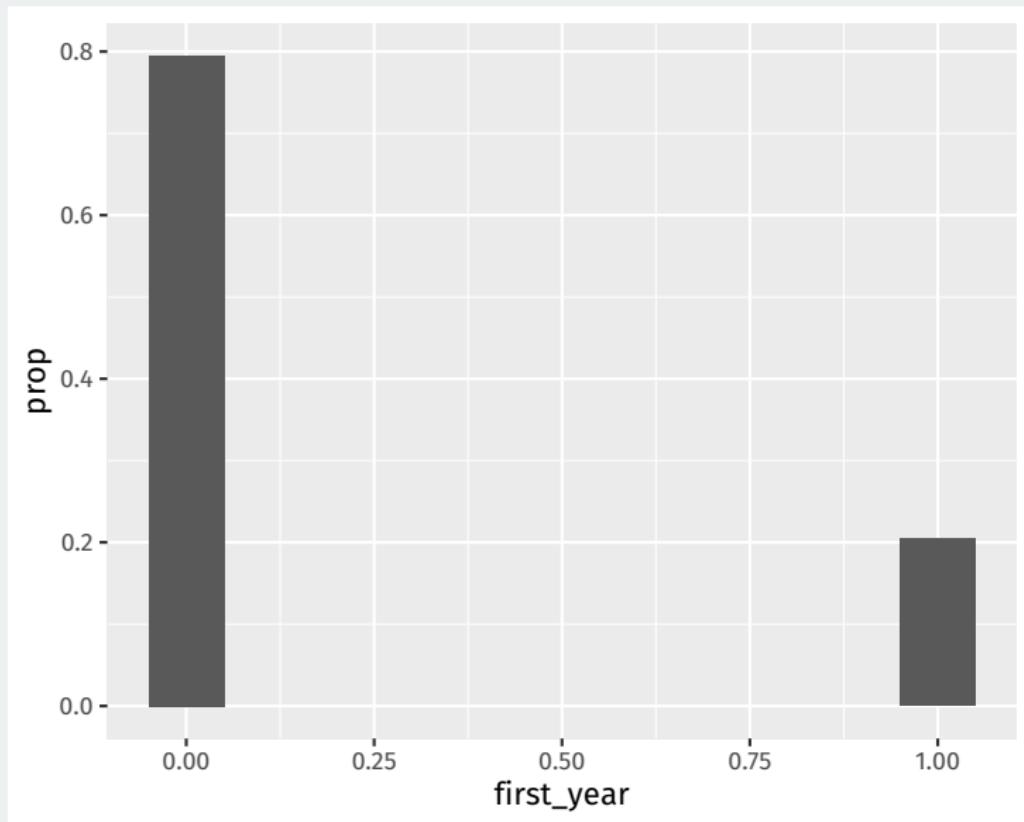
**Continuous variables:** like a continuous version of population histogram.

# Discrete probability distribution

We can use the `y = ..prop..` aesthetic to get a barplot with proportions instead of count to show us the chance/probability of selecting a first-year student:

```
library(gov50data)
class_years |>
  mutate(first_year = as.numeric(year == "First-Year")) |>
  ggplot(aes(x = first_year)) +
  geom_bar(mapping = aes(y = ..prop..), width = 0.1)
```

# Discrete probability distribution



# Midwest data

```
library(ggplot2)
midwest

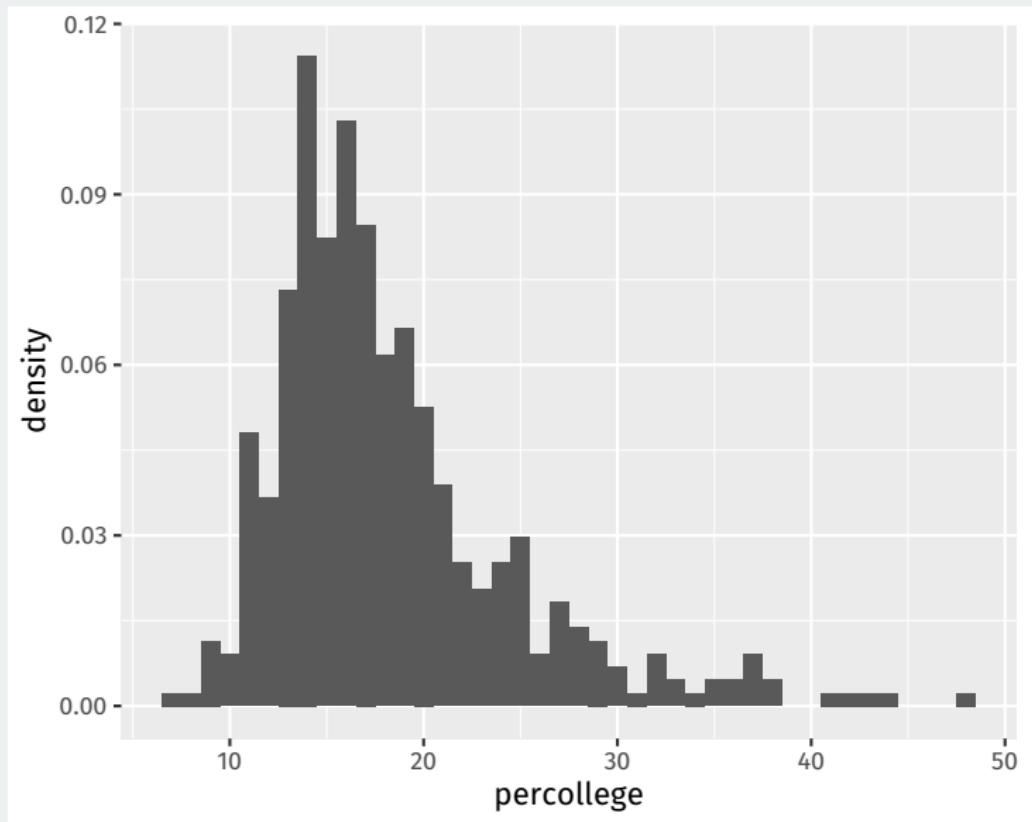
## # A tibble: 437 x 28
##       PID county state   area popto~1 popde~2 popwh~3 popbl~4
##       <int> <chr>   <chr> <dbl>    <int>    <dbl>    <int>    <int>
## 1      561 ADAMS    IL  0.052    66090    1271.   63917    1702
## 2      562 ALEXAN~  IL  0.014    10626     759    7054    3496
## 3      563 BOND     IL  0.022    14991    681.   14477    429
## 4      564 BOONE    IL  0.017    30806    1812.   29344    127
## 5      565 BROWN    IL  0.018    5836     324.   5264    547
## 6      566 BUREAU   IL  0.05     35688    714.   35157     50
## 7      567 CALHOUN  IL  0.017    5322     313.   5298     1
## 8      568 CARROLL  IL  0.027    16805    622.   16519    111
## 9      569 CASS     IL  0.024    13437    560.   13384    16
## 10     570 CHAMPA~  IL  0.058    173025   2983.  146506   16559
## # ... with 427 more rows, 20 more variables:
## #   popamerindian <int>, popasian <int>, popother <int>,
## #   percwhite <dbl>, percblack <dbl>, percamerindan <dbl>,
## #   percasian <dbl>, percother <dbl>, popadults <int>,
## #   perchsd <dbl>, percollege <dbl>, percprof <dbl>,
## #   poppovertyknown <int>, percpovertyknown <dbl>,
## #   ... and 10 more variables:
```

# Continuous probability distribution

We can use the `y = ..density..` to create a **density histogram** instead of a count histogram so that the area of the histogram boxes are equal to the chance of randomly selecting a unit in that bin:

```
midwest |>  
  ggplot(aes(x = percollege)) +  
  geom_histogram(aes(y = ..density..), binwidth = 1)
```

# Continuous probability distribution



# Why density?

Histograms with **density** on the y-axis are drawn so that the area of each box is equal to the proportion of units in the sample in that horizontal bin.

Easier to compare distributions across sample sizes.

Sum up all the area = 1 (but heights can go above 1)

# 3/ Sampling distribution

# Key properties of sums and means

Suppose  $X_1, X_2, \dots, X_n$  is a simple random sample from a population distribution with mean  $\mu$  ("mu") and variance  $\sigma^2$  ("sigma squared")

**Sample mean:**  $\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$

$$\bar{X}_n = \frac{X_1 + X_2 + \cdots + X_n}{n}$$

...

$\bar{X}_n$  is a random variable with a distribution!!

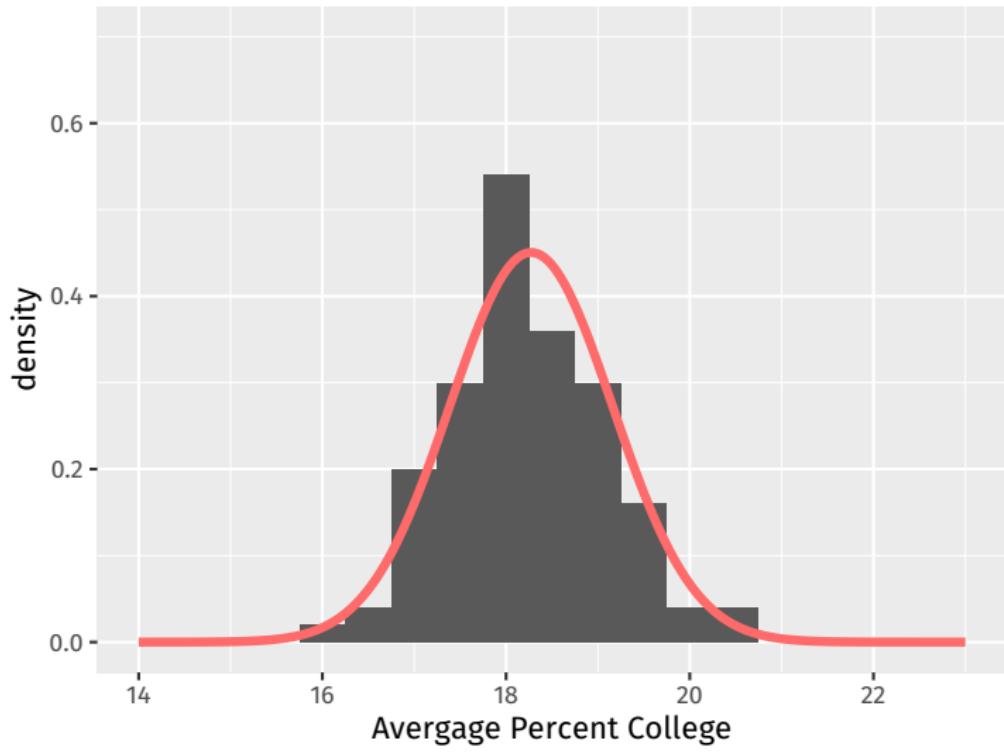
# Sample means/proportions distribution

**Sampling distributions** are the probability distributions of an estimator like  $\bar{X}_n$

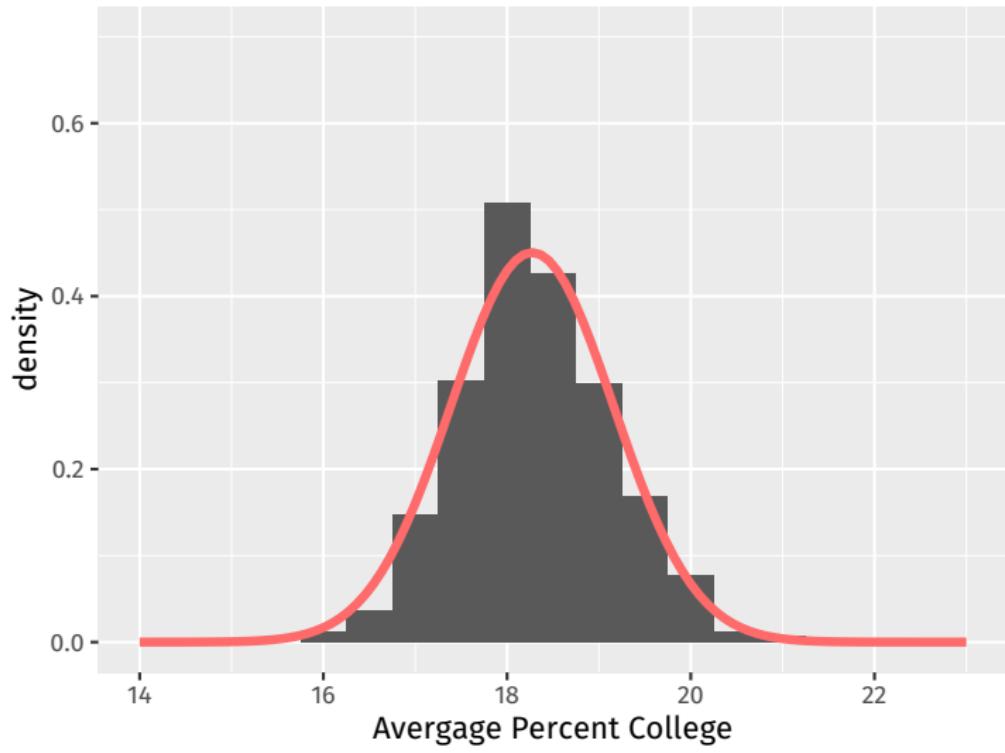
When we have access to the full population, we can approximate the sampling distribution with repeated sampling.

```
library(infer)
midwest |>
  rep_slice_sample(n = 50, reps = 100) |>
  group_by(replicate) |>
  summarize(`Avergage Percent College` = mean(percollege)) |>
  ggplot(aes(x = `Avergage Percent College`)) +
  geom_histogram(mapping = aes(y = ..density..), binwidth = 0.5) +
  coord_cartesian(xlim = c(14, 23), ylim = c(0, 0.7)) +
  labs(title = "100 Repititions") +
  stat_function(fun = dnorm, args = c(mean(midwest$percollege), sd(midwest$
```

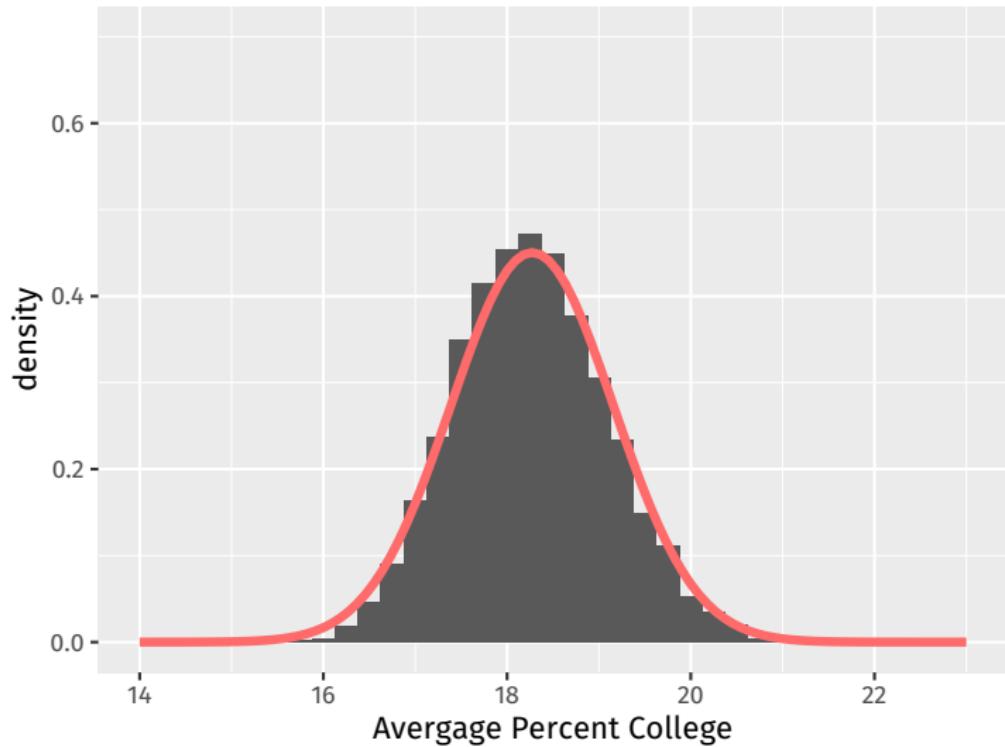
## 100 Repitions



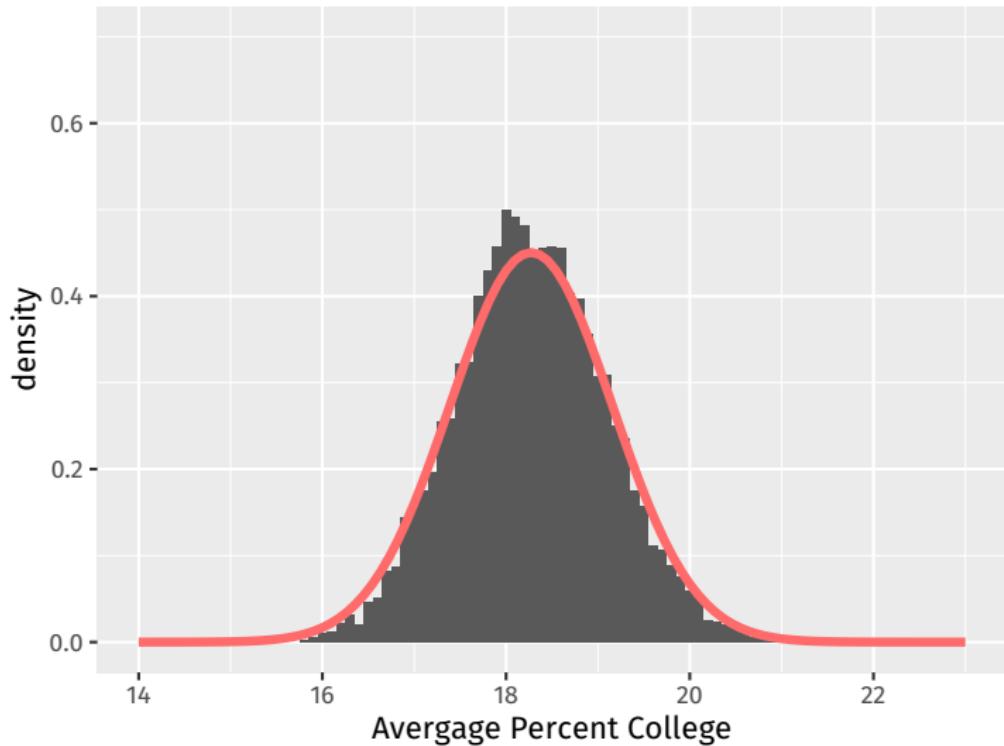
## 1,000 Repitions



10,000 Repitions



10,000 Repitions



# Sampling distribution of the sample mean

Suppose  $X_1, X_2, \dots, X_n$  is a simple random sample from a population distribution with mean  $\mu$  and variance  $\sigma^2$ .

**Expected value** of the distribution of  $\bar{X}_n$  is the population mean,  $\mu$ .

**Standard error** of the distribution of  $\bar{X}_n$  is approximately  $\sigma/\sqrt{n}$ :

$$SE \approx \frac{\text{population standard deviation}}{\sqrt{\text{sample size}}}$$

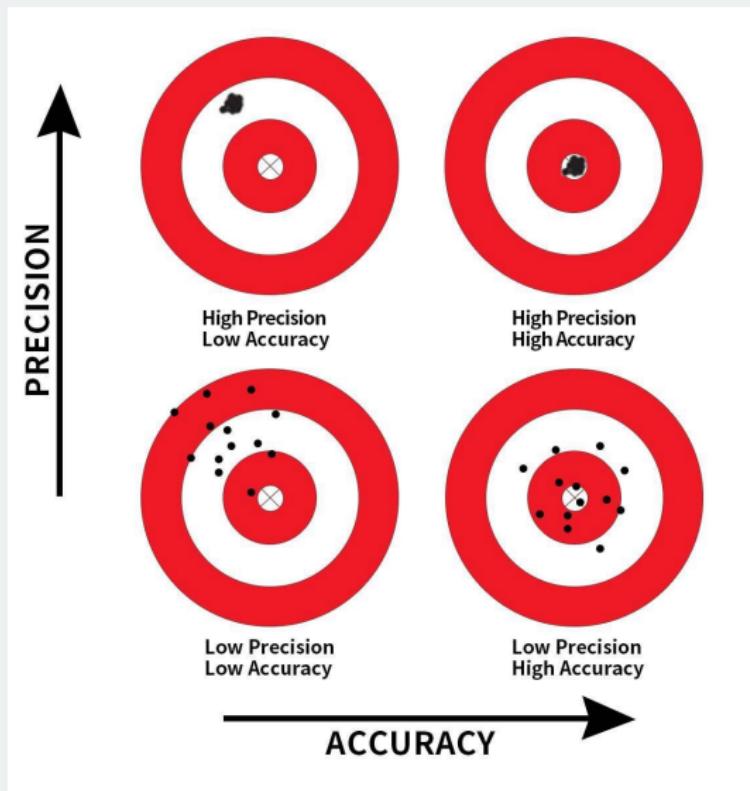
# Unbiasedness

An estimator is **unbiased** when its expected value across repeated samples equals the population parameter of interest.

Sample mean of a simple random sample is **unbiased** for the population mean,  $\mathbb{E}[\bar{X}_n] = \mu$

An estimator that isn't unbiased is called **biased**.

# Precision vs accuracy



# Law of large numbers

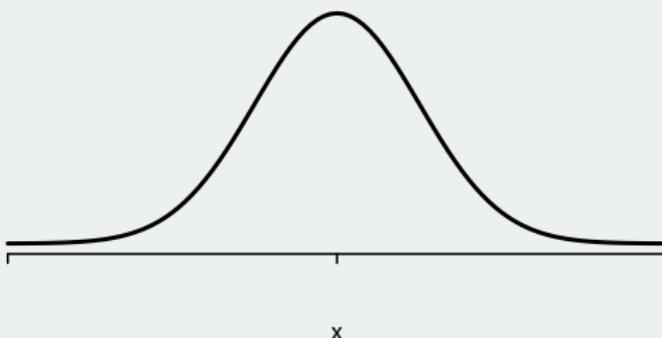
## Law of large numbers

Let  $X_1, \dots, X_n$  be a simple random sample from a population with mean  $\mu$  and finite variance  $\sigma^2$ . Then,  $\bar{X}_n$  converges to  $\mu$  as  $n$  gets large.

- Probability of  $\bar{X}_n$  being “far away” from  $\mu$  goes to 0 as  $n$  gets big.
- The distribution of sample mean “collapses” to population mean.
- Can see this from the SE of  $\bar{X}_n$ :  $SE = \sigma/\sqrt{n}$ .
- Not necessarily true with a biased sample!

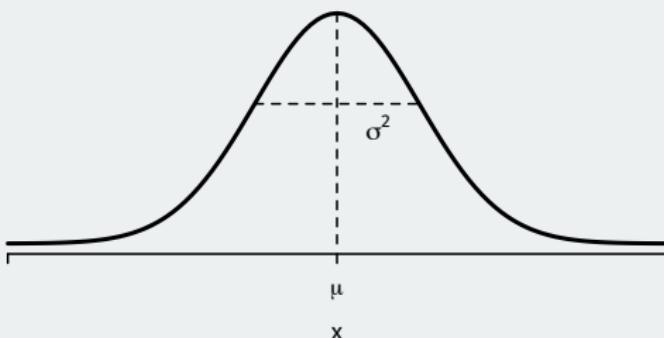
# 4| Normal variables and the Central Limit Theorem

# Normal random variable



- A **normal distribution** has a PDF that is the classic “bell-shaped” curve.
  - Extremely ubiquitous in statistics.
  - An r.v. is more likely to be in the center, rather than the tails.
- Three key properties of this PDF:
  - **Unimodal:** one peak at the mean.
  - **Symmetric** around the mean.
  - **Everywhere positive:** any real value can possibly occur.

# Normal distribution



- A normal distribution can be affected by two values:
  - **mean/expected value** usually written as  $\mu$
  - **variance** written as  $\sigma^2$  (standard deviation is  $\sigma$ )
  - Written  $X \sim N(\mu, \sigma^2)$ .
- **Standard normal distribution:** mean 0 and standard deviation 1.

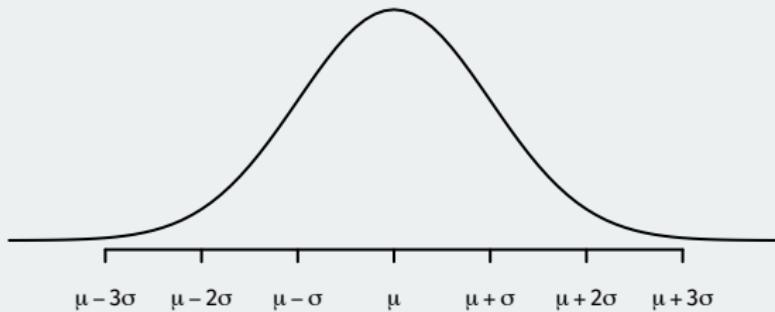
# Central limit theorem

## Central limit theorem

Let  $X_1, \dots, X_n$  be a simple random sample from a population with mean  $\mu$  and finite variance  $\sigma^2$ . Then,  $\bar{X}_n$  will be approximately distributed  $N(\mu, \sigma^2/n)$  in large samples.

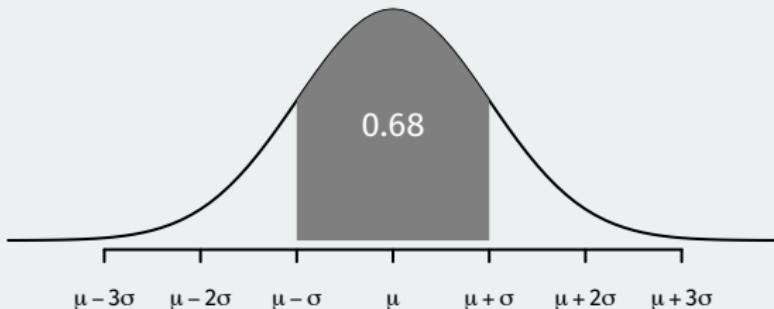
- “Sample means tend to be normally distributed as samples get large.”
- $\rightsquigarrow$  we know (an approx. of) the entire probability distribution of  $\bar{X}_n$ 
  - Approximation is better as  $n$  goes up.
  - Does not depend on the distribution of  $X_i$ !

# Empirical Rule for the Normal Distribution



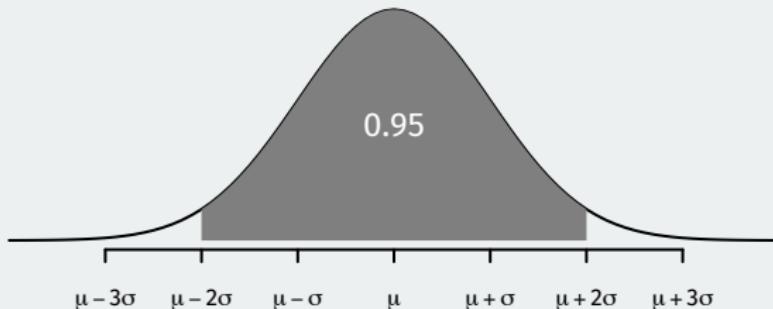
- If  $X \sim N(\mu, \sigma^2)$ , then:

# Empirical Rule for the Normal Distribution



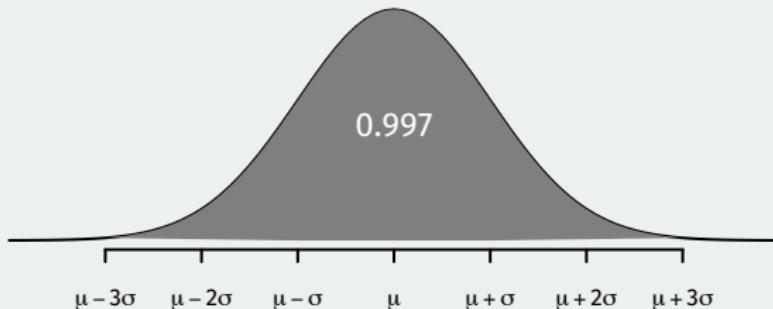
- If  $X \sim N(\mu, \sigma^2)$ , then:
  - $\approx 68\%$  of the distribution of  $X$  is within 1 SD of the mean.

# Empirical Rule for the Normal Distribution



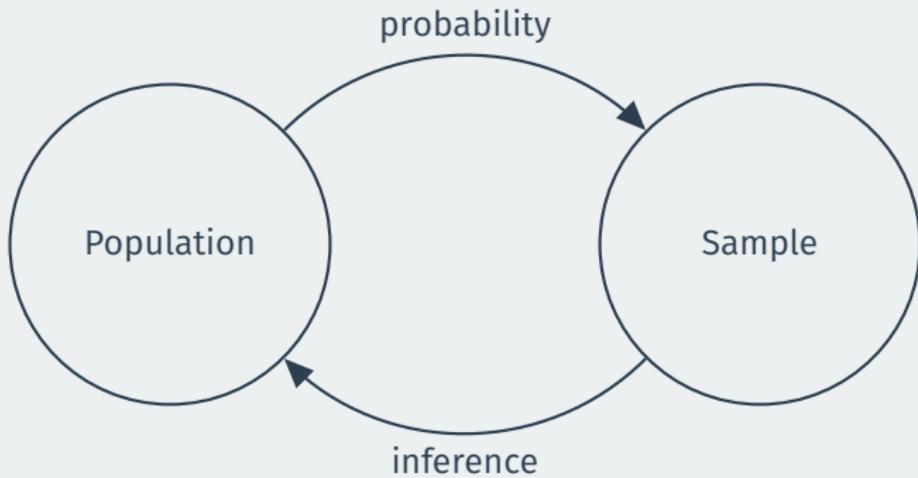
- If  $X \sim N(\mu, \sigma^2)$ , then:
  - $\approx 68\%$  of the distribution of  $X$  is within 1 SD of the mean.
  - $\approx 95\%$  of the distribution of  $X$  is within 2 SDs of the mean.

# Empirical Rule for the Normal Distribution



- If  $X \sim N(\mu, \sigma^2)$ , then:
  - $\approx 68\%$  of the distribution of  $X$  is within 1 SD of the mean.
  - $\approx 95\%$  of the distribution of  $X$  is within 2 SDs of the mean.
  - $\approx 99.7\%$  of the distribution of  $X$  is within 3 SDs of the mean.
- CLT + empirical rule: we'll know the rough distribution of estimation errors we should expect.

# Where are we going?



We only get 1 sample. Can we learn about the population from that sample?

# Gov 50: 18. The Bootstrap

Matthew Blackwell

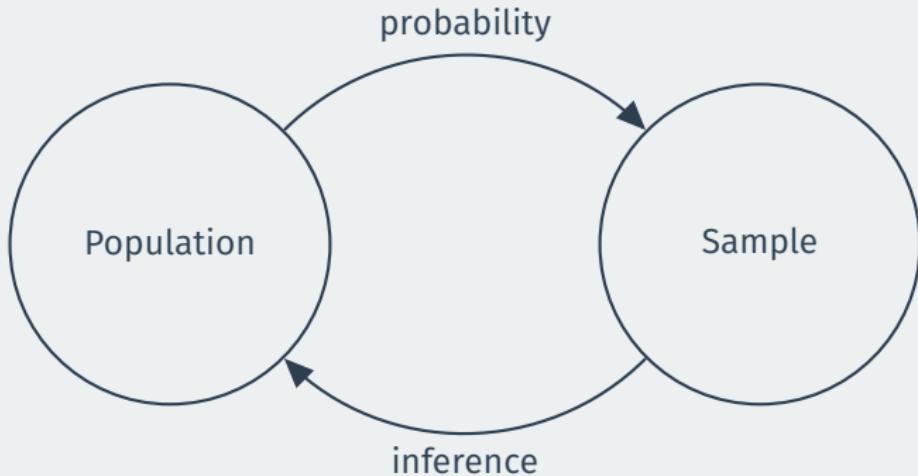
Harvard University

# Roadmap

1. Resampling from our sample
2. Confidence intervals
3. Calculating confidence intervals

# 1/ Resampling from our sample

# Where are we?



Can we approximate the **sampling distribution** with our single sample?

# American National Election Survey data

| Name             | Description  |
|------------------|--|
| state            | State of respondent  |
| district         | Congressional district of respondent                                   |
| pid7             | Party ID (1=Strong D, 7=Strong R)                                      |
| pres_vote        | Self reported vote in 2020   |
| sci_therm        | 0-100 therm score for scientists                                       |
| rural_therm      | 0-100 therm score for rural Americans                                  |
| favor_voter_id   | 1 if respondent thinks voter ID should be required                     |
| envir_doing_more | 1 if respondent thinks gov't should be doing more about climate change |

# ANES data

```
library(gov50data)
anes

## # A tibble: 5,162 x 8
##   state district pid7 pres_vote sci_therm rural_~1 favor~2
##   <chr>     <dbl> <dbl> <chr>          <dbl>    <dbl>    <dbl>
## 1 ID         2     4 Other           70      60      1
## 2 VA         2     3 Biden          100      75      0
## 3 CO         4     4 Trump          60      90      1
## 4 TX         5     3 Biden          85      85      1
## 5 WI         6     6 Trump          85      70      1
## 6 CA        40     2 Biden          50      50      1
## 7 WI         5     2 Biden          100     70      1
## 8 OR         4     7 Trump          70      50      0
## 9 MA         5     3 Biden          80      70      0
## 10 NV        3     1 Biden          85      40      0
## # ... with 5,152 more rows, 1 more variable:
## #   envir_doing_more <dbl>, and abbreviated variable names
## #   1: rural_therm, 2: favor_voter_id
```

# Sample statistic

What is the average thermometer score for scientists?

```
anes |>  
  summarize(mean(sci_therm))
```

```
## # A tibble: 1 x 1  
##   `mean(sci_therm)`  
##       <dbl>  
## 1      80.6
```

What is the sampling distribution of this average? We only have this 1 draw!

# Notation review

**Population:** all US adults.

**Population parameter:** average feeling thermometer score for scientists among all US adults.

**Sample:** (complicated) random sample of all US adults.

**Sample statistic/point estimate:** sample average of thermometer scores.

Roughly how far our point estimate is likely to be from the truth?

# The bootstrap

**Mimic** sampling from the population by **resampling** many times from the sample itself.

Bootstrap resampling done **with replacement** (same row can appear more than once)

# One bootstrap resample

```
boot_1 <- anes |>
  slice_sample(prop = 1, replace = TRUE)
boot_1

## # A tibble: 5,162 x 8
##   state district pid7 pres_vote sci_therm rural_~1 favor~2
##   <chr>     <dbl> <dbl> <chr>          <dbl>     <dbl>     <dbl>
## 1 CO         6     1 Biden        85      70       0
## 2 NY         8     1 Biden        85      70       1
## 3 SC         7     1 Biden       100     100       0
## 4 CO         3     4 Trump       85      85       1
## 5 CA        39     2 Biden       100      60       0
## 6 CA        37     3 Biden       90      65       0
## 7 AR         2     1 Biden       85      70       0
## 8 CO         6     1 Biden       90      70       0
## 9 WA         5     3 Biden       70      85       0
## 10 MI        7     3 Other       60      70       0
## # ... with 5,152 more rows, 1 more variable:
## #   envir_doing_more <dbl>, and abbreviated variable names
## #   1: rural_therm, 2: favor_voter_id
```

# Sample mean in the bootstrap sample

```
boot_1 |>  
  summarize(mean(sci_therm))
```

```
## # A tibble: 1 x 1  
##   `mean(sci_therm)`  
##             <dbl>  
## 1             81.0
```

# Many bootstrap samples

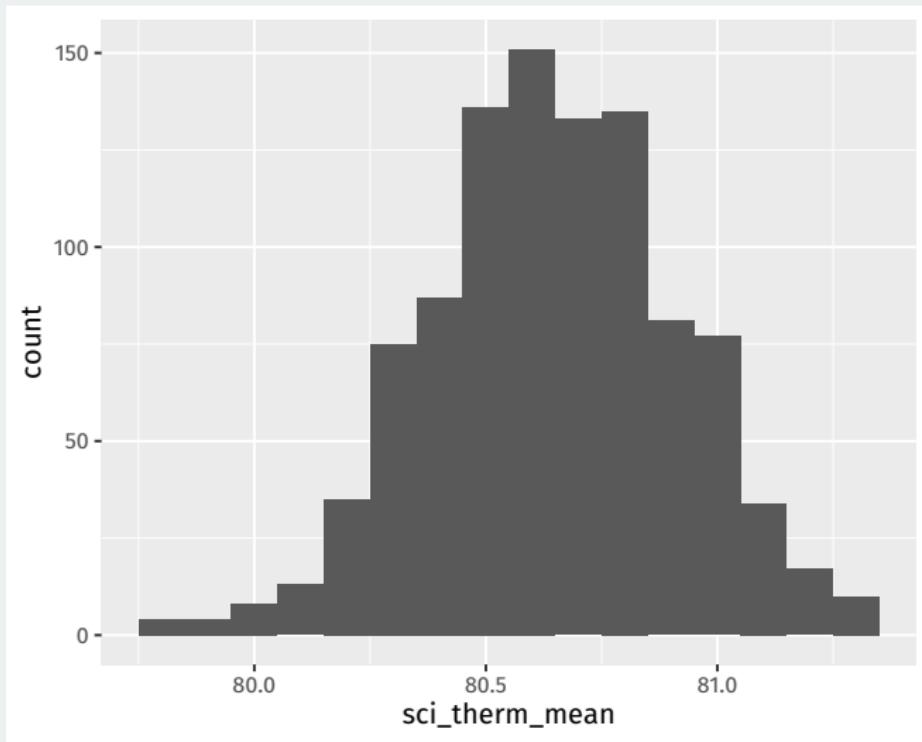
```
library(infer)
bootstrap_dist <- anes |>
  rep_slice_sample(prop = 1, reps = 1000, replace = TRUE) |>
  group_by(replicate) |>
  summarize(sci_therm_mean = mean(sci_therm))
bootstrap_dist
```

# Many bootstrap samples

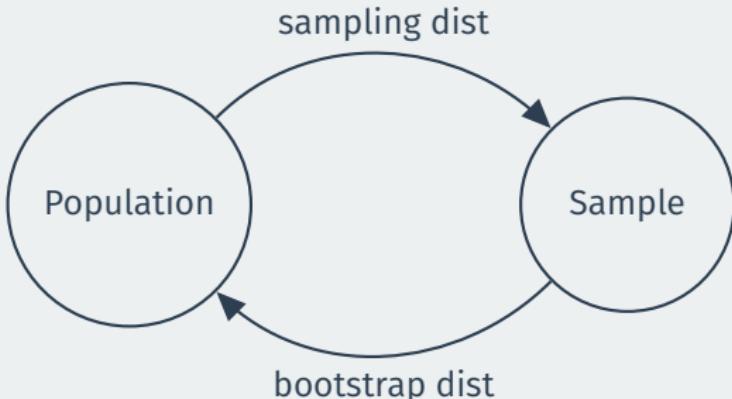
```
## # A tibble: 1,000 x 2
##   replicate sci_therm_mean
##       <int>          <dbl>
## 1        1          80.6
## 2        2          80.2
## 3        3          80.1
## 4        4          80.8
## 5        5          80.7
## 6        6          80.3
## 7        7          80.5
## 8        8          81.1
## 9        9          80.9
## 10      10          80.8
## # ... with 990 more rows
```

# Visualizing the bootstrap distribution

```
bootstrap_dist |>  
  ggplot(aes(x = sci_therm_mean)) + geom_histogram(binwidth = 0.1)
```



# Bootstrap distribution



Bootstrap distribution **approximates** the sampling distribution of the estimator.

Both should have a **similar shape and spread** if sampling from the distribution  $\approx$  bootstrap resampling.

Approximation gets better as sample gets bigger.

# Comparing to the point estimate

Given the sampling, not surprising that bootstrap distribution is centered on the point estimate:

```
bootstrap_dist |>  
  summarize(mean(sci_therm_mean))
```

```
## # A tibble: 1 × 1  
##   `mean(sci_therm_mean)`  
##                 <dbl>  
## 1                  80.6
```

```
anes |>  
  summarize(mean(sci_therm))
```

```
## # A tibble: 1 × 1  
##   `mean(sci_therm)`  
##                 <dbl>  
## 1                  80.6
```

# 2/ Confidence intervals

# What is a confidence interval?



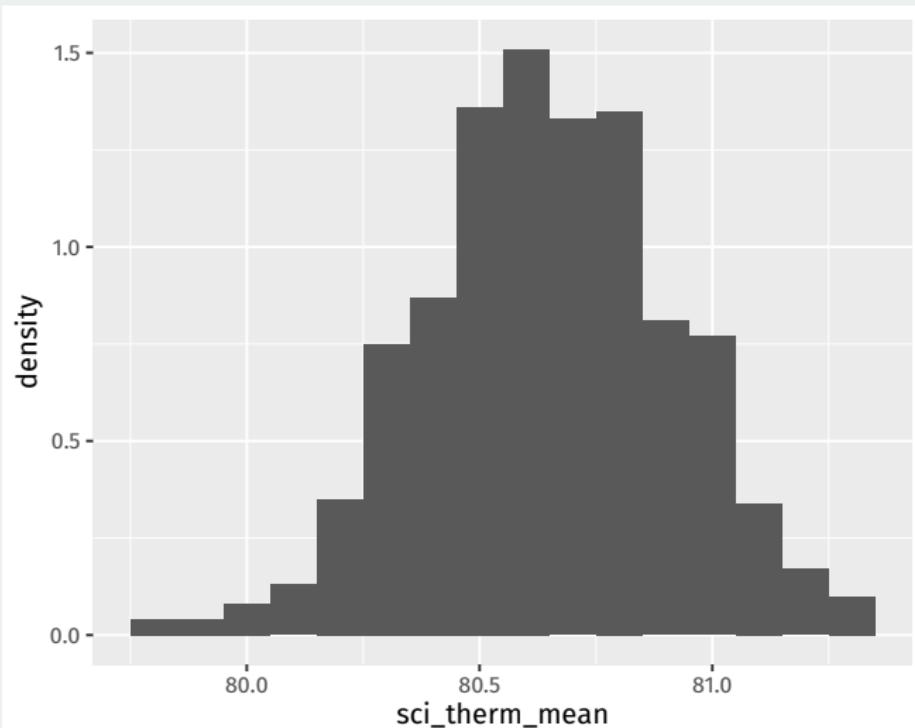
**Point estimate:** best single guess about the population parameter.  
Unlikely to be exactly correct.



**Confidence interval:** a range of plausible values of the population parameter.

# Where is most of the bootstrap distribution?

```
bootstrap_dist |>  
  ggplot(aes(x = sci_therm_mean)) +  
  geom_histogram(aes(y= ..density..), binwidth = 0.1)
```



# Confidence intervals



- Each sample gives a different CI or toss of the ring.
- Some samples the ring will contain the target (the CI will contain the truth) other times it won't.
  - We don't know if the CI for our sample contains the truth!
- **Confidence level:** percent of the time our CI will contain the population parameter.
  - Number of ring tosses that will hit the target.
  - We get to choose, but typical values are 90%, 95%, and 99%

# Confidence intervals as occasional liars

The **confidence level** of a CI determine how often the CI will be wrong.

A 95% confidence interval will:

- Tell you the truth in 95% of repeated samples (contain the population parameter 95% of the time)
- Lie to you in 5% of repeated sample (not contain the population parameter 5% of the time)

Can you tell if your particular confidence interval is telling the truth? No!

# Percentile method

**Percentile method:** find the middle 95% of the bootstrap distribution.

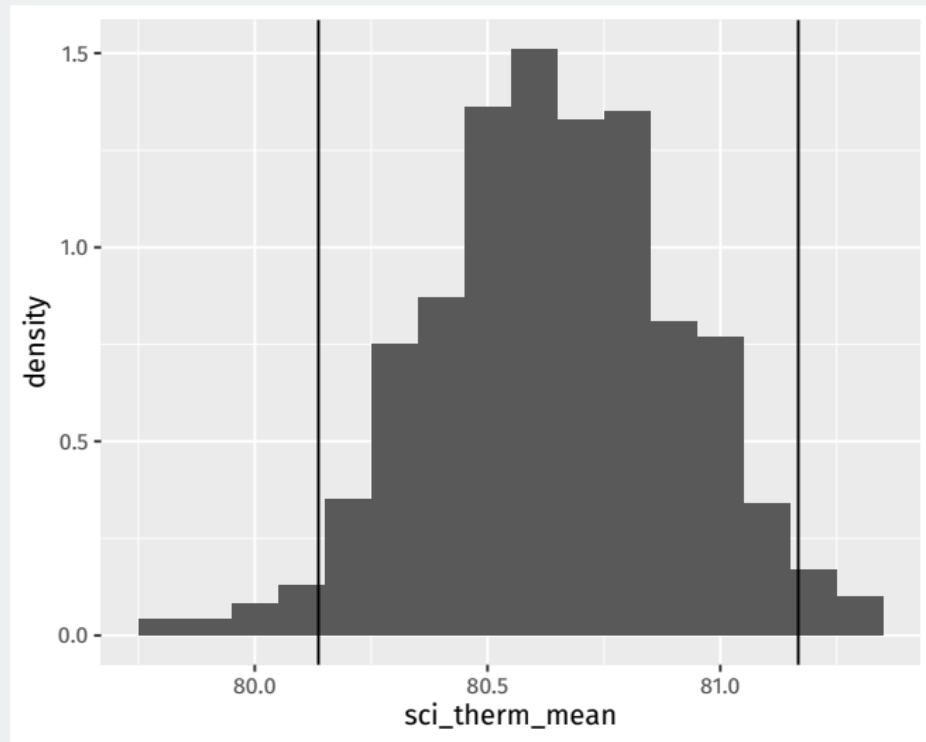
We can do this by finding the points that the 2.5th percentile and the 97.5th percentile.

```
perc_ci95 <- quantile(bootstrap_dist$sci_therm_mean,  
                      probs = c(0.025, 0.975))
```

```
perc_ci95
```

```
## 2.5% 97.5%  
## 80.1 81.2
```

# Visualizing the CI



# Width of the interval

What happens if we want the CI to be right more often? Will the width of a 99% confidence interval be wider or narrower?

# 99% confidence interval

For 99% CI we need to find the middle 99% of the bootstrap distribution.

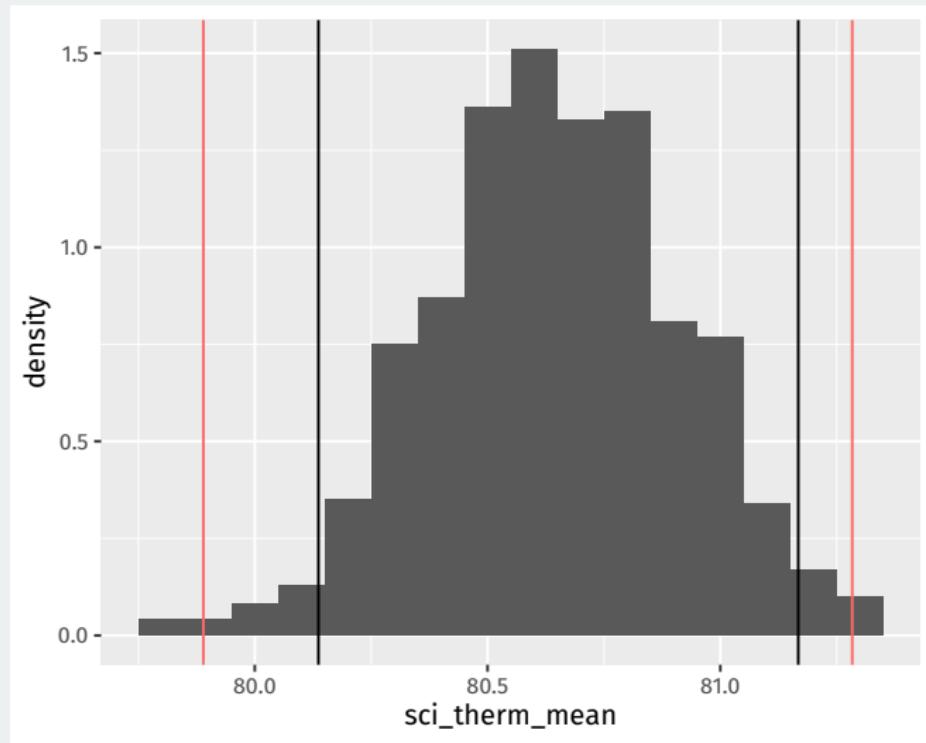
We can do this by finding the points that the 0.5th percentile and the 99.5th percentile.

```
perc_ci99 <- quantile(bootstrap_dist$sci_therm_mean,  
                      probs = c(0.005, 0.995))
```

```
perc_ci99
```

```
## 0.5% 99.5%  
## 79.9 81.3
```

# Visualizing the CIs



# **3/** Calculating confidence intervals

# infer package

Possible to use `quantile` to calculate CIs, but `infer` package is a more unified framework for CIs and hypothesis tests.

We'll use a `dplyr`-like approach of chained calls.

# Step 1: define an outcome of interest

Start with defining the variable of interest:

```
anes |>  
  specify(response = sci_therm)
```

```
## Response: sci_therm (numeric)  
## # A tibble: 5,162 x 1  
##   sci_therm  
##   <dbl>  
## 1 70  
## 2 100  
## 3 60  
## 4 85  
## 5 85  
## 6 50  
## 7 100  
## 8 70  
## 9 80  
## 10 85  
## # ... with 5,152 more rows
```

## Step 2: generate bootstraps

Next `infer` can generate bootstraps with the `generate()` function (similar to `rep_slice_sample()`):

```
anes |>
  specify(response = sci_therm) |>
  generate(reps = 1000, type = "bootstrap")
```

```
## Response: sci_therm (numeric)
## # A tibble: 5,162,000 x 2
## # Groups:   replicate [1,000]
##       replicate sci_therm
##   <int>     <dbl>
## 1 1         85
## 2 1         85
## 3 1         60
## 4 1         70
## 5 1         70
## 6 1         85
## 7 1         90
## 8 1        100
## 9 1         50
## 10 1        100
## # ... with 5,161,990 more rows
```

## Step 3: calculate sample statistics

Use `calculate()` to do the `group_by(replicate)` and `summarize` commands in one:

```
boot_dist_infer <- anes |>  
  specify(response = sci_therm) |>  
  generate(reps = 1000, type = "bootstrap") |>  
  calculate(stat = "mean")
```

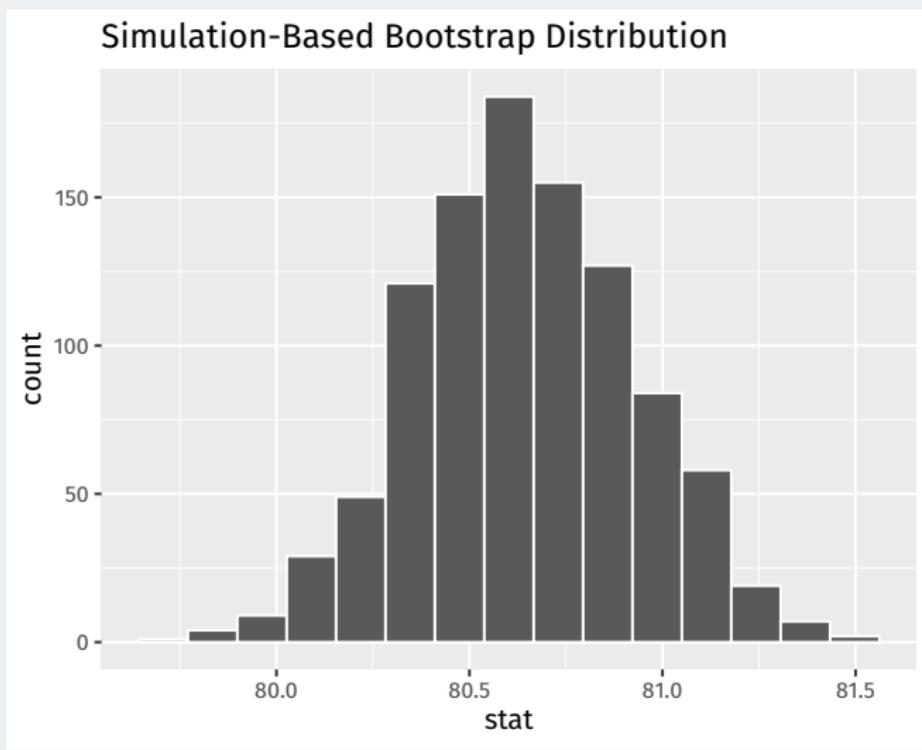
## boot\_dist\_infer

```
## Response: sci_therm (numeric)
## # A tibble: 1,000 x 2
##       replicate   stat
##           <int> <dbl>
## 1             1 80.7
## 2             2 80.8
## 3             3 80.5
## 4             4 80.9
## 5             5 80.4
## 6             6 81.2
## 7             7 81.0
## 8             8 80.7
## 9             9 80.5
## 10            10 80.4
## # ... with 990 more rows
```

## Step 3(b): visualize the bootstrap distribution

infer also has a shortcut for plotting called `visualize()`:

```
visualize(boot_dist_infer)
```



## Step 4: calculate CIs

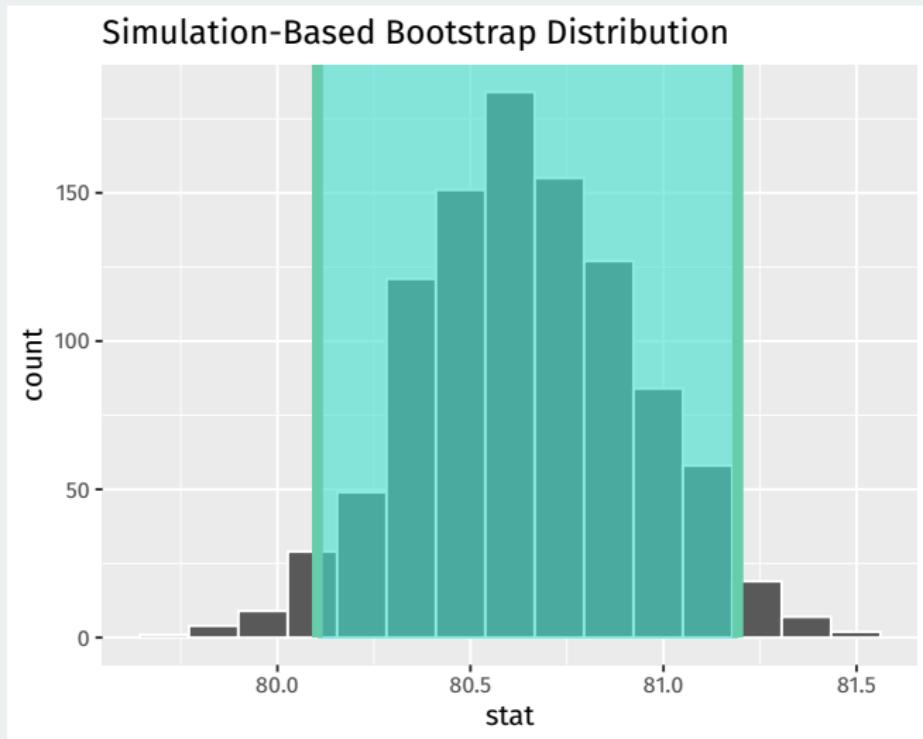
Finally we can calculate the CI using the percentile method with `get_confidence_interval()`:

```
perc_ci_95 <- boot_dist_infer |>  
  get_confidence_interval(level = 0.95, type = "percentile")  
perc_ci_95
```

```
## # A tibble: 1 x 2  
##   lower_ci upper_ci  
##       <dbl>    <dbl>  
## 1     80.1     81.2
```

## Step 4(b): visualize Cls

```
visualize(boot_dist_infer) +  
  shade_confidence_interval(endpoints = perc_ci_95)
```



# Gov 50: 19. More Confidence Intervals

Matthew Blackwell

Harvard University

# Roadmap

1. Bootstrap CIs for a difference in means
2. Bootstrap CIs for a difference in ATEs
3. Interpreting confidence intervals

# 1/ Bootstrap CIs for a difference in means

# Comparison between groups

- Last time: confidence intervals for means.
- More interesting to compare across groups.
  - Differences in public opinion across groups
  - Difference between treatment and control groups.
- Bedrock of causal inference!

# Trains experiment

- Back to the Boston trains example.
  - Boston commuter rail platform setting.
- Treatment group: presence of native Spanish-speaking confederates.
- Control group: no confederates.
- Outcome:  $X_i$ , change in views on immigration.
  - Sample average in the treated group,  $\bar{X}_T$
  - Sample average in the control group,  $\bar{X}_C$
- Estimated **average treatment effect**

$$\widehat{\text{ATE}} = \bar{X}_T - \bar{X}_C$$

# Inference for the difference

- Parameter: **population ATE**  $\mu_T - \mu_C$ 
  - $\mu_T$ : Average outcome in the population if everyone received treatment.
  - $\mu_C$ : Average outcome in the population if everyone received control.
- Difference-in-means estimator:  $\widehat{\text{ATE}} = \bar{X}_T - \bar{X}_C$
- $\bar{X}_T$  is a r.v. with mean  $\mathbb{E}[\bar{X}_T] = \mu_T$
- $\bar{X}_C$  is a r.v. with mean  $\mathbb{E}[\bar{X}_C] = \mu_C$
- $\rightsquigarrow \bar{X}_T - \bar{X}_C$  is a r.v. with mean  $\mu_T - \mu_C$ 
  - Sample difference in means is on average equal to the population difference in means.

# Trains data

```
library(gov50data)
trains

## # A tibble: 115 x 14
##       age   male income white college usborn treatment ideol~1
##     <dbl> <dbl> <dbl> <dbl>   <dbl>   <dbl>      <dbl>   <dbl>
## 1     31     0 135000     1     1     1         1     3
## 2     34     0 105000     1     1     0         1     4
## 3     63     1 135000     1     1     1         1     2
## 4     45     1 300000     1     1     1         1     4
## 5     55     1 135000     1     1     1         0     2
## 6     37     0  87500     1     1     1         1     5
## 7     53     0  87500     1     0     1         0     5
## 8     36     1 135000     1     1     1         1     4
## 9     54     0 105000     1     0     1         0     3
## 10    42     1 135000     1     1     1         1     4
## # ... with 105 more rows, 6 more variables:
## #   numberim.pre <dbl>, numberim.post <dbl>,
## #   remain.pre <dbl>, remain.post <dbl>, english.pre <dbl>,
## #   english.post <dbl>, and abbreviated variable name
## #   1: ideology
```

# Estimating the difference in means

```
diff_in_means <- trains |>
  group_by(treatment) |>
  summarize(post_mean = mean(numberim.post)) |>
  pivot_wider(names_from = treatment, values_from = post_mean) |>
  mutate(ATE = `1` - `0`)
diff_in_means

## # A tibble: 1 x 3
##       `0`     `1`   ATE
##     <dbl> <dbl> <dbl>
## 1    2.73   3.12  0.383
```

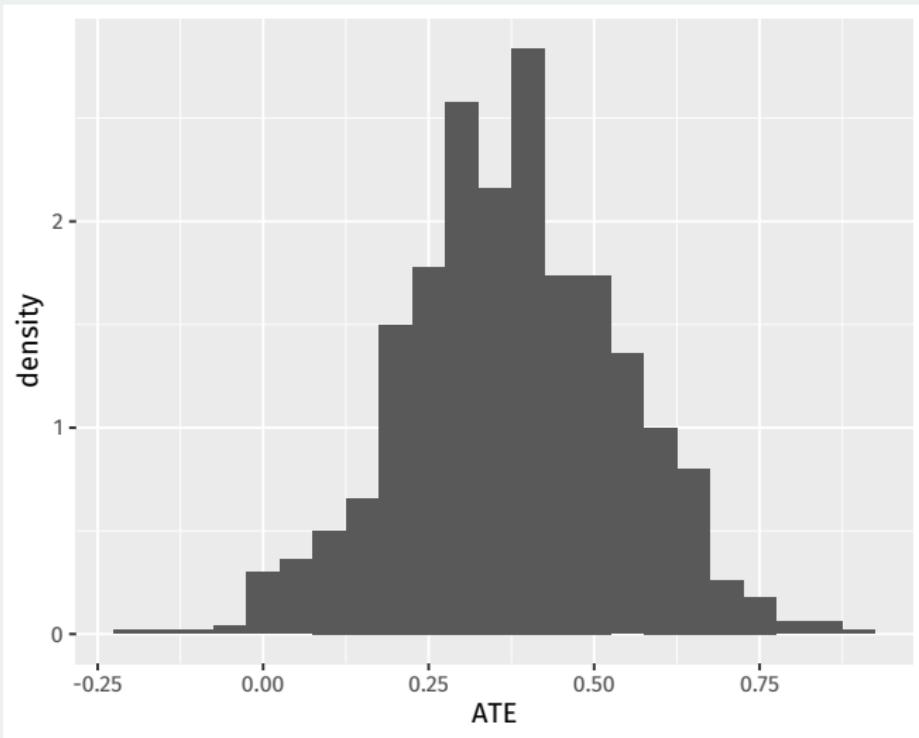
# Bootstrap for the difference in means

```
library(infer)
dim_boots <- trains |>
  rep_slice_sample(prop = 1, replace = TRUE, reps = 1000) |>
  group_by(replicate, treatment) |>
  summarize(post_mean = mean(numberim.post)) |>
  pivot_wider(names_from = treatment, values_from = post_mean) |>
  mutate(ATE = `1` - `0`)
dim_boots
```

```
## # A tibble: 1,000 x 4
## # Groups:   replicate [1,000]
##   replicate `0` `1`   ATE
##       <int> <dbl> <dbl> <dbl>
## 1     1    2.83  3.02  0.194
## 2     2    2.67  3.07  0.406
## 3     3    2.74  3.09  0.346
## 4     4    2.79  3.19  0.398
## 5     5    2.76  3.13  0.376
## 6     6    2.62  3.14  0.520
## 7     7    2.87  3.27  0.395
## 8     8    2.71  3.07  0.360
## 9     9    3.03  3.26  0.229
```

# Visualizing the bootstraps

```
dim_boots |>  
  ggplot(aes(x = ATE)) +  
  geom_histogram(aes(y = ..density..), binwidth = 0.05)
```



# Calculating the percentile CI

You can use `get_confidence_interval()` with your “hand-rolled” bootstraps, but you have to make sure you only pass it the variable of interest using `select`:

```
dim_ci_95 <- dim_boots |>  
  select(replicate, ATE) |>  
  get_confidence_interval(level = 0.95, type = "percentile")
```

```
dim_ci_95
```

```
## # A tibble: 1 x 2  
##   lower_ci upper_ci  
##       <dbl>    <dbl>  
## 1     0.0514    0.685
```

# What about change in views as the outcome?

```
change_ci_95 <- trains |>
  rep_slice_sample(prop = 1, replace = TRUE, reps = 1000) |>
  group_by(replicate, treatment) |>
  summarize(change_mean = mean(numberim.post - numberim.pre)) |>
  pivot_wider(names_from = treatment, values_from = change_mean) |>
  mutate(ATE = `1` - `0`) |>
  select(replicate, ATE) |>
  get_confidence_interval(level = 0.95, type = "percentile")
change_ci_95
```

```
## # A tibble: 1 x 2
##   lower_ci upper_ci
##       <dbl>    <dbl>
## 1     0.0157    0.613
```

# What's different?

Let's look at the width of the two confidence intervals:

```
## Post outcome width  
dim_ci_95[2]-dim_ci_95[1]
```

```
##      upper_ci  
## 1    0.634
```

```
## Change outcome width  
change_ci_95[2] - change_ci_95[1]
```

```
##      upper_ci  
## 1    0.597
```

# Width of CI depends on outcome variability

Change CI is narrower! Why? Because the change is less variable than the post outcome:

```
trains |> summarize(sd_post = sd(numberim.post),  
                     sd_change = sd(numberim.post - numberim.pre))
```

```
## # A tibble: 1 x 2  
##   sd_post  sd_change  
##     <dbl>      <dbl>  
## 1    0.917      0.826
```

# infer workflow

For `infer`, we have to do a bit of massaging. It wants the treatment variable to be a vector and we have to tell it what order we take the difference:

```
dim_boots_infer <- trains |>
  mutate(treatment = if_else(treatment == 1, "Treated", "Control")) |>
  specify(numberim.post ~ treatment) |>
  generate(reps = 1000, type = "bootstrap") |>
  calculate(stat = "diff in means", order = c("Treated", "Control"))
dim_boots_infer |>
  get_confidence_interval(level = 0.95, type = "percentile")
```

```
## # A tibble: 1 x 2
##   lower_ci upper_ci
##       <dbl>    <dbl>
## 1     0.0569    0.720
```

# 2/ Bootstrap CIs for a difference in ATEs

# Interactions

We have also estimated conditional ATEs:

$$ATE_{\text{college}} = \bar{X}_{T,\text{college}} - \bar{X}_{C,\text{college}}$$

$$ATE_{\text{noncollege}} = \bar{X}_{T,\text{noncollege}} - \bar{X}_{C,\text{noncollege}}$$

An **interaction** between treatment and college is the difference between these two effects:

$$ATE_{\text{college}} - ATE_{\text{noncollege}}$$

This is a random variable and has a **sampling distribution**.

# Estimating the interaction

To estimate the interaction, we need to pivot both treatment and college to the columns.

```
trains |>
  mutate(
    treatment = if_else(treatment == 1, "Treated", "Control"),
    college = if_else(college == 1, "College", "Noncollege")
  ) |>
  group_by(treatment, college) |>
  summarize(post_mean = mean(numberim$post)) |>
  pivot_wider(
    names_from = c(treatment, college),
    values_from = post_mean
  )
```

```
## # A tibble: 1 x 4
##   Control_College Control_Noncollege Treated_College Treat~1
##                 <dbl>              <dbl>            <dbl>    <dbl>
## 1             2.63              3.57            3.11    3.14
## # ... with abbreviated variable name 1: Treated_Noncollege
```

# Estimating the interaction

```
trains |>
  mutate(
    treatment = if_else(treatment == 1, "Treated", "Control"),
    college = if_else(college == 1, "College", "Noncollege")
  ) |>
  group_by(treatment, college) |>
  summarize(post_mean = mean(numberim.post)) |>
  pivot_wider(
    names_from = c(treatment, college),
    values_from = post_mean
  ) |>
  mutate(
    ATE_c = Treated_College - Control_College,
    ATE_nc = Treated_Noncollege - Control_Noncollege,
    interaction = ATE_c - ATE_nc
  ) |>
  select(ATE_c, ATE_nc, interaction)
```

```
## # A tibble: 1 x 3
##   ATE_c ATE_nc interaction
##   <dbl>  <dbl>      <dbl>
## 1  0.482 -0.429     0.911
```

# Bootstrapping the interaction

```
int_boots <- trains |>
  mutate(
    treatment = if_else(treatment == 1, "Treated", "Control"),
    college = if_else(college == 1, "College", "Noncollege")
  ) |>
  rep_slice_sample(prop = 1, replace = TRUE, reps = 1000) |>
  group_by(replicate, treatment, college) |>
  summarize(post_mean = mean(numberim.post)) |>
  pivot_wider(
    names_from = c(treatment, college),
    values_from = post_mean
  ) |>
  mutate(
    ATE_c = Treated_College - Control_College,
    ATE_nc = Treated_Noncollege - Control_Noncollege,
    interaction = ATE_c - ATE_nc
  ) |>
  select(replicate, ATE_c, ATE_nc, interaction)
```

```
int_boots
```

```
## # A tibble: 1,000 x 4
## # Groups:   replicate [1,000]
##   replicate ATE_c ATE_nc interaction
##   <int>    <dbl>    <dbl>      <dbl>
## 1 1        0.580 -0.175     0.755
## 2 2        0.515 -0.458     0.973
## 3 3        0.753 -0.812     1.57
## 4 4        0.339  0.125     0.214
## 5 5        0.355  0         0.355
## 6 6        0.465 -0.568     1.03
## 7 7        0.492 -0.75      1.24
## 8 8        0.382 -0.5       0.882
## 9 9        0.277  0.125     0.152
## 10 10      0.449 -0.633     1.08
## # ... with 990 more rows
```

# Getting the confidence interval

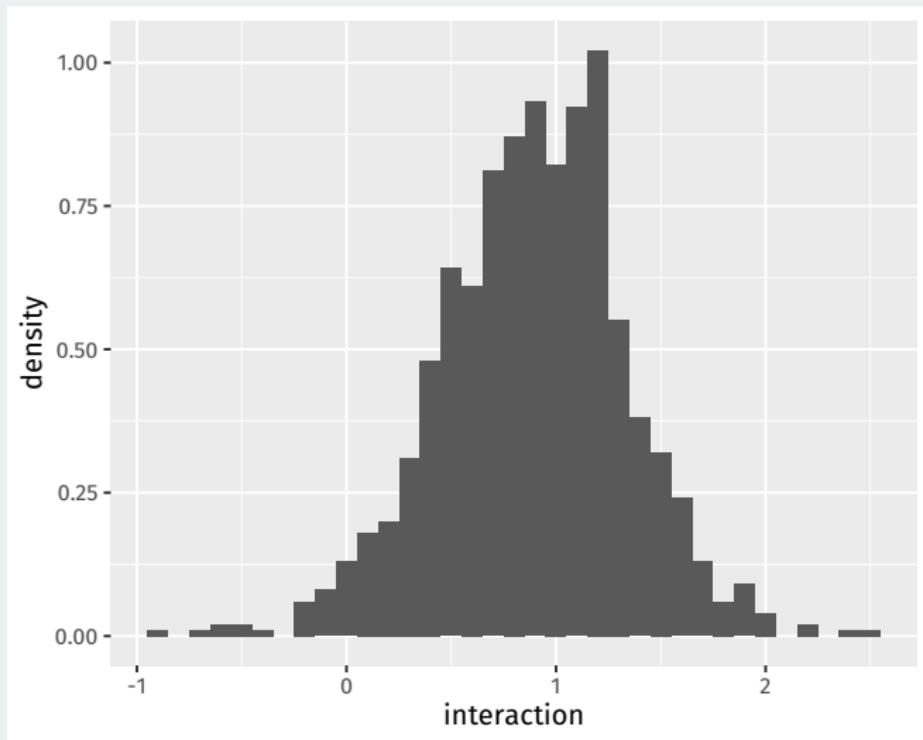
We have to drop NA values because sometimes the bootstrap gets a draw of all college or all noncollege and we can't calculate the interaction:

```
int_boots |>
  select(replicate, interaction) |>
  drop_na() |>
  get_confidence_interval(level = 0.95)
```

```
## # A tibble: 1 x 2
##   lower_ci upper_ci
##       <dbl>    <dbl>
## 1 -0.00805     1.72
```

# Visualizing the bootstrap

```
int_boots |>  
  ggplot(aes(x = interaction)) +  
  geom_histogram(aes(y = ..density..), binwidth = 0.1)
```



# **3/** Interpreting confidence intervals

# Interpretation and simulation

- Be careful about interpretation:
  - A 95% confidence interval will contain the true value in 95% of repeated samples.
  - For a particular calculated confidence interval, truth is either in it or not.
- A simulation can help our understanding:
  - Draw samples of size 1500 assuming population approval for Trump of  $p = 0.4$ .
  - Calculate 95% confidence intervals in each sample.
  - See how many overlap with the true population approval.

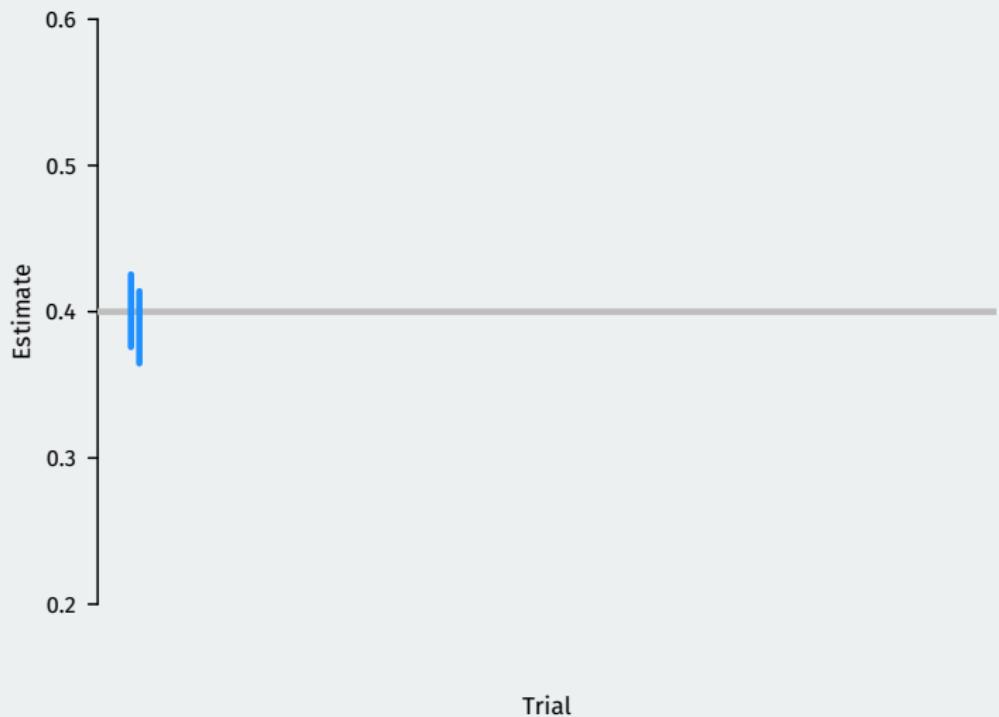
# Plotting the CIs



# Plotting the CIs



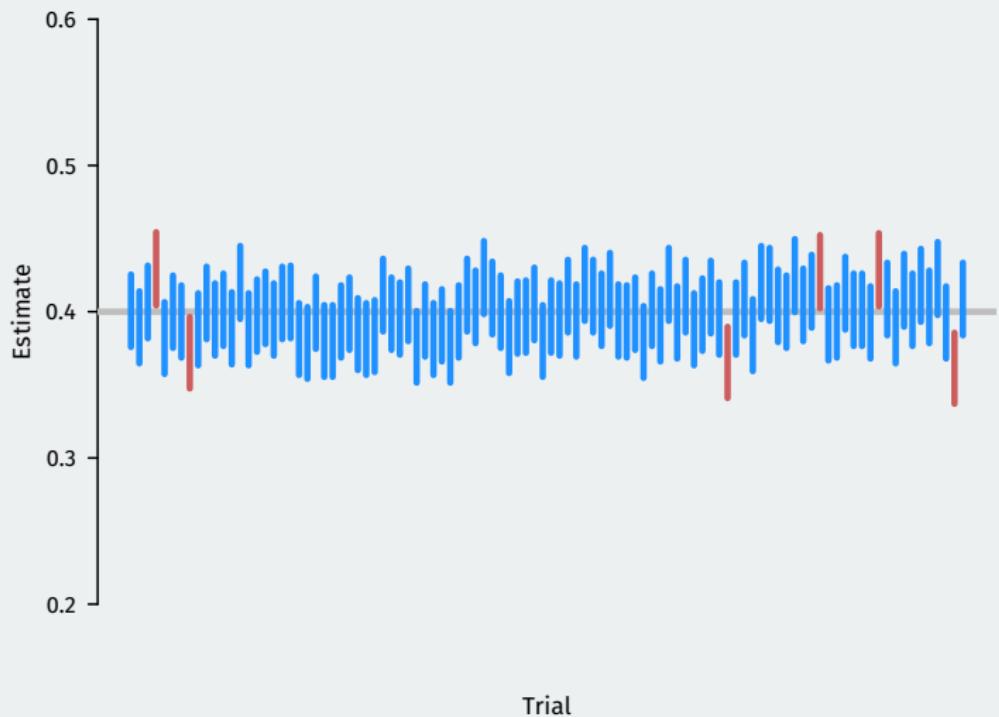
# Plotting the CIs



# Plotting the CIs



# Plotting the CIs



# Gov 50: 20. Hypothesis testing

Matthew Blackwell

Harvard University

# Roadmap

1. The lady tasting tea
2. Hypothesis tests
3. Hypothesis testing using infer

# 1/ The lady tasting tea

# The lady tasting tea

*Your friend asks you to grab a tea with milk for her before meeting up and she says that she prefers tea poured before the milk. You stop by a local tea shop and ask for a tea with milk. When you bring it to her, she complains that it was prepared milk-first.*

- You're skeptical that she can tell the difference, so you devise a test:
  - Prepare 8 cups of tea, 4 milk-first, 4 tea-first
  - Present cups to friend in a **random** order
  - Ask friend to pick which 4 of the 8 were milk-first.

# Lady Tasting Tea data

Friend picks out all 4 milk-first cups correctly!

```
library(gov50data)
tea

## # A tibble: 8 x 2
##   truth     guess
##   <chr>     <chr>
## 1 tea-first tea-first
## 2 milk-first milk-first
## 3 milk-first milk-first
## 4 tea-first tea-first
## 5 tea-first tea-first
## 6 milk-first milk-first
## 7 tea-first tea-first
## 8 milk-first milk-first
```

# Thought experiment

Could she have been guessing at random? What would guessing look like?

```
set.seed(02138)
one_guess <- tea |>
  mutate(random_guess = sample(guess))
one_guess
```

```
## # A tibble: 8 x 3
##   truth     guess   random_guess
##   <chr>     <chr>    <chr>
## 1 tea-first tea-first milk-first
## 2 milk-first milk-first tea-first
## 3 milk-first milk-first tea-first
## 4 tea-first  tea-first  milk-first
## 5 tea-first  tea-first  tea-first
## 6 milk-first milk-first milk-first
## 7 tea-first  tea-first  tea-first
## 8 milk-first milk-first milk-first
```

4 correct in this random guess!

# Another guess

```
another_guess <- tea |>
  mutate(random_guess = sample(guess))
another_guess

## # A tibble: 8 x 3
##   truth     guess   random_guess
##   <chr>     <chr>     <chr>
## 1 tea-first tea-first tea-first
## 2 milk-first milk-first tea-first
## 3 milk-first milk-first milk-first
## 4 tea-first tea-first tea-first
## 5 tea-first tea-first milk-first
## 6 milk-first milk-first milk-first
## 7 tea-first tea-first tea-first
## 8 milk-first milk-first milk-first
```

6 correct in this random guess!

# All possible guesses

We could enumerate all possible guesses. “Guessing” would mean choosing one of these at random:

```
##   Cup 1 Cup 2 Cup 3 Cup 4 Cup 5 Cup 6 Cup 7 Cup 8  
## 1 milk milk milk milk tea tea tea tea  
## 2 milk milk milk tea milk tea tea tea  
## 3 milk milk tea milk milk tea tea tea  
## 4 milk tea milk milk milk tea tea tea  
## 5 tea milk milk milk milk tea tea tea  
## 6 milk milk milk tea tea milk tea tea
```

[snip]

```
##   Cup 1 Cup 2 Cup 3 Cup 4 Cup 5 Cup 6 Cup 7 Cup 8  
## 65 tea tea tea milk milk tea milk milk  
## 66 milk tea tea tea tea milk milk milk  
## 67 tea milk tea tea tea milk milk milk  
## 68 tea tea milk tea tea milk milk milk  
## 69 tea tea tea milk tea milk milk milk  
## 70 tea tea tea tea milk milk milk milk
```

# Statistical thought experiment

- Statistical thought experiment: how often would she get all 4 correct **if she were guessing randomly?**
  - Only one way to choose all 4 correct cups.
  - But 70 ways of choosing 4 cups among 8.
  - Choosing at random: picking each of these 70 with equal probability.
- Chances of guessing all 4 correct is  $\frac{1}{70} \approx 0.014$  or 1.4%.
- → the guessing hypothesis might be implausible.
  - Impossible? No, because of random chance!

# 2/ Hypothesis tests

# Statistical hypothesis testing

- Statistical hypothesis testing is a **thought experiment**.
  - Could our results just be due to random chance?
- What would the world look like **if we knew the truth?**
- Example 1:
  - An analyst claims that 20% of Boston households are in poverty.
  - You take a sample of 900 households and find that 23% of the sample is under the poverty line.
  - Should you conclude that the analyst is wrong?
- Example 2:
  - Trump won 47.5% of the vote in the 2020 election.
  - Last YouGov poll of 1,363 likely voters said 44% planned to vote for Trump.
  - Could the difference between the poll and the outcome be just due to random chance?

# Null and alternative hypothesis

- **Null hypothesis:** Some statement about the population parameters.
  - “Devil’s advocate” position  $\rightsquigarrow$  assumes what you seek to prove wrong.
  - Usually that an observed difference is due to chance.
  - Ex: poll drawn from the same population as all voters.
  - Denoted  $H_0$
- **Alternative hypothesis:** The statement we hope or suspect is true instead of  $H_0$ .
  - It is the opposite of the null hypothesis.
  - An observed difference is real, not just due to chance.
  - Ex: polling for Trump is systematically wrong.
  - Denoted  $H_1$  or  $H_a$
- **Probabilistic** proof by contradiction: try to “disprove” the null.

# Hypothesis testing example

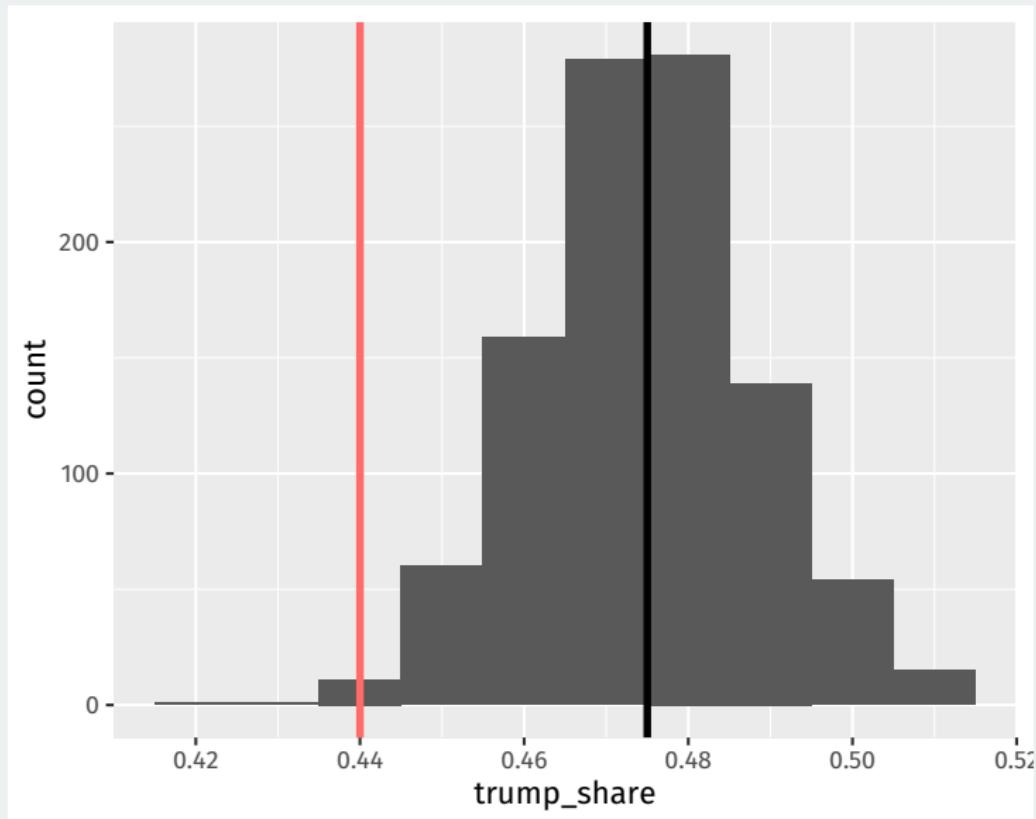
- Are we polling the same population as the actual voters?
  - If so, how likely are we to see polling error this big by chance?
- What is the parameter we want to learn about?
  - True population mean of the surveys,  $p$ .
  - Null hypothesis:  $H_0 : p = 0.475$  (surveys drawing from same population)
  - Alternative hypothesis:  $H_1 : p \neq 0.475$
- Data: poll has  $\bar{X} = 0.44$  with  $n = 1363$ .

# Statistical thought experiment

- If the null were true, what should the distribution of the data be?
  - $X_i$  is 1 if respondent  $i$  will vote for Trump.
  - Under null,  $X_i$  is a coin flip with probability  $p = 0.475$  of landing on “Trump”.
  - $\sum_{i=1}^n X_i$  is the number in sample that will vote for Trump.
- We can simulate sums of coin flips using a function called `rbinom()`
- Compare the distribution of proportions under the null to the observed proportion.

```
null_dist <- tibble(  
  trump_share = rbinom(n = 1000, size = 1363, prob = 0.475) / 1363  
)  
ggplot(null_dist, aes(x = trump_share)) +  
  geom_histogram(binwidth = 0.01) +  
  geom_vline(xintercept = 0.44, color = "indianred1", size = 1.25) +  
  geom_vline(xintercept = 0.475, size = 1.25)
```

# Simulations of the reference distribution



# p-value

## p-value

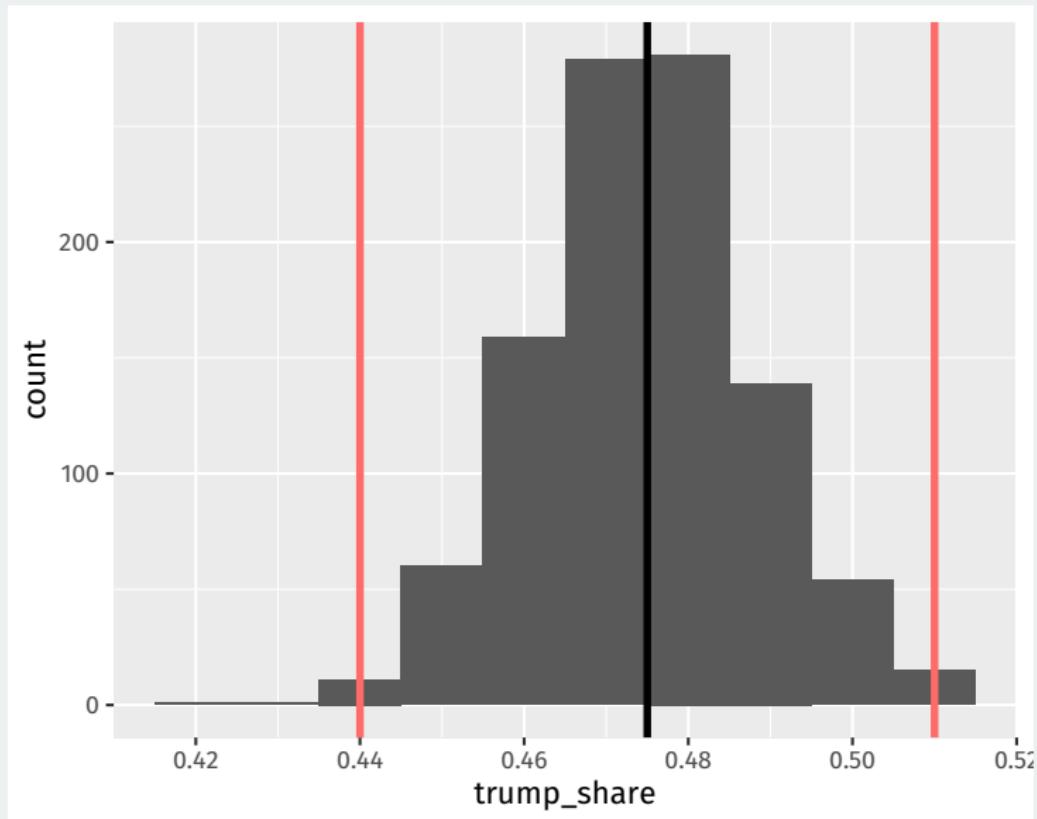
The **p-value** is the probability of observing data as or more extreme as our data under the null.

- If the null is true, how often would we expect polling errors this big?
  - Smaller p-value  $\rightsquigarrow$  stronger evidence against the null
  - **NOT** the probability that the null is true!
- p-values are usually **two-sided**:
  - Observed error of  $0.44 - 0.475 = -0.035$  under the null.
  - p-value is probability of sample proportions being less than 0.44 **plus**
  - Probability of sample proportions being greater than  $0.475 + 0.035 = 0.51$ .

```
mean(null_dist$trump_share < 0.44) + mean(null_dist$trump_share > 0.51)
```

```
## [1] 0.01
```

# Two-sided p-value



# One-sided tests

- Sometimes our hypothesis is directional.
  - We only consider evidence against the null from one direction.
- Null: our polls are from the same population as actual voters
  - $H_0 : p = 0.475$
- **One-sided alternative:** polls underestimate Trump support.
  - $H_1 : p < 0.475$
- Makes the p-value one-sided:
  - What's the probability of a random sample underestimating Trump support by as much as we see in the sample?
  - Always smaller than a two-sided p-value.

```
mean(null_dist$trump_share < 0.44)
```

```
## [1] 0.005
```

# Rejecting the null

- Tests usually end with a decision to reject the null or not.
- Choose a threshold below which you'll reject the null.
  - **Test level**  $\alpha$ : the threshold for a test.
  - Decision rule: “reject the null if the p-value is below  $\alpha$ ”
  - Otherwise “fail to reject” or “retain”, not “accept the null”
- Common (arbitrary) thresholds:
  - $p \geq 0.1$  “not statistically significant”
  - $p < 0.05$  “statistically significant”
  - $p < 0.01$  “highly significant”

# Testing errors

- A p-value of 0.05 says that data this extreme would only happen in 5% of repeated samples if the null were true.
  - $\rightsquigarrow$  5% of the time we'll reject the null when it is actually true.
- Test errors:

|              | $H_0$ True   | $H_0$ False   |
|--------------|--------------|---------------|
| Retain $H_0$ | Awesome!     | Type II error |
| Reject $H_0$ | Type I error | Good stuff!   |

- Type I error because it's the worst
  - “Convicting” an innocent null hypothesis
- Type II error less serious
  - Missed out on an awesome finding

# **3/** Hypothesis testing using infer

# GSS data from infer

```
library(infer)
gss

## # A tibble: 500 x 11
##   year   age sex college partyid hompop hours income
##   <dbl> <dbl> <fct>  <fct>    <fct>    <dbl> <dbl> <ord>
## 1 2014   36 male   degree   ind        3     50 $25000~
## 2 1994   34 female no degree rep        4     31 $20000~
## 3 1998   24 male   degree   ind        1     40 $25000~
## 4 1996   42 male   no degree ind       4     40 $25000~
## 5 1994   31 male   degree   rep        2     40 $25000~
## 6 1996   32 female no degree rep       4     53 $25000~
## 7 1990   48 female no degree dem       2     32 $25000~
## 8 2016   36 female degree   ind        1     20 $25000~
## 9 2000   30 female degree   rep        5     40 $25000~
## 10 1998  33 female no degree dem      2     40 $15000~
## # ... with 490 more rows, and 3 more variables:
## #   class <fct>, finrela <fct>, weight <dbl>
```

# What is the average hours worked?

dplyr way:

```
gss |>  
  summarize(mean(hours))
```

```
## # A tibble: 1 x 1  
##   `mean(hours)`  
##       <dbl>  
## 1        41.4
```

infer way:

```
observed_mean <- gss |>  
  specify(response = hours) |>  
  calculate(stat = "mean")  
observed_mean
```

```
## Response: hours (numeric)  
## # A tibble: 1 x 1  
##   stat  
##   <dbl>  
## 1 41.4
```

# Hypothesis test

Could we get a mean this different from 40 hours if that was the true population average of hours worked?

Null and alternative:

$$H_0 : \mu_{\text{hours}} = 40$$

$$H_1 : \mu_{\text{hours}} \neq 40$$

How do we perform this test using infer? The **bootstrap!**

# Specifying the hypotheses

```
gss |>  
  specify(response = hours) |>  
  hypothesize(null = "point", mu = 40)
```

```
## Response: hours (numeric)  
## Null Hypothesis: point  
## # A tibble: 500 x 1  
##       hours  
##   <dbl>  
## 1     50  
## 2     31  
## 3     40  
## 4     40  
## 5     40  
## 6     53  
## 7     32  
## 8     20  
## 9     40  
## 10    40  
## # ... with 490 more rows
```

# Generating the null distribution

We can use the bootstrap to determine how much variation there will be around 40 in the null distribution.

```
null_dist <- gss |>  
  specify(response = hours) |>  
  hypothesize(null = "point", mu = 40) |>  
  generate(reps = 1000, type = "bootstrap") |>  
  calculate(stat = "mean")  
null_dist
```

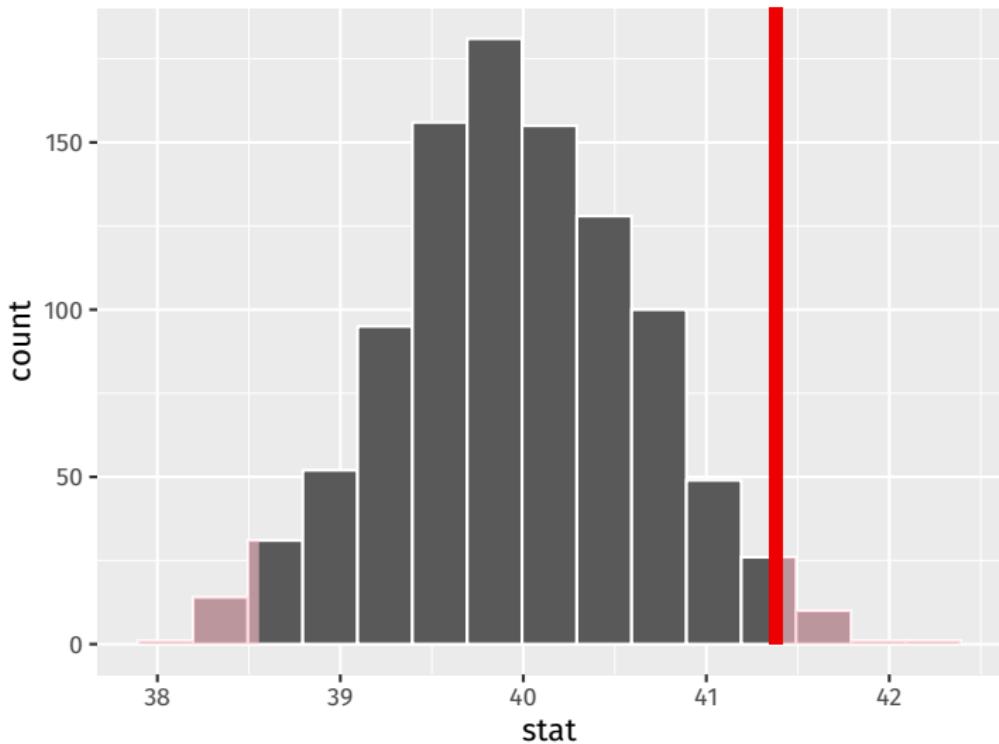
```
## Response: hours (numeric)  
## Null Hypothesis: point  
## # A tibble: 1,000 x 2  
##       replicate   stat  
##           <int> <dbl>  
## 1             1  40.3  
## 2             2  39.8  
## 3             3  40.0  
## 4             4  39.2  
## 5             5  40.3  
## 6             6  40.2  
## 7             7  40.4
```

# Visualizing the p-value

We can visualize our bootstrapped null distribution and the p-value as a shaded region:

```
null_dist |>  
  visualize() +  
  shade_p_value(observed_mean,  
                 direction = "two-sided")
```

## Simulation-Based Null Distribution



# Gov 50: 21. More Hypothesis testing

Matthew Blackwell

Harvard University

# Roadmap

1. Hypothesis testing using infer
2. Two-sample tests
3. Two-sample permutation tests with infer

# 1/ Hypothesis testing using infer

# Statistical hypothesis testing

- Statistical hypothesis testing is a **thought experiment**.
- What would the world look like **if we knew the truth?**
- Conducted with several steps:
  1. Specify your **null** and **alternative hypotheses**
  2. Choose an appropriate **test statistic** and level of test  $\alpha$
  3. Derive the **reference distribution** of the test statistic under the null.
  4. Use this distribution to calculate the **p-value**.
  5. Use p-value to decide whether to reject the null hypothesis or not

# GSS data from infer

```
library(infer)
gss

## # A tibble: 500 x 11
##   year   age sex college partyid hompop hours income
##   <dbl> <dbl> <fct>  <fct>    <fct>    <dbl> <dbl> <ord>
## 1 2014   36 male   degree   ind        3     50 $25000~
## 2 1994   34 female no degree rep        4     31 $20000~
## 3 1998   24 male   degree   ind        1     40 $25000~
## 4 1996   42 male   no degree ind       4     40 $25000~
## 5 1994   31 male   degree   rep        2     40 $25000~
## 6 1996   32 female no degree rep       4     53 $25000~
## 7 1990   48 female no degree dem       2     32 $25000~
## 8 2016   36 female degree   ind        1     20 $25000~
## 9 2000   30 female degree   rep        5     40 $25000~
## 10 1998  33 female no degree dem      2     40 $15000~
## # ... with 490 more rows, and 3 more variables:
## #   class <fct>, finrela <fct>, weight <dbl>
```

# What is the average hours worked?

dplyr way:

```
gss |>  
  summarize(mean(hours))
```

```
## # A tibble: 1 x 1  
##   `mean(hours)`  
##       <dbl>  
## 1        41.4
```

infer way:

```
observed_mean <- gss |>  
  specify(response = hours) |>  
  calculate(stat = "mean")  
observed_mean
```

```
## Response: hours (numeric)  
## # A tibble: 1 x 1  
##   stat  
##   <dbl>  
## 1 41.4
```

# Hypothesis test

Could we get a mean this different from 40 hours if that was the true population average of hours worked?

Null and alternative:

$$H_0 : \mu_{\text{hours}} = 40$$

$$H_1 : \mu_{\text{hours}} \neq 40$$

How do we perform this test using infer? The **bootstrap!**

# Specifying the hypotheses

```
gss |>  
  specify(response = hours) |>  
  hypothesize(null = "point", mu = 40)
```

```
## Response: hours (numeric)  
## Null Hypothesis: point  
## # A tibble: 500 x 1  
##       hours  
##   <dbl>  
## 1     50  
## 2     31  
## 3     40  
## 4     40  
## 5     40  
## 6     53  
## 7     32  
## 8     20  
## 9     40  
## 10    40  
## # ... with 490 more rows
```

# Generating the null distribution

We can use the bootstrap to determine how much variation there will be around 40 in the null distribution.

```
null_dist <- gss |>
  specify(response = hours) |>
  hypothesize(null = "point", mu = 40) |>
  generate(reps = 1000, type = "bootstrap") |>
  calculate(stat = "mean")
null_dist
```

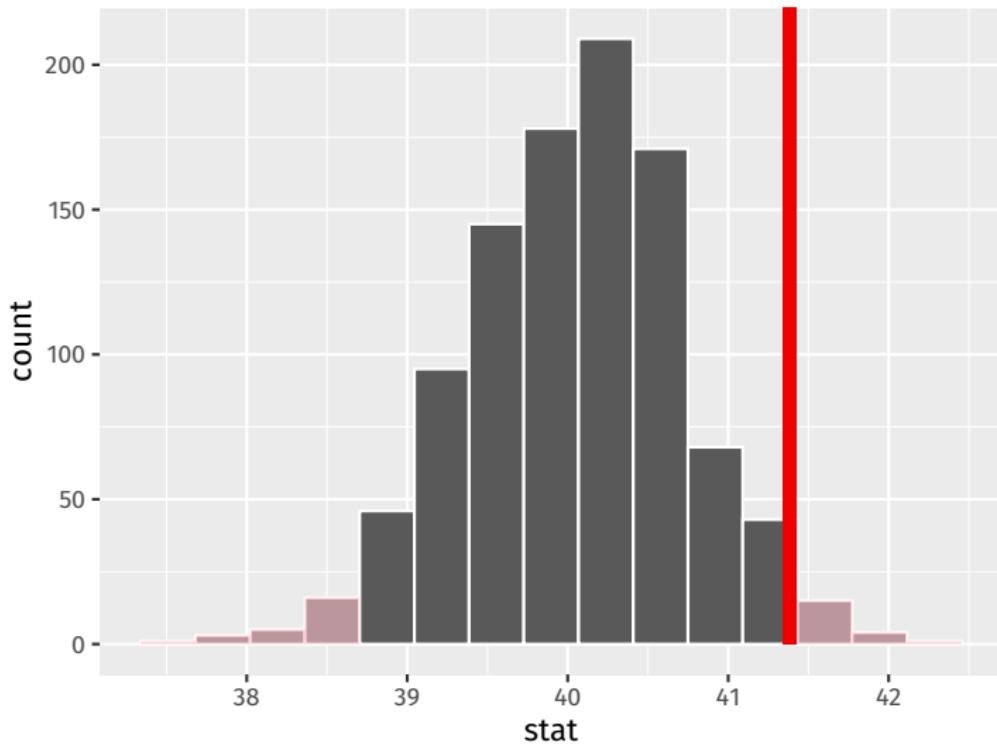
```
## Response: hours (numeric)
## Null Hypothesis: point
## # A tibble: 1,000 x 2
##       replicate   stat
##           <int> <dbl>
## 1             1 40.3
## 2             2 39.6
## 3             3 40.8
## 4             4 39.6
## 5             5 39.8
## 6             6 39.8
## 7             7 40.6
```

# Visualizing the p-value

We can visualize our bootstrapped null distribution and the p-value as a shaded region:

```
null_dist |>  
  visualize() +  
  shade_p_value(observed_mean,  
                 direction = "two-sided")
```

## Simulation-Based Null Distribution



# 2/ Two-sample tests

# Social pressure experiment

- Experimental study where each household for 2006 MI primary was randomly assigned to one of 4 conditions:
  - Control: no mailer
  - Civic Duty: mailer saying voting is your civic duty.
  - Hawthorne: a “we’re watching you” message.
  - Neighbors: naming-and-shaming social pressure mailer.
- Outcome: whether household members voted or not.
- We’ll focus on Neighbors vs Control
- Randomized implies samples are **independent**

# Neighbors mailer

Dear Registered Voter:

## WHAT IF YOUR NEIGHBORS KNEW WHETHER YOU VOTED?

Why do so many people fail to vote? We've been talking about the problem for years, but it only seems to get worse. This year, we're taking a new approach. We're sending this mailing to you and your neighbors to publicize who does and does not vote.

The chart shows the names of some of your neighbors, showing which have voted in the past. After the August 8 election, we intend to mail an updated chart. You and your neighbors will all know who voted and who did not.

## DO YOUR CIVIC DUTY — VOTE!

|                          |  | Aug 04 | Nov 04 | Aug 06 |
|--------------------------|--|--------|--------|--------|
| MAPLE DR                 |  | Voted  | Voted  | _____  |
| 9995 JOSEPH JAMES SMITH  |  |        | Voted  | _____  |
| 9995 JENNIFER KAY SMITH  |  |        | Voted  | _____  |
| 9997 RICHARD B JACKSON   |  |        | Voted  | _____  |
| 9999 KATHY MARIE JACKSON |  |        | Voted  | _____  |

# Social pressure data

```
data(social, package = "qss")
social <- as_tibble(social)
social

## # A tibble: 305,866 x 6
##   sex     yearofbirth primary2004 messages primar~1 hhsiz
##   <chr>      <int>      <int> <chr>      <int>    <int>
## 1 male        1941        0 Civic Duty      0        2
## 2 female      1947        0 Civic Duty      0        2
## 3 male        1951        0 Hawthorne       1        3
## 4 female      1950        0 Hawthorne       1        3
## 5 female      1982        0 Hawthorne       1        3
## 6 male        1981        0 Control         0        3
## 7 female      1959        0 Control         1        3
## 8 male        1956        0 Control         1        3
## 9 female      1968        0 Control         0        2
## 10 male       1967        0 Control         0        2
## # ... with 305,856 more rows, and abbreviated variable name
## #   1: primary2006
```

# Two-sample hypotheses

- Parameter: **population ATE**  $\mu_T - \mu_C$ 
  - $\mu_T$ : Turnout rate in the population if everyone received treatment.
  - $\mu_C$ : Turnout rate in the population if everyone received control.
- Goal: learn about the population difference in means
- Usual null hypothesis: no difference in population means (ATE = 0)
  - Null:  $H_0 : \mu_T - \mu_C = 0$
  - Two-sided alternative:  $H_1 : \mu_T - \mu_C \neq 0$
- In words: are the differences in sample means just due to chance?

# Permutation test

How do we generate draws of the difference in means under the null?

$$H_0 : \mu_T - \mu_C = 0$$

If the voting distribution is the same in the treatment and control groups, we could randomly swap who is labelled as treated and who is labelled as control and it shouldn't matter.

**Permutation test:** generate the null distribution by permuting the group labels and see the resulting distribution of differences in proportions

# Permuting the labels

```
social <- social |>  
  filter(messages %in% c("Neighbors", "Control"))  
  
social |>  
  mutate(messages_permute = sample(messages)) |>  
  select(primary2006, messages, messages_permute)
```

```
## # A tibble: 229,444 x 3  
##   primary2006 messages messages_permute  
##       <int> <chr>    <chr>  
## 1          0 Control   Control  
## 2          1 Control   Control  
## 3          1 Control   Neighbors  
## 4          0 Control   Control  
## 5          0 Control   Control  
## 6          1 Control   Neighbors  
## 7          0 Control   Control  
## 8          1 Control   Control  
## 9          1 Control   Control  
## 10         1 Control   Control  
## # ... with 229,434 more rows
```

# **3/** Two-sample permutation tests with infer

# Calculating the difference in proportion

infer functions with binary outcomes work best with factor variables:

```
social <- social |>
  mutate(turnout = if_else(primary2006 == 1, "Voted", "Didn't Vote"))

est_ate <- social |>
  specify(turnout ~ messages, success = "Voted") |>
  calculate(stat = "diff in props", order = c("Neighbors", "Control"))
est_ate

## Response: turnout (factor)
## Explanatory: messages (factor)
## # A tibble: 1 x 1
##       stat
##   <dbl>
## 1 0.0813
```

# Specifying the relationship of interest

infer functions with binary outcomes work best with factor variables:

```
social |>  
  specify(turnout ~ messages, success = "Voted")
```

```
## Response: turnout (factor)  
## Explanatory: messages (factor)  
## # A tibble: 229,444 x 2  
##   turnout     messages  
##   <fct>      <fct>  
## 1 Didn't Vote Control  
## 2 Voted       Control  
## 3 Voted       Control  
## 4 Didn't Vote Control  
## 5 Didn't Vote Control  
## 6 Voted       Control  
## 7 Didn't Vote Control  
## 8 Voted       Control  
## 9 Voted       Control  
## 10 Voted      Control  
## # ... with 229,434 more rows
```

# Setting the hypotheses

The null for these two-sample tests is called "independence" for the `infer` package because the assumption is that the two variables are statistically independent.

```
social |>  
  specify(turnout ~ messages, success = "Voted") |>  
  hypothesize(null = "independence")
```

```
## Response: turnout (factor)  
## Explanatory: messages (factor)  
## Null Hypothesis: independence  
## # A tibble: 229,444 x 2  
##       turnout     messages  
##       <fct>      <fct>  
## 1 Didn't Vote Control  
## 2 Voted      Control  
## 3 Voted      Control  
## 4 Didn't Vote Control  
## 5 Didn't Vote Control  
## 6 Voted      Control  
## 7 Didn't Vote Control  
## 8 Voted      Control
```

# Generating the permutations

We can tell `infer` to do our permutation test by using the argument `type = "permute"` to `generate()`:

```
social |>  
  specify(turnout ~ messages, success = "Voted") |>  
  hypothesize(null = "independence") |>  
  generate(reps = 1000, type = "permute")
```

```
## Response: turnout (factor)  
## Explanatory: messages (factor)  
## Null Hypothesis: independence  
## # A tibble: 229,444,000 x 3  
## # Groups:   replicate [1,000]  
##       turnout     messages replicate  
##       <fct>      <fct>      <int>  
## 1 Voted        Control        1  
## 2 Didn't Vote Control        1  
## 3 Voted        Control        1  
## 4 Didn't Vote Control        1  
## 5 Didn't Vote Control        1  
## 6 Voted        Control        1  
## 7 Voted        Control        1
```

# Calculating the diff in proportions in each sample

```
null_dist <- social |>  
  specify(turnout ~ messages, success = "Voted") |>  
  hypothesize(null = "independence") |>  
  generate(reps = 1000, type = "permute") |>  
  calculate(stat = "diff in props", order = c("Neighbors", "Control"))
```

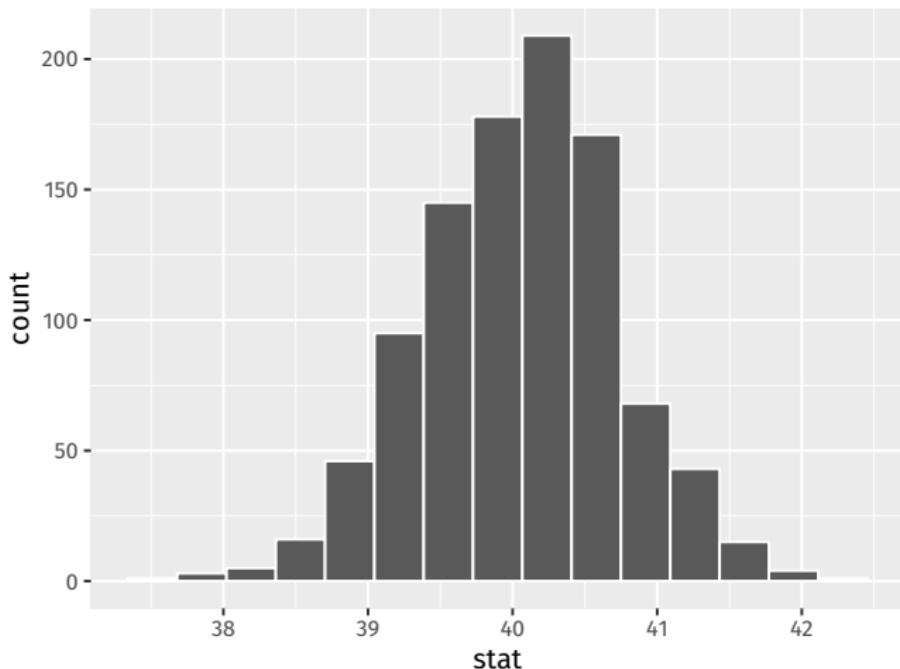
```
null_dist
```

```
## Response: hours (numeric)
## Null Hypothesis: point
## # A tibble: 1,000 x 2
##       replicate   stat
##       <int> <dbl>
## 1          1  40.3
## 2          2  39.6
## 3          3  40.8
## 4          4  39.6
## 5          5  39.8
## 6          6  39.8
## 7          7  40.6
## 8          8  40.5
## 9          9  38.6
## 10        10  41.2
## # ... with 990 more rows
```

# Visualizing

```
null_dist |>  
visualize()
```

Simulation-Based Null Distribution



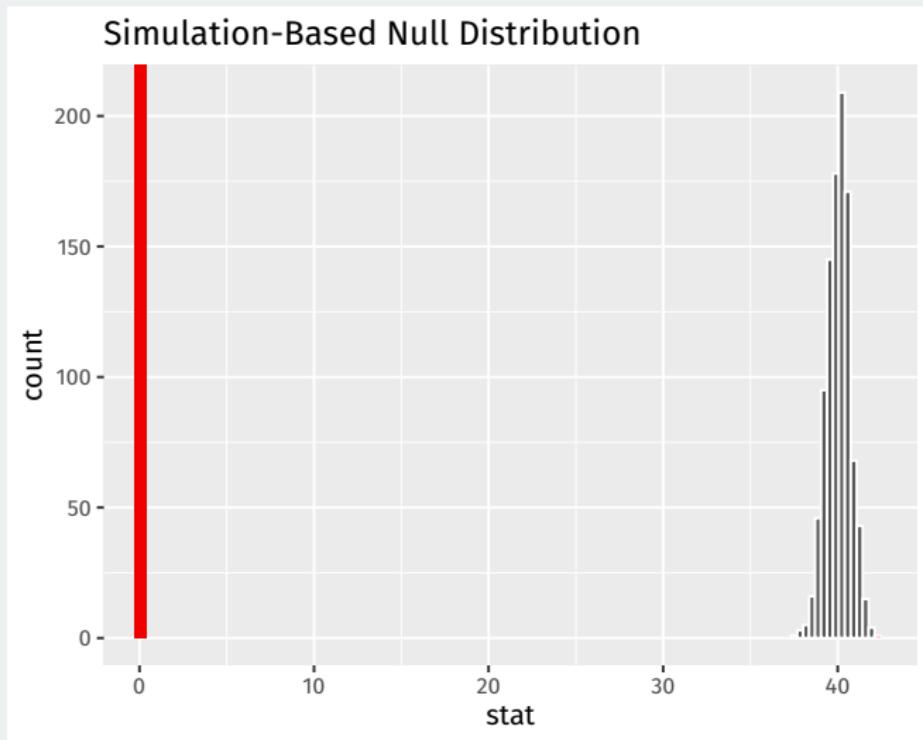
# Calculating p-values

```
ate_pval <- null_dist |>
  get_p_value(obs_stat = est_ate, direction = "both")
ate_pval
```

```
## # A tibble: 1 x 1
##   p_value
##   <dbl>
## 1 0
```

# Visualizing p-values

```
null_dist |>  
  visualize() +  
  shade_p_value(obs_stat = est_ate, direction = "both")
```



# Gov 50: 22. More Hypothesis testing

Matthew Blackwell

Harvard University

# Roadmap

1. Reviewing hypothesis testing
2. Issues with hypothesis testing
3. Power Analyses

# 1/ Reviewing hypothesis testing

# Difference-in-means

```
library(gov50data)
trains <- trains |>
  mutate(treated = if_else(treatment == 1, "Treated", "Untreated"))
trains
```

```
## # A tibble: 115 x 15
##       age   male income white college usborn treatment ideol~1
##       <dbl> <dbl> <dbl> <dbl>   <dbl>   <dbl>     <dbl>   <dbl>
## 1     31     0 135000     1      1      1        1      3
## 2     34     0 105000     1      1      0        1      4
## 3     63     1 135000     1      1      1        1      2
## 4     45     1 300000     1      1      1        1      4
## 5     55     1 135000     1      1      1        0      2
## 6     37     0  87500     1      1      1        1      5
## 7     53     0  87500     1      0      1        0      5
## 8     36     1 135000     1      1      1        1      4
## 9     54     0 105000     1      0      1        0      3
## 10    42     1 135000     1      1      1        1      4
## # ... with 105 more rows, 7 more variables:
## #   numberim.pre <dbl>, numberim.post <dbl>,
## #   remain.pre <dbl>, remain.post <dbl>, english.pre <dbl>,
## #   english.post <dbl>, treated <chr>, and abbreviated
## #   titles <chr>
```

# Calculating the ATE

```
library(infer)
ate <- trains |>
  specify(numberim.post ~ treated) |>
  calculate(stat = "diff in means",
             order = c("Treated", "Untreated"))
ate

## Response: numberim.post (numeric)
## Explanatory: treated (factor)
## # A tibble: 1 x 1
##       stat
##   <dbl>
## 1 0.383
```

# Difference in means hypotheses

Hypotheses:

$$H_0 : \mu_T - \mu_C = 0$$
$$H_1 : \mu_T - \mu_C \neq 0$$

Observed difference in means:

$$\widehat{ATE} = \bar{Y}_T - \bar{Y}_C$$

How can we approximate the **null distribution?** **Permute** the outcome/treatment variables.

# Permuting the treatment

Let's do 2 permutations to see how things vary:

```
set.seed(02138)
perm <- trains |>
  specify(numberim.post ~ treated) |>
  hypothesize(null = "independence") |>
  generate(reps = 1000,
            type = "permute")
```

`generate(type = "permute")` shuffles to the outcomes, keeping treatment the same:

```
perm |> filter(replicate == 1)
```

```
## # A tibble: 115 x 3
## # Groups:   replicate [1]
##   numberim.post treated  replicate
##   <dbl> <fct>     <int>
## 1 3 Treated      1
## 2 2 Treated      1
## 3 5 Treated      1
## 4 3 Treated      1
## 5 3 Untreated    1
## 6 3 Treated      1
## 7 2 Untreated    1
## 8 2 Treated      1
## 9 3 Untreated    1
## 10 3 Treated     1
## # ... with 105 more rows
```

```
perm |> filter(replicate == 2)
```

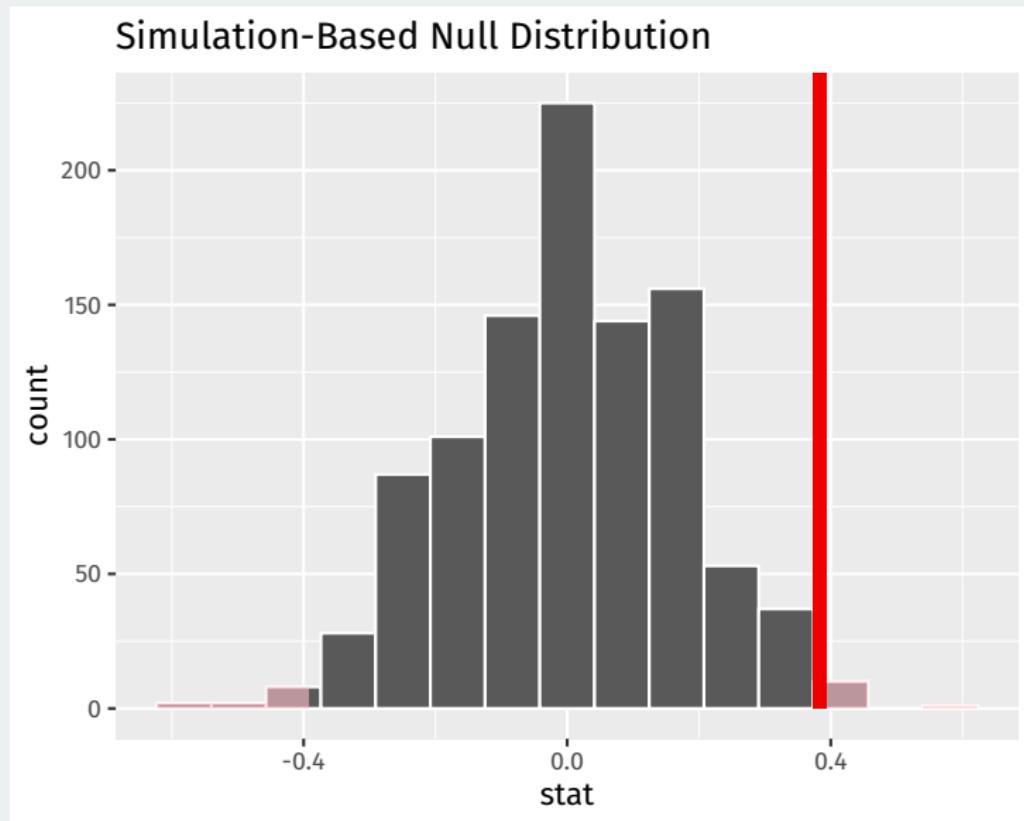
```
## # A tibble: 115 x 3
## # Groups:   replicate [1]
##   numberim.post treated  replicate
##   <dbl> <fct>     <int>
## 1 2 Treated      2
## 2 3 Treated      2
## 3 3 Treated      2
## 4 3 Treated      2
## 5 3 Untreated    2
## 6 4 Treated      2
## 7 2 Untreated    2
## 8 3 Treated      2
## 9 3 Untreated    2
## 10 2 Treated     2
## # ... with 105 more rows
```

# Null distribution

The distribution of the differences-in-means under permutation will be mean 0 because shuffling the outcomes means that the outcomes in each permutation's treated and control group are coming from the same distribution.

```
null_dist <- trains |>  
  specify(numberim.post ~ treated) |>  
  hypothesize(null = "independence") |>  
  generate(reps = 1000,  
            type = "permute") |>  
  calculate(stat = "diff in means", order = c("Treated", "Untreated"))
```

```
null_dist |>  
  visualize() +  
  shade_p_value(obs_stat = ate, direction = "both")
```



# Interpreting p-values

```
get_p_value(null_dist, obs_stat = ate, direction = "both")
```

```
## # A tibble: 1 x 1
##   p_value
##   <dbl>
## 1 0.022
```

Hypotheses:

$$H_0 : \mu_T - \mu_C = 0$$

$$H_1 : \mu_T - \mu_C \neq 0$$

Observed difference in means:

$$\widehat{ATE} = \bar{Y}_T - \bar{Y}_C$$

**p-value:** probability of an estimated ATE as big as  $|\widehat{ATE}|$  by random chance if there is no treatment effect.

# Rejecting the null

Decision rule: “reject the null if the p-value is below the **test level  $\alpha$** ”

Rejecting the null in two-sample tests: there is a true difference in means.

Test level  $\alpha$  controls the amount of false positives:

---

|             | Null False (True difference)   | Null True (No true difference) |
|-------------|--------------------------------|--------------------------------|
| Reject Null | True Positive                  | False Positive (Type I error)  |
| Retain Null | False Negative (Type II error) | True Negative                  |

---

# Tests and confidence intervals

- There is a deep connection between confidence intervals and tests.
- Any value outside of a  $100 \times (1 - \alpha)\%$  confidence interval would have a p-value less than  $\alpha$  if we tested it as the null hypothesis.
  - 95% CI for social pressure experiment: [0.016, 0.124]
  - $\rightsquigarrow$  p-value for  $H_0 : \mu_T - \mu_C = 0$  less than 0.05.
- Confidence intervals are all of the null hypotheses we **can't reject** with a test.

# CI in the trains example

```
trains |>
  specify(numberim.post ~ treated) |>
  generate(reps = 1000, type = "bootstrap") |>
  calculate(stat = "diff in means",
             order = c("Treated", "Untreated")) |>
  get_ci(level = 0.95)
```

```
## # A tibble: 1 x 2
##   lower_ci upper_ci
##       <dbl>     <dbl>
## 1     0.0893     0.698
```

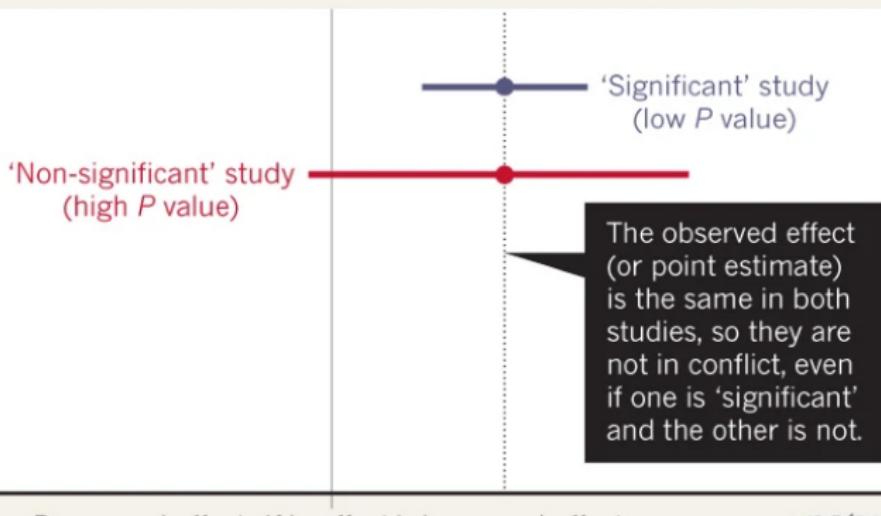
# **2/** Issues with hypothesis testing

# Significant vs not significant

The difference between statistically significant and not statistically significant is itself not statistically significant:

## BEWARE FALSE CONCLUSIONS

Studies currently dubbed ‘statistically significant’ and ‘statistically non-significant’ need not be contradictory, and such designations might cause genuine effects to be dismissed.

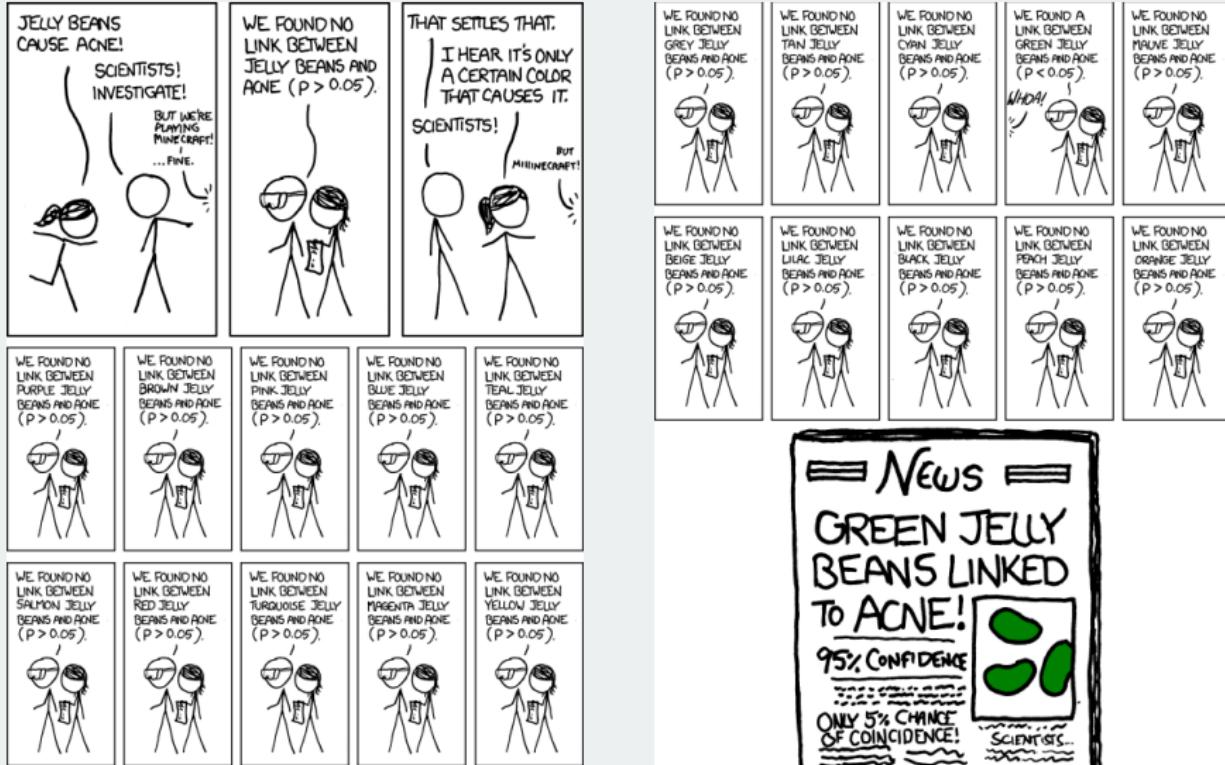


# What kind of significance

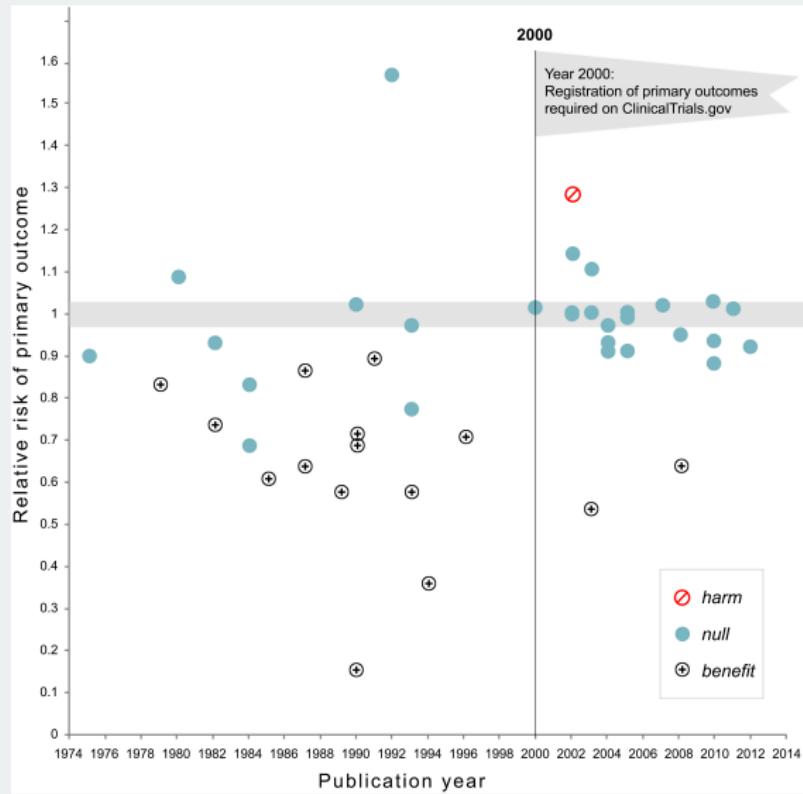
There are different types of significance that don't all have to be true together:

1. **Statistical significance:** we can reject the null of no effect.
2. **Causal significance:** we can interpret our estimated difference in means as a causal effect.
3. **Practical significance:** the estimated effect is meaningfully large.

# p-hacking



# p-hacking



# 3/ Power Analyses

# Effect sizes

**TABLE 2. Effects of Four Mail Treatments on Voter Turnout in the August 2006 Primary Election**

|                   | Experimental Group |            |           |        |           |
|-------------------|--------------------|------------|-----------|--------|-----------|
|                   | Control            | Civic Duty | Hawthorne | Self   | Neighbors |
| Percentage Voting | 29.7%              | 31.5%      | 32.2%     | 34.5%  | 37.8%     |
| N of Individuals  | 191,243            | 38,218     | 38,204    | 38,218 | 38,201    |

- Why did Gerber, Green, and Larimer use sample sizes of 38,000 for each treatment condition?
- Choose the sample size to ensure that you can *detect* what you think might be the true treatment effect:
  - Small effect sizes (half percentage point) will require huge  $n$
  - Large effect sizes (10 percentage points) will require smaller  $n$
- **Detect** here means “reject the null of no effect”

# Power of a test

- **Definition** The **power** of a test is the probability that a test rejects the null.
  - Probability that we reject given some specific value of the parameter
  - Power =  $1 - \mathbb{P}(\text{Type II error})$
  - Better tests = higher power.
- If we fail to reject a null hypothesis, two possible states of the world:
  - Null is true (no treatment effect)
  - Null is false (there is a treatment effect), but test had low power.

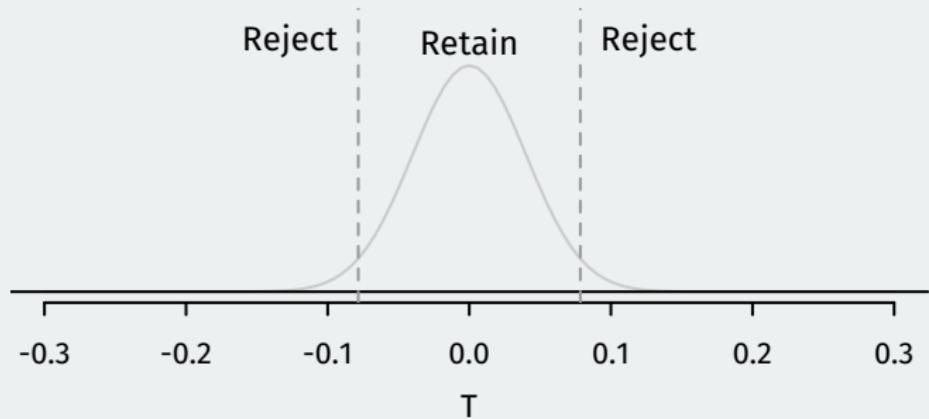
# Why care about power?

- Imagine you are a company being sued for racial discrimination in hiring.
- Judge forces you to conduct hypothesis test:
  - Null hypothesis is that hiring rates for white and black people are equal,  
 $H_0 : \mu_w - \mu_b = 0$
  - You sample 10 hiring records of each race, conduct hypothesis test and fail to reject null.
- Say to judge, “look we don’t have any racial discrimination”! What’s the problem?

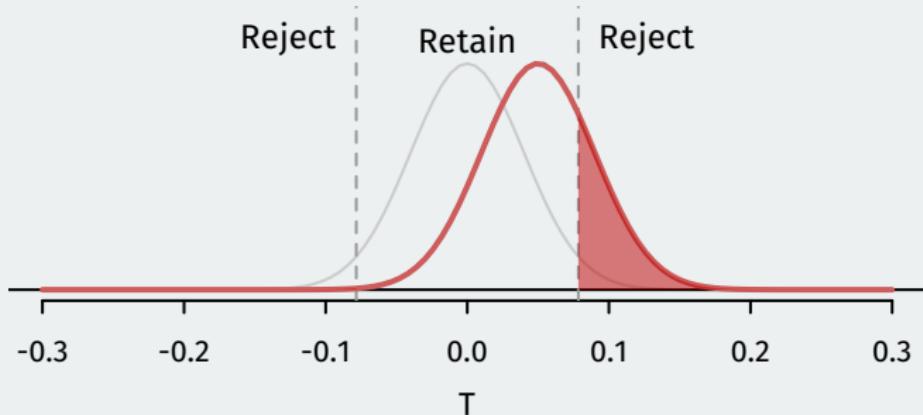
# Power analysis procedure

- Power can help guide the choice of sample size through a **power analysis**.
  - Calculate how likely we are to reject different possible treatment effects at different sample sizes.
  - **Can be done before the experiment:** which effects will I be able to detect with high probability at my  $n$ ?
- Steps to a power analysis:
  - Pick some hypothetical effect size,  $\mu_T - \mu_C = 0.05$
  - Calculate the distribution of  $T$  under that effect size.
  - Calculate the probability of rejecting the null under that distribution.
  - Repeat for different effect sizes.

# Power graph

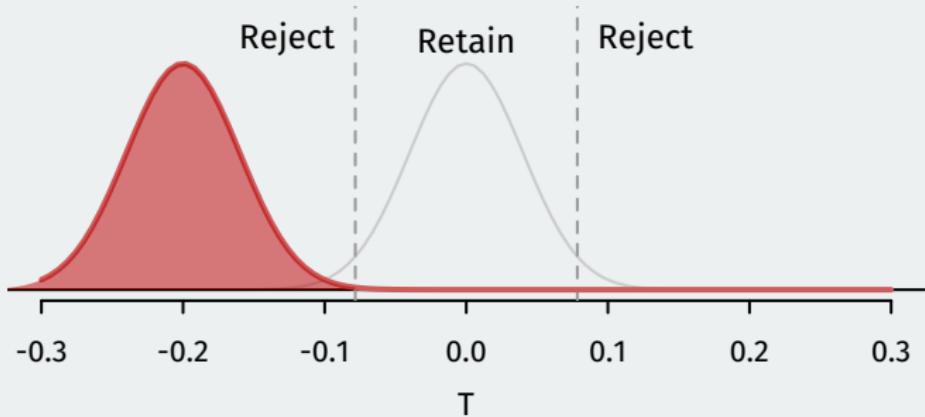


# Power graph



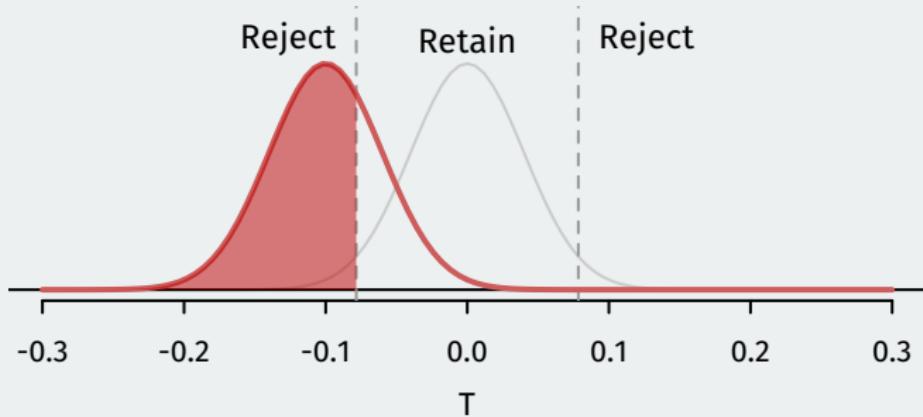
Assumed treatment effect = 0.05 and power = 0.24.

# Power graph



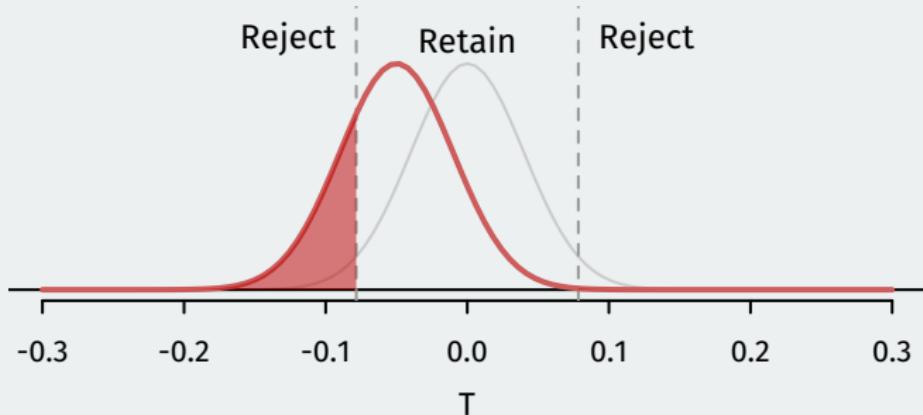
Assumed treatment effect = -0.2 and power = 0.999.

# Power graph



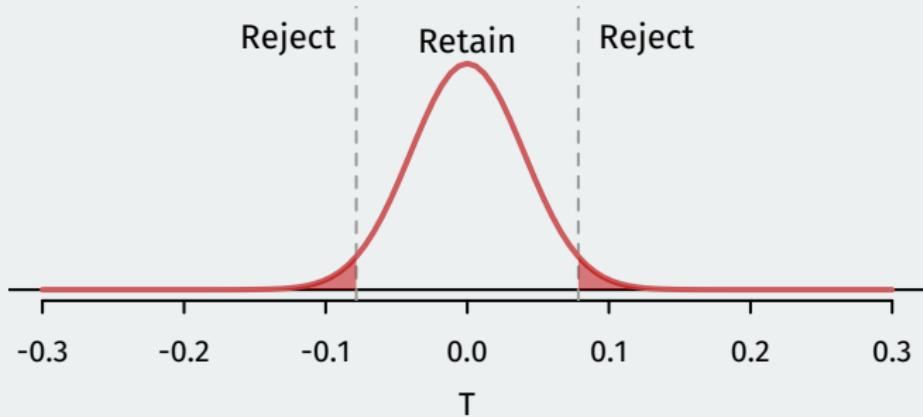
Assumed treatment effect = -0.1 and power = 0.705.

# Power graph



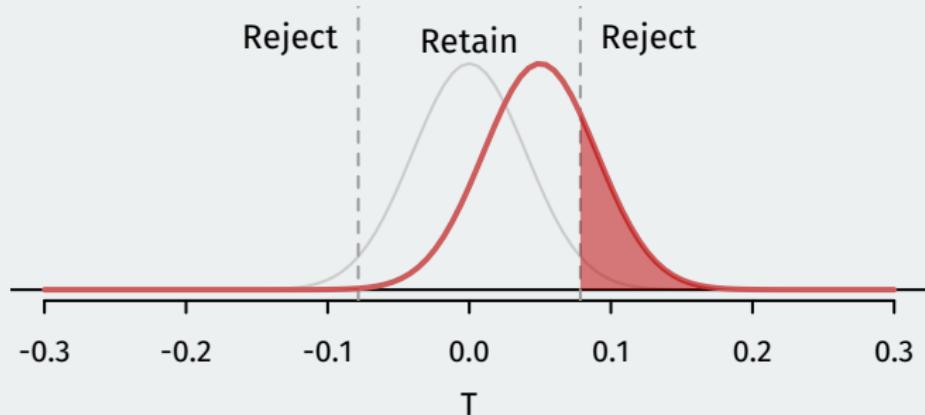
Assumed treatment effect = -0.05 and power = 0.24.

# Power graph



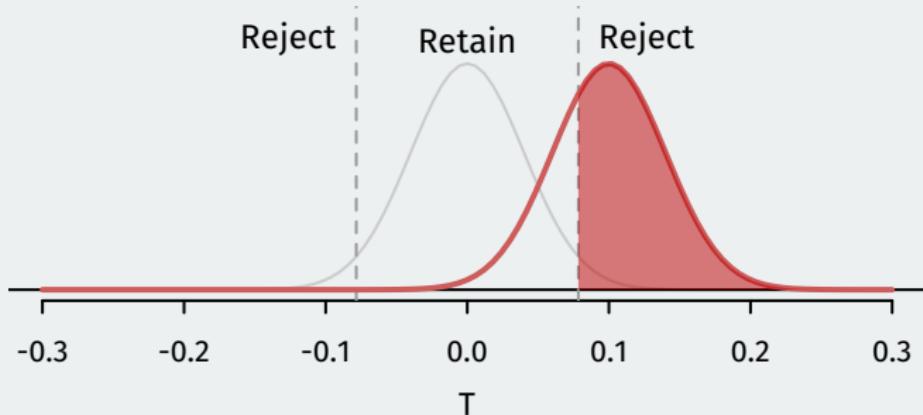
Assumed treatment effect = 0 and power = 0.05.

# Power graph



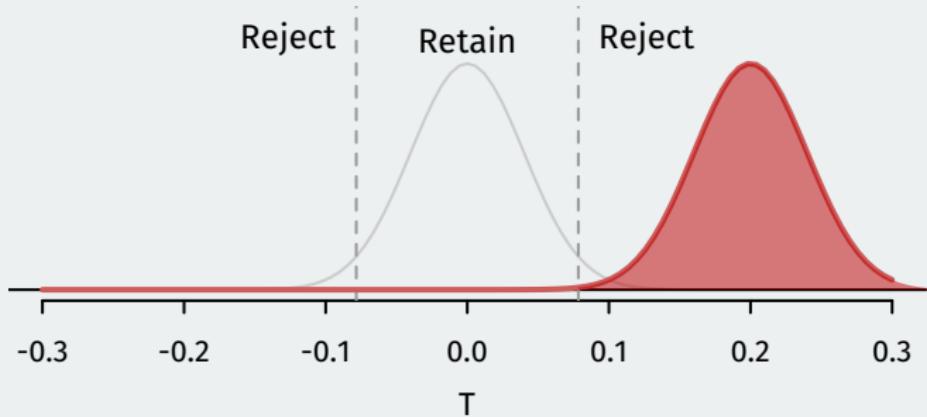
Assumed treatment effect = 0.05 and power = 0.24.

# Power graph



Assumed treatment effect = 0.1 and power = 0.705.

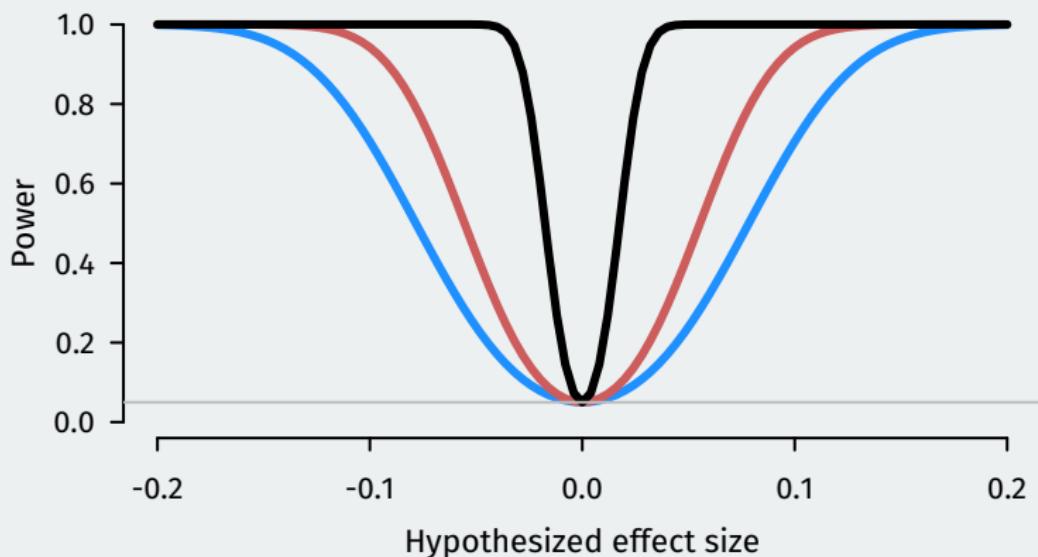
# Power graph



Assumed treatment effect = 0.2 and power = 0.999.

# A power analysis

- We can calculate the power for every possible effect size and plot the resulting **power curve**:
  - $n = 500$  (blue), 1000 (red), 10000 (black)



# Gov 50: 23. Inference with Mathematical Models

Matthew Blackwell

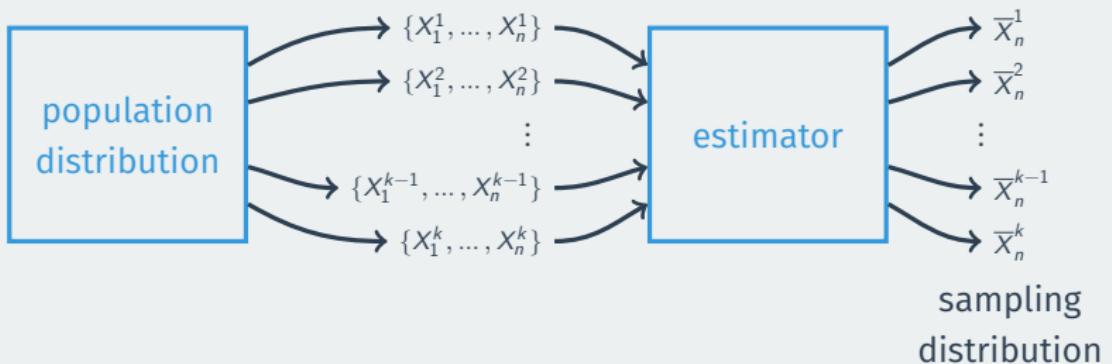
Harvard University

# Roadmap

1. Central limit theorem
2. Normal distribution
3. Using the Normal for inference

# 1/ Central limit theorem

# Sampling distribution, in pictures



# Sampling distribution of the sample proportion

sample mean = population mean + chance error

$$\bar{X} = \mu + \text{chance error}$$

Then  $\bar{X}$  centered at  $\mu$ .

Spread: standard deviation of the sampling distribution is the **standard error**

# Spread of the sample mean

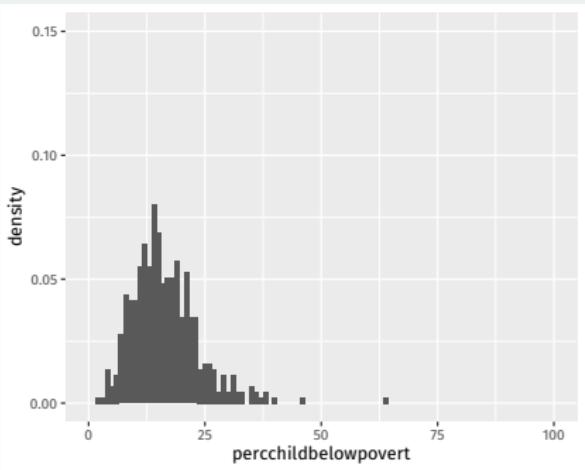
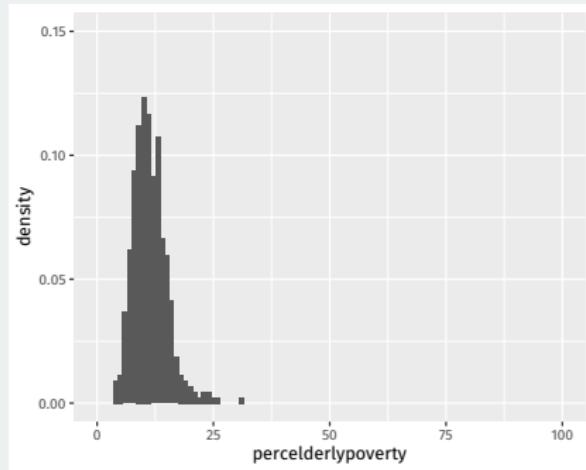
- **Standard error:** how big is the chance error on average?
  - This is the standard deviation of the estimator **across repeated samples**.
  - With random samples, we can get a formula for the SE for many estimators.
- Standard error for the sample mean:

$$SE = \frac{\sigma}{\sqrt{n}} = \frac{\text{population standard deviation}}{\sqrt{\text{sample size}}}$$

- Two components:
  - Population SD: more spread of the variable in the population → more spread of sample means
  - Size of the sample: larger sample → smaller spread of the sample means

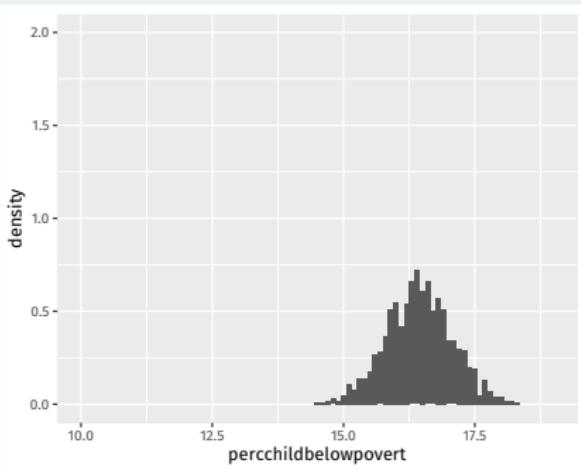
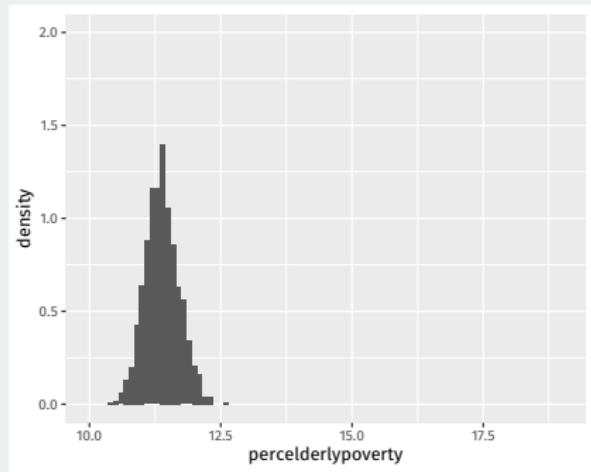
# Midwest counties

Population distributions:



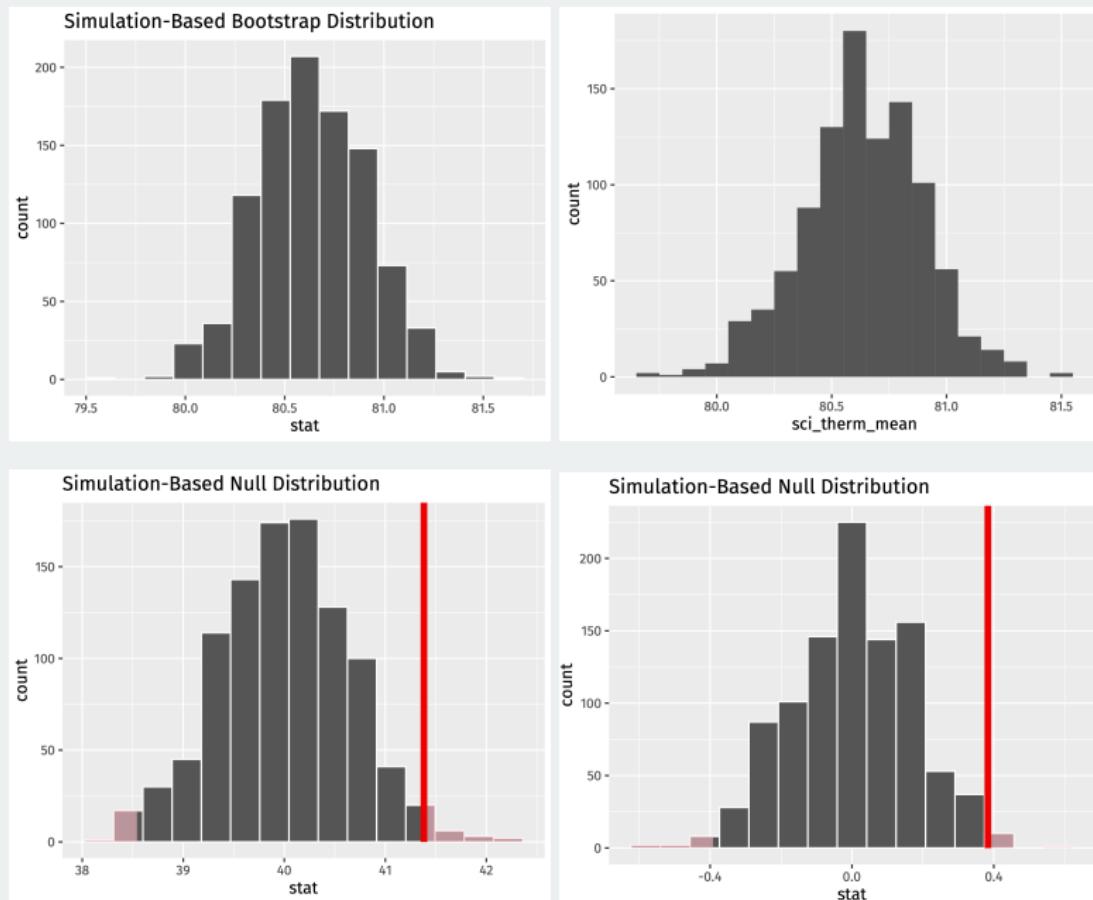
# Midwest counties

Sampling distributions with  $n = 100$

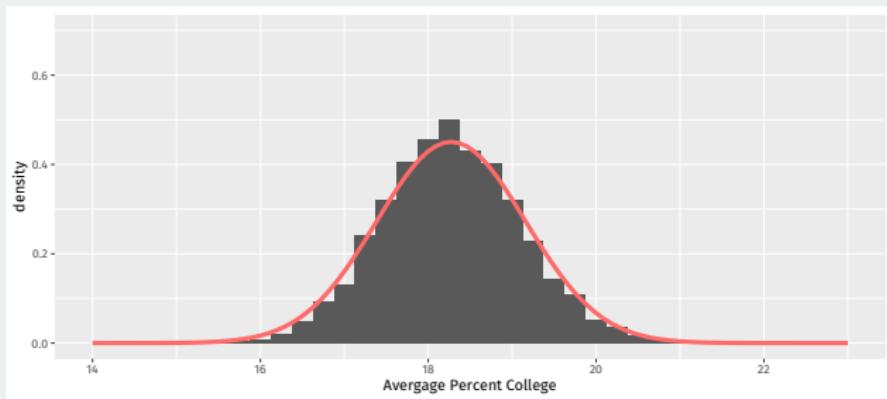


More population spread → higher SE

# Similarity in the bootstrap/null distributions



# Conditions for the CLT

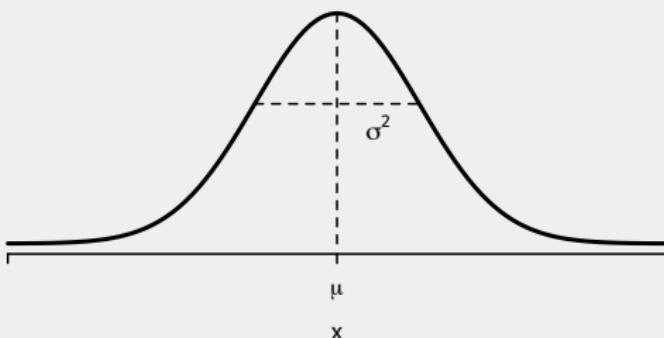


**Central limit theorem:** sums and means of **random samples** tend to be normally distributed as the **sample size grows**.

Many, many estimators will follow the CLT and have a normal distribution and will be easier to use this to do inference rather than doing increasingly complicated simulations.

# **2/** Normal distribution

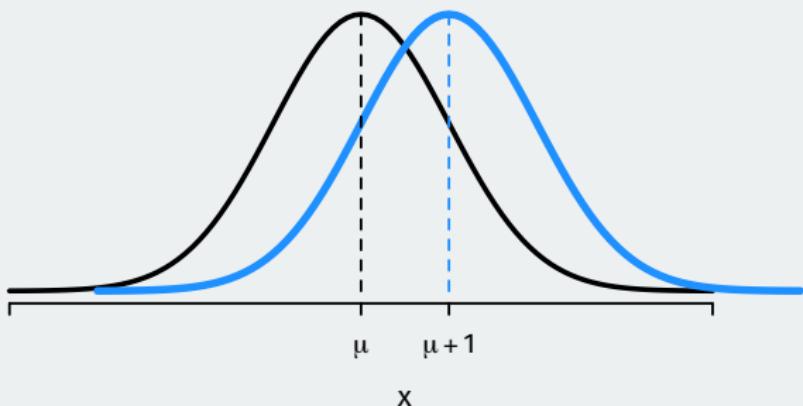
# Normal distribution



- A normal distribution can be affected by two values:
  - **mean/expected value** usually written as  $\mu$
  - **variance** written as  $\sigma^2$  (standard deviation is  $\sigma$ )
  - Written  $X \sim N(\mu, \sigma^2)$ .
- **Standard normal distribution:** mean 0 and standard deviation 1.

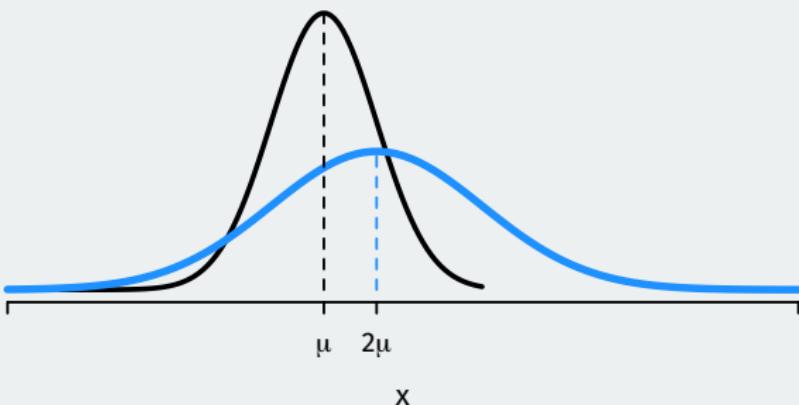
# Reentering and scaling the normal

- How do transformations of a normal work?
- Let  $X \sim N(\mu, \sigma^2)$  and  $c$  be a constant.
- If  $Z = X + c$ , then  $Z \sim N(\mu + c, \sigma^2)$ .
- Intuition: adding a constant to a normal shifts the distribution by that constant.



# Recentering and scaling the normal

- Let  $X \sim N(\mu, \sigma^2)$  and  $c$  be a constant.
- If  $Z = cX$ , then  $Z \sim N(c\mu, (c\sigma)^2)$ .
- Intuition: multiplying a normal by a constant scales the mean and the variance.



# Z-scores of normals

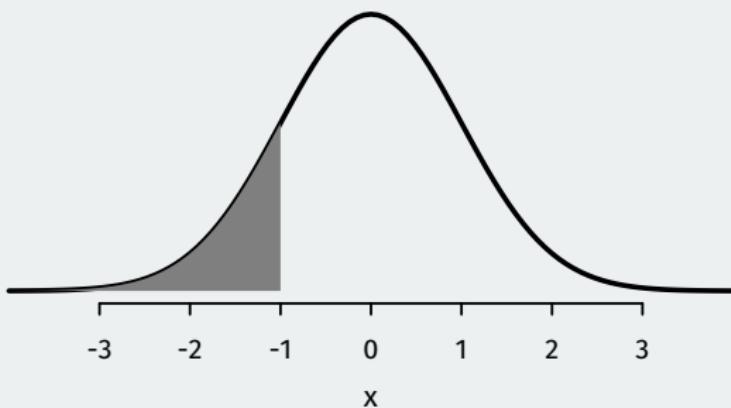
- These facts imply the **z-score** of a normal variable is a standard normal:

$$z = \frac{X - \mu}{\sigma} \sim N(0, 1)$$

- Subtract the mean and divide by the SD  $\rightsquigarrow$  standard normal.
- z-score measures how many SDs away from the mean a value of  $X$  is.

# Normal probability calculations

What's the probability of being below -1 for a standard normal?



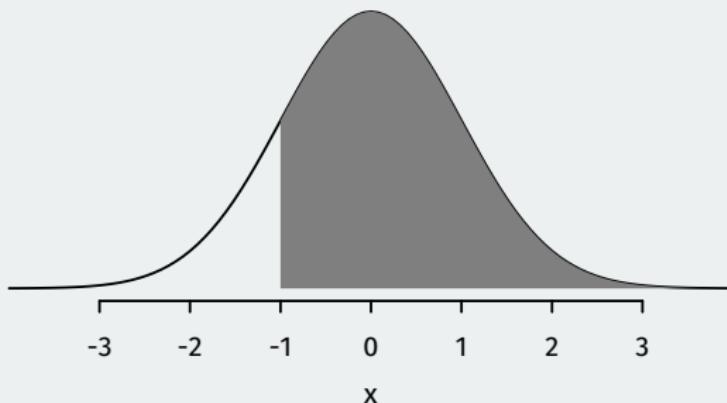
This is the area under the normal curve, which `pnorm( )` function gives us this:

```
pnorm(-1, mean = 0, sd = 1)
```

```
## [1] 0.159
```

# Normal probability calculations

What's the probability of being **above** -1 for a standard normal?



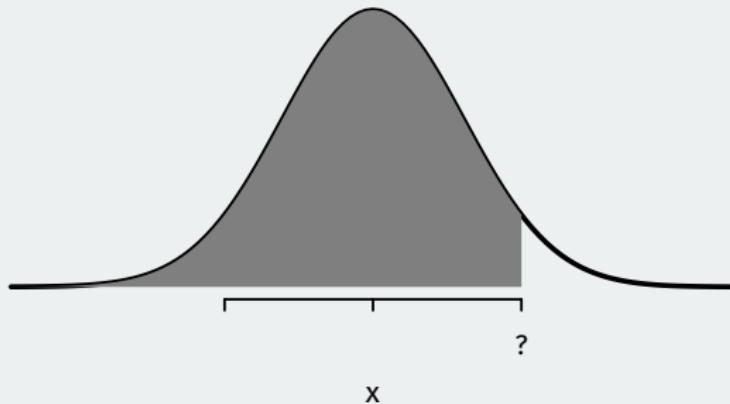
Total area under the curve (1) minus the area below -1:

```
1 - pnorm(-1, mean = 0, sd = 1)
```

```
## [1] 0.841
```

# Normal quantiles

What if we want to know the opposite? What value of the normal distribution puts 95% of the distribution below it?



This is a **quantile** and we can get it using `qnorm()`:

```
qnorm(0.95, mean = 0, sd = 1)
```

```
## [1] 1.64
```

# **3/** Using the Normal for inference

# How popular is Joe Biden?



- What proportion of the public approves of Biden's job as president?
- Latest Gallup poll:
  - Sept 1st-16th
  - 812 adult Americans
  - Telephone interviews
  - Approve (42%), Disapprove (56%)
- Define r.v.  $Y_i$  for Biden approval:
  - $Y_i = 1 \rightsquigarrow$  respondent  $i$  approves of Biden, 0 otherwise.
  - $p = \mathbb{P}(Y_i = 1)$  the population proportion of Biden approvers.
  - $\bar{Y} = 0.42$  is the sample proportion.

# Standard errors for sample proportions

How variable will our sample proportion be? Depends on the **standard error**.

Special rule for SEs of sample proportion  $\bar{Y}$ :

$$SE \text{ for } \bar{Y} = \sqrt{\frac{p(1-p)}{n}} = \sqrt{\frac{(\text{pop. proportion}) \times (1 - \text{pop. proportion})}{\text{sample size}}}$$

Because we don't know  $p$ , we replace it with our best guess,  $\bar{Y}$ :

$$\widehat{SE} = \sqrt{\frac{\bar{Y}(1-\bar{Y})}{n}}$$

# CLT for confidence intervals

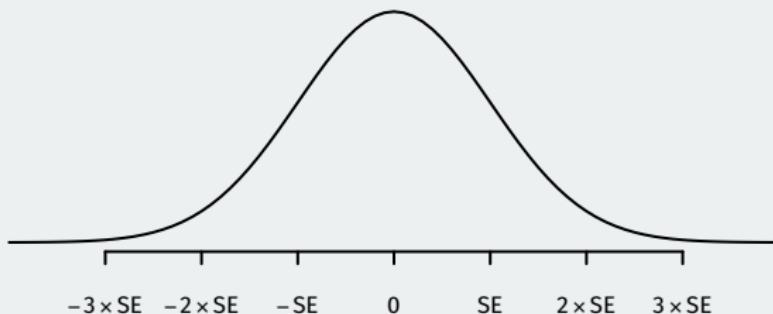
$$\bar{Y} - p = \text{chance error}$$

- How can we figure out a range of plausible chance errors?
  - Find a range of plausible chance errors and add them to  $\bar{Y}$
  - With **bootstrap**, we used resampling to simulate chance error.
- Central limit theorem implies

$$\bar{Y} \approx N\left(p, \frac{p(1-p)}{n}\right)$$

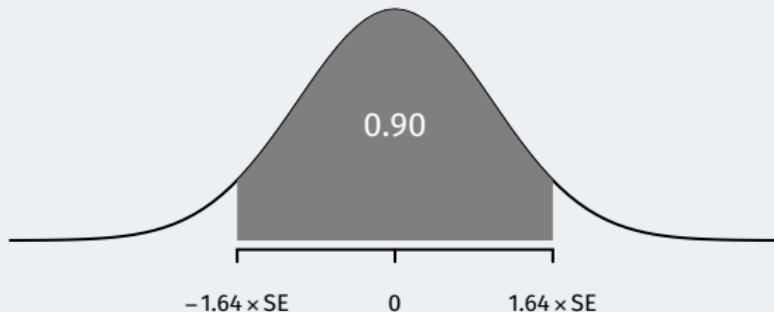
Chance error:  $\bar{Y} - p$  is approximately normal with mean 0 and SE equal to  $\sqrt{p(1-p)/n}$

# Chance errors



If  $\bar{Y} \sim N(p, SE^2)$ , then chance errors are  $\bar{Y} - p \sim N(0, SE^2)$  so:

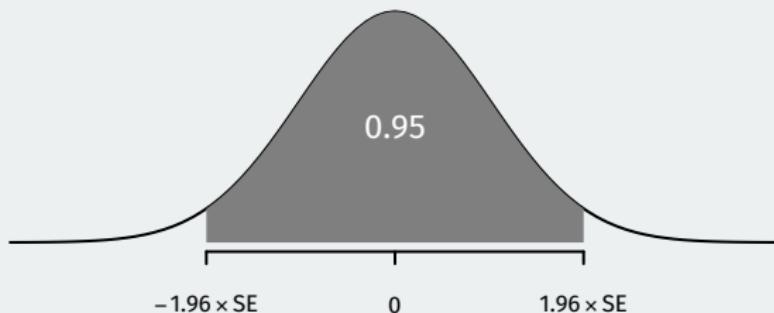
# Chance errors



If  $\bar{Y} \sim N(p, SE^2)$ , then chance errors are  $\bar{Y} - p \sim N(0, SE^2)$  so:

- $\approx 90\%$  of chance errors  $\bar{Y} - p$  are within 1.64 SEs of the mean.

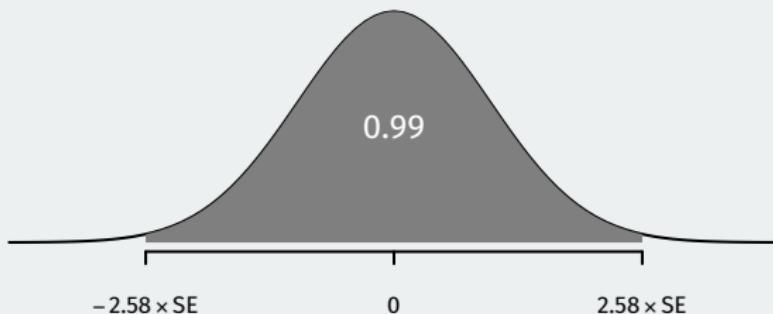
# Chance errors



If  $\bar{Y} \sim N(p, SE^2)$ , then chance errors are  $\bar{Y} - p \sim N(0, SE^2)$  so:

- $\approx 90\%$  of chance errors  $\bar{Y} - p$  are within 1.64 SEs of the mean.
- $\approx 95\%$  of chance errors  $\bar{Y} - p$  are within 1.96 SEs of the mean.

# Chance errors



If  $\bar{Y} \sim N(p, SE^2)$ , then chance errors are  $\bar{Y} - p \sim N(0, SE^2)$  so:

- $\approx 90\%$  of chance errors  $\bar{Y} - p$  are within 1.64 SEs of the mean.
- $\approx 95\%$  of chance errors  $\bar{Y} - p$  are within 1.96 SEs of the mean.
- $\approx 99\%$  of chance errors  $\bar{Y} - p$  are within 2.58 SEs of the mean.

This implies we can build a 95% confidence interval with  $\bar{Y} \pm 1.96 \times SE$

# How did we get those values?

- First, choose a **confidence level**.
  - What percent of chance errors do you want to count as “plausible”?
  - Convention is 95%.
- $100 \times (1 - \alpha)\%$  confidence interval:

$$CI = \bar{Y} \pm z_{\alpha/2} \times SE$$

- In polling,  $\pm z_{\alpha/2} \times SE$  is called the **margin of error**
- $z_{\alpha/2}$  is the  $N(0, 1)$  z-score that would put  $\alpha/2$  in the upper tail:
  - $\mathbb{P}(-z_{\alpha/2} < Z < z_{\alpha/2}) = \alpha$
  - 90% CI  $\rightsquigarrow \alpha = 0.1 \rightsquigarrow z_{\alpha/2} = 1.64$
  - 95% CI  $\rightsquigarrow \alpha = 0.05 \rightsquigarrow z_{\alpha/2} = 1.96$
  - 99% CI  $\rightsquigarrow \alpha = 0.01 \rightsquigarrow z_{\alpha/2} = 2.58$

# Standard normal z-scores in R

`qnorm(x, lower.tail = FALSE)` will find the quantile of  $N(0, 1)$  that puts  $x$  in the upper tail:

```
qnorm(0.05, lower.tail = FALSE)
```

```
## [1] 1.64
```

```
qnorm(0.025, lower.tail = FALSE)
```

```
## [1] 1.96
```

```
qnorm(0.005, lower.tail = FALSE)
```

```
## [1] 2.58
```

# Gov 50: 24. More Inference with Mathematical Models

Matthew Blackwell

Harvard University

# Roadmap

1. Confidence intervals for experiments
2. Hypothesis testing with the CLT
3. Two-sample tests

# 1/ Confidence intervals for experiments

# Comparison between groups

- More interesting to compare across groups.
  - Differences in public opinion across groups
  - Difference between treatment and control groups.
- Bedrock of causal inference!

# Social pressure experiment

- Back to the Social Pressure Mailer GOTV example.
  - Primary election in MI 2006
- Treatment group: postcards showing their own and their neighbors' voting records.
  - Sample size of treated group,  $n_T = 360$  (artificially reducing sample size to highlight the math)
- Control group: received nothing.
  - Sample size of the control group,  $n_C = 1890$

# Outcomes

- Outcome:  $Y_i = 1$  if  $i$  voted, 0 otherwise.
- Turnout rate (sample mean) in treated group,  $\bar{Y}_T = 0.37$
- Turnout rate (sample mean) in control group,  $\bar{Y}_C = 0.30$
- Estimated **average treatment effect**

$$\widehat{\text{ATE}} = \bar{Y}_T - \bar{Y}_C = 0.07$$

# Inference for the difference

- Parameter: **population ATE**  $\mu_T - \mu_C$ 
  - $\mu_T$ : Turnout rate in the population if everyone received treatment.
  - $\mu_C$ : Turnout rate in the population if everyone received control.
- Estimator:  $\widehat{\text{ATE}} = \bar{Y}_T - \bar{Y}_C$

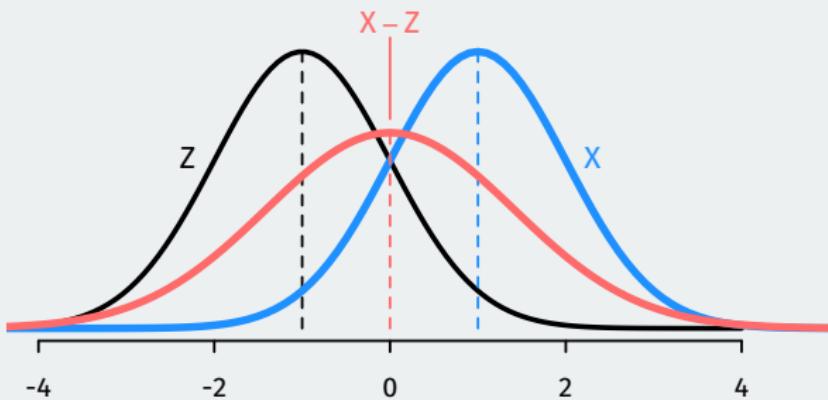
By the CLT in large samples, we know that:

- $\bar{Y}_T \approx N\left(\mu_T, \frac{\mu_T(1-\mu_T)}{n_c}\right)$
- $\bar{Y}_C \approx N\left(\mu_C, \frac{\mu_C(1-\mu_C)}{n_c}\right)$
- $\rightsquigarrow \bar{Y}_T - \bar{Y}_C \approx N(\mu_T - \mu_C, SE_{\text{diff}}^2)$

But what is the  $SE_{\text{diff}}$  in this case?

# Spread of a difference in normals

If we take the difference between two independent normal r.v.s, what happens to the spread?



The spread of the difference is **larger** than the spread of the two variables being differenced!

# Standard error for the estimated ATE

- SE of a difference in means **adds** the SEs for each group

$$SE_{\text{diff}} = \sqrt{SE_T^2 + SE_C^2}$$

- Using what we know about SEs with binary outcomes:

$$SE_{\text{diff}} = \sqrt{\frac{\mu_T(1 - \mu_T)}{n_t} + \frac{\mu_C(1 - \mu_C)}{n_C}}$$

- Chance errors  $\bar{Y}_T - \bar{Y}_C - (\mu_T - \mu_C) \approx N(0, SE_{\text{diff}})$ 
  - We can construct a 95% CI with  $\widehat{\text{ATE}} \pm 1.96 \times SE_{\text{diff}}$

# Confidence intervals

But we don't know  $\mu_T$  or  $\mu_C$ ! Plug in our sample proportions to estimate the SE:

$$\begin{aligned}\widehat{SE}_{\text{diff}} &= \sqrt{\frac{\bar{Y}_T(1 - \bar{Y}_T)}{n_t} + \frac{\bar{Y}_C(1 - \bar{Y}_C)}{n_C}} \\ &= \sqrt{\frac{0.37 \times 0.63}{360} + \frac{0.3 \times 0.7}{1890}} = 0.028\end{aligned}$$

Now we can construct confidence intervals based on the CLT like last time:

$$\begin{aligned}CI_{95} &= \widehat{ATE} \pm 1.96 \times \widehat{SE}_{\text{diff}} \\ &= 0.07 \pm 1.96 \times 0.028 \\ &= 0.07 \pm 0.054 \\ &= [0.016, 0.124]\end{aligned}$$

Range of possibilities taking into account plausible chance errors.

# **2/** Hypothesis testing with the CLT

# Statistical hypothesis testing

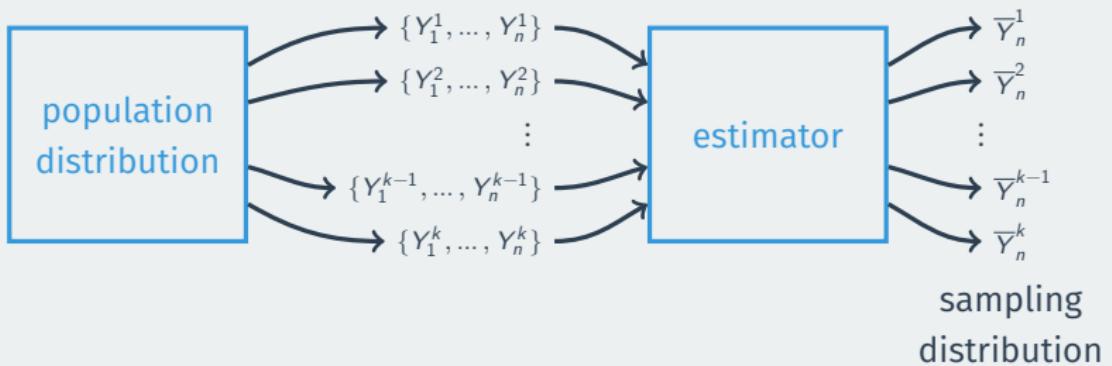
- Statistical hypothesis testing is a **thought experiment**.
- What would the world look like **if we knew the truth?**
- Conducted with several steps:
  1. Specify your **null** and **alternative hypotheses**
  2. Choose an appropriate **test statistic** and level of test  $\alpha$
  3. Derive the **reference distribution** of the test statistic under the null.
  4. Use this distribution to calculate the **p-value**.
  5. Use p-value to decide whether to reject the null hypothesis or not

# How popular is Joe Biden?



- What proportion of the public approves of Biden's job as president?
- Latest Gallup poll:  $\bar{Y} = 0.42$  with  $n = 812$
- Could we reject the null that Biden's national support is 50%?
  - Null:  $H_0 : p = 0.5$
  - Alternative:  $H_1 : p \neq 0.5$

# Sampling distribution, in pictures



# CLT for hypothesis testing

Under the null, we know the distribution of  $\bar{Y}$ :

$$\bar{Y} \approx N\left(p, \frac{p(1-p)}{n}\right) = N\left(0.5, \frac{0.5 \times 0.5}{812}\right)$$

Using the rules of normal transformations if  $X \sim N(\mu, \sigma^2)$ :

$$\frac{X - \mu}{\sigma} \sim N(0, 1)$$

Then under the null, know the distribution of the following test statistic:

$$Z = \frac{\bar{Y} - 0.5}{0.5/\sqrt{812}} \approx N(0, 1)$$

# p-values

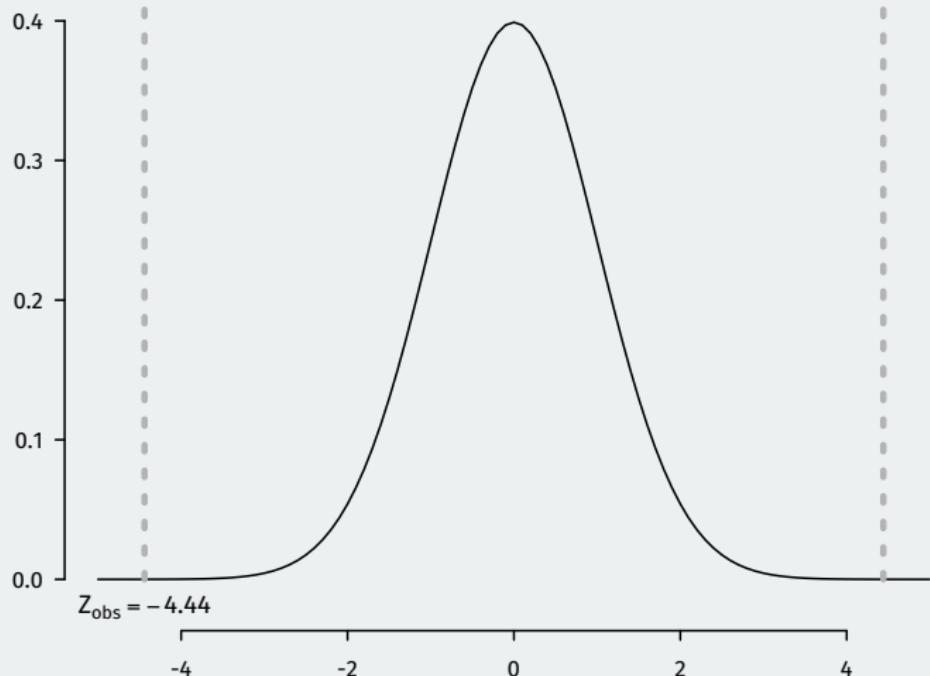
What we observe:

$$\begin{aligned} Z_{\text{obs}} &= \frac{\bar{Y} - 0.5}{0.5/\sqrt{812}} = \frac{0.42 - 0.5}{0.5/\sqrt{812}} \\ &= -\frac{0.08}{0.018} = -4.44 \end{aligned}$$

Our observed sample proportion is 4.44 SEs away from 0.5 under the null.

What's the probability of being that far away? (**p-value**)

```
pnorm(-4.44, mean = 0, sd = 1) + ## prob being below -4.44  
(1 - pnorm(4.44, mean = 0, sd = 1)) ## prob being above 4.44  
  
## [1] 0.000009
```



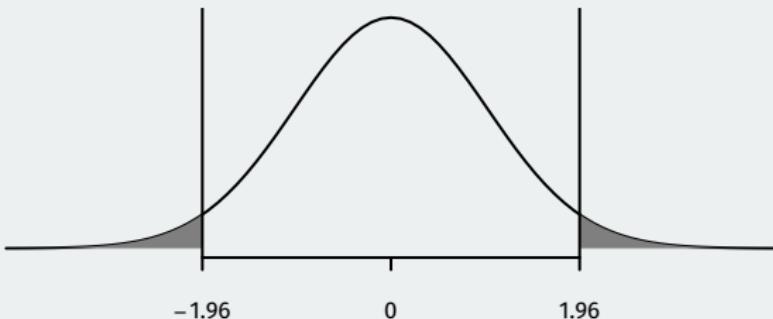
# Generalizing hypothesis tests

- Hypothesis testing using the CLT pretty much takes this general form no matter what the estimator of interest is.
- Hypotheses:  $H_0 : \mu = \mu_0$  (null guess),  $H_1 : \mu \neq \mu_0$
- Test statistic:

$$Z = \frac{\text{observed value} - \text{null guess}}{\widehat{SE}} = \frac{\bar{Y} - \mu_0}{\widehat{SE}}$$

- The exact estimator for the standard error  $\widehat{SE}$  will depend on the estimator of interest.
- Null distribution:  $Z \approx N(0, 1)$  by the CLT
- p-value: probability of a standard normal being bigger than  $|Z_{\text{obs}}|$

# Rejecting regions



- Reject if p-value is below  $\alpha$  (usually 0.05).
  - We know 5% of the time  $Z$  will be bigger than 1.96.
  - If  $Z_{\text{obs}} > 1.96$  or  $Z_{\text{obs}} < -1.96$ , then the p-value must be below 0.05
  - We can reject if  $|Z_{\text{obs}}| > 1.96$

# 3/ Two-sample tests

# Two-sample hypotheses

- Parameter: **population ATE**  $\mu_T - \mu_C$
- Goal: learn about the population difference in means
- Usual null hypothesis: no difference in population means (ATE = 0)
  - Null:  $H_0 : \mu_T - \mu_C = 0$
  - Two-sided alternative:  $H_1 : \mu_T - \mu_C \neq 0$
- In words: are the differences in sample means just due to chance?

# Difference-in-means review

- Sample turnout rates:  $\bar{Y}_T = 0.37$ ,  $\bar{Y}_C = 0.30$
- Sample sizes:  $n_T = 360$ ,  $n_C = 1890$
- Estimator is the **sample difference-in-means**:

$$\widehat{\text{ATE}} = \bar{Y}_T - \bar{Y}_C = 0.07$$

- Estimated SE for the difference in means:

$$\widehat{\text{SE}}_{\text{diff}} = \sqrt{\frac{\bar{Y}_T(1 - \bar{Y}_T)}{n_T} + \frac{\bar{Y}_C(1 - \bar{Y}_C)}{n_C}} = 0.028$$

# CLT again and again

Earlier we saw that by the CLT we have:

$$\bar{Y}_T - \bar{Y}_C \approx N(\mu_T - \mu_C, SE_{\text{diff}}^2)$$

We can use Z-scores to get a test statistic:

$$Z = \frac{(\bar{Y}_T - \bar{Y}_C) - (\mu_T - \mu_C)}{SE_{\text{diff}}} \sim N(0, 1)$$

Same general form of the test statistic as with one sample mean/proportion:

$$\frac{\text{observed} - \text{null guess}}{\text{SE}}$$

# The usual null of no difference

- Null hypothesis:  $H_0 : \mu_T - \mu_C = 0$
- Test statistic:

$$Z = \frac{(\bar{Y}_T - \bar{Y}_C) - (\mu_T - \mu_C)}{\text{SE}_{\text{diff}}} = \frac{(\bar{Y}_T - \bar{Y}_C) - 0}{\text{SE}_{\text{diff}}}$$

- In large samples, we can replace true SE with an estimate:

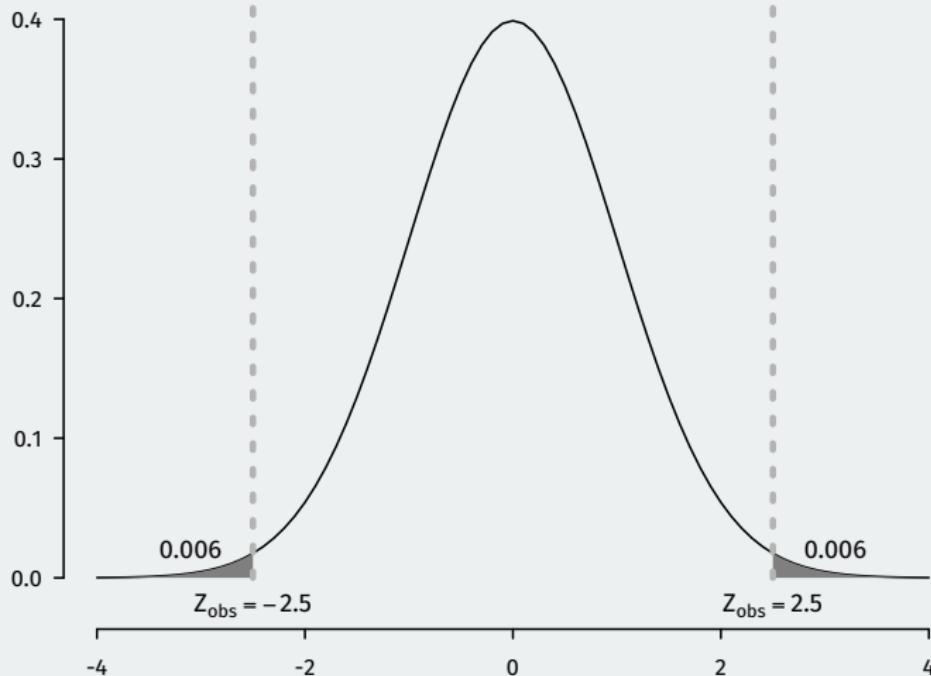
$$\widehat{\text{SE}}_{\text{diff}} = \sqrt{\widehat{\text{SE}}_T^2 + \widehat{\text{SE}}_C^2}$$

# Calculating p-values

- Finally! Our test statistic in this sample:

$$Z = \frac{\bar{Y}_T - \bar{Y}_C}{\widehat{SE}_{\text{diff}}} = \frac{0.07}{0.028} = 2.5$$

- p-value based on a two-sided test: probability of getting a difference in means this big (or bigger) if the null hypothesis were true
  - Lower p-values  $\rightsquigarrow$  stronger evidence against the null.



```
2 * pnorm(2.5, lower.tail = FALSE)
```

```
## [1] 0.0124
```

# Gov 50: 25. Inference for Linear Regression

Matthew Blackwell

Harvard University

# Roadmap

1. Inference for linear regression
2. Presenting OLS regressions
3. Wrapping up the class

# 1/ Inference for linear regression

# Data

- Do political institutions promote economic development?
  - Famous paper on this: Acemoglu, Johnson, and Robinson (2001)
  - Relationship between strength of property rights in a country and GDP.
- Data:

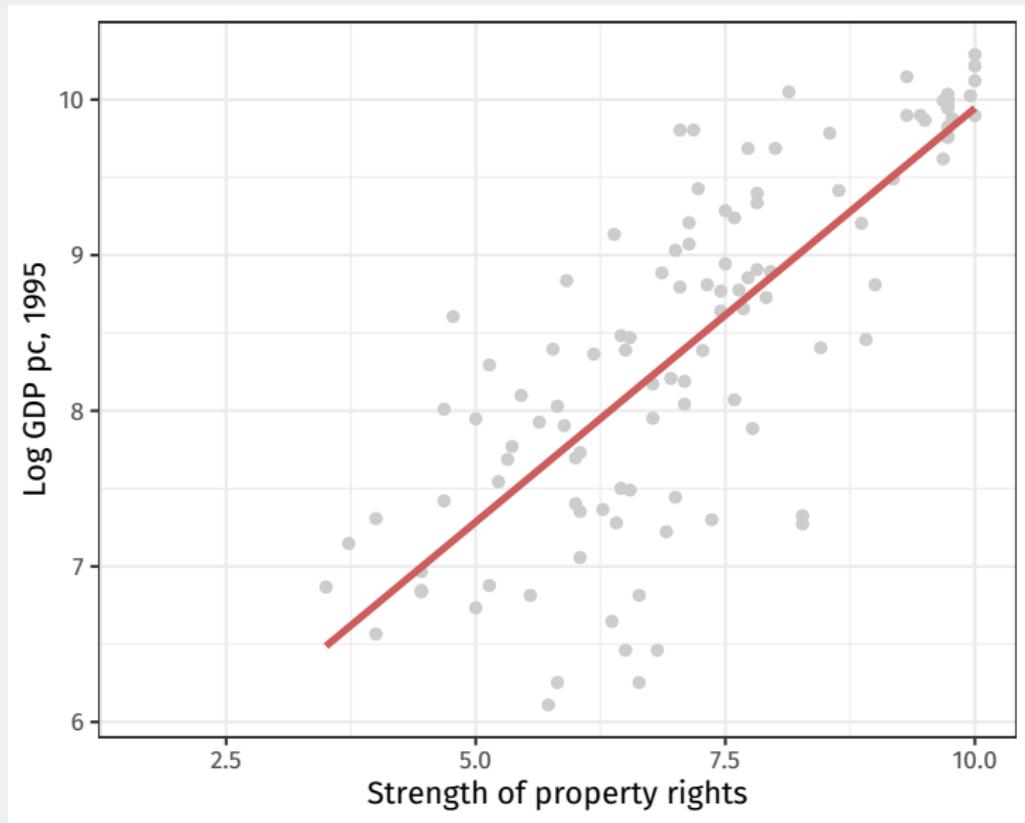
| Name                  | Description  |
|-----------------------|--|
| <code>shortnam</code> | three-letter country code                                      |
| <code>africa</code>   | indicator for if the country is in Africa                      |
| <code>asia</code>     | indicator for if country is in Asia                            |
| <code>avexpr</code>   | strength of property rights (protection against expropriation) |
| <code>logpgp95</code> | log GDP per capita   |

# Loading the data

```
library(gov50data)
head(ajr)
```

```
## # A tibble: 6 x 15
##   short~1 africa lat_a~2 malfa~3 avexpr logpg~4 logem4  asia
##   <chr>     <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 AFG        0     0.367  0.00372    NA      NA     4.54     1
## 2 AGO        1     0.137  0.950      5.36    7.77    5.63     0
## 3 ARE        0     0.267  0.0123     7.18    9.80    NA      1
## 4 ARG        0     0.378  0          6.39    9.13    4.23     0
## 5 ARM        0     0.444  0          NA      7.68    NA      1
## 6 AUS        0     0.300  0          9.32    9.90    2.15     0
## # ... with 7 more variables: yellow <dbl>, baseco <dbl>,
## #   leb95 <dbl>, imr95 <dbl>, meantemp <dbl>,
## #   lt100km <dbl>, latabs <dbl>, and abbreviated variable
## #   names 1: shortnam, 2: lat_abst, 3: malfal94,
## #   4: logpgp95
```

# AJR scatterplot



# Simple linear regression model

- We are going to assume a linear model:

$$Y_i = \beta_0 + \beta_1 X_i + \varepsilon_i$$

- Data:
  - Dependent variable:  $Y_i$
  - Independent variable:  $X_i$
- Population parameters:
  - Population intercept:  $\beta_0$
  - Population slope:  $\beta_1$
- Error/disturbance:  $\varepsilon_i$ 
  - Represents all unobserved error factors influencing  $Y_i$  other than  $X_i$ .

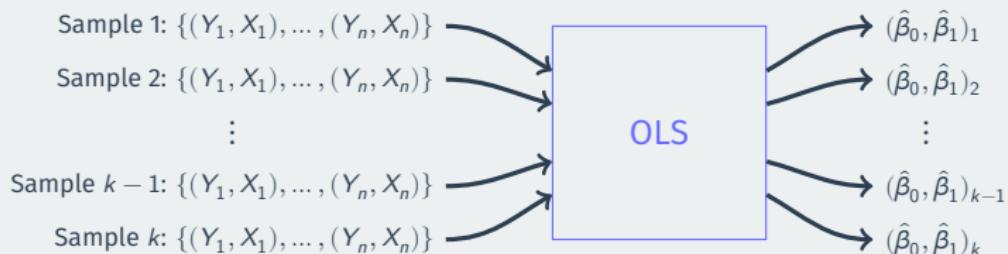
# Least squares

- How do we figure out the best line to draw?
  - Alt question: how do we figure out  $\beta_0$  and  $\beta_1$ ?
  - $(\hat{\beta}_0, \hat{\beta}_1)$ : estimated coefficients.
  - $\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_i$ : predicted/fitted value.
  - $\hat{\epsilon}_i = Y_i - \hat{Y}_i$ : residual.
- Get these estimates by the **least squares method**.
- Minimize the **sum of the squared residuals** (SSR):

$$SSR = \sum_{i=1}^n \hat{\epsilon}_i^2 = \sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i)^2$$

# Estimators

- Least squares is an **estimator**
  - it's a machine that we plug data into and we get out estimates.

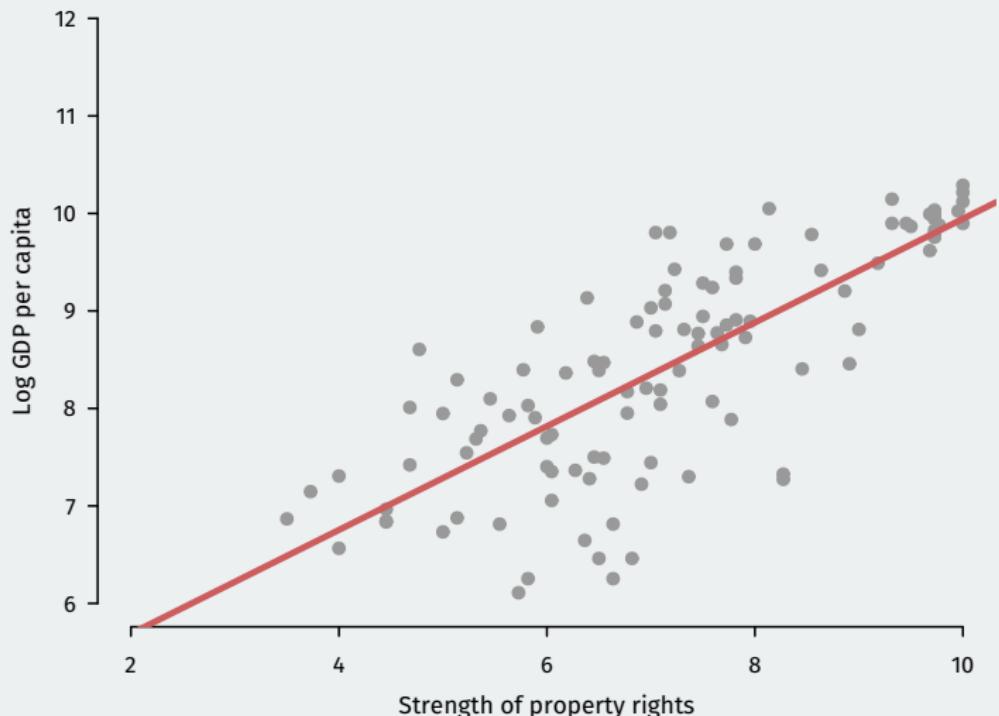


- Just like the sample mean or difference in sample means
- $\rightsquigarrow$  sampling distribution with a standard error, etc.

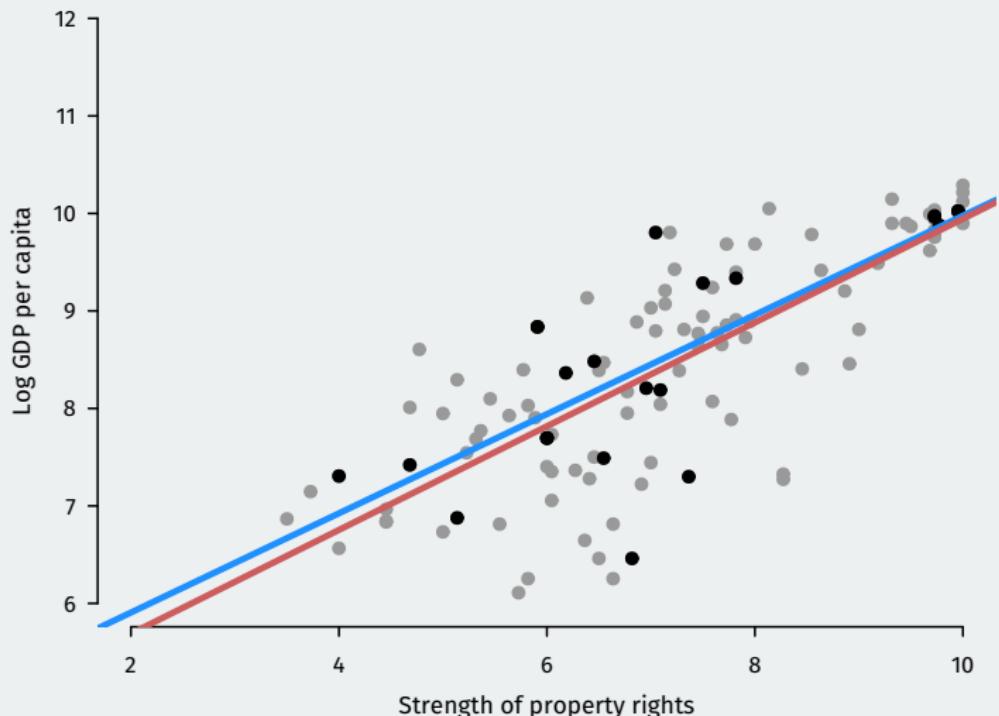
# Simulation procedure

- Let's take a simulation approach to demonstrate:
    - Pretend that the AJR data represents the population of interest
    - See how the line varies from sample to sample
1. Randomly sample  $n = 30$  countries w/ replacement using `sample()`
  2. Use `lm()` to calculate the OLS estimates of the slope and intercept
  3. Plot the estimated regression line

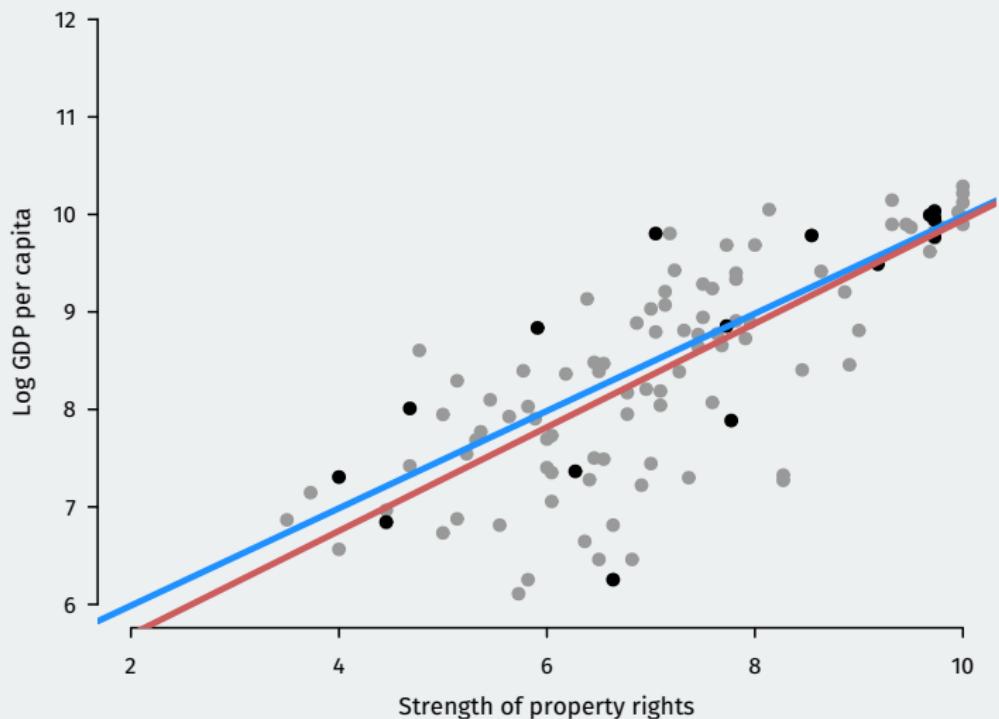
# Population regression



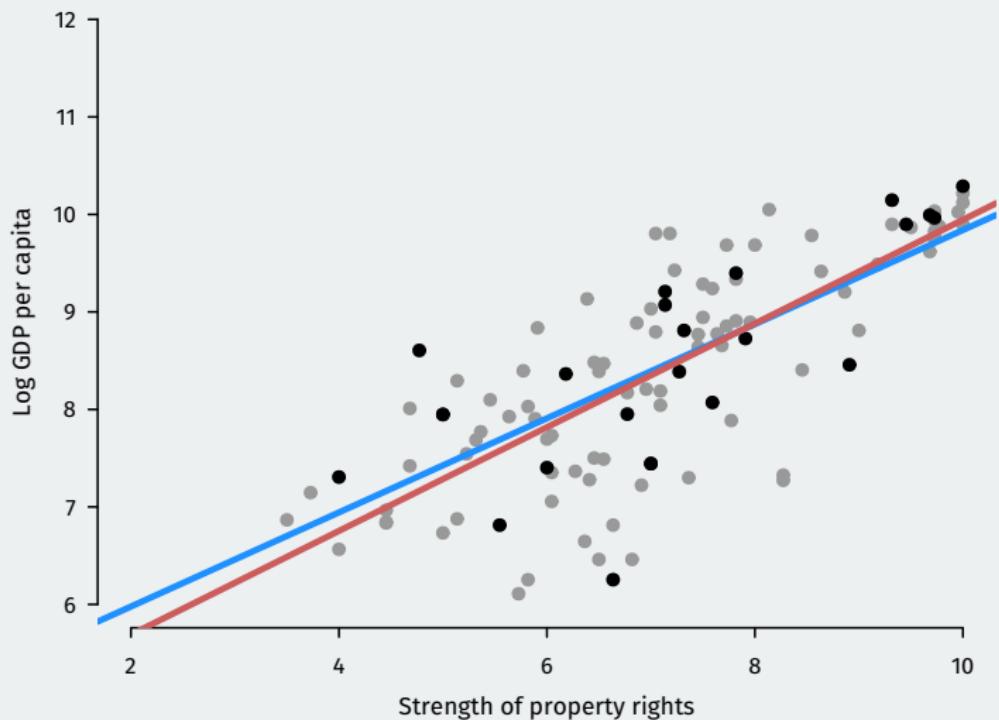
# Randomly sample from AJR



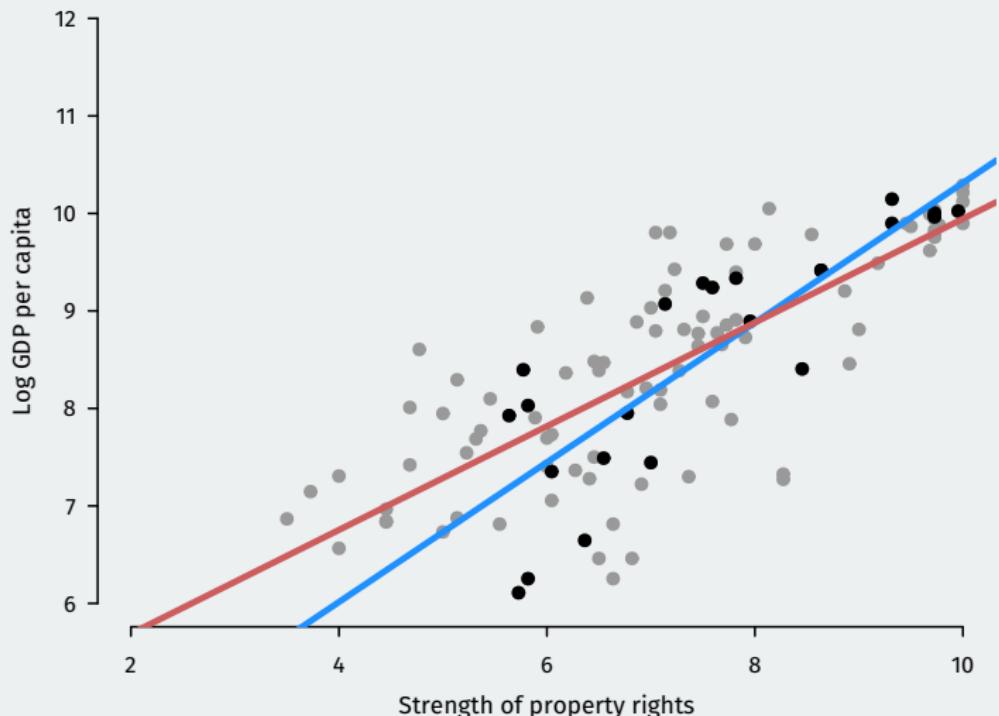
# Randomly sample from AJR



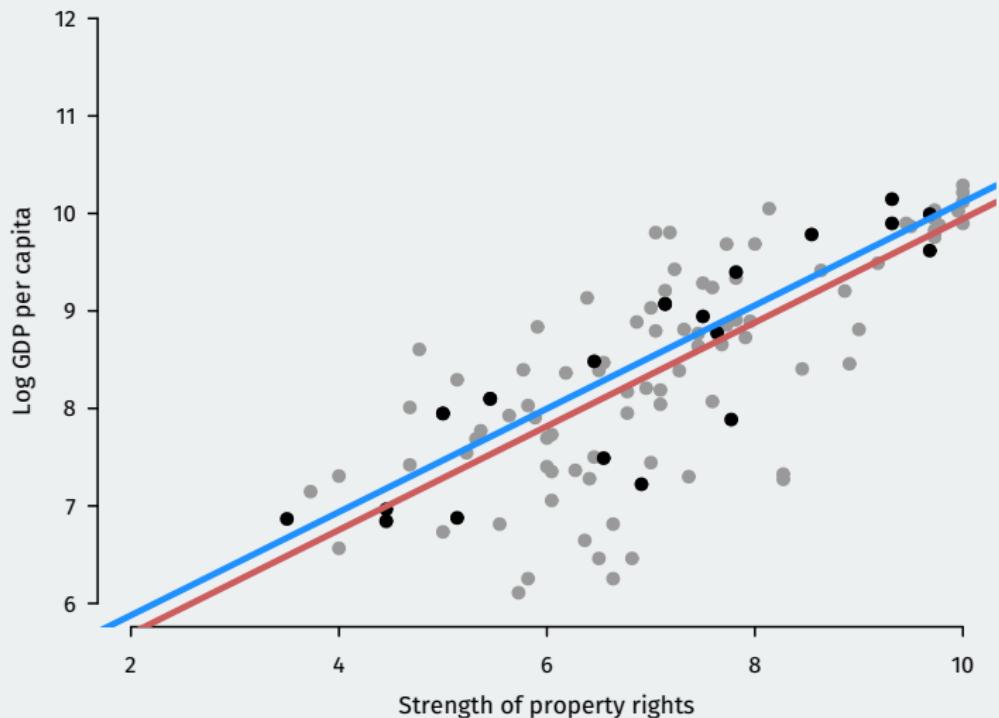
# Randomly sample from AJR



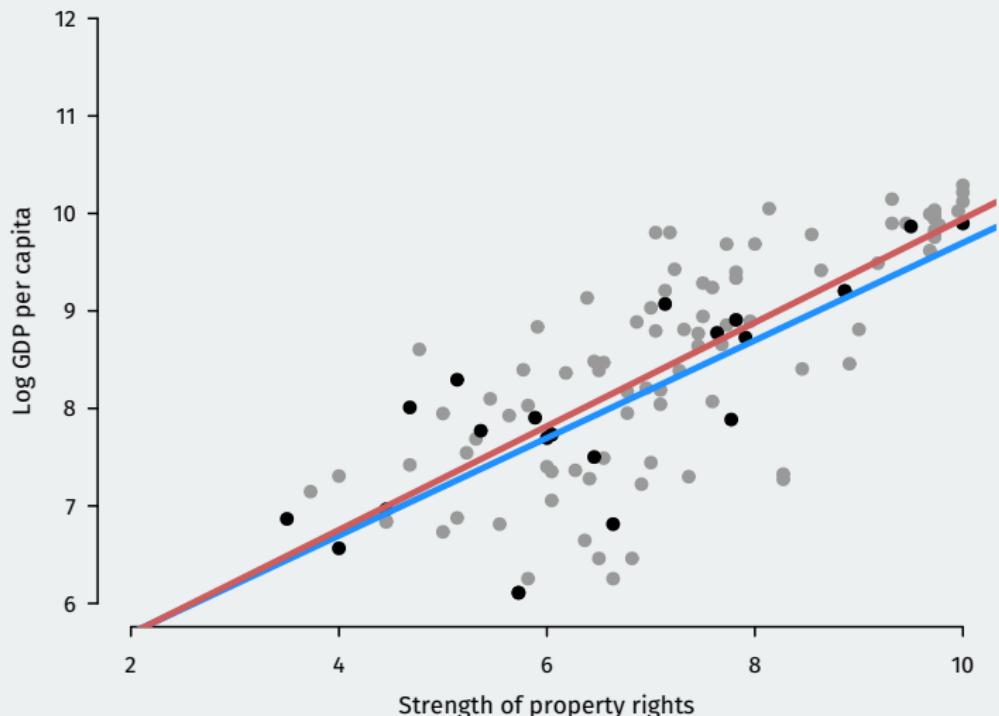
# Randomly sample from AJR



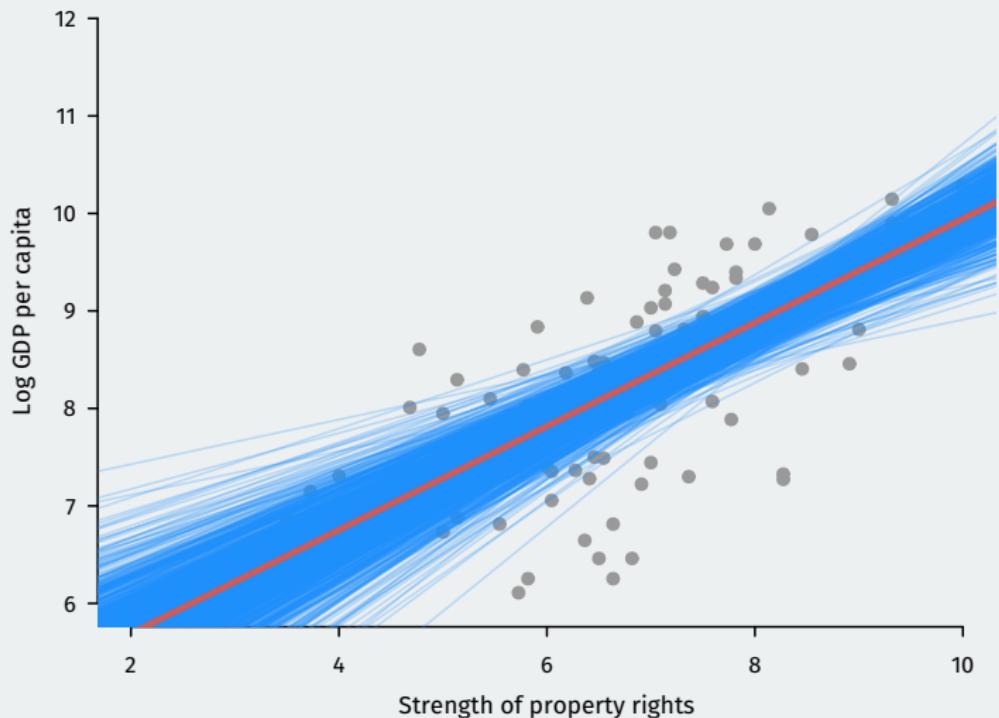
# Randomly sample from AJR



# Randomly sample from AJR

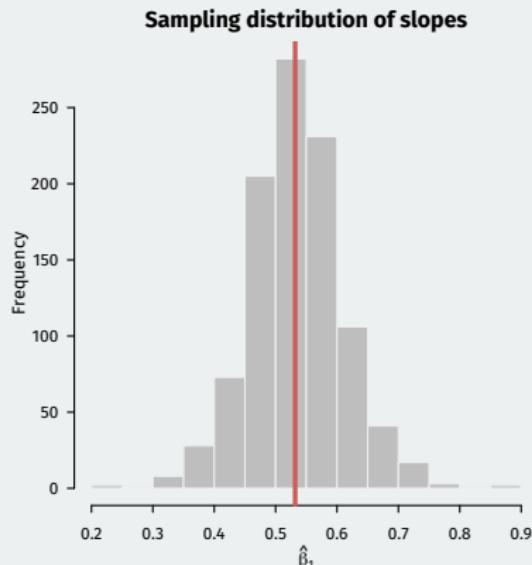
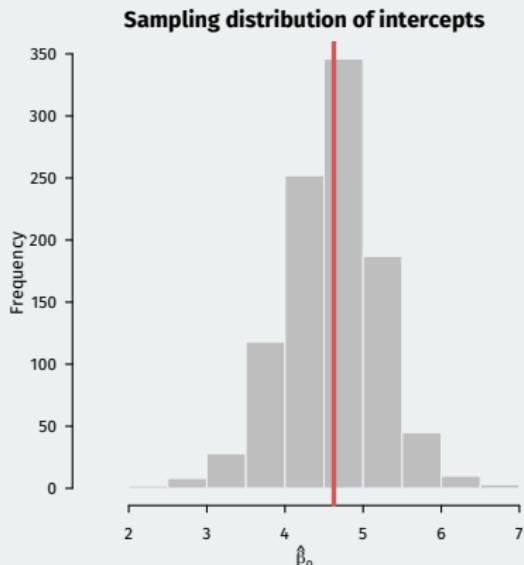


# Randomly sample from AJR



# Sampling distribution of OLS

- Estimated slope and intercept vary between samples, centered on truth.



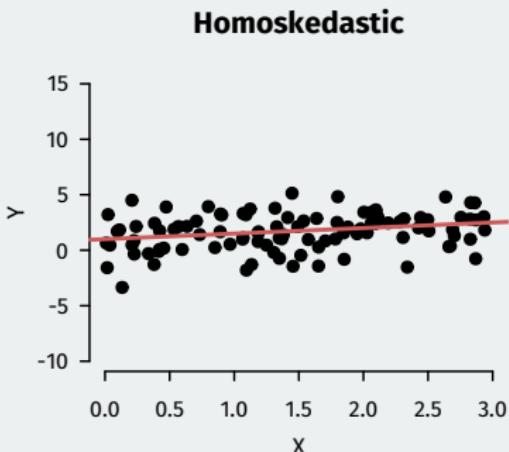
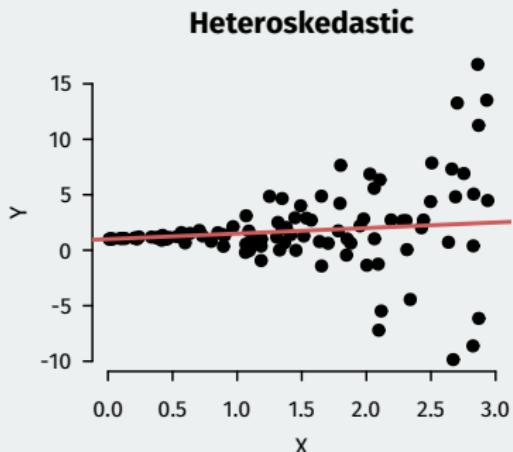
# Properties of OLS

- $\hat{\beta}_0$  and  $\hat{\beta}_1$  are random variables
  - Are they on average equal to the true values (bias)?
  - How spread out are they around their center (variance)?
- Under minimal conditions,  $\hat{\beta}_0$  and  $\hat{\beta}_1$  are unbiased for the population line of best fit, but...
  - This might be misleading if the true relationship is nonlinear.
  - May not represent a causal effect unless causal assumptions hold.

# Standard errors of OLS

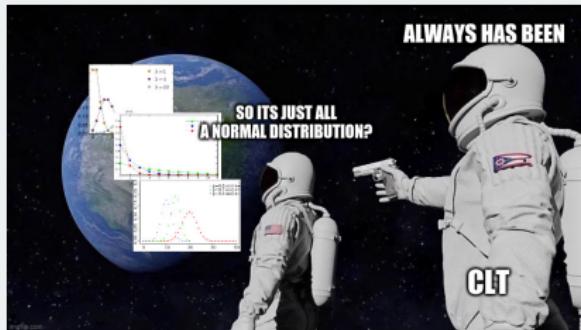
R will also calculate an estimate of the standard error:  $\widehat{SE}(\hat{\beta}_1)$

Default estimators for the SEs assume **homoskedasticity** or that the spread around the regression line is the same for all values of the independent variables.



Relatively easy fixes exist, but beyond the scope of this class.

# Tests and CIs for regression



- $(\hat{\beta}_0, \hat{\beta}_1)$  can be written as weighted averages of the outcome...
  - Which means they follow the Central Limit Theorem!
- BAM! 95% confidence intervals:  $\hat{\beta}_1 \pm 1.96 \times \widehat{SE}(\hat{\beta}_1)$
- BOOM! Hypothesis tests:
  - Null hypothesis:  $H_0 : \beta_1 = \beta_1^*$
  - Test statistic:  $\frac{\hat{\beta}_1 - \beta_1^*}{\widehat{SE}(\hat{\beta}_1)} \sim N(0, 1)$
  - Usual test is of  $\beta_1 = 0$ .
  - $\hat{\beta}_1$  is **statistically significant** if its p-value from this test is below some threshold (usually 0.05)

```
ajr.reg <- lm(logpgp95 ~ avexpr, data = ajr)
summary(ajr.reg)

##
## Call:
## lm(formula = logpgp95 ~ avexpr, data = ajr)
##
## Residuals:
##    Min      1Q  Median      3Q     Max 
## -1.902 -0.316  0.138  0.422  1.441 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept)  4.6261    0.3006   15.4   <2e-16 ***
## avexpr       0.5319    0.0406   13.1   <2e-16 ***
## ---
## Signif. codes:
## 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.718 on 109 degrees of freedom
##   (52 observations deleted due to missingness)
## Multiple R-squared:  0.611, Adjusted R-squared:  0.608 
## F-statistic: 171 on 1 and 109 DF,  p-value: <2e-16
```

# Using broom with regression

```
library(broom)
tidy(ajr.reg)

## # A tibble: 2 x 5
##   term      estimate std.error statistic p.value
##   <chr>      <dbl>     <dbl>      <dbl>    <dbl>
## 1 (Intercept)  4.63     0.301     15.4 4.28e-29
## 2 avexpr       0.532    0.0406    13.1 4.16e-24
```

# Multiple regression

- Correlation doesn't imply causation
- Omitted variables  $\rightsquigarrow$  violation of exogeneity
- You can adjust for multiple confounding variables:

$$Y_i = \beta_0 + \beta_1 X_{i1} + \beta_2 X_{i2} + \cdots + \beta_p X_{ip} + \epsilon_i$$

- Interpretation of  $\beta_j$ : an increase in the outcome associated with a one-unit increase in  $X_{ij}$  when other variables don't change their values
- Inference:
  - Confidence intervals constructed exactly the same for  $\hat{\beta}_j$
  - Hypothesis tests done exactly the same for  $\hat{\beta}_j$
  - $\rightsquigarrow$  interpret p-values the same as before.

# Using knitr::kable to produce tables

```
ajr.multreg <- lm(logpgp95 ~ avexpr + lat_abst + asia + africa, data = ajr)
tidy(ajr.multreg) |>
  knitr::kable(digits = 3)
```

| term        | estimate | std.error | statistic | p.value |
|-------------|----------|-----------|-----------|---------|
| (Intercept) | 5.840    | 0.339     | 17.239    | 0.000   |
| avexpr      | 0.394    | 0.050     | 7.843     | 0.000   |
| lat_abst    | 0.312    | 0.444     | 0.703     | 0.484   |
| asia        | -0.170   | 0.153     | -1.108    | 0.270   |
| africa      | -0.930   | 0.165     | -5.628    | 0.000   |

# **2|** Presenting OLS regressions

# Regression tables

- In papers, you'll often find regression tables that have several models.
- Each column is a different regression:
  - Might differ by independent variables, dependent variables, sample, etc.
- Standard errors, p-values, sample size, and  $R^2$  may be reported as well.

# AJR regression table

VOL. 91 NO. 5

ACEMOGLU ET AL.: THE COLONIAL ORIGINS OF DEVELOPMENT

1379

TABLE 2—OLS REGRESSIONS

|  | Whole world<br>(1) | Base sample<br>(2) | Whole world<br>(3) | Whole world<br>(4) | Base sample<br>(5) | Base sample<br>(6) | Whole world<br>(7) | Base sample<br>(8) |
|--|--------------------|--------------------|--------------------|--------------------|--------------------|--------------------|--------------------|--------------------|
| Dependent variable is log GDP per capita in 1995         |                    |                    |                    |                    |                    |                    |                    |                    |
| Dependent variable is log output per worker in 1988      |                    |                    |                    |                    |                    |                    |                    |                    |
| Average protection against expropriation risk, 1985–1995 | 0.54<br>(0.04)     | 0.52<br>(0.06)     | 0.47<br>(0.06)     | 0.43<br>(0.05)     | 0.47<br>(0.06)     | 0.41<br>(0.06)     | 0.45<br>(0.04)     | 0.46<br>(0.06)     |
| Latitude   |                    |                    | 0.89<br>(0.49)     | 0.37<br>(0.51)     | 1.60<br>(0.70)     | 0.92<br>(0.63)     |                    |                    |
| Asia dummy   |                    |                    |                    | -0.62<br>(0.19)    |                    | -0.60<br>(0.23)    |                    |                    |
| Africa dummy   |                    |                    |                    | -1.00<br>(0.15)    |                    | -0.90<br>(0.17)    |                    |                    |
| “Other” continent dummy                                  |                    |                    |                    | -0.25<br>(0.20)    |                    | -0.04<br>(0.32)    |                    |                    |
| <i>R</i> <sup>2</sup>                                    | 0.62               | 0.54               | 0.63               | 0.73               | 0.56               | 0.69               | 0.55               | 0.49               |
| Number of observations                                   | 110                | 64                 | 110                | 110                | 64                 | 64                 | 108                | 61                 |

# `modelsummary()` to produce tables

We can use `modelsummary()` to produce a table. It takes a list of outputs from `lm` and aligns them in the correct way.

```
modelsummary::modelsummary(list(ajr.reg, ajr.multreg))
```

# Output

```
modelsummary::modelsummary(list(ajr.reg, ajr.multreg))
```

|             | Model 1          | Model 2           |
|-------------|------------------|-------------------|
| (Intercept) | 4.626<br>(0.301) | 5.840<br>(0.339)  |
| avexpr      | 0.532<br>(0.041) | 0.394<br>(0.050)  |
| lat_abst    |                  | 0.312<br>(0.444)  |
| asia        |                  | -0.170<br>(0.153) |
| africa      |                  | -0.930<br>(0.165) |
| Num.Obs.    | 111              | 111               |
| R2          | 0.611            | 0.713             |
| R2 Adj.     | 0.608            | 0.703             |
| AIC         | 245.4            | 217.6             |
| BIC         | 253.5            | 233.8             |
| Log.Lik.    | -119.709         | -102.795          |
| RMSE        | 0.71             | 0.61              |

# Cleaning up the goodness of fit statistics

```
modelsummary::modelsummary(  
  list(ajr.reg, ajr.multreg),  
  gof_map = c("nobs", "r.squared", "adj.r.squared"))
```

|             | Model 1          | Model 2           |
|-------------|------------------|-------------------|
| (Intercept) | 4.626<br>(0.301) | 5.840<br>(0.339)  |
| avexpr      | 0.532<br>(0.041) | 0.394<br>(0.050)  |
| lat_abst    |                  | 0.312<br>(0.444)  |
| asia        |                  | -0.170<br>(0.153) |
| africa      |                  | -0.930<br>(0.165) |
| Num.Obs.    | 111              | 111               |
| R2          | 0.611            | 0.713             |
| R2 Adj.     | 0.608            | 0.703             |

# Cleaning up the variable names

We can also map the variable names to more readable names using the `coef_map` argument. But first, we should do the mapping in a vector. Any term omitted from this vector will be omitted from the table

```
var_labels <- c(  
  "avexpr" = "Avg. Expropriation Risk",  
  "lat_abst" = "Abs. Value of Latitude",  
  "asia" = "Asian country",  
  "africa" = "African country"  
)  
var_labels
```

```
##                               avexpr                  lat_abst  
## "Avg. Expropriation Risk"  "Abs. Value of Latitude"  
##                               asia                   africa  
## "Asian country"            "African country"
```

# Nice table

```
modelsummary::modelsummary(  
  list(ajr.reg, ajr.multreg),  
  coef_map = var_labels,  
  gof_map = c("nobs", "r.squared", "adj.r.squared"))
```

|                         | Model 1          | Model 2           |
|-------------------------|------------------|-------------------|
| Avg. Expropriation Risk | 0.532<br>(0.041) | 0.394<br>(0.050)  |
| Abs. Value of Latitude  |                  | 0.312<br>(0.444)  |
| Asian country           |                  | -0.170<br>(0.153) |
| African country         |                  | -0.930<br>(0.165) |
| Num.Obs.                | 111              | 111               |
| R2                      | 0.611            | 0.713             |
| R2 Adj.                 | 0.608            | 0.703             |

# 3/ Wrapping up the class

# Big takeaways

Important takeaways from the course:

1. Data wrangling and data visualizations are really important skills that you now have!
2. Causality is hugely important in the world but difficult to establish.
3. Really important to understand and assess statistical uncertainty when working with data.

# I'm really proud of you!



You've come a long way! Hopefully the tools you learned in this course will help you throughout your life and career!

# What next?



- Gov 51 with Naijia Liu:
  - A more in-depth review of some ideas from Gov 50 including causality and regression plus new models (maybe some machine learning).
  - Really helpful for students looking to write senior theses.
- Only need 3 more classes to finish the data science track in Gov!
- More theoretical stats side: Stat 110/111
- More CS approach to data science: CS109 (Data Science 1)

# Thanks!



Fill out your evaluations!